



中国科学院计算技术研究所  
INSTITUTE OF COMPUTING TECHNOLOGY, CHINESE ACADEMY OF SCIENCES



北京开源芯片研究院  
BEIJING INSTITUTE OF OPEN SOURCE CHIP



香山开源处理器社区  
XiangShan Open-Source Processor Community

*At the 58th IEEE/ACM International Symposium on Microarchitecture (MICRO 2025)*

# DiffTest-H: Toward Semantic-Aware Communication in Hardware-Accelerated Processor Verification

---

**Kunlin You<sup>1,2</sup>, Yinan Xu<sup>1</sup>, Kehan Feng<sup>3</sup>, Luoshan Cai<sup>1,2</sup>, Yaoyang Zhou<sup>3</sup> and Yungang Bao<sup>1,2</sup>**

*<sup>1</sup> State Key Lab of Processors, Institute of Computing Technology, Chinese Academy of Sciences*


*<sup>2</sup> University of Chinese Academy of Sciences*

*<sup>3</sup> Beijing Institute of Open Source Chip*


October 21, 2025

# XiangShan: Open-source High Performance Processors

## • Top-Performing Open-Source Processor



## • Vision: “Linux of Processors” for Industry & Academia



### Research on XiangShan

XiangShan has been used by researchers as the underlying platform for their evaluations. We appreciate their contributions to enhancing XiangShan and strengthening the community.

- Imprecise Store Exceptions, EPFL, ISCA'23
- TEEsec: Pre-Silicon Vulnerability Discovery for Trusted Execution Environments, OSU & SUSTech, ISCA'23

### Research Platform

- of CPU, ICT-CAS, ICCAD'23
- A Distributed ATPG System Combining Test Compaction Based on Pure MaxSAT, ICT-CAS, ATS'23
- REMU: Enabling Cost-Effective Checkpointing and Deterministic Replay in FPGA-based Emulation, ICT-CAS, ICCD'23
- Asynchronous Memory Access Unit: Exploiting Massive Parallelism for Far Memory Access, ICT-CAS, TACO
- Single-Address-Space FaaS with Jord, EPFL, ISCA'25

# XiangShan: Powering Desktop, GPU SoC, Server ...


1st Gen  
XiangShan  
Chip



与开芯院南湖RISC-V CPU协同开发，开创RISC-V+AI新生态

**XiangShan on GPU  
@InnoSilicon**

- ✓ 计算数据调度预处理
- ✓ 芯片功耗控制
- ✓ 跨芯片数据通讯控制
- ✓ 温度监控和风扇管理
- ✓ 电源管理
- ✓ BMC控制
- ✓ 异常监控



The world's first laptop powered by  
a open-source RISC-V processor

**XiangShan Laptop  
@ISCAS**

Ruyi Book	
CPU	"XiangShan Nanhu" (RV64GCBK), up to 2.5GHz
Memory	8GB DDR5 4800MT/s
GPU	AMD RX 550
USB	2x USB3
Ethernet	2x 2.5Gbps Ethernet Port
Display	1x 14-inch LCD Display

```
hw_001: 0000000000000000: 0 bytes: 0x00000000
hw_002: 0000000000000000: 0 bytes: 0x00000000
hw_003: 0000000000000000: 0 bytes: 0x00000000
hw_004: 0000000000000000: 0 bytes: 0x00000000
hw_005: 0000000000000000: 0 bytes: 0x00000000
hw_006: 0000000000000000: 0 bytes: 0x00000000
hw_007: 0000000000000000: 0 bytes: 0x00000000
hw_008: 0000000000000000: 0 bytes: 0x00000000
hw_009: 0000000000000000: 0 bytes: 0x00000000
hw_010: 0000000000000000: 0 bytes: 0x00000000
hw_011: 0000000000000000: 0 bytes: 0x00000000
hw_012: 0000000000000000: 0 bytes: 0x00000000
hw_013: 0000000000000000: 0 bytes: 0x00000000
hw_014: 0000000000000000: 0 bytes: 0x00000000
hw_015: 0000000000000000: 0 bytes: 0x00000000
hw_016: 0000000000000000: 0 bytes: 0x00000000
hw_017: 0000000000000000: 0 bytes: 0x00000000
hw_018: 0000000000000000: 0 bytes: 0x00000000
hw_019: 0000000000000000: 0 bytes: 0x00000000
hw_020: 0000000000000000: 0 bytes: 0x00000000
hw_021: 0000000000000000: 0 bytes: 0x00000000
hw_022: 0000000000000000: 0 bytes: 0x00000000
hw_023: 0000000000000000: 0 bytes: 0x00000000
hw_024: 0000000000000000: 0 bytes: 0x00000000
hw_025: 0000000000000000: 0 bytes: 0x00000000
hw_026: 0000000000000000: 0 bytes: 0x00000000
hw_027: 0000000000000000: 0 bytes: 0x00000000
hw_028: 0000000000000000: 0 bytes: 0x00000000
hw_029: 0000000000000000: 0 bytes: 0x00000000
hw_030: 0000000000000000: 0 bytes: 0x00000000
hw_031: 0000000000000000: 0 bytes: 0x00000000
hw_032: 0000000000000000: 0 bytes: 0x00000000
hw_033: 0000000000000000: 0 bytes: 0x00000000
hw_034: 0000000000000000: 0 bytes: 0x00000000
hw_035: 0000000000000000: 0 bytes: 0x00000000
hw_036: 0000000000000000: 0 bytes: 0x00000000
hw_037: 0000000000000000: 0 bytes: 0x00000000
hw_038: 0000000000000000: 0 bytes: 0x00000000
hw_039: 0000000000000000: 0 bytes: 0x00000000
hw_040: 0000000000000000: 0 bytes: 0x00000000
hw_041: 0000000000000000: 0 bytes: 0x00000000
hw_042: 0000000000000000: 0 bytes: 0x00000000
hw_043: 0000000000000000: 0 bytes: 0x00000000
hw_044: 0000000000000000: 0 bytes: 0x00000000
hw_045: 0000000000000000: 0 bytes: 0x00000000
hw_046: 0000000000000000: 0 bytes: 0x00000000
hw_047: 0000000000000000: 0 bytes: 0x00000000
hw_048: 0000000000000000: 0 bytes: 0x00000000
hw_049: 0000000000000000: 0 bytes: 0x00000000
hw_050: 0000000000000000: 0 bytes: 0x00000000
hw_051: 0000000000000000: 0 bytes: 0x00000000
hw_052: 0000000000000000: 0 bytes: 0x00000000
hw_053: 0000000000000000: 0 bytes: 0x00000000
hw_054: 0000000000000000: 0 bytes: 0x00000000
hw_055: 0000000000000000: 0 bytes: 0x00000000
hw_056: 0000000000000000: 0 bytes: 0x00000000
hw_057: 0000000000000000: 0 bytes: 0x00000000
hw_058: 0000000000000000: 0 bytes: 0x00000000
hw_059: 0000000000000000: 0 bytes: 0x00000000
hw_060: 0000000000000000: 0 bytes: 0x00000000
hw_061: 0000000000000000: 0 bytes: 0x00000000
hw_062: 0000000000000000: 0 bytes: 0x00000000
hw_063: 0000000000000000: 0 bytes: 0x00000000
hw_064: 0000000000000000: 0 bytes: 0x00000000
hw_065: 0000000000000000: 0 bytes: 0x00000000
hw_066: 0000000000000000: 0 bytes: 0x00000000
hw_067: 0000000000000000: 0 bytes: 0x00000000
hw_068: 0000000000000000: 0 bytes: 0x00000000
hw_069: 0000000000000000: 0 bytes: 0x00000000
hw_070: 0000000000000000: 0 bytes: 0x00000000
hw_071: 0000000000000000: 0 bytes: 0x00000000
hw_072: 0000000000000000: 0 bytes: 0x00000000
hw_073: 0000000000000000: 0 bytes: 0x00000000
hw_074: 0000000000000000: 0 bytes: 0x00000000
hw_075: 0000000000000000: 0 bytes: 0x00000000
hw_076: 0000000000000000: 0 bytes: 0x00000000
hw_077: 0000000000000000: 0 bytes: 0x00000000
hw_078: 0000000000000000: 0 bytes: 0x00000000
hw_079: 0000000000000000: 0 bytes: 0x00000000
hw_080: 0000000000000000: 0 bytes: 0x00000000
hw_081: 0000000000000000: 0 bytes: 0x00000000
hw_082: 0000000000000000: 0 bytes: 0x00000000
hw_083: 0000000000000000: 0 bytes: 0x00000000
hw_084: 0000000000000000: 0 bytes: 0x00000000
hw_085: 0000000000000000: 0 bytes: 0x00000000
hw_086: 0000000000000000: 0 bytes: 0x00000000
hw_087: 0000000000000000: 0 bytes: 0x00000000
hw_088: 0000000000000000: 0 bytes: 0x00000000
hw_089: 0000000000000000: 0 bytes: 0x00000000
hw_090: 0000000000000000: 0 bytes: 0x00000000
hw_091: 0000000000000000: 0 bytes: 0x00000000
hw_092: 0000000000000000: 0 bytes: 0x00000000
hw_093: 0000000000000000: 0 bytes: 0x00000000
hw_094: 0000000000000000: 0 bytes: 0x00000000
hw_095: 0000000000000000: 0 bytes: 0x00000000
hw_096: 0000000000000000: 0 bytes: 0x00000000
hw_097: 0000000000000000: 0 bytes: 0x00000000
hw_098: 0000000000000000: 0 bytes: 0x00000000
hw_099: 0000000000000000: 0 bytes: 0x00000000
hw_100: 0000000000000000: 0 bytes: 0x00000000
```

**4-core XiangShan  
@SpacemiT**




2nd Gen  
XiangShan Chip

Gaming on 2nd Gen  
XiangShan Chip



**8-core XiangShan  
@LanXin Computing**




Fedora on XiangShan  
@Fedora-V Force




# Open-source Demands Agile Verification

- Highly configurable ISA
- Rapid code iterations
- **1,400+** bugs uncovered in XiangShan (2025)



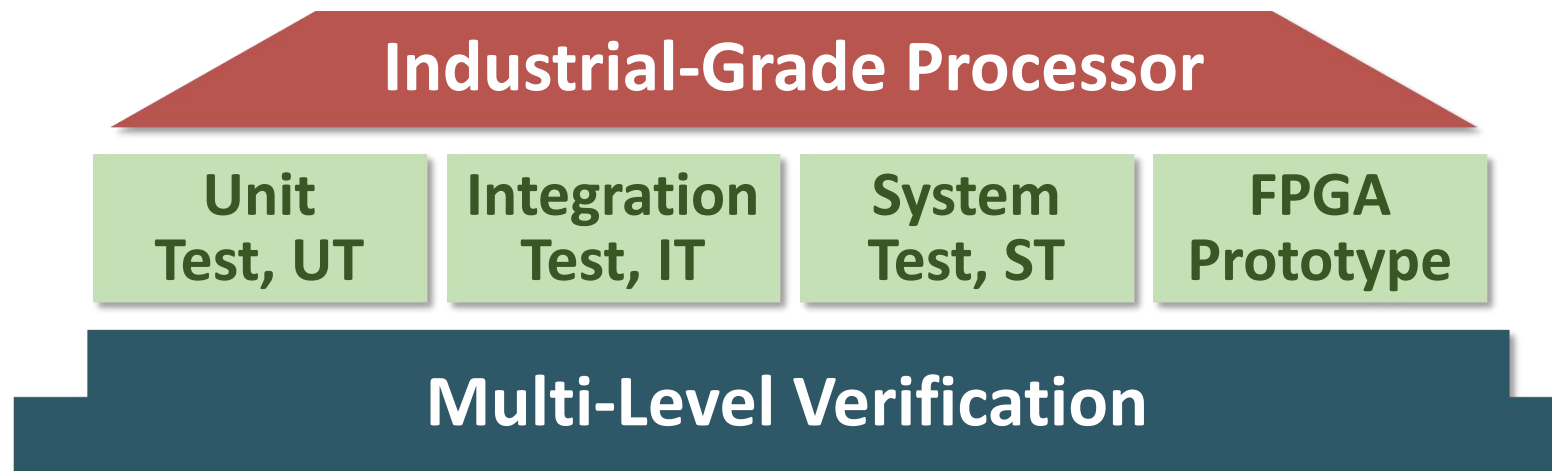
**Bugs Report  
(2025@XiangShan)**



■ XiangShan ■ Partners



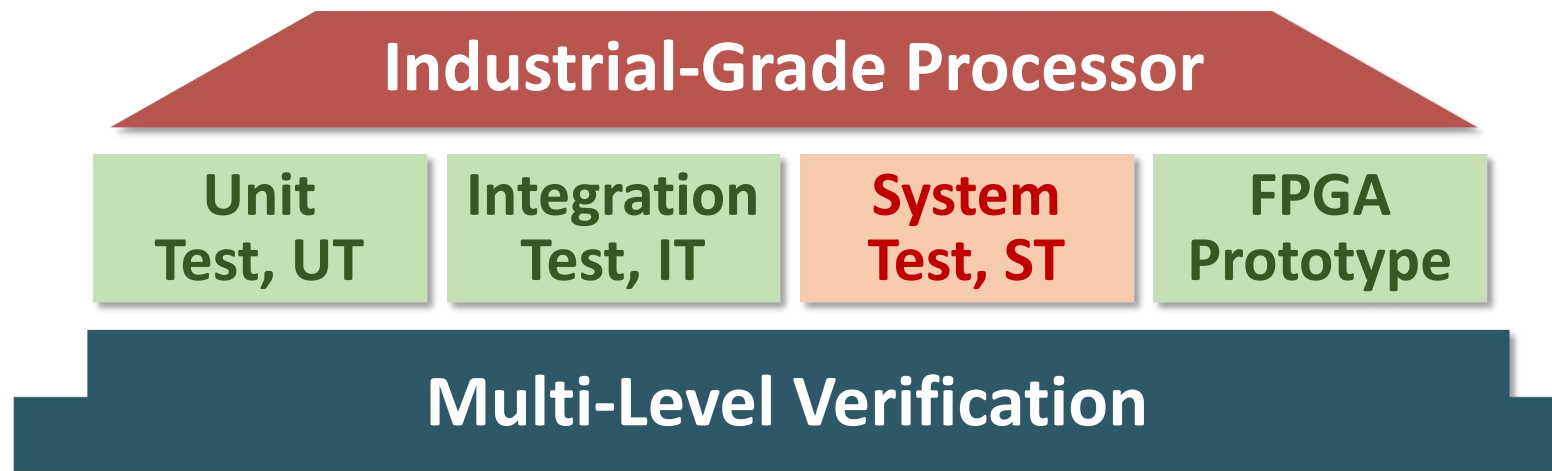
# The Era of Verification



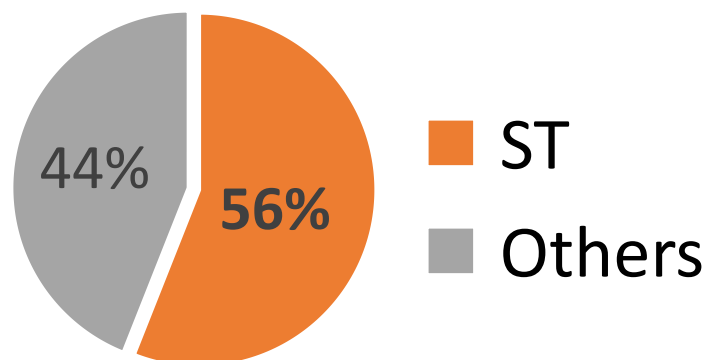




# The Era of Verification



Covering **56% bugs** of XiangShan

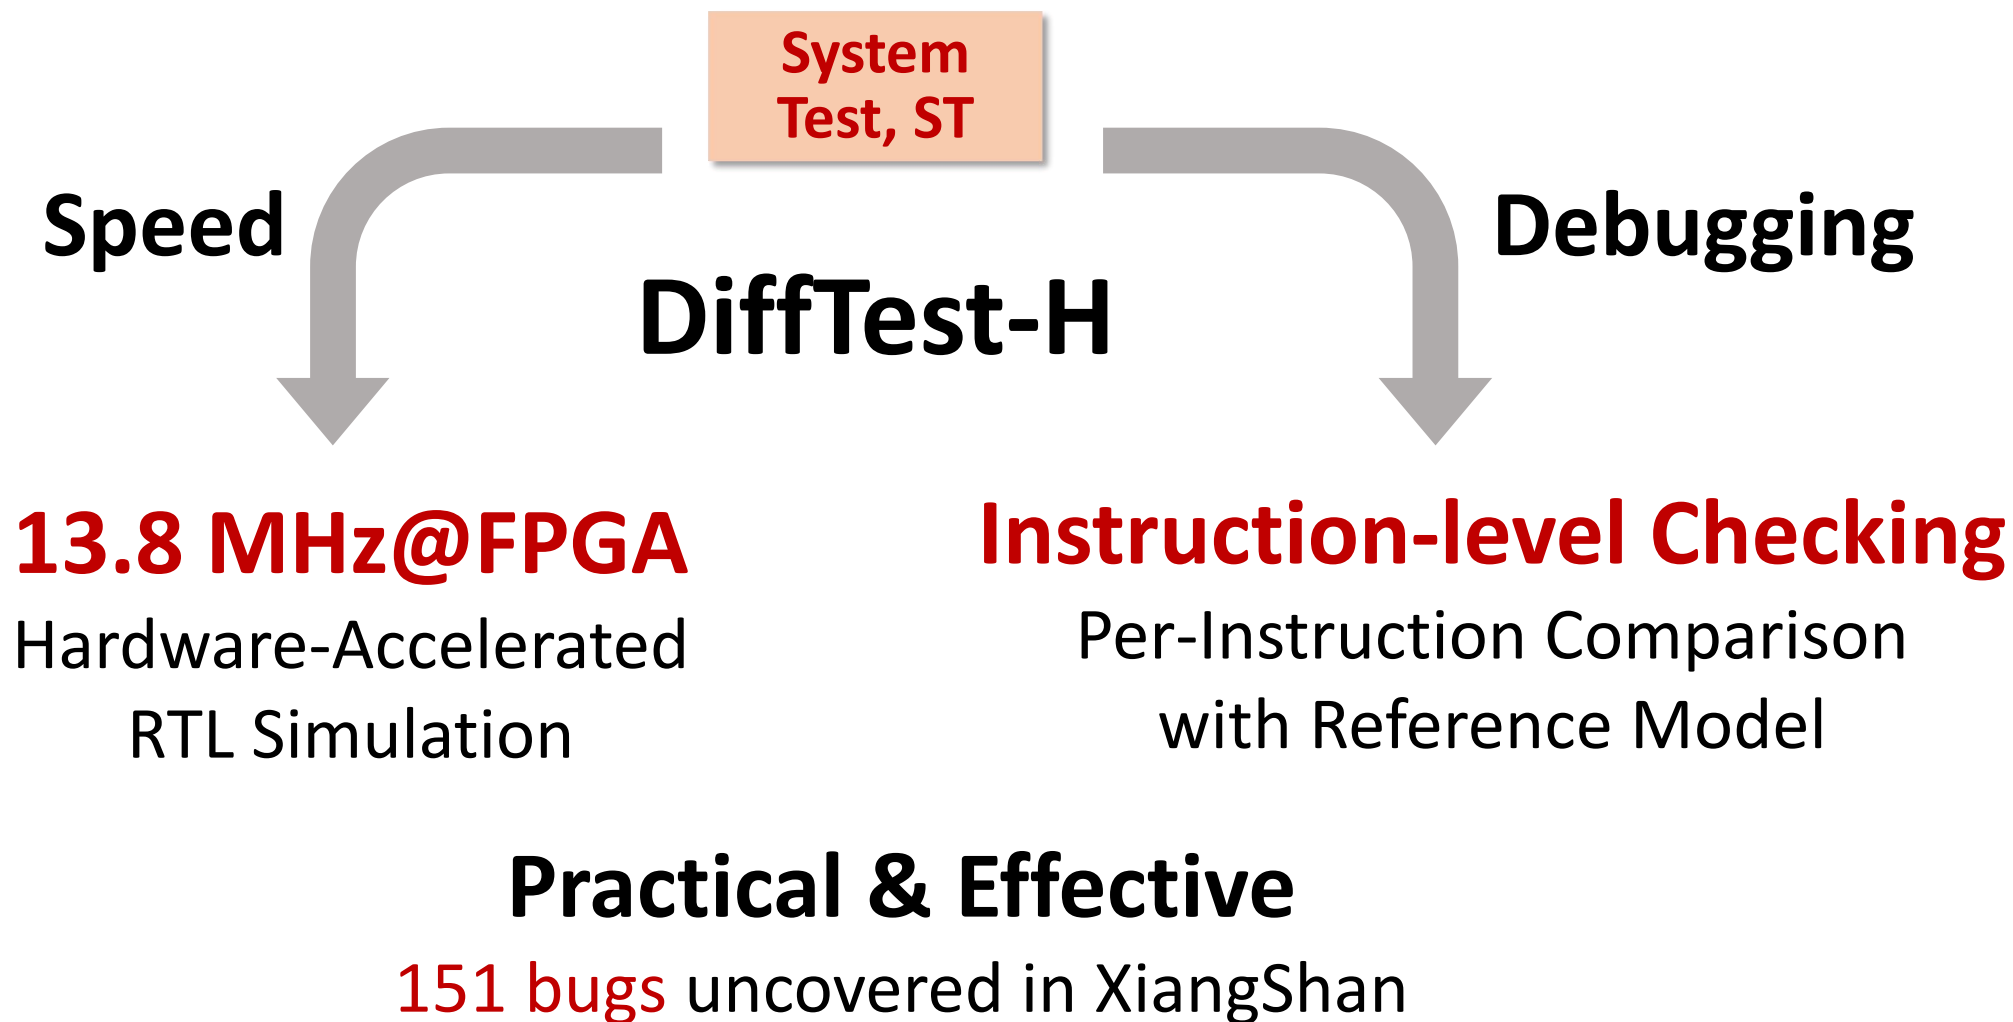


Checking under **real-world cases**





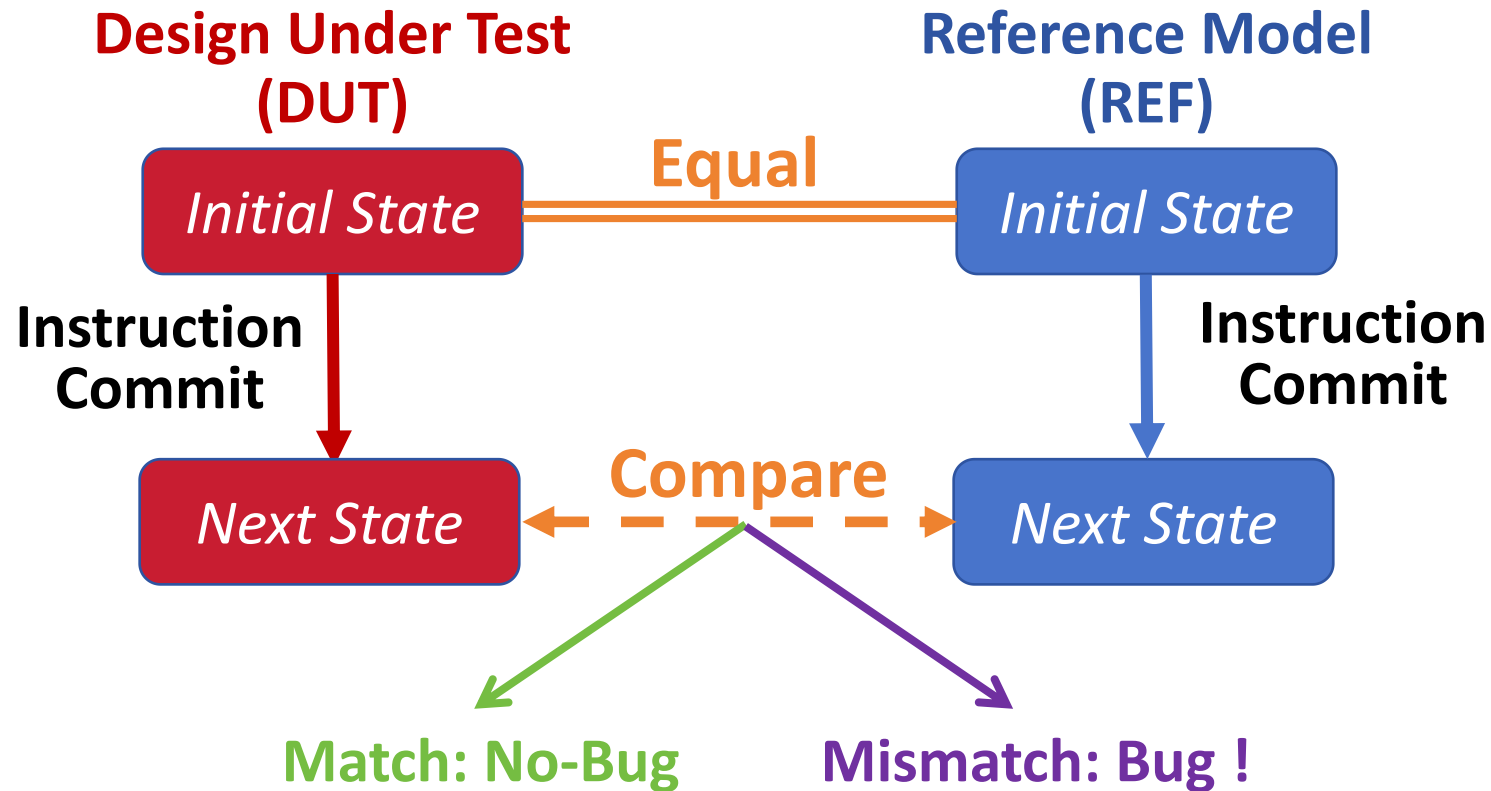
# This Work: Fast & Debug-friendly System Test





# Co-simulation: Instruction-level Debugging

- DiffTest<sup>[1]</sup>: Instr-by-Instr Comparison of Architectural State (DUT vs. REF)

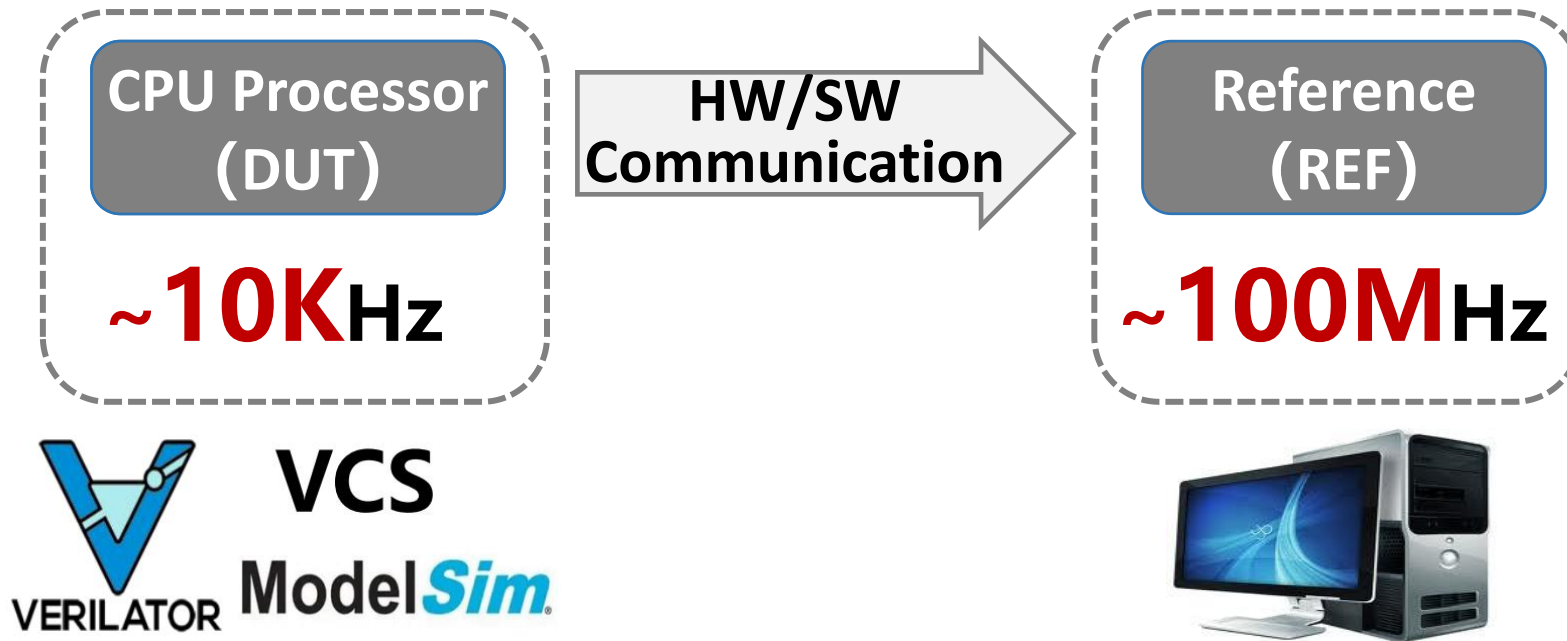


**Well debugging, but what about speed?**



# Software-based Co-simulation

- Based on **RTL Simulator**:
  - Verilator, VCS: **only KHz-scale speed** (limited by RTL simulation)
  - Even with optimizations<sup>[1,2,3]</sup>, still limited at **KHz-scale**



[1] Chen et al. GSIM: Accelerating RTL Simulation for Large-Scale Designs. (DAC 2025)

[2] Wang et al. Reput: Superlinear parallel rtl simulation with replication-aided partitioning. (ASPLOS 2023)

[3] Zhou et al. Khronos: Fusing memory access for improved hardware RTL simulation. (MICRO 2023)

# Hardware-accelerated Co-simulation

- Based on **Emulator/FPGA**:
  - Emulator/FPGA: DUT at 1~100 MHz, **up to 10,000× speedup**
  - Reference Model (software): ~100 MHz




# Speed GAP: HW/SW Communication

**Ideal: DUT Speed**

**Up to 10,000×  
Speedup**

Emulator: 1 MHz  
FPGA: 100 MHz



**Reality: Co-sim Speed**

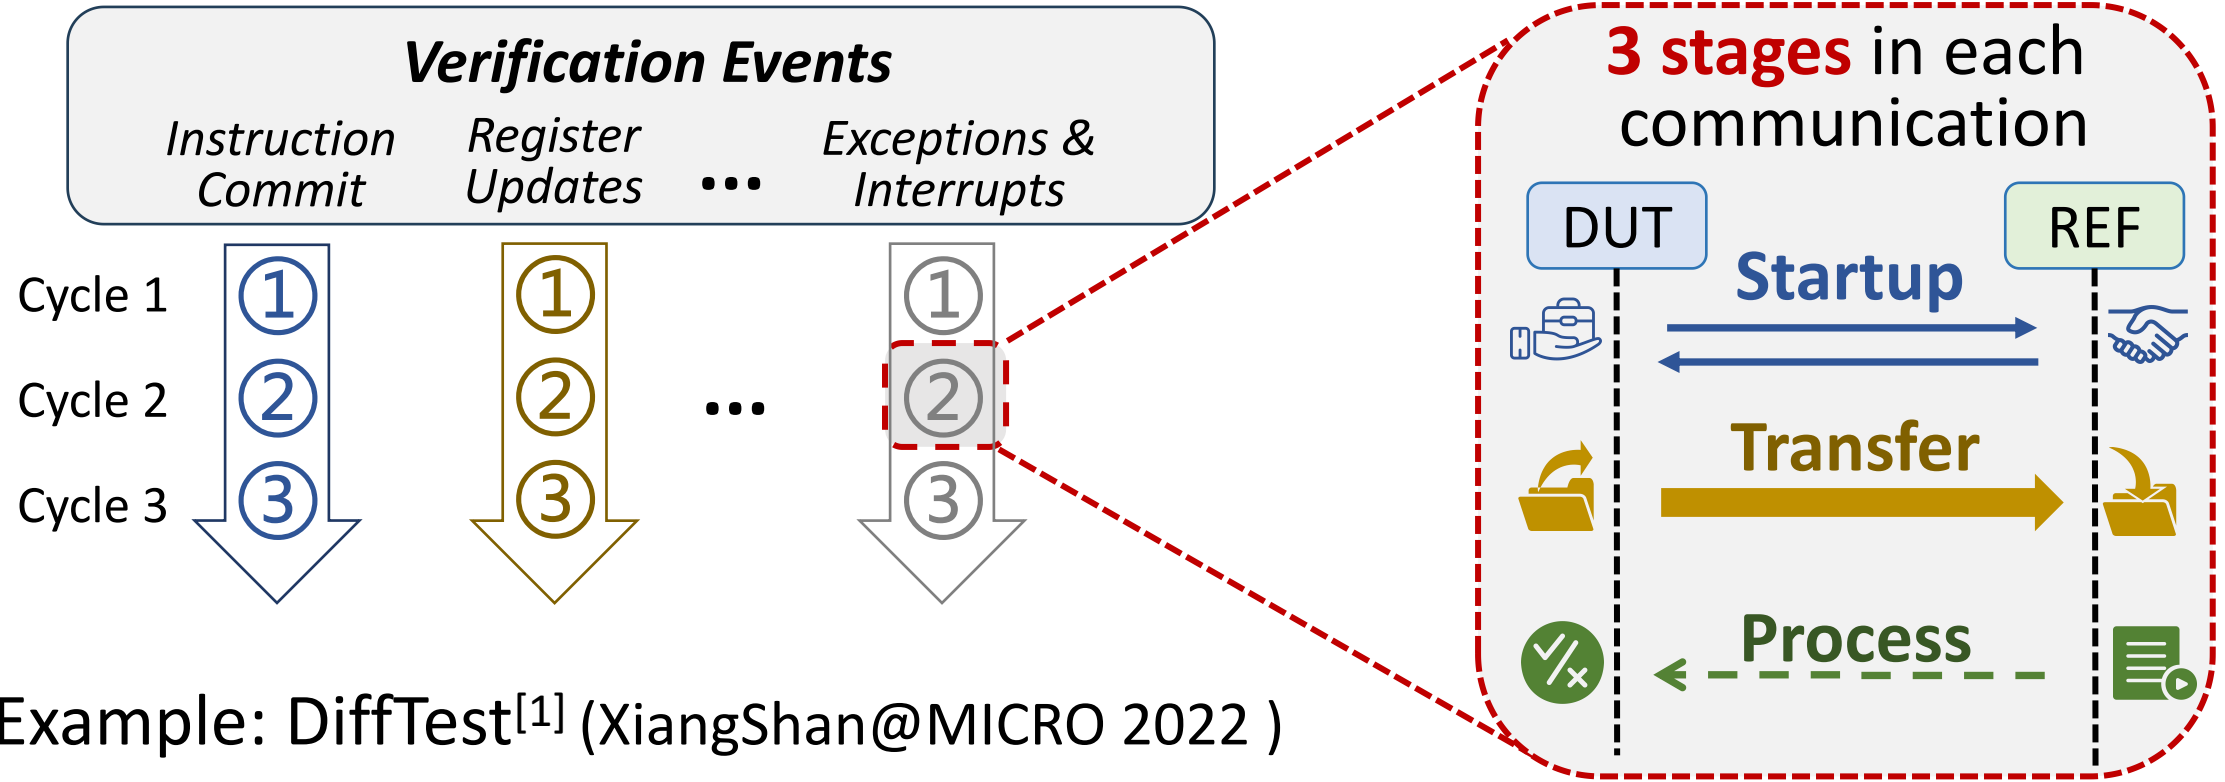
**Less than 20×  
Speedup**

Emulator: < 10 KHz  
FPGA: < 100 KHz

**GAP: HW/SW  
Communication**


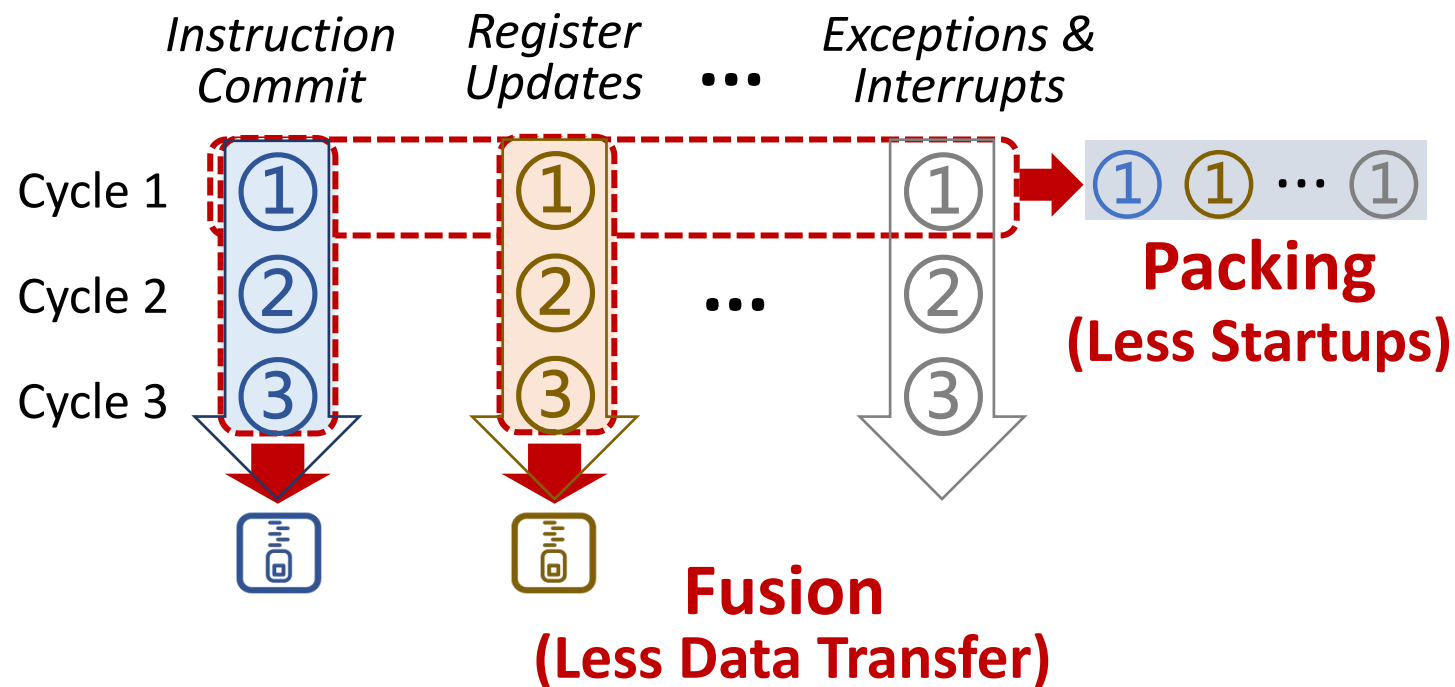
**98%~99.8% of co-sim time**

# What is HW/SW Communication?



- Example: DiffTest<sup>[1]</sup> (XiangShan@MICRO 2022 )
  - **32 types** of architectural events
  - 1 event trigger 1 communication
  - **~15 communications, ~1.2 KB data** per cycle

# How to optimize 3-stage communication?




Focus on **Startup & Transfer**

*Process latency is hidden by Async Transmission*




# Toward Fast & Debug-friendly Communication



- **Challenge 1:** Packing Under Structural Diversity
  - 170× Length Diversity + Per-cycle Event Variation




# Toward Fast & Debug-friendly Communication



- **Challenge 1:** Packing Under Structural Diversity
  - $170\times$  Length Diversity + Per-cycle Event Variation
- **Challenge 2:** Fusion Under Order Constraints
  - Same-type Fusion vs. Cross-type Ordering?



# Toward Fast & Debug-friendly Communication



- **Challenge 1:** Packing Under Structural Diversity
  - 170× Length Diversity + Per-cycle Event Variation
- **Challenge 2:** Fusion Under Order Constraints
  - Same-type Fusion vs. Cross-type Ordering?
- **Challenge 3:** Debugging after Fusion
  - Instr-level Debugging vs. Fusion Detail Loss?



# DiffTest-H Overview

A Semantic-aware, Hardware-accelerated Framework  
Toward **Fast & Debug-friendly** Communication

*Challenge 1 (Frequency)*  
Packing Under Structural Diversity

***Batch: Structure-wise Packing***

*Challenge 2 (Data Volume)*  
Fusion Under Order Constraints



***Squash: Order-decoupling Fusion***

*Challenge 3 (Debuggability)*  
Debugging after Fusion

***Replay: Instruction-level Debugging***



# Batch: Packing under Structural Diversity



HW Pack: Offset from length  
SW Unpack: Extract by type

## Prior: Fixed Space

Padding invalid space with bubbles

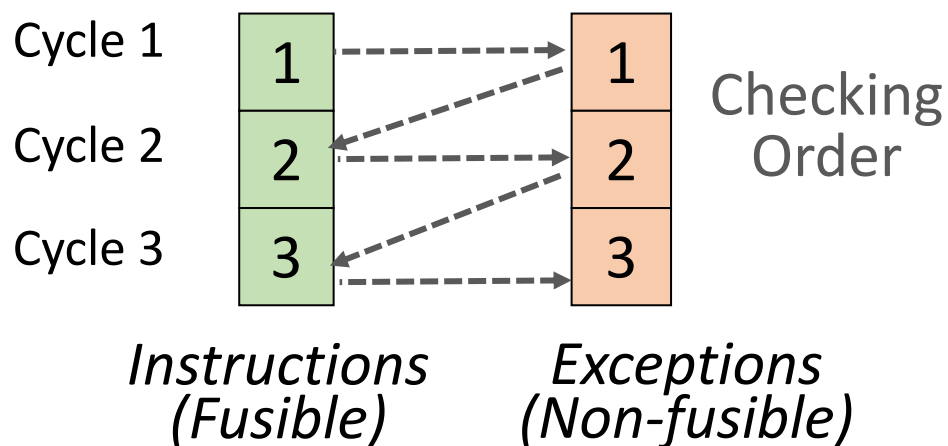
**60%** Bubbles, **1.67×** Frequency

## Now: Structure-wise Packing

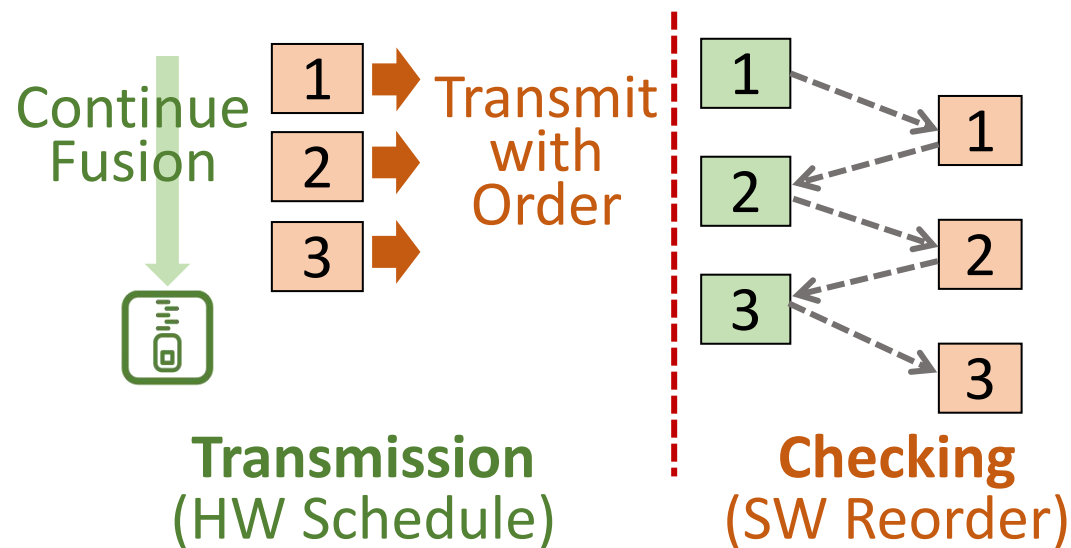
Packing with length & type

No bubble, Less frequency

# Squash: Order-decoupling Fusion

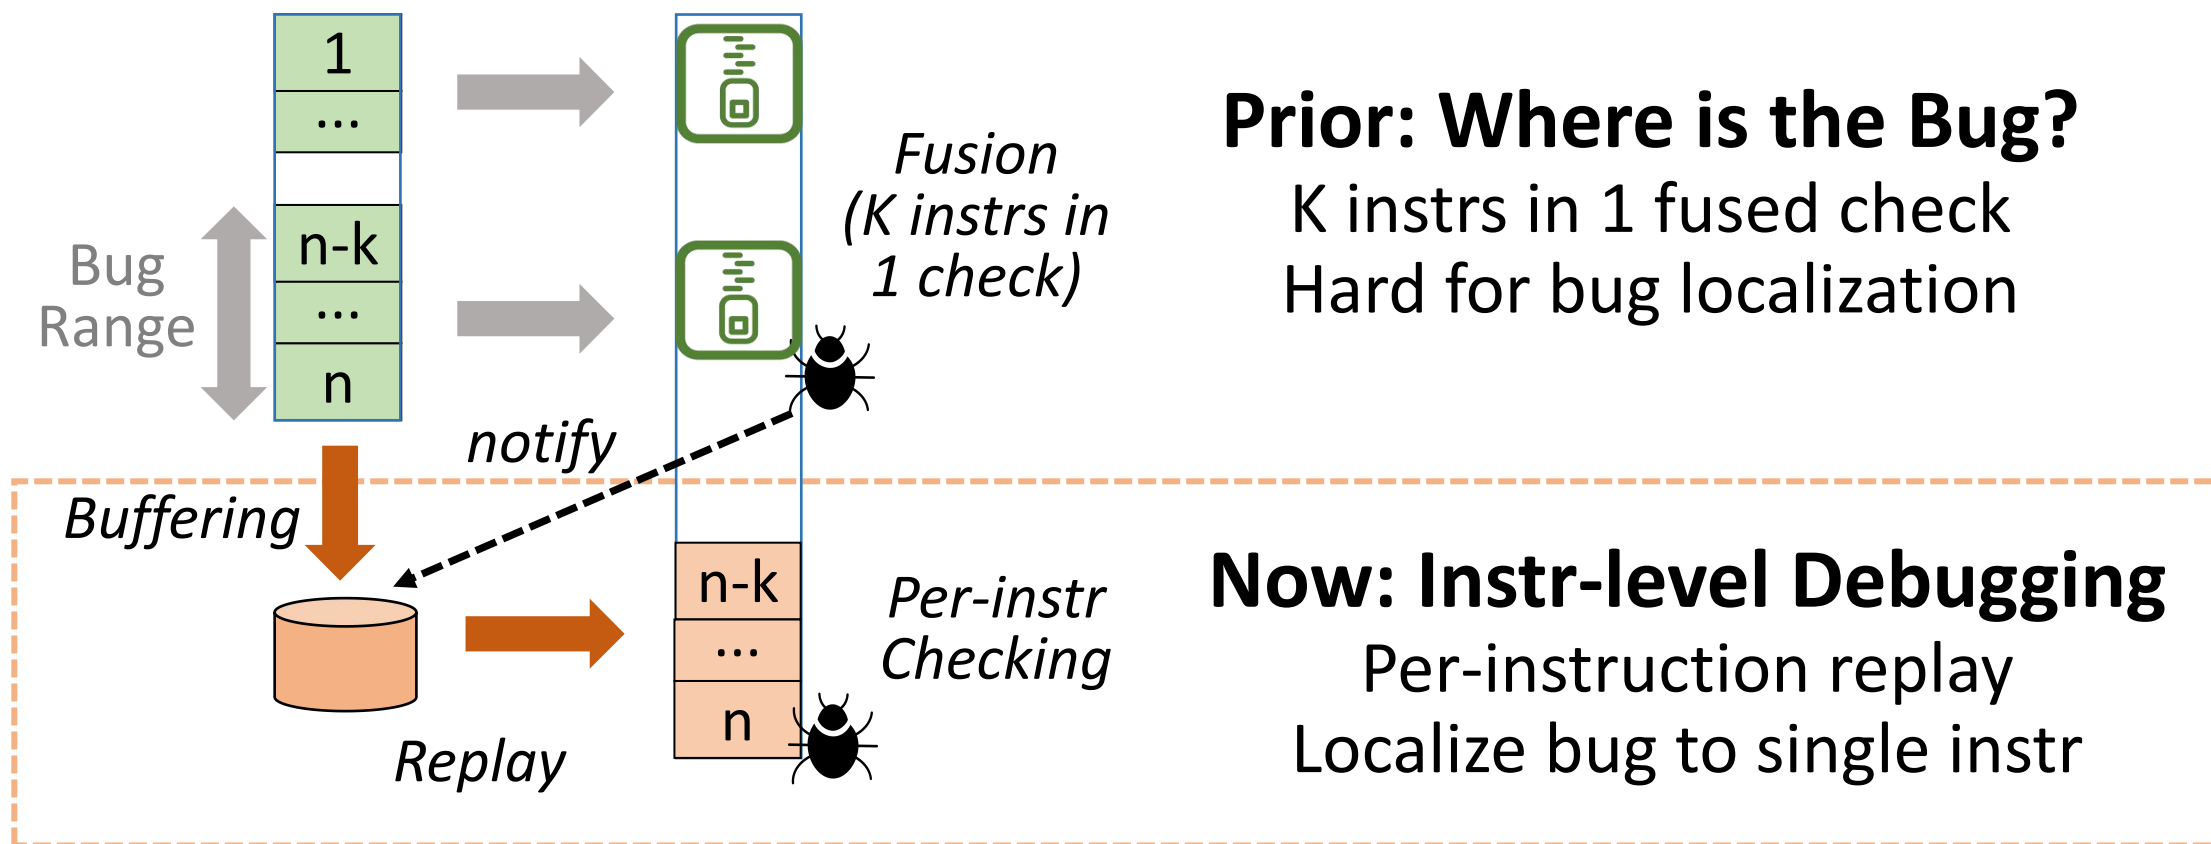


**Prior: Order Constraints**  
e.g. Exception after specific instrs  
Frequent Fusion Breaks



**Now: Order-decoupling Fusion**  
Decouple fusion & checking order  
Better fusion ratio, Less Data

# Replay: Instruction-level Debugging







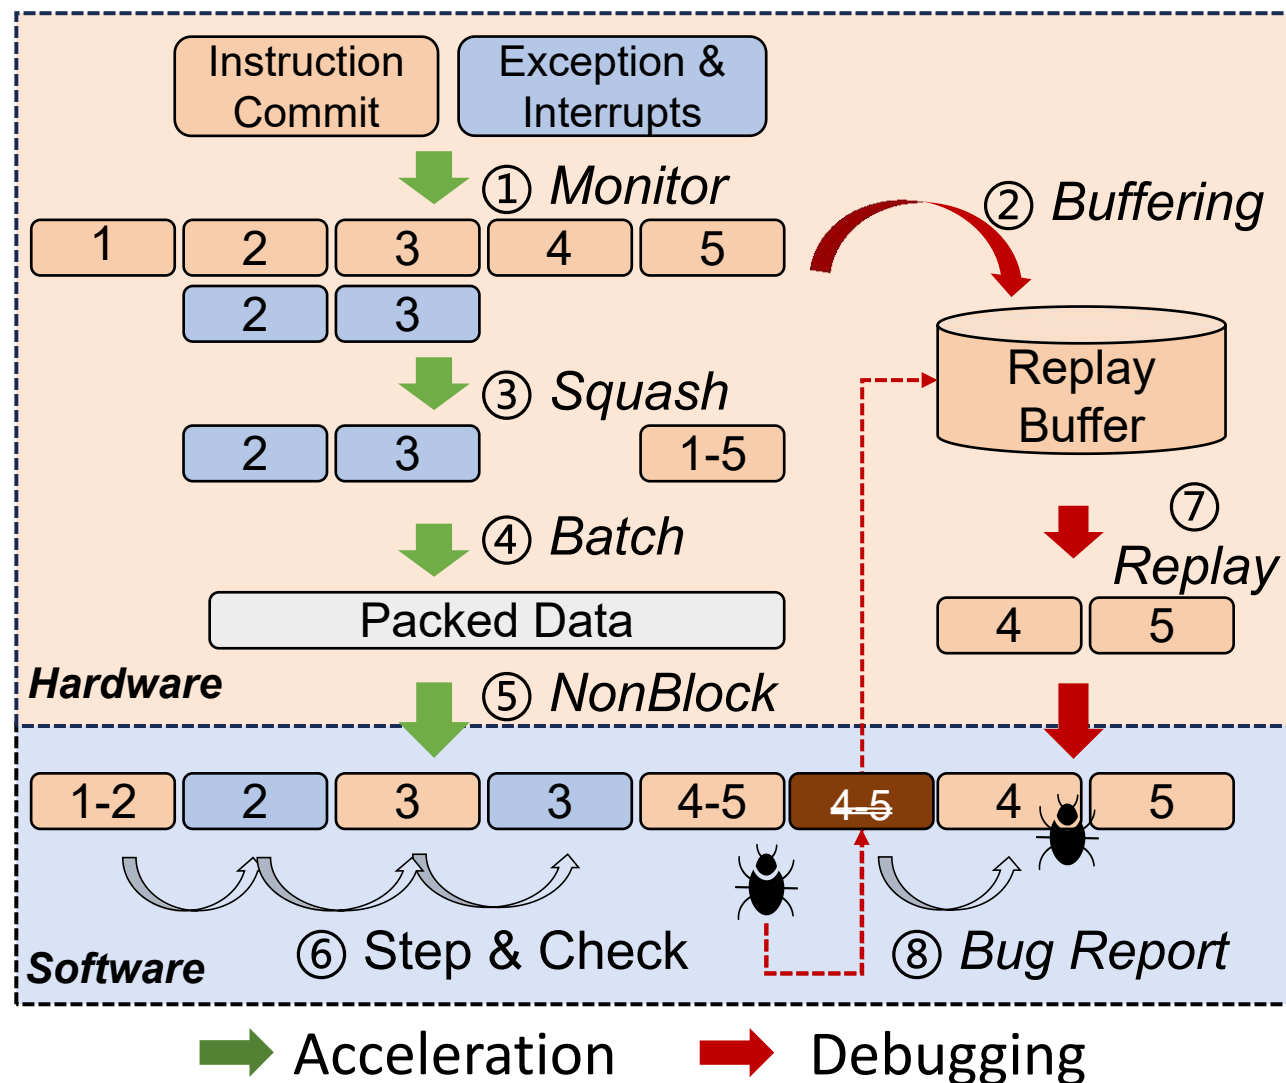
# DiffTest-H Workflow

Accelerated Co-simulation:

- ① **Monitor**: Capture DUT events
- ③ **Squash**: Fusion for less data
- ④ **Batch**: Packing for less frequency
- ⑤ **Nonblock**: Overlap SW latency
- ⑥ **Check**: Compare DUT vs. REF

Debugging on Error:


- ② **Buffering**: Backup unfused events
- ⑦ **Replay**: Recheck unfused events
- ⑧ **Report**: Bug localization



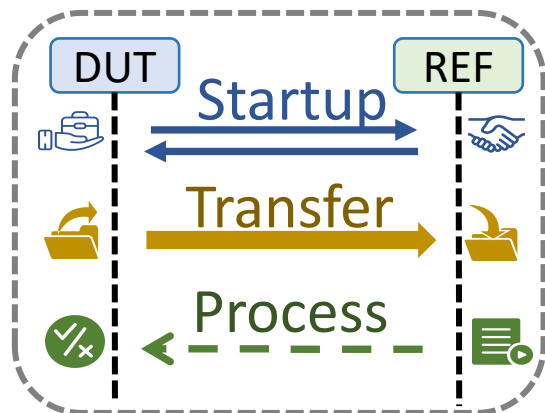


# Performance Evaluation


XiangShan: **13.8 MHz@FPGA**, **2,300× faster** than Verilator



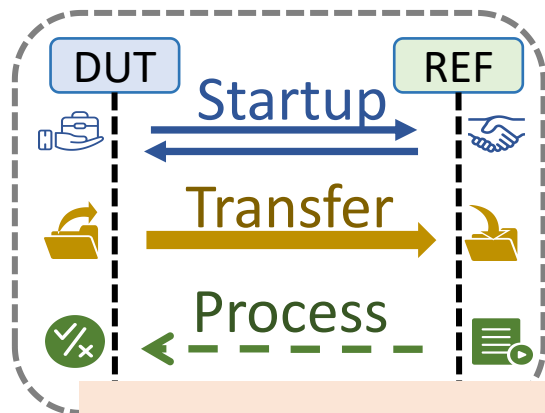
# Optimization Breakdown



- **Frequency:** Batch reduces by **43x**
- **Process Latency:** NonBlock overlaps with async
- **Data Volume:** Squash reduces by **47x**




# Optimization Breakdown



- **Frequency:** Batch reduces by **43x**
- **Process Latency:** NonBlock overlaps with async
- **Data Volume:** Squash reduces by **47x**

***74 x-138 x speedup, 98.0%-99.8% overhead reduction***



## Comparison: Fast & Debug-friendly

- **13.8× faster** than FPGA SOTA
- **32** verification event with **1.2KB** size, **Instruction-level** Debugging

Platform	Work	Checking Types	Checking Size B/cycle	Co-sim Speed
Emulator	IBI-Check <sup>[1]</sup>	2	7	80 KHz
	<b>This work</b>	32	1200	<b>478 KHz</b>
FPGA	Fromajo <sup>[2]</sup>	7	24	1 MHz
	<b>This work</b>	32	1200	<b>13.8 MHz</b>

[1] Chatterjee et al. Checking architectural outputs instruction-by-instruction on acceleration platforms. (DAC'12)

[2] Zhang et al. Integrating a High-Performance Instruction Set Simulator with FireSim to Co-simulate Operating System Boots. (ASPLOS '23 Workshops)



# Bug Discovery: 151 Bugs Uncovered in XiangShan

- In the past 6 months, uncover **151 Bugs** in XiangShan
- Over **780+** lines change with 19 Pull Request

fix(LoadUnit): fix misalign load wrong wakeup ✓

type: bug/confirmed

#4263 by cz4e was merged on Feb 16

fix(LSU): fix misalign store exception logic ✓

#4239 by Anzooooo was merged on Jan 26

fix(CSR): fix CSR misalign load/store ✓

#4157 by sinceforY was merged on Jan 10

fix(interrupt): vset should not respond to interrupts ✓

#3991 by Anzooooo was merged on Dec 7, 2024

fix(VecExcp): isEnqExcp should be set 0 when writeback has older exception ✓

#3778 by huxuan0307 was merged on Oct 24, 2024

fix(exception): fix exception vaddr generate logic ✓

#3639 by good-circle was merged on Sep 27, 2024

fix(LoadUnit): uncache should not be generated when page fault ✓

#4442 by Anzooooo was merged on Mar 20

fix(StoreUnit): no uncache store misalign of mmio ✓

#4441 by Anzooooo was merged on Mar 20

fix(StoreQueue): fix the vecExceptionFlag setting

fix(TLB): avoid flush when TLB hit

#3964 by cedrobot was merged on Dec 2, 2024

fix(ICache): block waylookup if there is a pending gpf ✓

#3719 by ngc7331 was merged on Oct 12, 2024

fix(TLB, RVH): delete the s1tagfix which maybe cause the tag check to fail ✓

#3685 by pxk27 was merged on Sep 30, 2024

submodule(CoupledL2): bump CoupledL2 ✓

#3621 by linjuanZ was merged on Sep 21, 2024

fix(vector): do not set vs.dirty for some type of vecInsts ✓

#3965 by Tang-Haojin was merged on Dec 1, 2024

fix(vstart): fix vstart wrong update when other instruction handling interrupt ✓

#3887 by Tang-Haojin was merged on Nov 18, 2024

fix(vstart): fix vstart wrong update when other instruction handling interrupt ✓

#3876 by Tang-Haojin was merged on Nov 18, 2024

feat(rv64v): fix exception for fof/non-fof load ✓

#3695 by huxuan0307 was merged on Dec 2, 2024

feat(rv64v): Added support for vector's vstart, first-only-fault, trigger ✓

feat(DiffTest): add multi-core vector load check ✓

#4361 by Anzooooo was merged on Mar 7

fix(vtypegen): fix initial condition after receive redirect ✓

#3664 by Ziyue-Zhang was merged on Sep 27, 2024

fix(VLSU): fix bug in flush of pipeline connect & skid buffer ✓

#3646 by weidingliu was merged on Sep 26, 2024

## Exception & Interrupts

## Memory hierarchy & coherence

## Vector Control Logic





# DiffTest-H Summary



- **DiffTest-H: Semantic-Aware, Hardware-Accelerated Framework**

- **13.8 MHz@FPGA, Instruction-level Debugging**
- Deployed in **XiangShan**, with **151 bugs uncovered**
- **Open-Source** @*github/OpenXiangShan/difftest*



**GitHub Repo**  
**DiffTest-H**

***Batch***

Structure-wise Packing

***Squash***

Order-decoupling Fusion

***Replay***

Instruction-level Debugging



**XIANGSHAN**

# Thanks for Your Attention