



XiangShan: an Open-Source High-Performance RISC-V Processor

Yungang Bao

Institute of Computing Technology (ICT)

Chinese Academy of Sciences (CAS)

Dec 6, 2021@RISC-V Summit

Outline

• Introduction of XiangShan

- Agile Performance Modelling with Software-based RTL-simulation
- Micro-architecture Design and Optimizations of XiangShan

🌟 Open-source Processor Ecosystem

① ISA + ② Microarch. design/impl. + ③ Workflow/Tools

Open Framework

microarch.

Open Flow

coding

Open Tools

EDA Tools

ISA Spec.

Docs

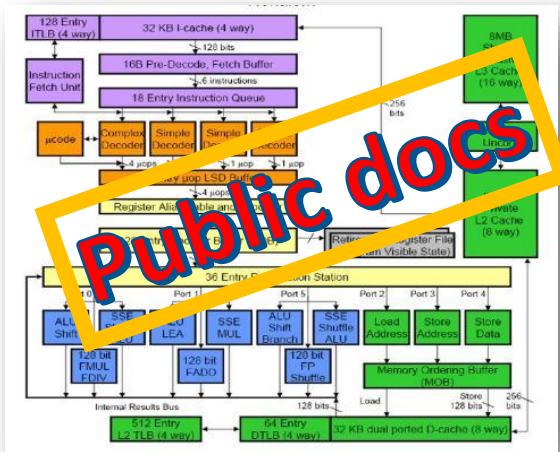
RTL

Layout


Open & Free

public docs

Open-source



```
component DebugCoreTop is
port (
  -- Trigger and Data
  cu_Clk       : in   std_logic_vector(1 downto 0) := (others => '0');
  cu0_Trig     : in   t_trig_0 := (others => (others => '0'));
  cu1_Trig     : in   t_trig_1 := (others => (others => '0'));
  cu2_Trig     : in   t_trig_2 := (others => (others => '0'));
  cu0_Data     : in   t_data_0 := (others => (others => '0'));
  cu1_Data     : in   t_data_1 := (others => (others => '0'));
  cu2_Data     : in   t_data_2 := (others => (others => '0'));
  -- Debug
  SDebugClk_n : in   std_logic := '0';
  SDebugInout : inout std_logic := '0';
  -- Upstream
  e_RefClk_p  : in   std_logic := '0';
  gt_RxClk_n  : in   std_logic := '0';
  gt_RX_p     : in   std_logic_vector(2 downto 0) := (others => '0');
  gt_RX_n     : in   std_logic_vector(2 downto 0) := (others => '0');
  gt_TX_p     : out  std_logic_vector(2 downto 0);
  gt_TX_n     : out  std_logic_vector(2 downto 0);
);
end component;
```



Open-source RISC-V processor

- Open-source RISC-V processor
- Chip agile development flow and tools

Open Framework

Open Flow

Open Tools

microarch.

coding

EDA Tools

ISA Spec.

Docs


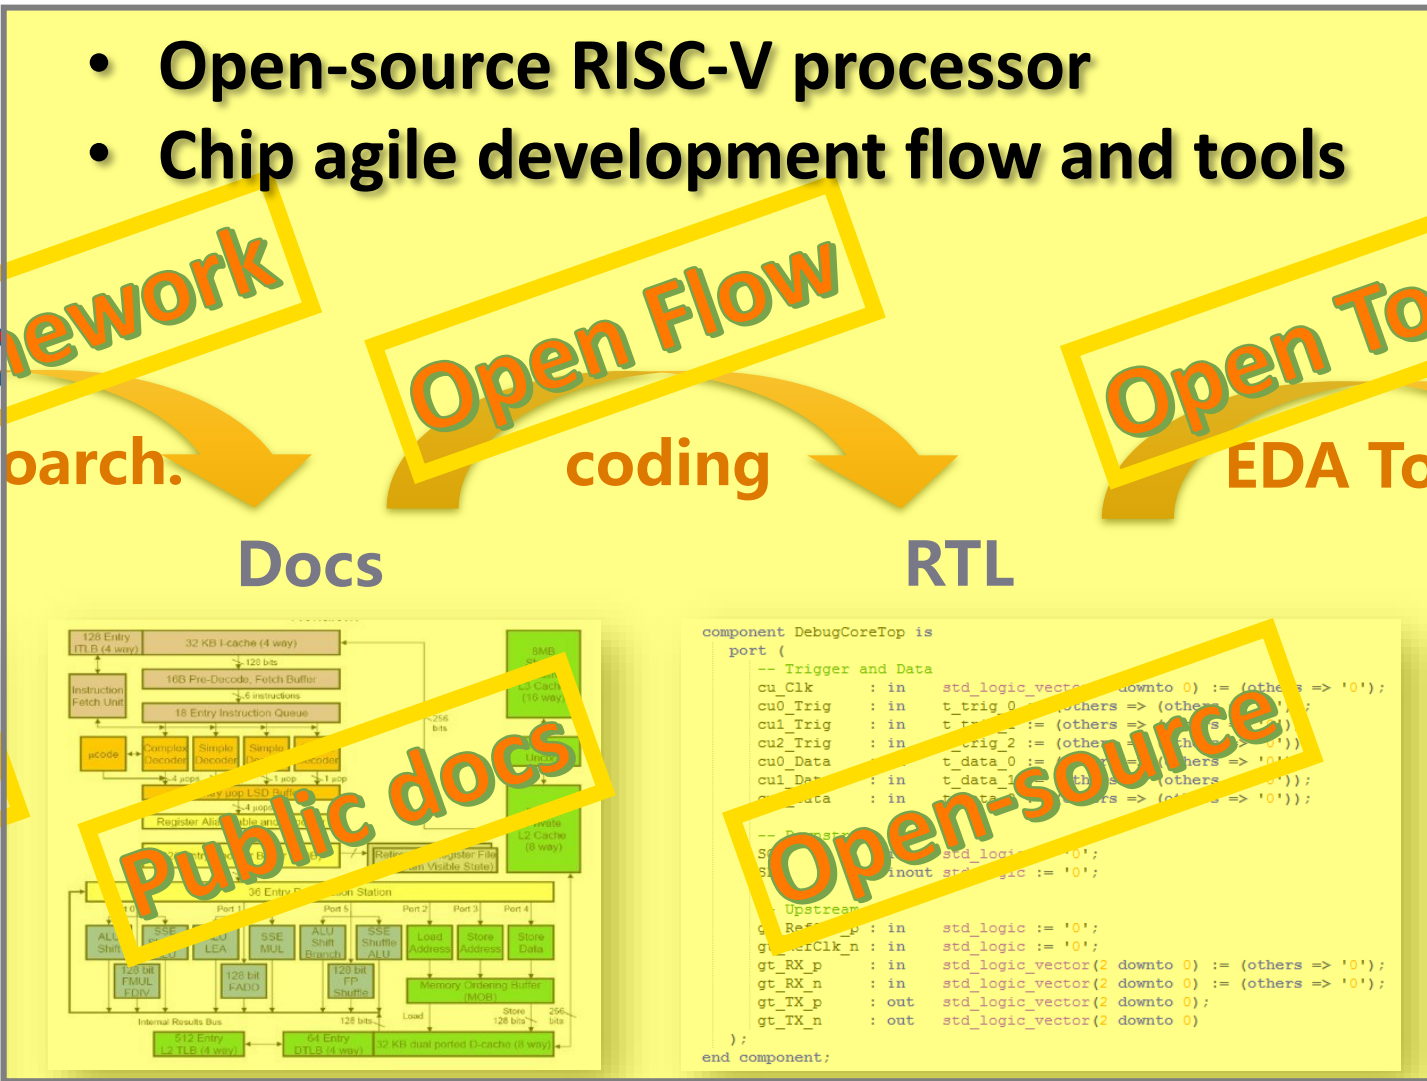
RTL

Layout

Open & Free

public docs

Open-source



❄️ Why open-source high-perf. RISC-V processor?

- **Why RISC-V:** Free and open ISA
- **Why high-perf :** Most RISC-V processors are for IoT/AI, but both academic and industrial community need high-performance RISC-V processors
- **Why open-source:** An open and innovative hardware platform, “hardware version of Linux”
- **Build a leading platform with chip development method, flow and tools**


		Designs (“Source”)			Products
		<i>Free & Open Designs</i>	<i>Licensable Designs</i>	<i>Closed Designs</i>	
Specifications	<i>Designs</i>				
	<i>Specifications</i>				
		Free & Open Spec	“Open Source”		
	Licensable Spec				Based on Licensed or Closed
	Closed Spec				Based on Closed Designs

Source: David Patterson, Keynote @ CRVF 2019, <https://crvf2019.github.io/pdf/keynote1.pdf> 5

XiangShan: an open-source high-performance processor

- **XiangShan**: an industrial-level processor written in Chisel
 - Named after a mountain in Beijing, China

- **Timeline**



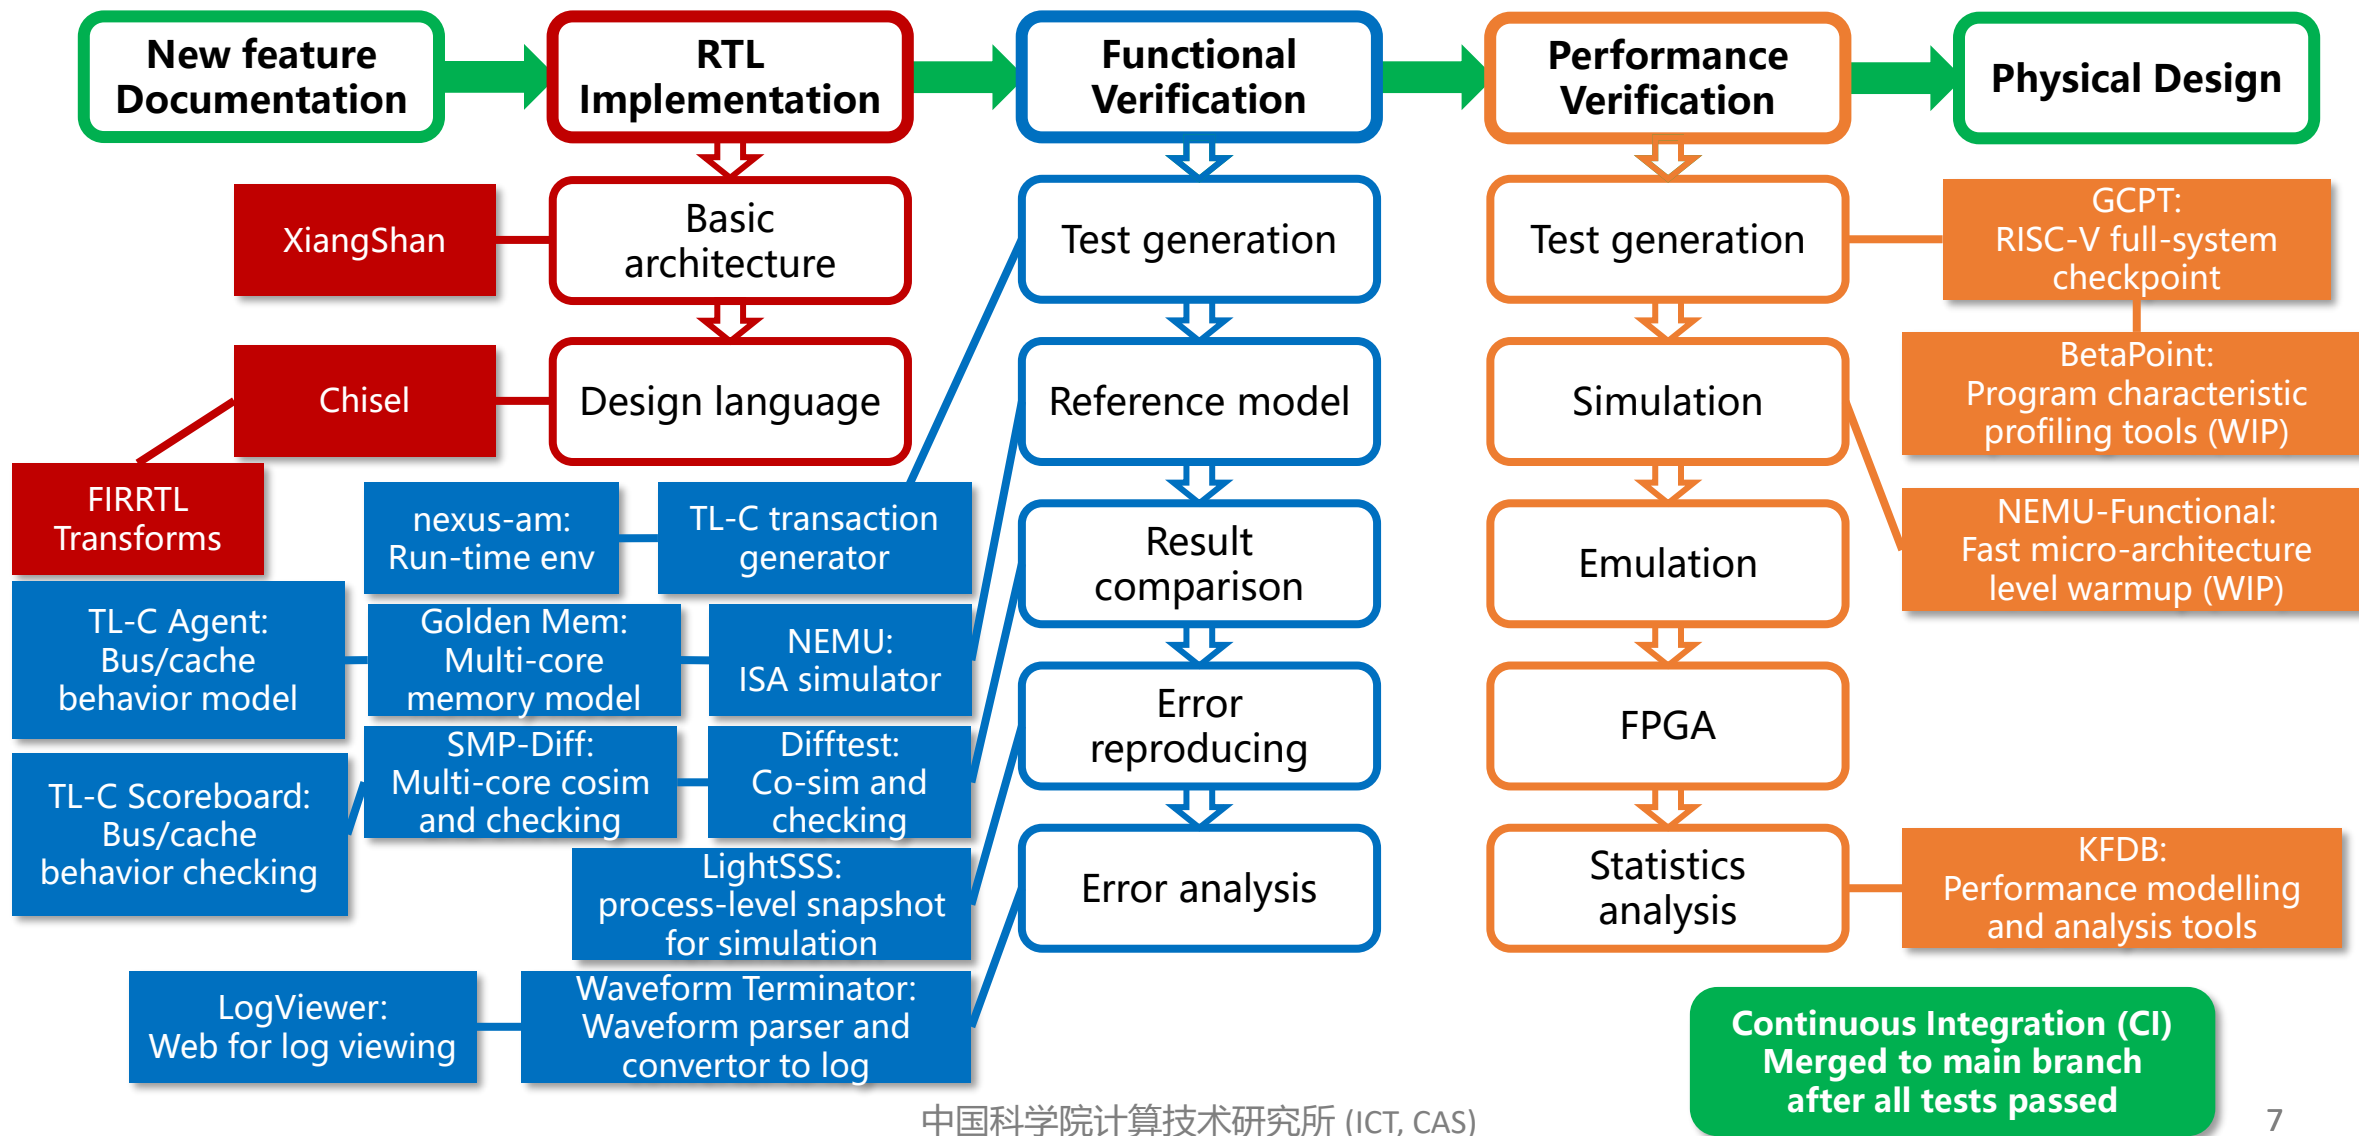
Fragrant Hills in Beijing

- Open-sourced at <https://github.com/OpenXiangShan/XiangShan>



Scan to follow us on GitHub

Agile development infrastructures



Outline

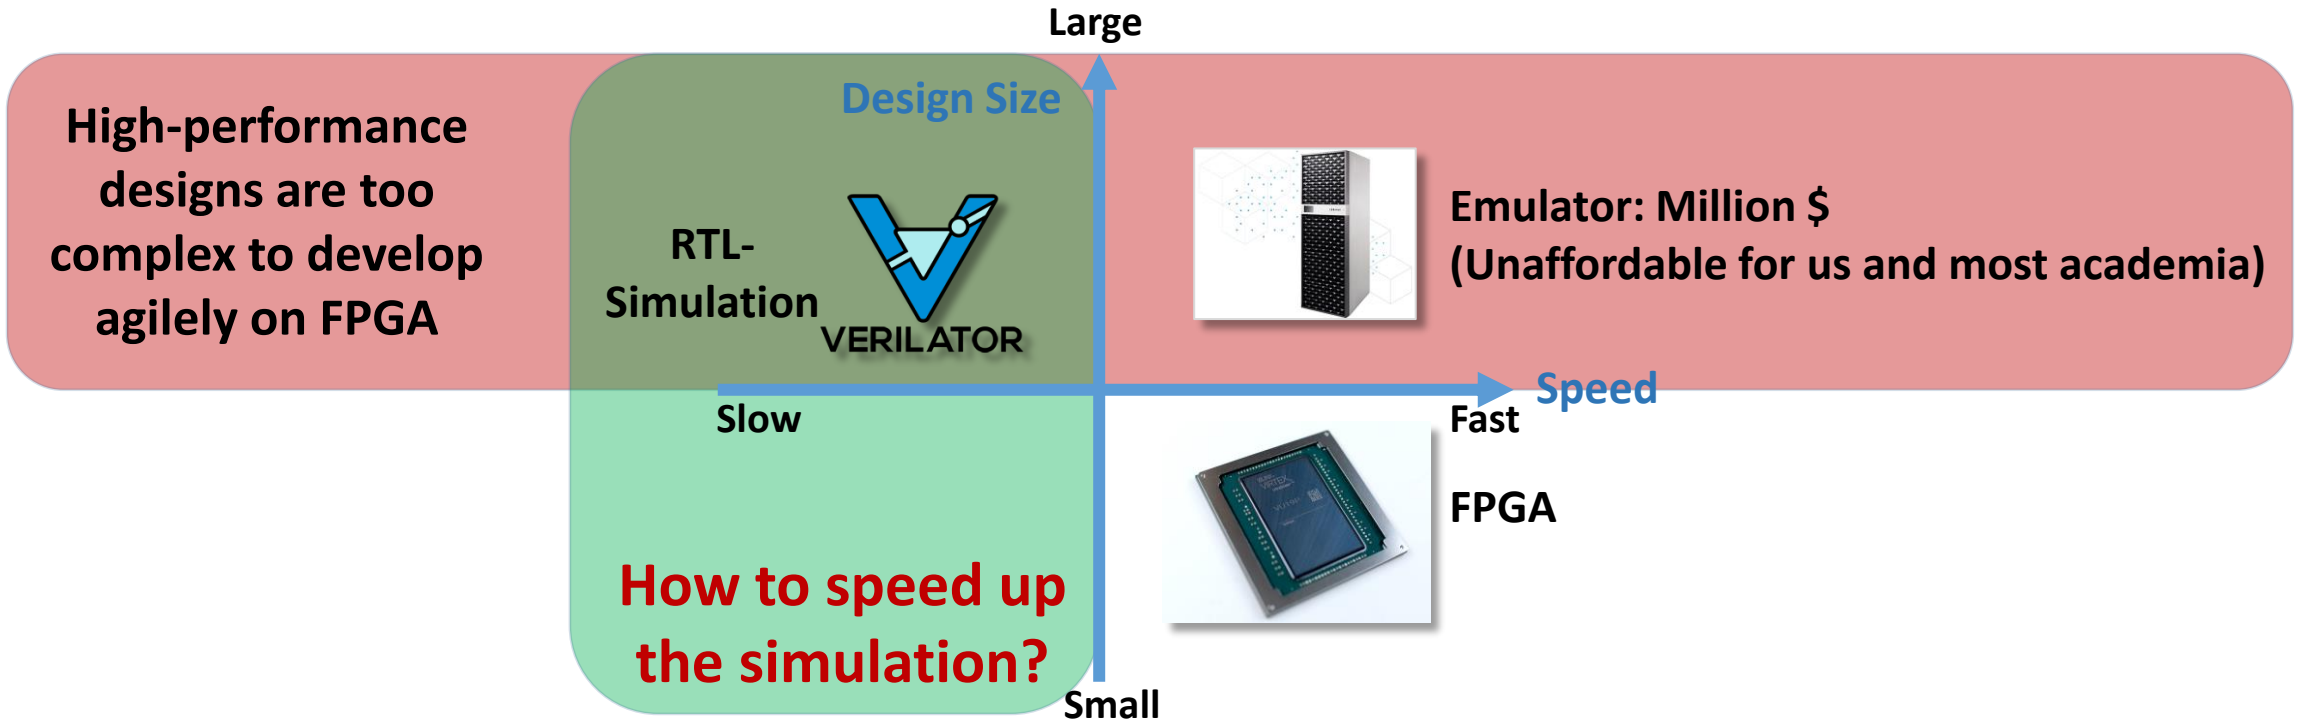
- Introduction of XiangShan

-  • **Agile Performance Modelling with Software-based RTL-simulation**

- Micro-architecture Design and Optimizations of XiangShan



Performance Modelling with Software-based RTL-simulation




Time for performance modelling of single-core XiangShan on SPEC CPU2006

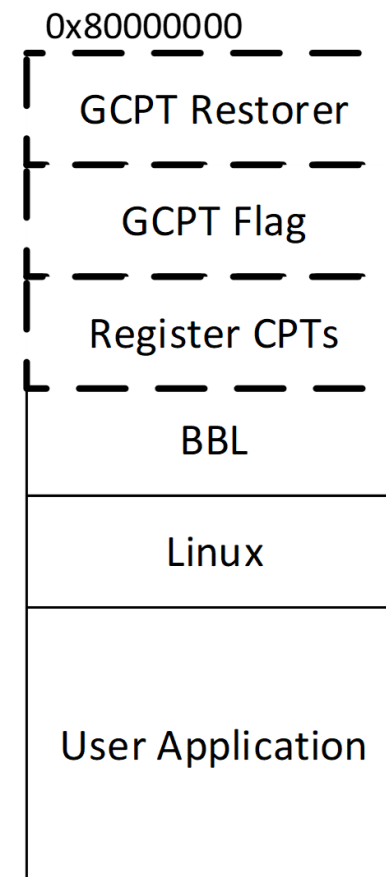
	RTL-simulation	FPGA	RTL-Simulation w/ Checkpoint
Compile	20 minutes	5 hours	20 minutes
Simulation	958 years@2KHz (2K CPS)	7 days@100MHz	5.5 hours with enough x86 servers
Our Approach			

Parallelism via RISC-V Checkpoint

- Cut a program into small segments, and simulate the segments simultaneously




- **Challenge ①: create checkpoints on a certain position of the program**
 - Based on NEMU, an instruction set simulator (ISS) running at ~300MIPS
 - Required information: program counter, architecture registers, memory, etc
- **Challenge ②: restore checkpoints when simulating XiangShan**
 - Key problem: the design is changing, and we need a generic approach for RISC-V
 - **Solution: use privilege instructions to initialize the registers**
 - Dromajo [MICRO'21] uses debug mode, which does not exist on some processors




**Format of RISC-V
Checkpoint**

Feature Clustering via Simpoint

- Simpoint[1]: finding and exploiting program phase behavior



- Generating representative RISC-V checkpoints from SPEC CPU2006
- Example: execution time estimation[2] based on clustering weights and simulated IPC



Estimated Yanqihu SPEC CPU2006 results@1.3GHz

[1] Greg Hamerly, Erez Perelman, Jeremy Lau, and Brad Calder, SimPoint 3.0: Faster and More Flexible Program Analysis, Journal of Instruction Level Parallel, September 2005.

[2] 90-cycle fixed latency. Simpoints with 80% coverage and maxK=30.

Outline

- Introduction of XiangShan
- Agile Performance Modelling with Software-based RTL-simulation

-  • **Micro-architecture Design and Optimizations of XiangShan**

Yanqihu: 1st generations of XiangShan

- **Yanqihu**: named after a lake in Beijing, China
 - RV64GC, 11-stage, superscalar, out-of-order
 - 5.3 CoreMark/MHz (gcc-9.3.0 -O2)
 - **Estimated SPEC CPU2006 ~9@1.3GHz**
- **Tape-out**: single XiangShan core (commit hash cbcac07) with 1MB L2 Cache



Yanqi Lake in Beijing

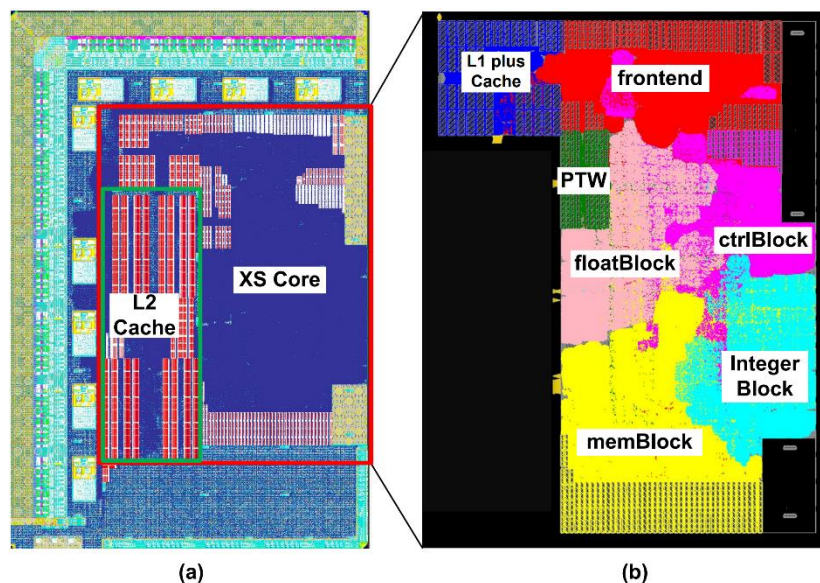
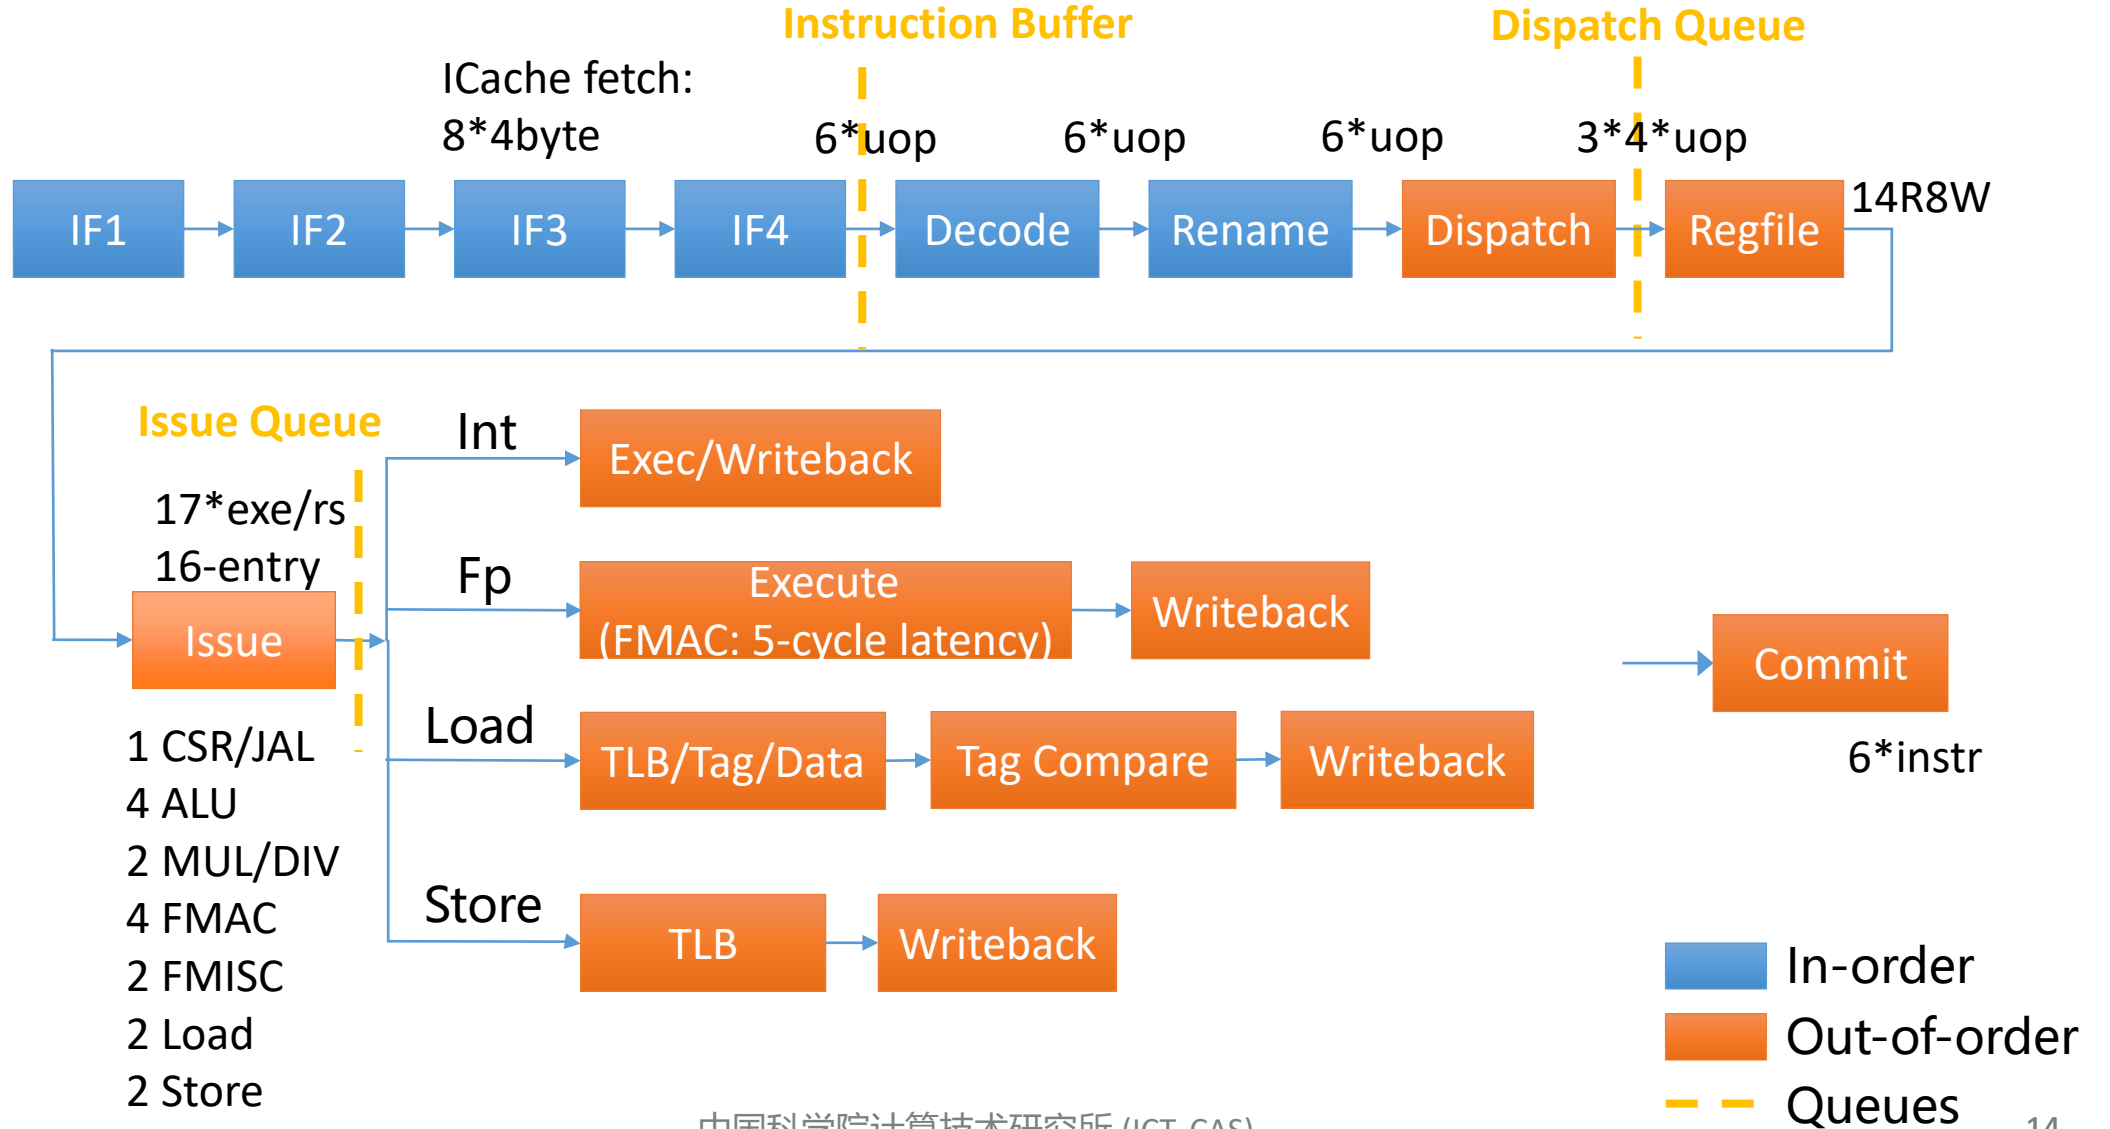


Figure. Layout of (a) the entire chip; (b) the core

Tape-out information for the processor core	
Process Node	28nm
Die Size	8.6 mm ²
Std Cell	5.05M, 4.27 mm ²
Mem	261, 1.7mm ²
Density	66%
Cell	ULVT 1.04%, LVT 19.32%, SVT 25.19%, HVT 53.67%
Estimated Power	5W
Frequency	1.3GHz, TT85C



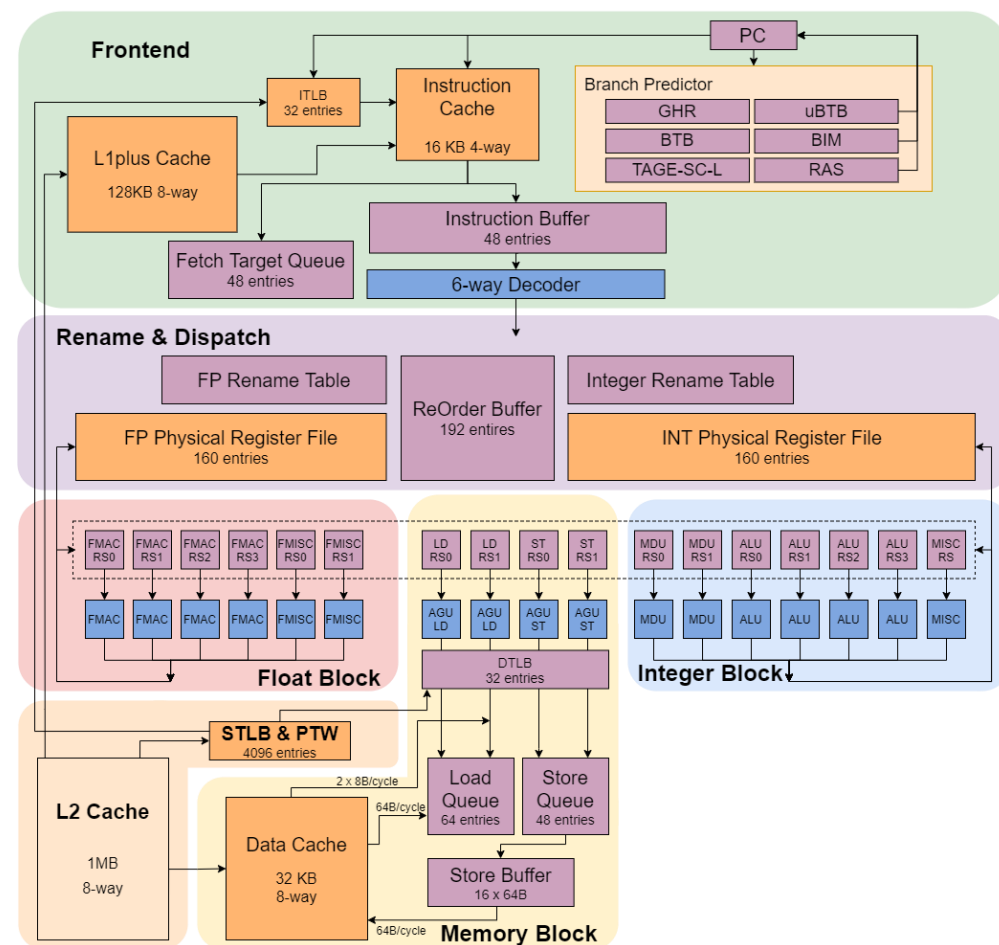
Yanqihu Pipelines



XiangShan microarchitecture (Yanqihu)




- **11**-stage, **6**-wide decode/rename
- **TAGE-SC-L** branch prediction
- **160** Int PRF + **160** FP PRF
- **192**-entry ROB, **64**-entry LQ, **48**-entry SQ
- **16**-entry RS for each FU
- **16KB** L1 Cache, **128KB** L1plus Cache for instruction
- **32KB** L1 Data Cache
- **32**-entry ITLB/DTLB, **4K**-entry STLB
- **1MB** inclusive L2 Cache





Estimated SPEC CPU2006 results of Yanqihu

- **RTL simulation** (commit hash adf9fcb, yanqihu)
 - DDR4-2400 under DRAMsim3
 - Sampling with SimPoint simulation points of SPEC CPU2006
 - maxK = 30, benchmarks with coverage > 80%
 - Each point with 50M warmup + 50M sampling
 - Customized RISC-V checkpoint format



Nanhu: 2nd generation microarchitecture

- Named after a lake in Jiaxing, Zhejiang, China
- Target: 2GHz@14nm, SPEC CPU2006 20 marks
- **Major changes**
 - New frontend design: decoupled BP and instruction fetch
 - Improved backend: better scheduler, instruction fusions, and more
 - New L2 cache: designed for high frequency and high performance
 - Tape-out with dual cores (RV64GCBK), more devices support (PCIe, USB, ...)
- **Continuously optimize the performance, frequency, ..., in an agile way**



Estimated performance of Nanhu

- Estimated **SPECint 2006 18.41, SPECfp 2006 20.94@2GHz**
 - RTL simulation, DDR4-2400 under DRAMsim3
 - SimPoint simulation points of SPEC CPU2006 (RV64GCB, O2)
 - maxK = 30, each point with 20M warmup + 20M sampling

```
Running CoreMark for 100 iterations
2K performance run parameters for coremark.
CoreMark Size : 666
Total time (ms) : 128
Iterations : 100
Compiler version : GCC10.2.0
seedcrc : 0xe9f5
[0]crclist : 0xe714
[0]crcmatrix : 0x1fd7
[0]crcstate : 0x8e3a
[0]crcfinal : 0x988c
Finised in 128 ms.

=====
CoreMark Iterations/Sec 781
Core 0: HIT GOOD TRAP at pc = 0x80001c76
total guest instructions = 31,217,108
instrCnt = 31,217,108, cycleCnt = 12,979,299, IPC = 2.405146
```

7.81 CoreMark/MHz @ (RV64GCB, O2)

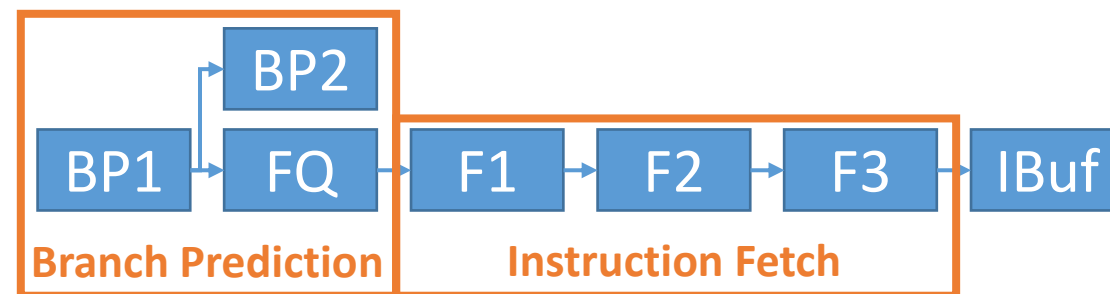
```
***** SPECINT 2006 *****
400.perlbench: 20.25, 10.12
401.bzip2: 11.89, 5.94
403.gcc: 20.51, 10.26
429.mcf: 11.81, 5.90
445.gobmk: 19.49, 9.74
456.hmmer: 20.09, 10.04
458.sjeng: 18.91, 9.45
462.libquantum: 41.80, 20.90
464.h264ref: 29.46, 14.73
471.omnetpp: 12.14, 6.07
473.astar: 14.33, 7.17
483.xalancbmk: 16.40, 8.20
SPECint2006@2GHz: 18.41
SPECint2006/GHz: 9.21
```

Estimated SPECint 2006 18.41, SPECfp 2006 20.94@2GHz

```
***** SPECFP 2006 *****
410.bwaves: 28.29, 14.14
416.gamess: 22.17, 11.08
433.milc: 20.01, 10.00
434.zeusmp: 21.95, 10.98
435.gromacs: 18.81, 9.41
436.cactusADM: 14.70, 7.35
437.leslie3d: 18.11, 9.05
444.namd: 26.86, 13.43
447.dealII: 33.21, 16.61
450.soplex: 18.24, 9.12
453.povray: 27.96, 13.98
454.Calculix: 9.45, 4.72
459.GemsFDTD: 17.48, 8.74
465.tonto: 17.85, 8.93
470.lbm: 39.93, 19.97
481.wrf: 18.42, 9.21
482.sphinx3: 21.03, 10.51
SPECfp2006@2GHz: 20.94
SPECfp2006/GHz: 10.47
```

🏔️ Opt. ①: Instruction Fetch and Branch Prediction

- **Decoupled IF and BP**
- **Higher Branch Prediction Accuracy**
 - More accurate branch history
 - One-cycle earlier prediction latency of TAGE
 - ITTAGE indirect branch prediction
 - Larger RAS for the call-return instruction pair
- **Higher Fetch Throughput**
 - Hide the instruction fetch bubbles
 - BP continues when IF is blocked



Opt. ②: Execution Units

- **More instructions supported**
 - RISC-V bit-manipulation extension (RVB)
 - RISC-V scalar crypto extension (RBK)
 - Fine-Grained Address-Translation Cache Invalidation (Svinval)
- **Instruction fusion** for common cases
- **IEEE754-compatible high-performance FPU**
 - Open-source at <https://github.com/OpenXiangShan/fudian>
 - Dual-path floating-point adder
 - Cascade FMA (5-cycle latency)
 - Improved Int/Fp Division with SRT-16




Figure. Normalized Dynamic Instruction Count for CoreMark (-O2)

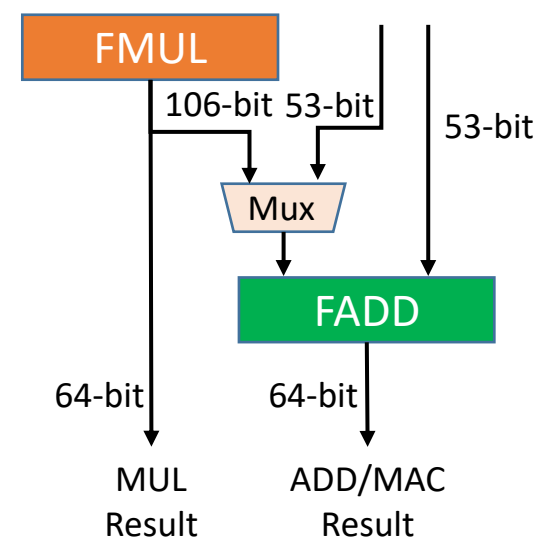


Figure. Cascade FMA

Opt. ③: Load Store Unit

• New Features Supported

- Customized Configurable Physical Memory Attributes (PMA)
- RISC-V Physical Memory Protection (PMP)
- ECC for L1/L2/L3
- Cache coherency between L1I and L1D (eliminated overhead of fence.i)

• L1 TLB/STLB

- Larger size (136 entries for DTLB) and better replacement (PLRU)
- Max. 8 in-flight PTW and memory requests supported
- Next-line prefetcher for STLB

• Load/Store Pipelines

- Virtual-address based fast data bypassing
- Memory dependence prediction using store sets
- Separated store address and data

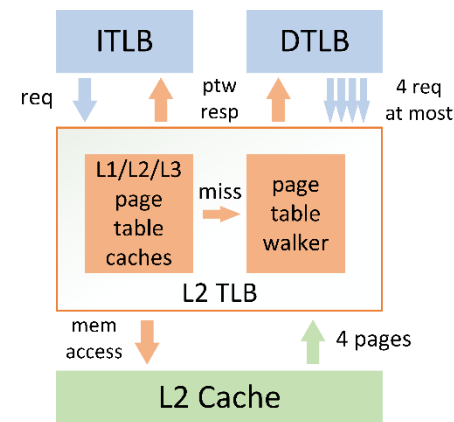


Figure. TLB Hierarchy

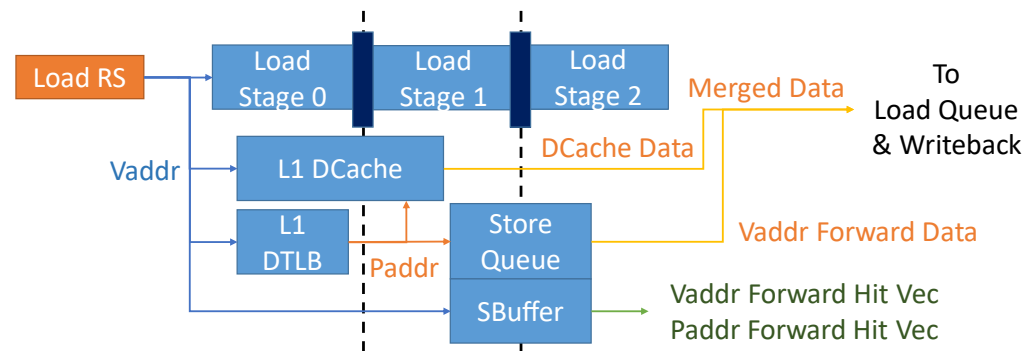
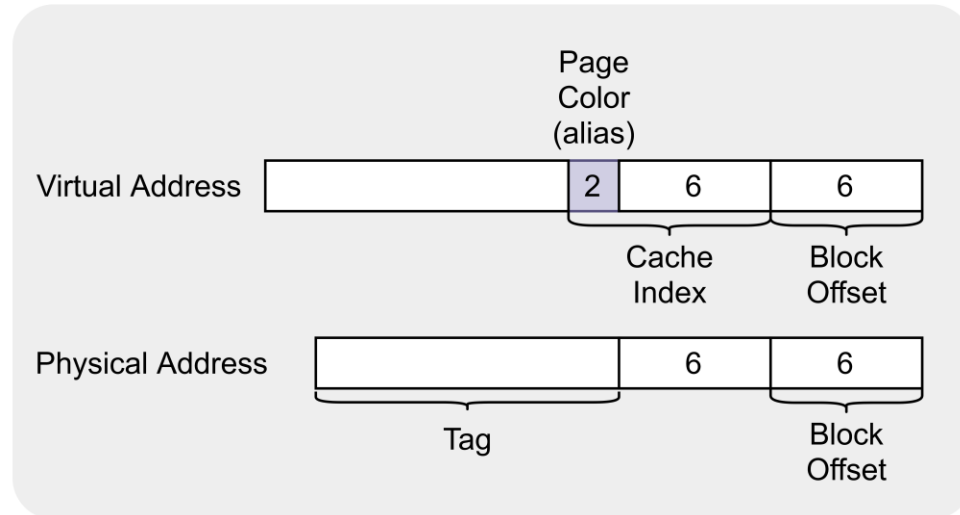
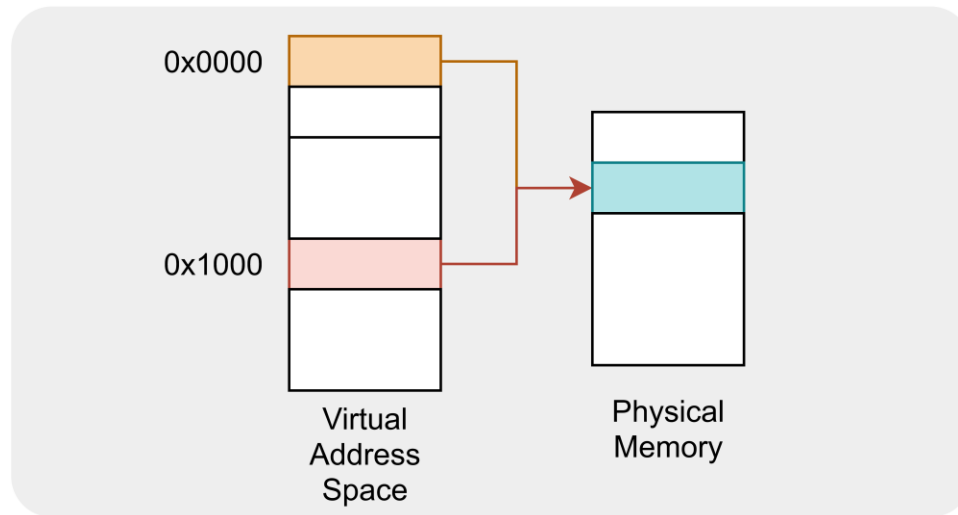


Figure. 3-cycle Load Latency

Opt. ④: Data Cache

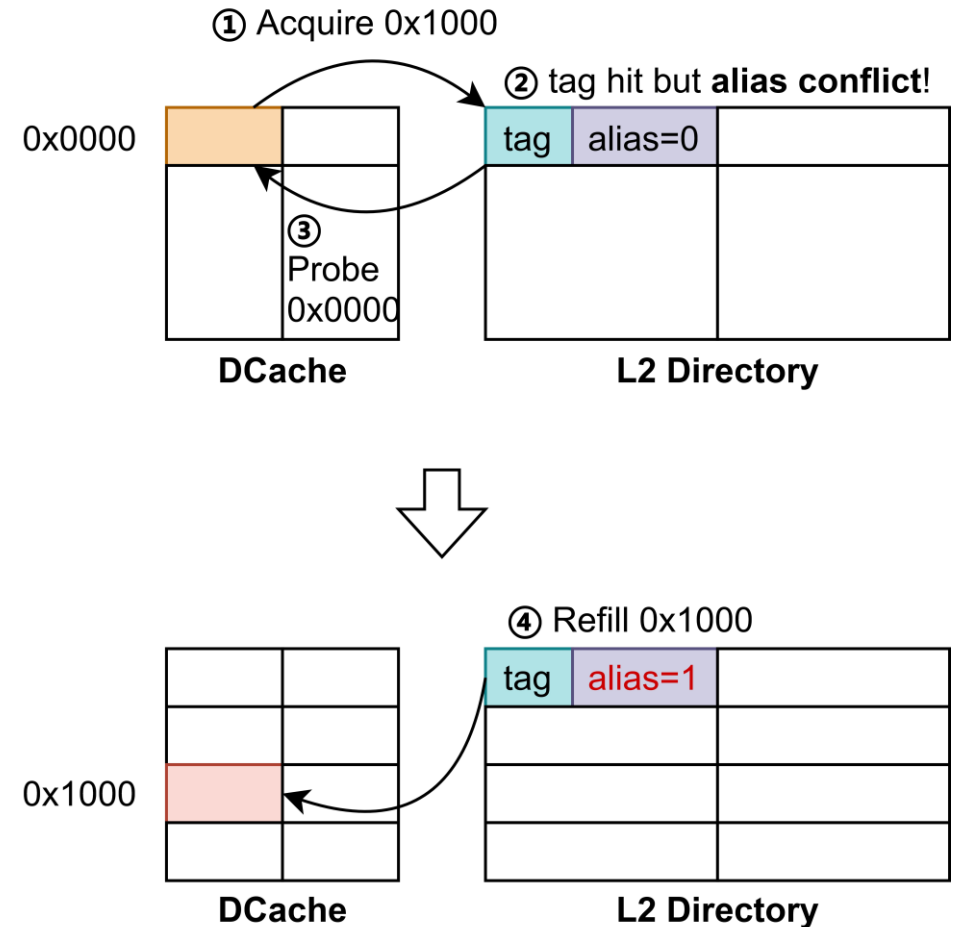
- 128KB DCache with hardware-based solution to cache alias
 - Key idea: disallow aliases (different page colors) to co-exist in a VIPT cache



Opt. ④: Data Cache

- **Hardware solution for cache alias**

- Virtual-indexed DCache
- Physical-indexed L2
- L2 directory saves alias bits (page color) of the block in DCache
- L2 will ensure only one alias exists in L1





Open-source plan

- XiangShan has been open-sourced on GitHub since June 2020
 - MulanPSLv2 License (compatible with Apache v2.0)
 - <https://github.com/OpenXiangShan/XiangShan>
 - Mailing list: xiangshan-all@ict.ac.cn
- Every commit, every generation, every tool, ..., **ALL on GitHub**
- Any discussions, features, patches, ..., **ALL are welcome**
 - After the verification and review processes, PRs will be merged
- **Contribute to XiangShan and realize your ideas on real chips!**
 - The open-source XiangShan will be **taped-out every ~6 months**



Scan to follow us on GitHub

Thanks!