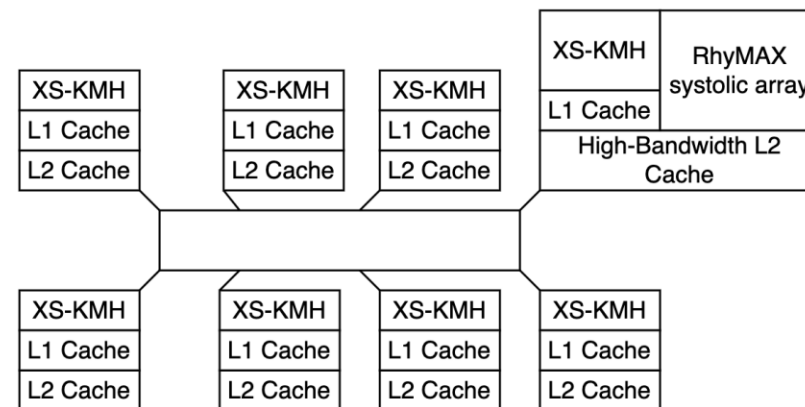
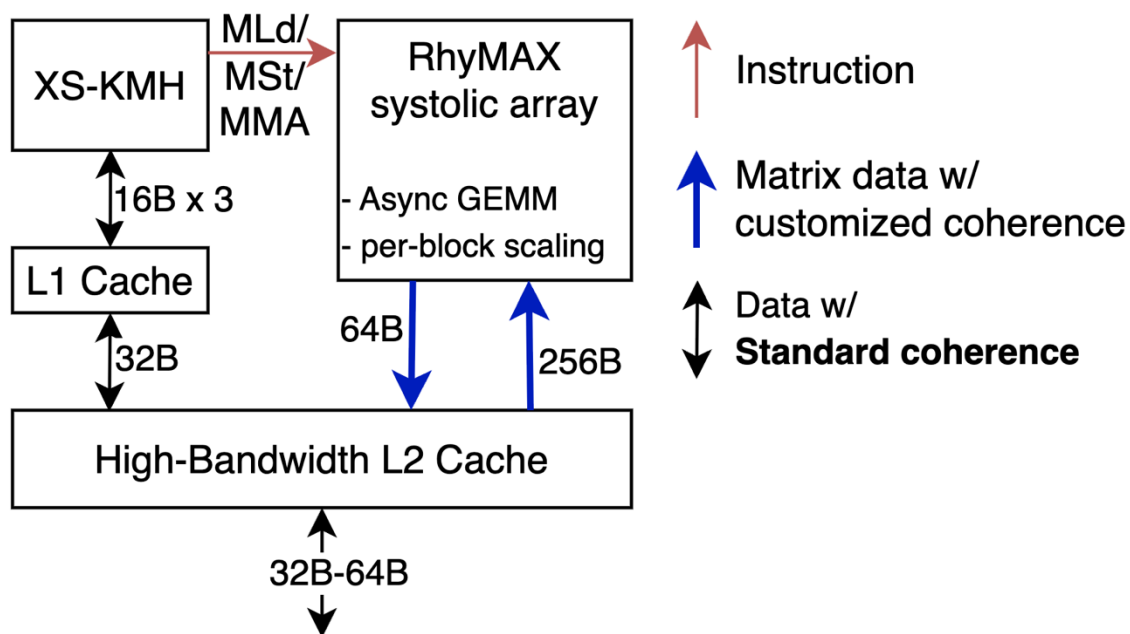


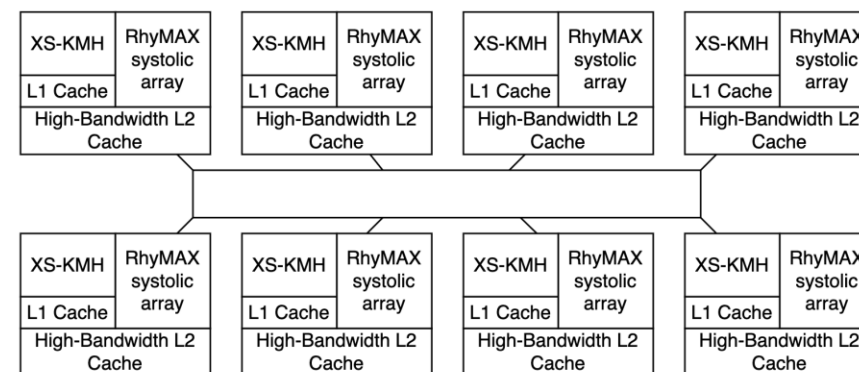
XSAI (ξ) : Hardware Support for Modern LLM Kernels in a CPU Paradigm

XSAI (ξ) : Xiangshan-KMH + RhyMAX

• Consumer SoC (AIPC)



• Cloud SoC



Agenda

Motivation of RISC-V + GEMM

- Why GEMM
- Why RISC-V CPU

Current status

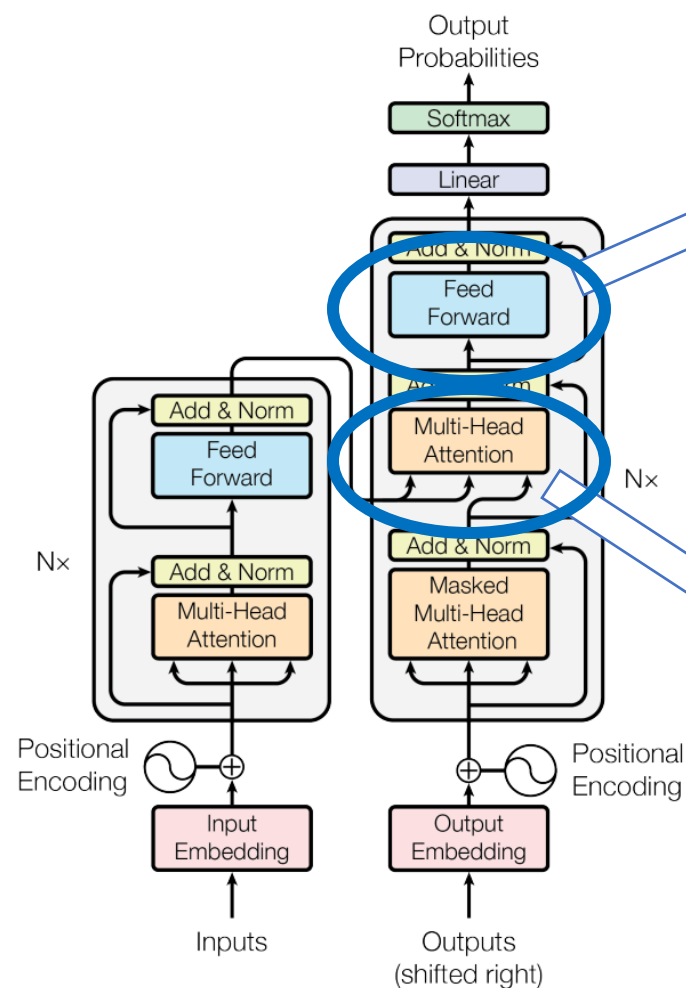
- High-bandwidth L2-cache
- Asynchronous GEMM
- MXFP8

GEMM is important, even for decoding & consumer-side

FAQ: But decoding is GEMV/mem-bound?

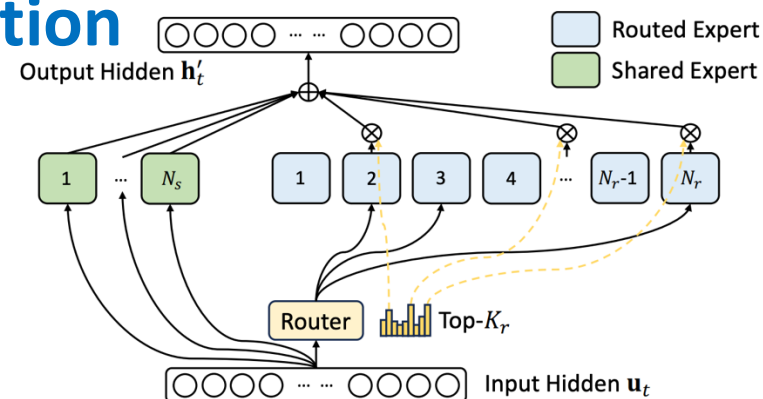
- Takeaway1: Decoding in cloud is not memory-bound
- Takeaway2: Decoding in consumer-side can be GEMM
- Takeaway3: Speculative decoding is cheaper than memory channel

Evolution of transformer block



Dense FFN -> MoE FFN

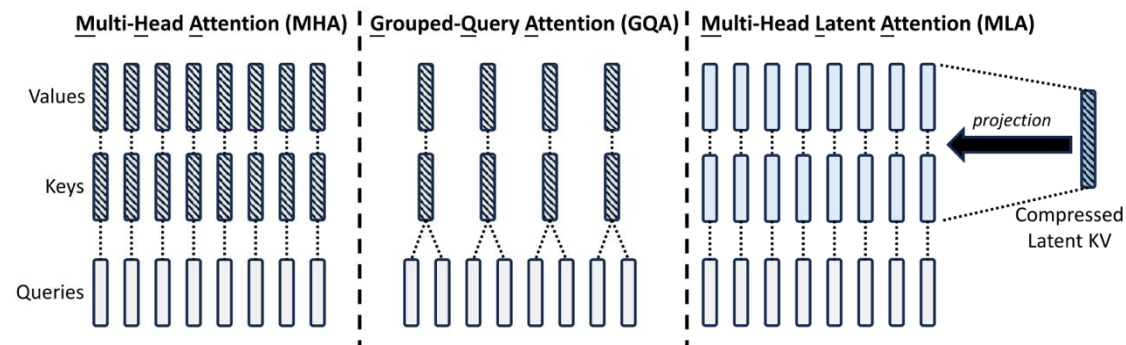
Less Communication



MHA -> GQA -> MLA

Less KV Cache occupation

Higher computation/mem ratio

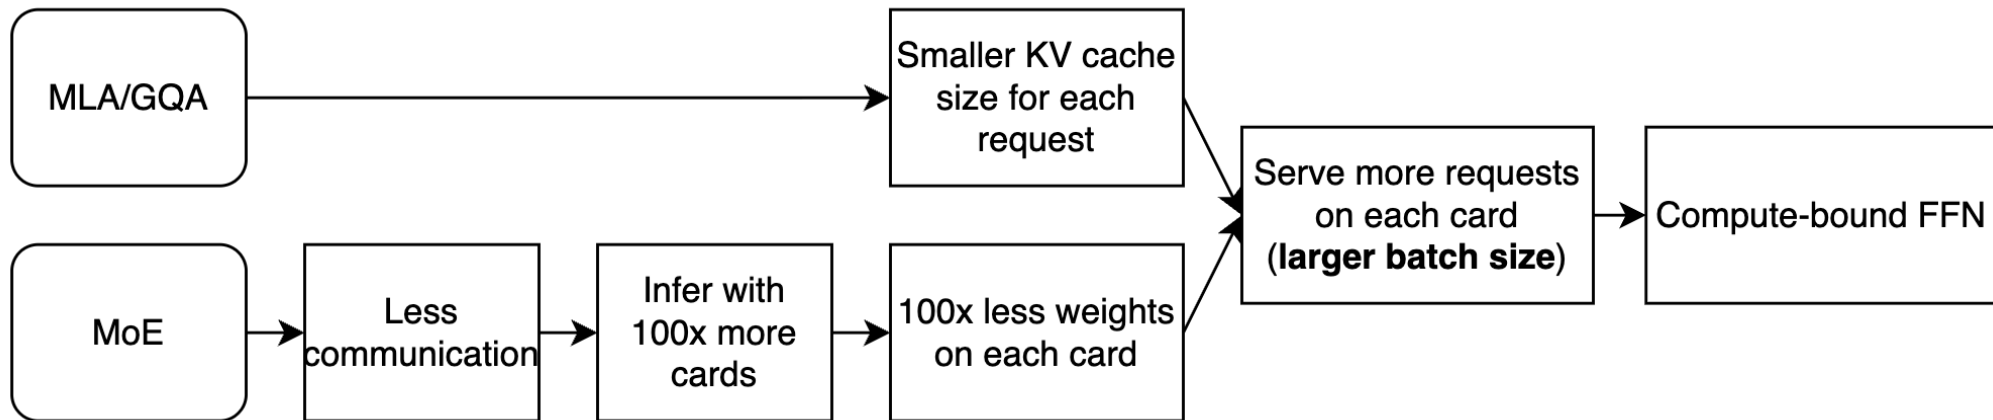


<https://arxiv.org/pdf/1706.03762>

<https://arxiv.org/abs/2405.04434>

Takeaway1: Decoding (MoE) in cloud is not mem-bound


- With EP, DSV3 is either **compute-bound or network-bound**
 - **NOT memory bound**
- Hard to make profit without EP
- Flag-ship open-weight models are becoming MoE
 - Deepseek V3, Qwen3, LLAMA4




Takeaway2: Decoding in consumer-side can be GEMM

Speculative decoding

- Guess multiple tokens with a draft model
 - **Verify** them in parallel (load weights once)
 - **Similar with prefill (GEMM)**
-
- For QWEN3 8B + speculation tree size=64
 - GQA: $QO_heads = 64 \times 8 = 512$, head_dim=128
 - Up proj: $M=64$, $K=5k$, $N=50k$
 - Down proj: $M=64$, $K=25k$, $N=5k$
 - **Main ops are ~~GEMV~~ GEMM during decoding**



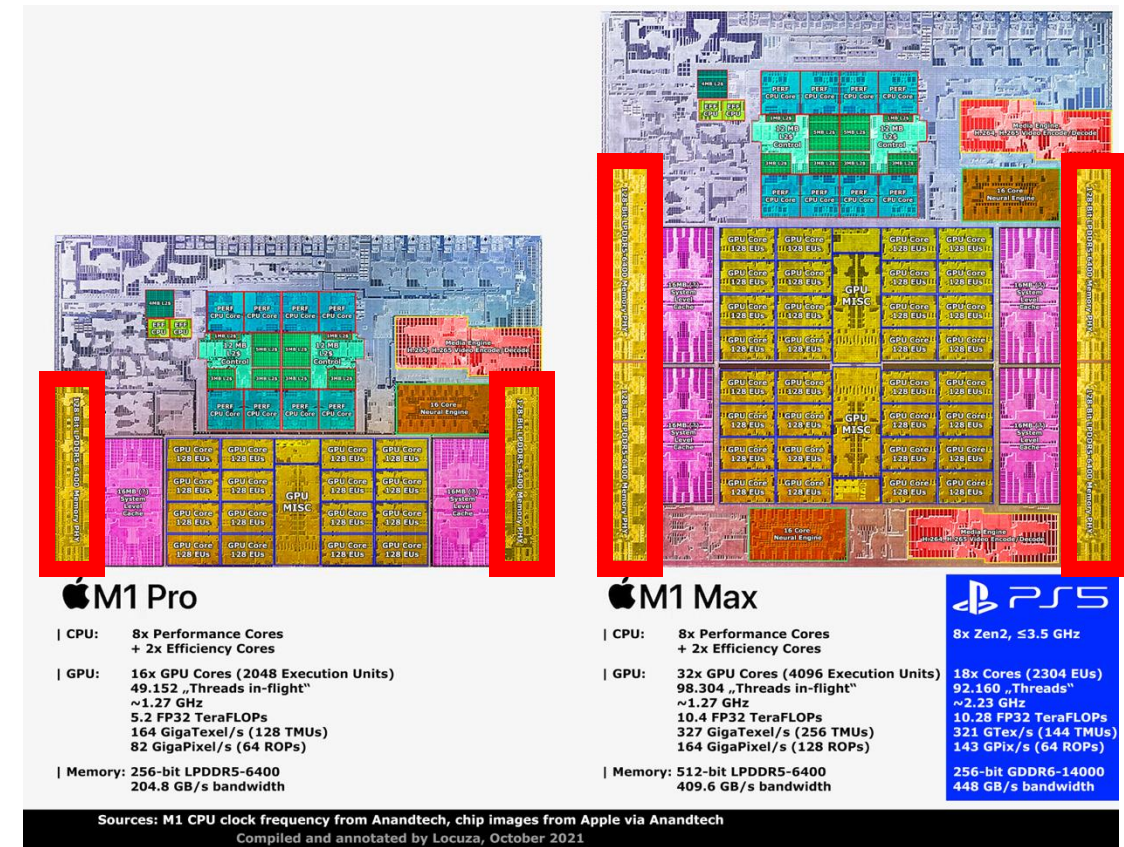
Step1: Autoregressive generation



Step2: Parallel verification

Takeaway3: Spec-decoding is cheaper than memory

- **Speculative decoding** VS. larger mem bandwidth
- Increasing memory channels is expensive
 - Double memory channels
 - Impact on total area
 - Double particles
- Increase GEMM FLOPS is cheaper
 - Double height of systolic array
 - Limited impact on SRAM/datapath


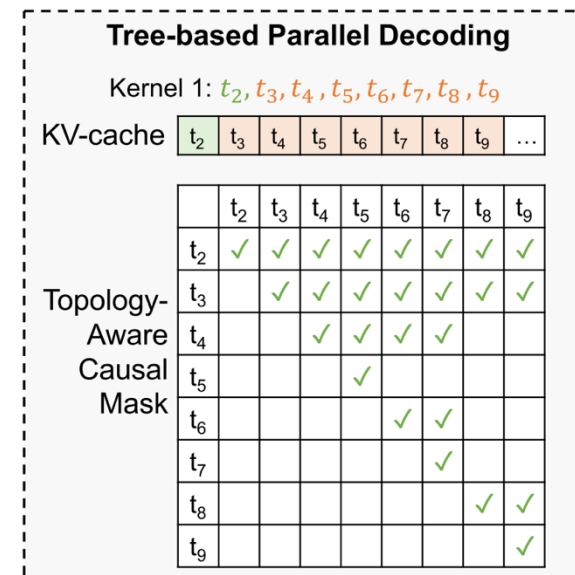


<https://www.anandtech.com/show/17024/apple-m1-max-performance-review>

<https://zhuanlan.zhihu.com/p/12618700803> 北京开源芯片研究院 (BOSC)

Why RISC-V CPU + GEMM?

- Software eco-system
 - Scalar compilers
 - Auto-vectorization for non-critical operations
 - OpenMP
- General computing
 - variable tree attention mask (SpecInfer and EAGLE-2)
 - Jagged attention



<https://arxiv.org/abs/2409.15373>
<https://arxiv.org/abs/2305.09781>

CPU VS. SIMT

- CUTE is trying to express GEMM in SIMT-paradigm (Warp)
 - Obscure (at least for me)
- Mainstream languages are Shifting from Warp-based to tile-based
 - OpenAI Triton
 - Tile-lang
- Support tile-based primitives in CPU is straight-forward
 - Intel AMX
 - RISC-V Attached matrix unit

How to satisfy computing demand in XSAI

- Better **Model FLOPS utilization (MFU)**
 - High-bandwidth L2 cache
 - Asynchronous GEMM / matrix load / matrix store
- Higher **Tflops**: Hardware support for quantization precision
 - Float: mxfp8
 - Int: Hadamard transform (Future works)

High-bandwidth L2 cache

- Demands of matrix unit conflicts with high-perf CPU

	Primary demand	what can be traded	Best buffer in existing memory hierarchy
CPU	latency	capacity	L1-cache
Matrix unit	capacity and hit bandwidth	Latency	?

- High-bandwidth L2 cache (HBL2)
 - Capacity: 1-2 MB
 - Bandwidth: 256-512 Bytes per cycle
 - Latency: ~12 cycle hit latency

	Effective hit bandwidth/Bytes/cycle
HBL2 (8-bank)	480
Xiangshan Coupled L2 cache (4-bank)	17
XuanTie C910 L2 cache	5.5


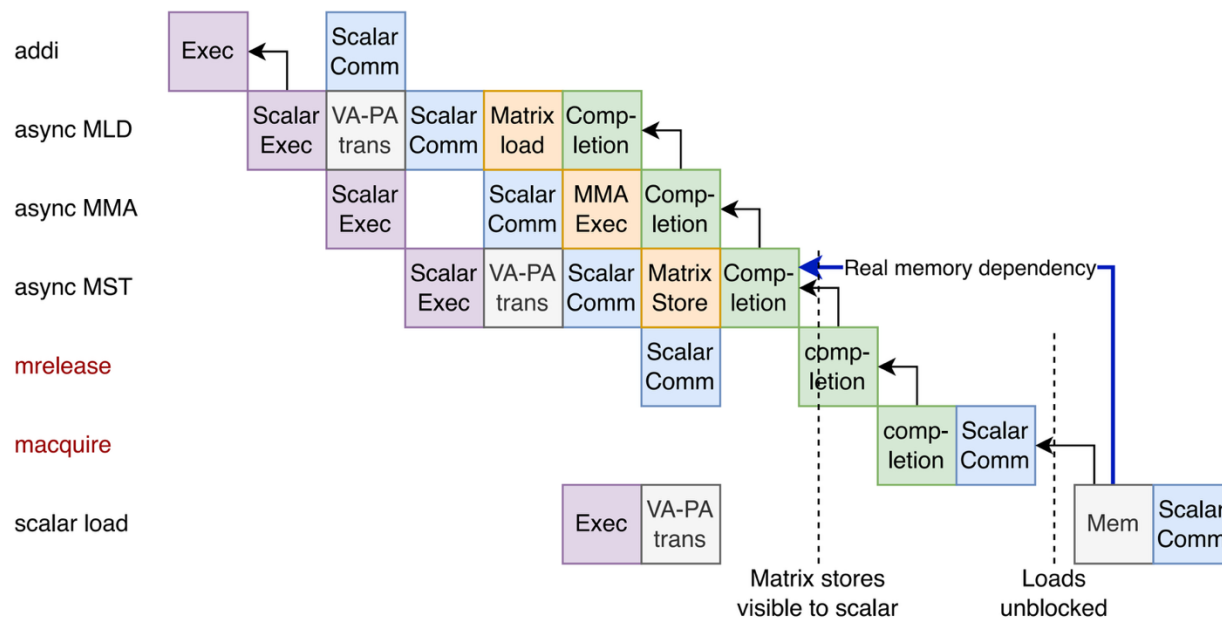
<https://chipsandcheese.com/p/alibabat-heads-xuantie-c910>

High-bandwidth L2 cache

- When coherent cache meets GEMM
 - Get/Put VS. Acquire/Release
 - Block invalidation from HNF
 - First-fetching permission of weights
 - Block insertion policy
- HBL2
 - Avoid crossbar: Direct datapath to matrix unit
 - Avoid Read/Modify/Write pattern of Acquire by using Get/Put semantic
 - 17 Bytes/cycle → 32 Bytes/cycle
 - Avoid block invalidation from HNF with smart permission speculation

Support async MMA in XSAI

- Async mma/mld/mst commit in scalar core without execution
 - Precise exception: VA-PA translation & permission check in scalar core
 - Explicit synchronization:
 - Matrix stores are not visible to scalar pipeline/vector pipeline before
 - M.release/M.acquire to enforce dependency explicitly



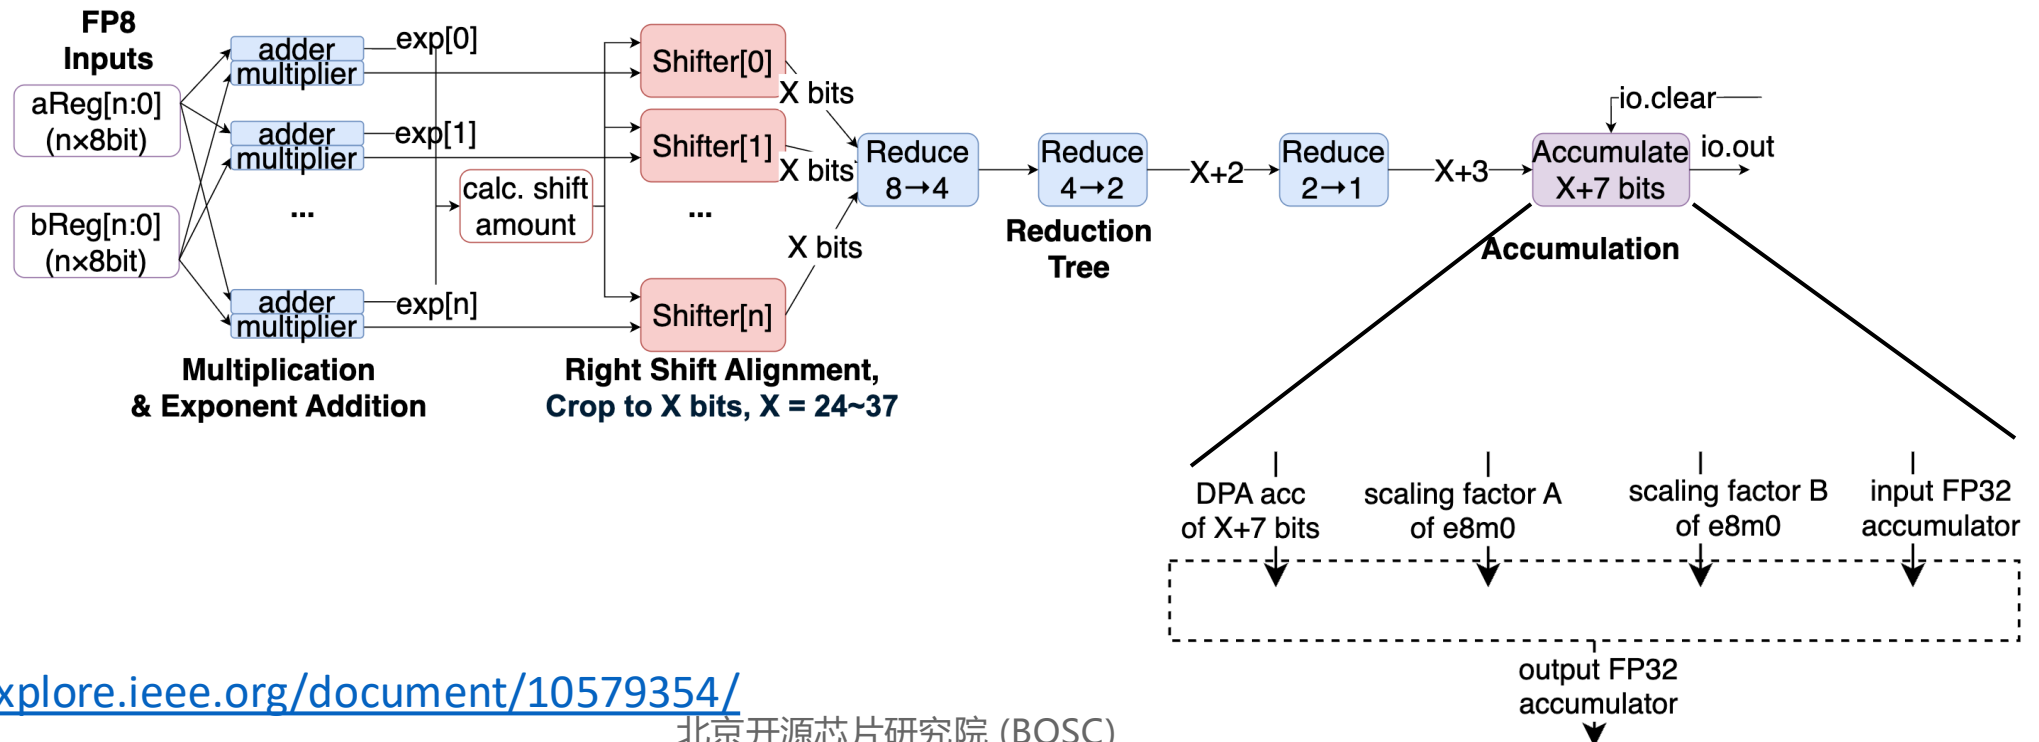
3-stage flash attention 3
Overlap GEMM with softmax

Achieve higher TFLOPS with quantization

- 8-bit tensor core
 - INT8/INT4
 - MXFP8
- Quantization enhancement
 - High accumulation precision
 - Fine-grained scaling factor

Mx8 DPA (Dot-Product Accelerator)

- 8-way fp8 DPA
- Every 32 elements share a scaling factor at K dimension
- Optional full-mantissa



Future works

- Hardware support for Hadamard transform
- Hardware assisted multi-cast
- Pre-silicon end-2-end tuning framework for LLM