

# XiangShan: An Open-Source High-Performance RISC-V Processor and Infrastructure for Architecture Research

---

*The XiangShan Team*

Institute of Computing Technology (ICT)  
Chinese Academy of Sciences (CAS)

MICRO'23@Toronto, Canada

October 28, 2023



# Tutorial@MICRO'23 Schedule

Time (PM)	Topic
1:00 - 1:25	Introduction of the XiangShan Project
1:25 - 1:30	Tutorial Overview and Highlights
<b>1:30 - 2:30</b>	<b><i>Microarchitecture Design and Implementation</i></b>
2:30 - 3:00	Hands-on Development
	Coffee Break
3:30 - 4:00	Hands-on Development & Discussions



# XiangShan: Open-Source High-Performance Processor

- **1<sup>st</sup> generation: YANQIHU**
  - 2020/6: first commit of design RTL
  - 2021/7: **28nm tape-out, 1.3GHz**
  - Performance: **SPEC CPU2006 7.01@1GHz, DDR4-1600**
- **2<sup>nd</sup> generation: NANHU**
  - 2021/5: starting design exploration and RTL design
  - 2022/12: RTL Freeze
  - Plan: **14-nm tape-out soon, estimated SPEC CPU2006 20@2GHz**
- **3<sup>rd</sup> generation: KUNMINGHU**
  - ISA feature: **Vector (V) extension, Hypervisor (H) extension**
  - Comprehensive performance improvement
- Open-sourced at <https://github.com/OpenXiangShan/XiangShan>



Fragrant Hills in Beijing

The screenshot shows the GitHub repository page for "OpenXiangShan / XiangShan". The repository is public and has 7,306 commits. It features tabs for Code, Issues (33), Pull requests (3), Discussions, Actions, Projects, Wiki, and Security. The repository description is "Open-source high-performance RISC-V processor" and it includes tags for "chisel3", "risc-v", and "microarchitecture". The repository has 3.3k stars, 75 watching, and 409 forks.

>3.2K stars, >400 forks on GitHub



# Yanqihu: 1<sup>st</sup> generation of XiangShan

- Yanqihu: named after a lake in Beijing, China
  - RV64GC, 11-stage, superscalar, out-of-order
  - 5.3 CoreMark/MHz (gcc-9.3.0 –O2)
  - Real chip: SPEC CPU2006 7@1GHz with DDR4-1600 (DDR not fully optimized)
- Tape-out: single XiangShan core (commit hash ccbca07) with 1MB L2 Cache



Yanqi Lake in Beijing

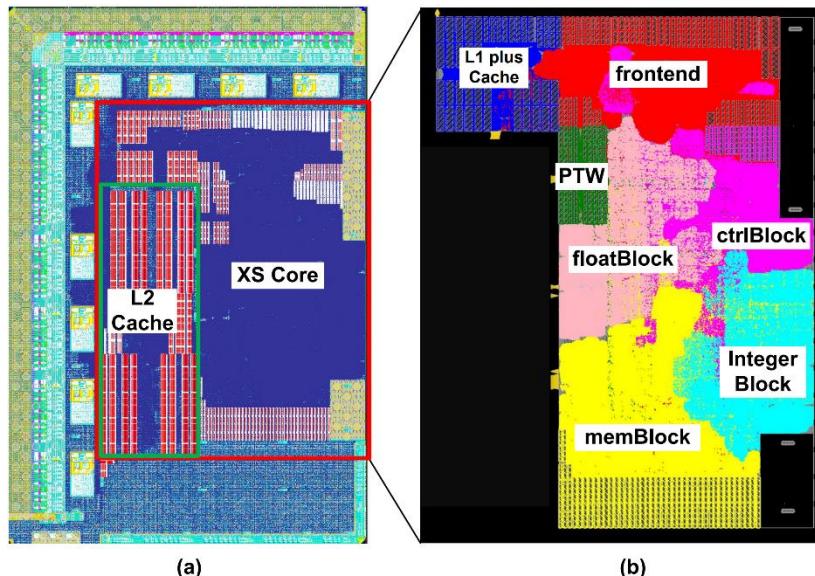
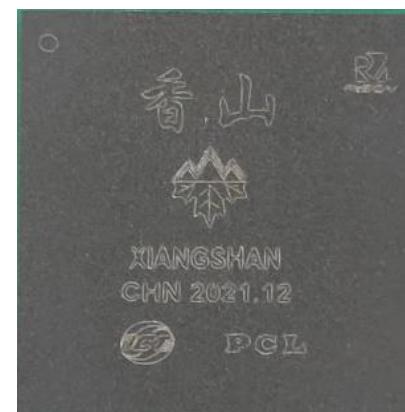


Figure. Layout of (a) the entire chip; (b) the core

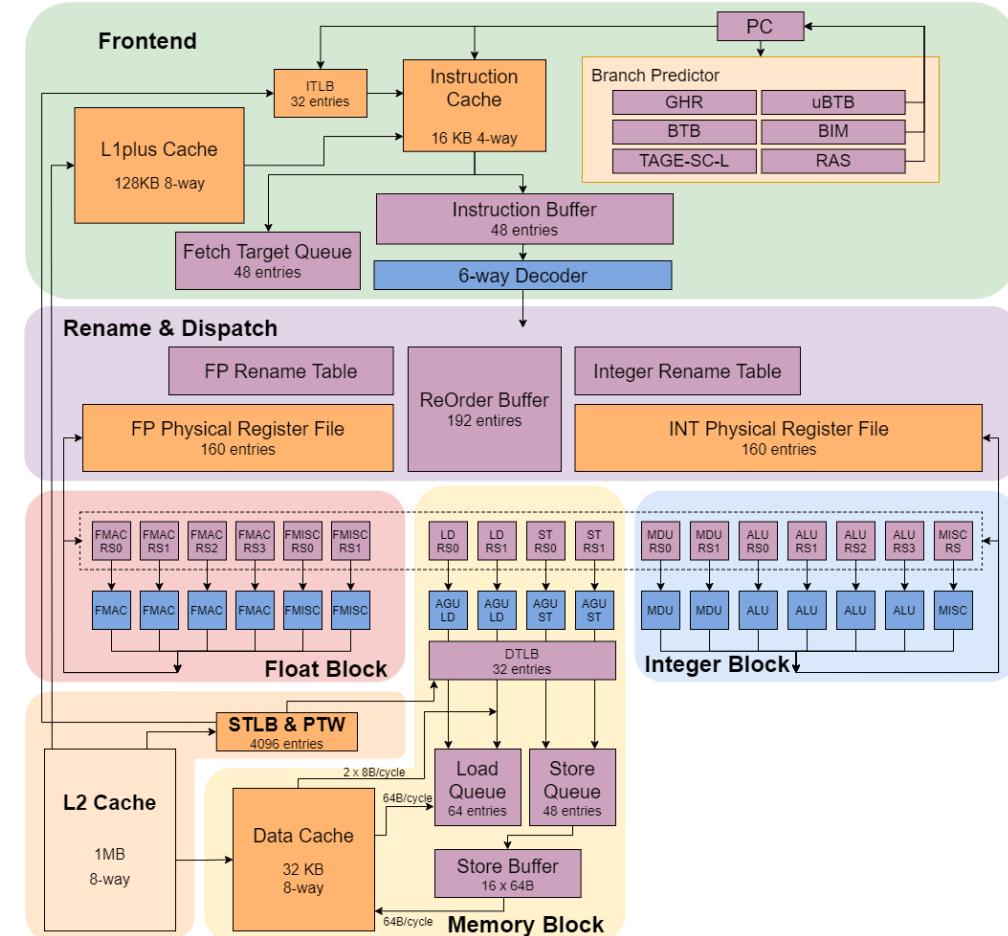


Tape-out information for the processor core	
Process Node	28nm
Die Size	8.6 mm <sup>2</sup>
Std Cell	5.05M, 4.27 mm <sup>2</sup>
Mem	261, 1.7mm <sup>2</sup>
Density	66%
Cell	ULVT 1.04%, LVT 19.32%, SVT 25.19%, HVT 53.67%
Estimated Power	5W
Frequency	1.3GHz, TT85C

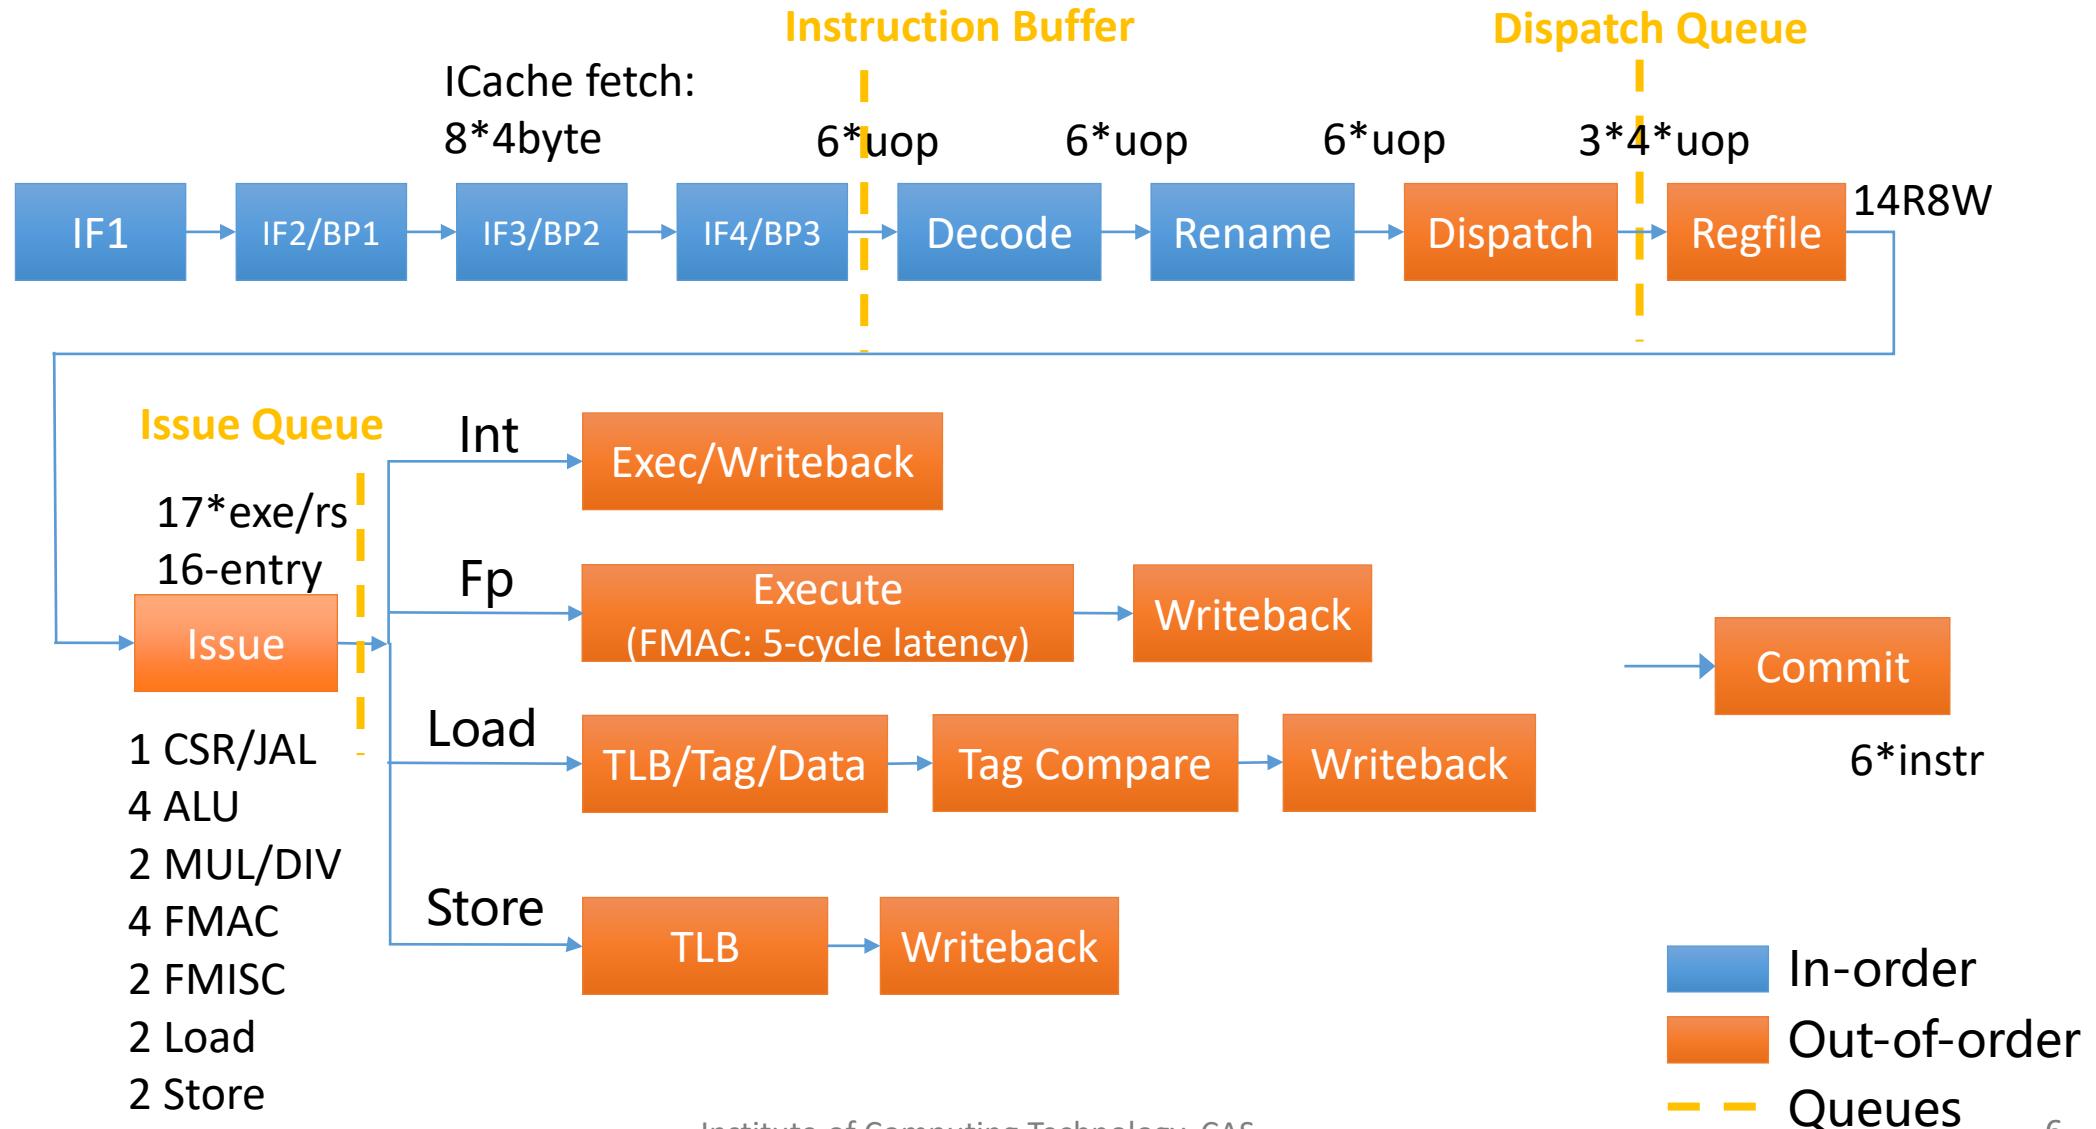


# XiangShan Microarchitecture (Yanqihu)

- **11**-stage, **6**-wide decode/ rename
- **TAGE-SC-L** branch prediction
- **160** Int PRF + **160** FP PRF
- **192**-entry ROB, **64**-entry LQ, **48**-entry SQ
- **16**-entry RS for each FU
- **16KB** L1 Cache, **128KB** L1plus Cache for instruction
- **32KB** L1 Data Cache
- **32**-entry ITLB/DTLB, **4K**-entry STLB
- **1MB** inclusive L2 Cache



# XiangShan Microarchitecture (Yanqihu)





# NANHU: 2<sup>nd</sup> generation microarchitecture

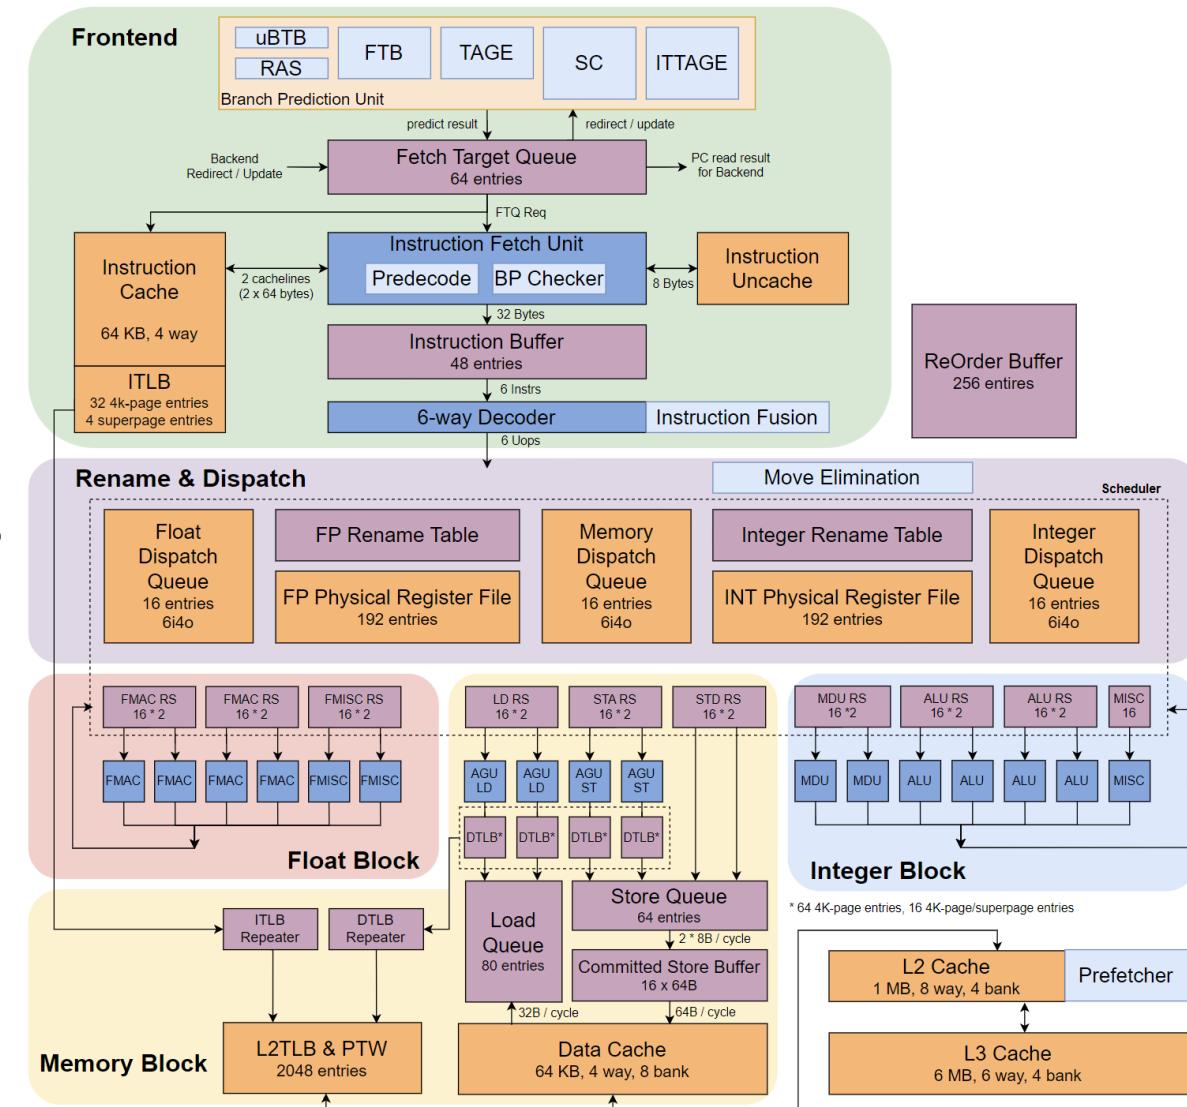
- Named after a lake in Jiaxing, Zhejiang, China
- **2GHz@14nm, taped out**
- Estimated **SPECint 2006 19.10, SPECfp 2006 22.18**
- **Major changes**
  - New frontend design: decoupled BP and instruction fetch
  - Improved backend: better scheduler, instruction fusions, and more
  - New L2 cache: designed for high frequency and high performance
  - Tape-out with dual cores (RV64GCBK), more devices support in SoC (PCIe, USB, ...)





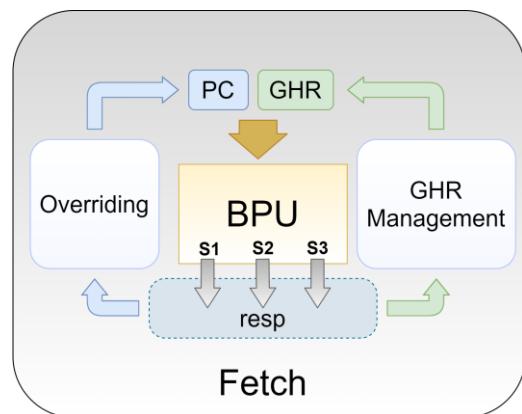
# XiangShan Microarchitecture (Nanhу)

- **192** Int PRF + **192** FP PRF
- **256**-entry ROB, **80**-entry LQ, **64**-entry SQ
- **16**-entry RS for each FU (32-entry as an RS)
- **64KB** L1 Cache, **64KB** L1 Data Cache
- **80**-entry DTLB, **36**-entry ITLB, **2K**-entry STLB
- **1MB** non-inclusive L2 Cache
- **6MB** non-inclusive L3 Cache (LLC)

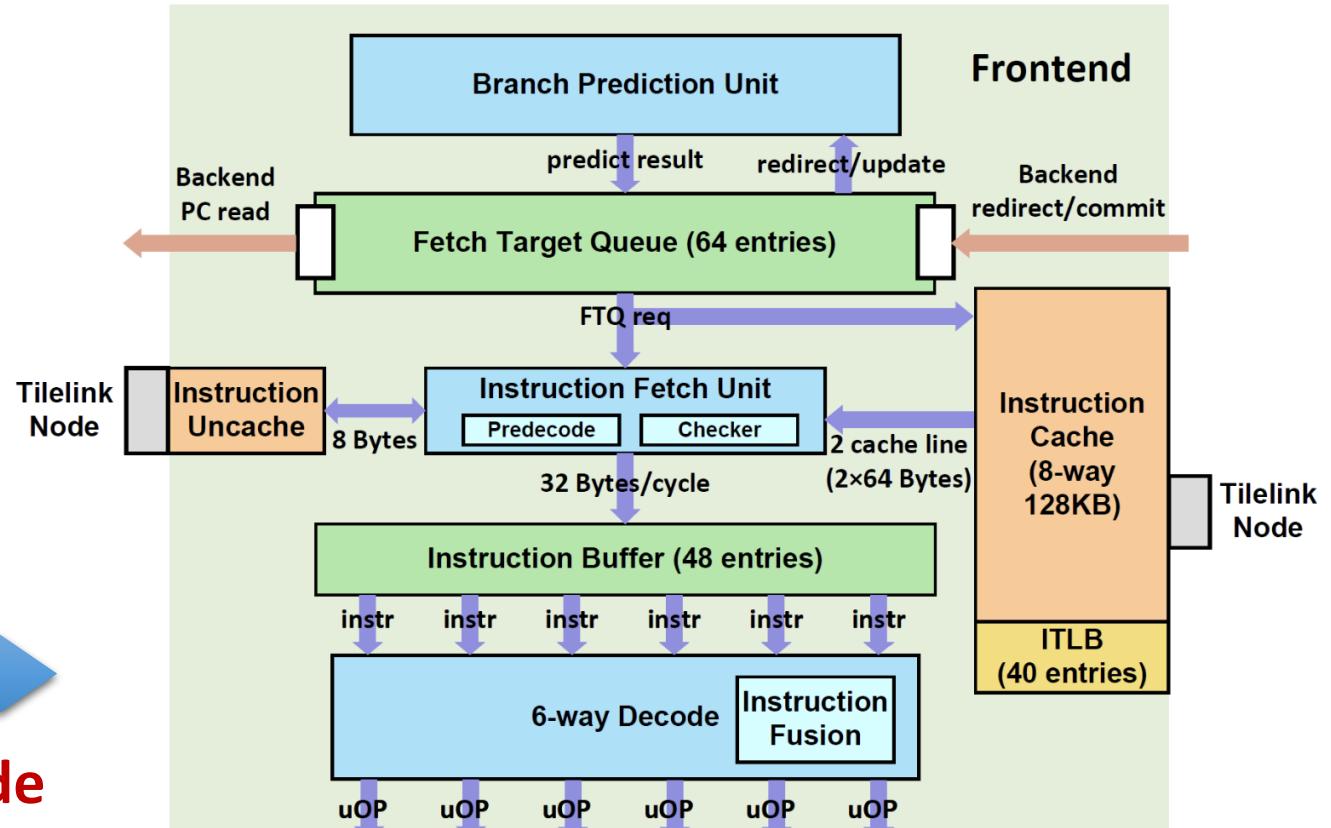


# Frontend

- Decoupled frontend
  - **Producer**: branch predictor unit
  - **Consumer**: instruction fetch unit
- Merit
  - Fewer fetch bubbles
  - Fetch directed instr. prefetching



Upgrade



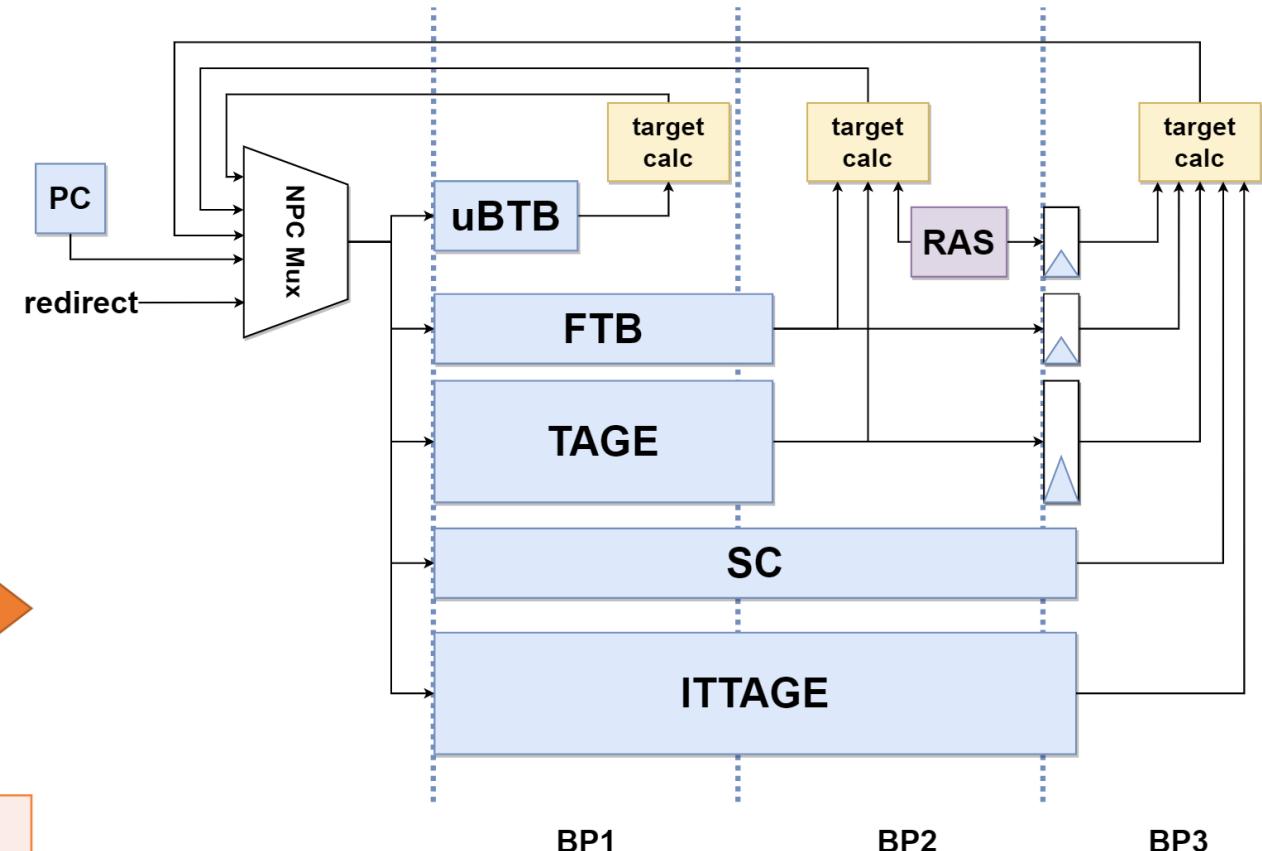
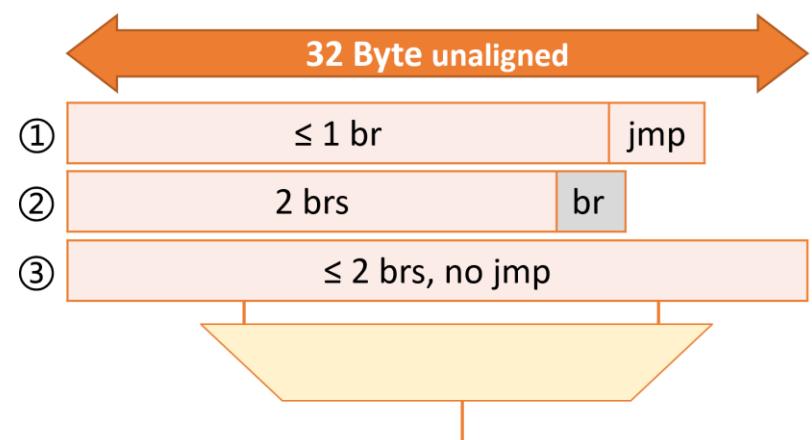
Yanqihu

Nanhu

# Branch Predictor

Three-stage overriding prediction

- Stage 1: uBTB
- Stage 2: FTB + TAGE
- Stage 3: SC + ITTAGE + RAS
- FTB: Fetch Target Buffer

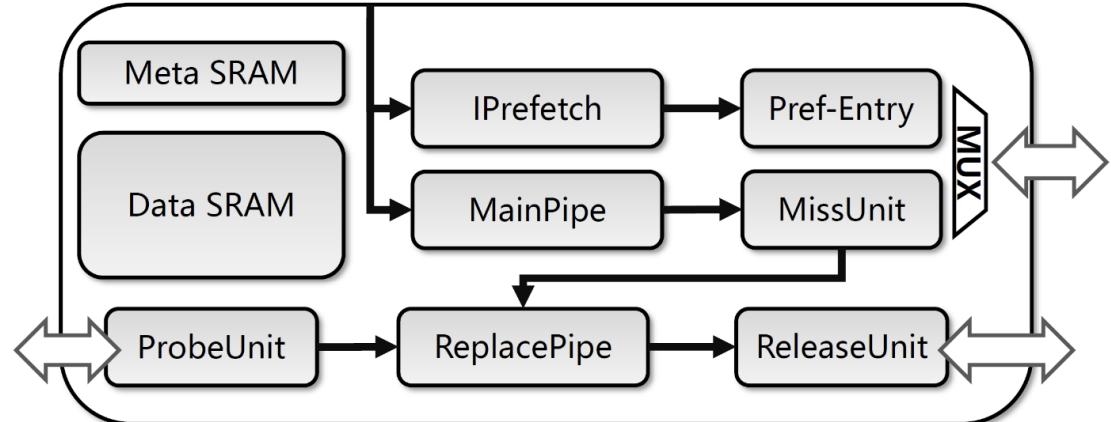




# Instruction Cache

- **Main feature**

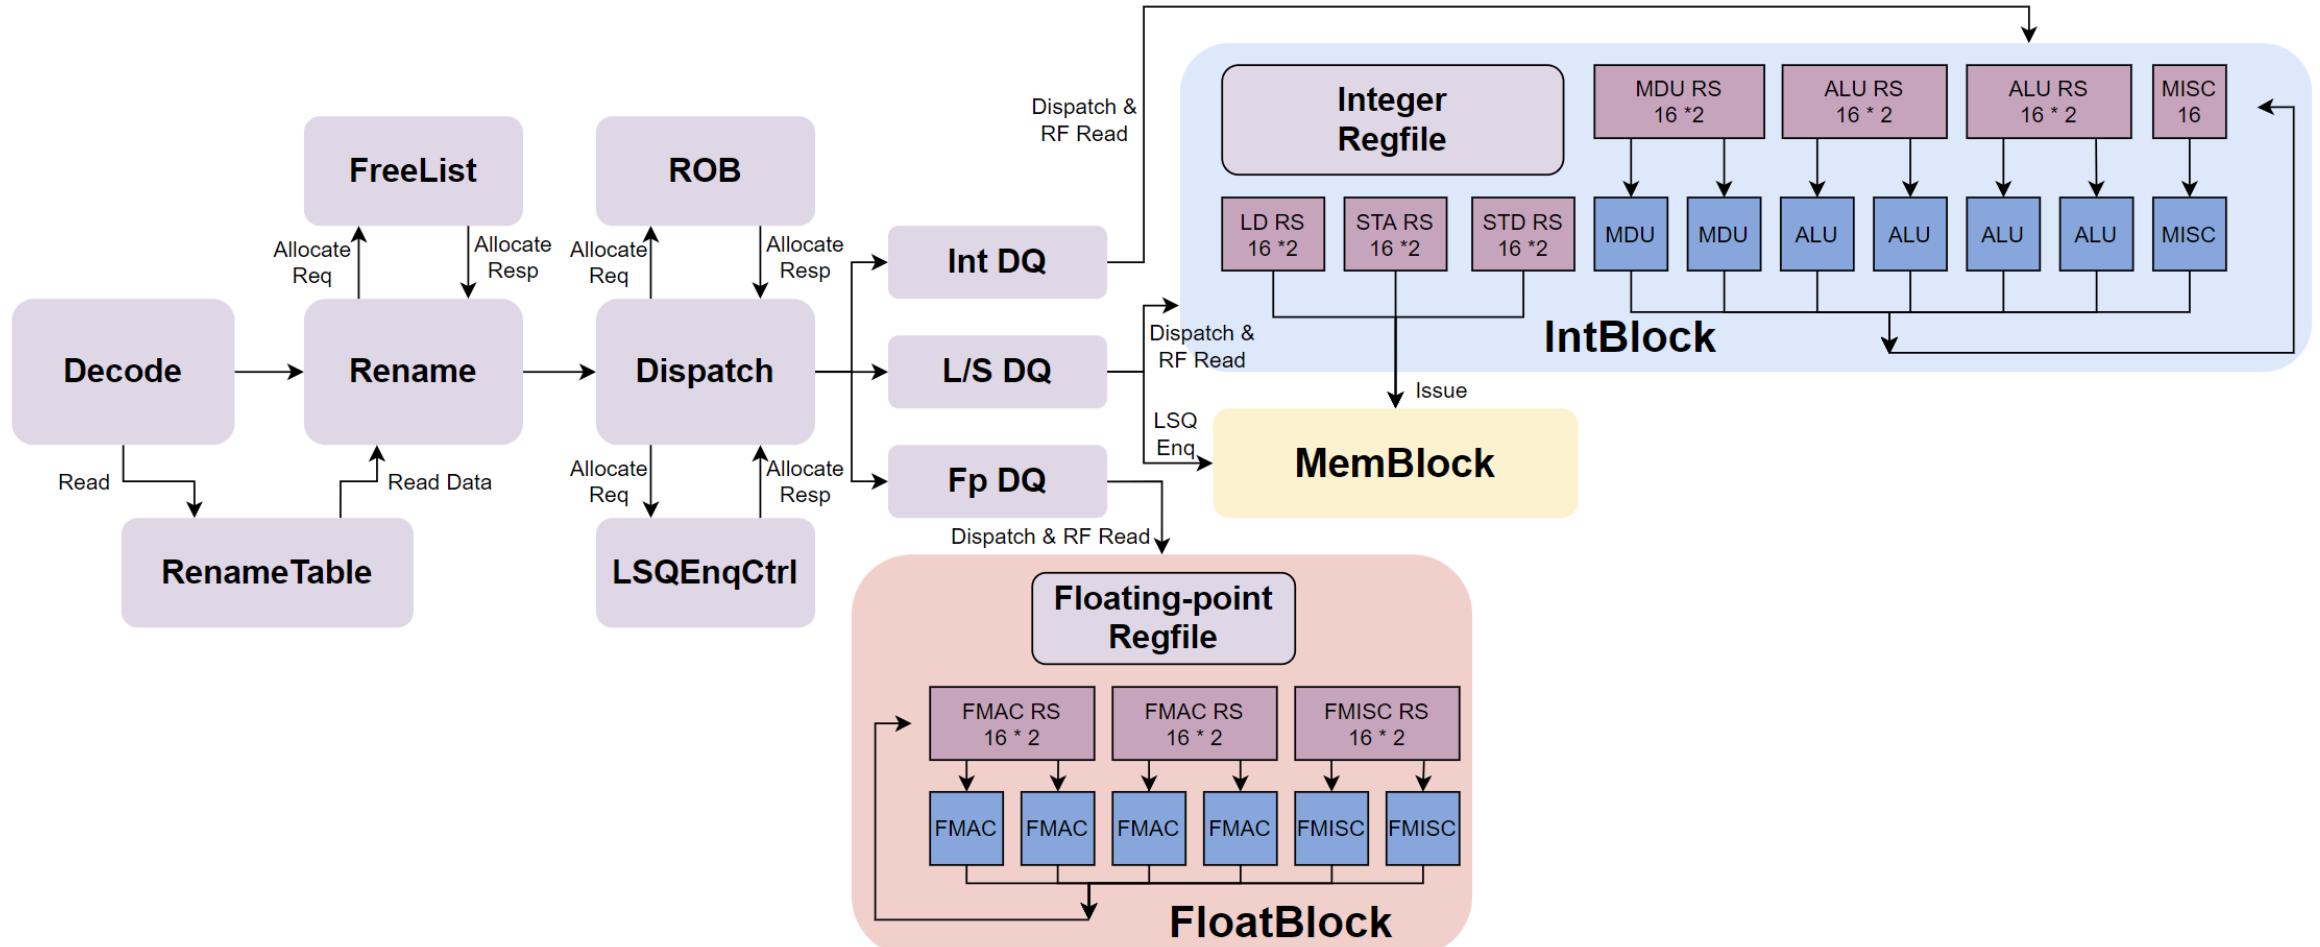
- 64KB 4-w VIPT blocking cache
- Even/odd bank interleaving read
- **Support TileLink protocol**
- PLRU replacement



	Yanqihu	Nanhu	
ICache level	Two	One	Reduce
Size	16KB, 4-w	64KB, 4-w	<b>4×</b>
L1plus Cache	128KB, 8-w	-	Reduce
Read bandwidth	64 B	128 B	<b>2×</b>
TileLink Support	No	Yes	New
Instruction prefetch	Stream	L2 FDIP	New



# Backend





# Instruction Fusion

- Fusion of adjacent UOPs
  - Improve backend efficiency
  - Reduce data bypass delay
- How can we do that?
  1. Reuse of the original calculation
  2. Designing new types of calculations
- **Reduce dynamic instr. greatly**

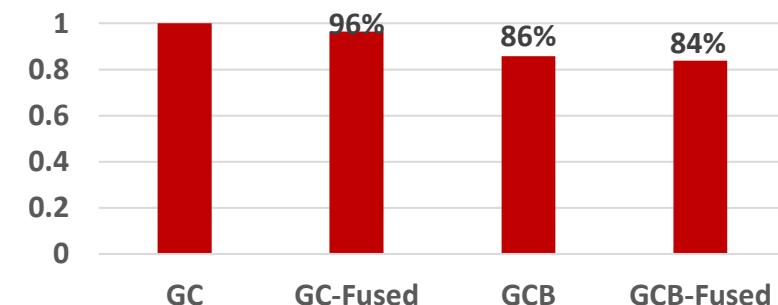
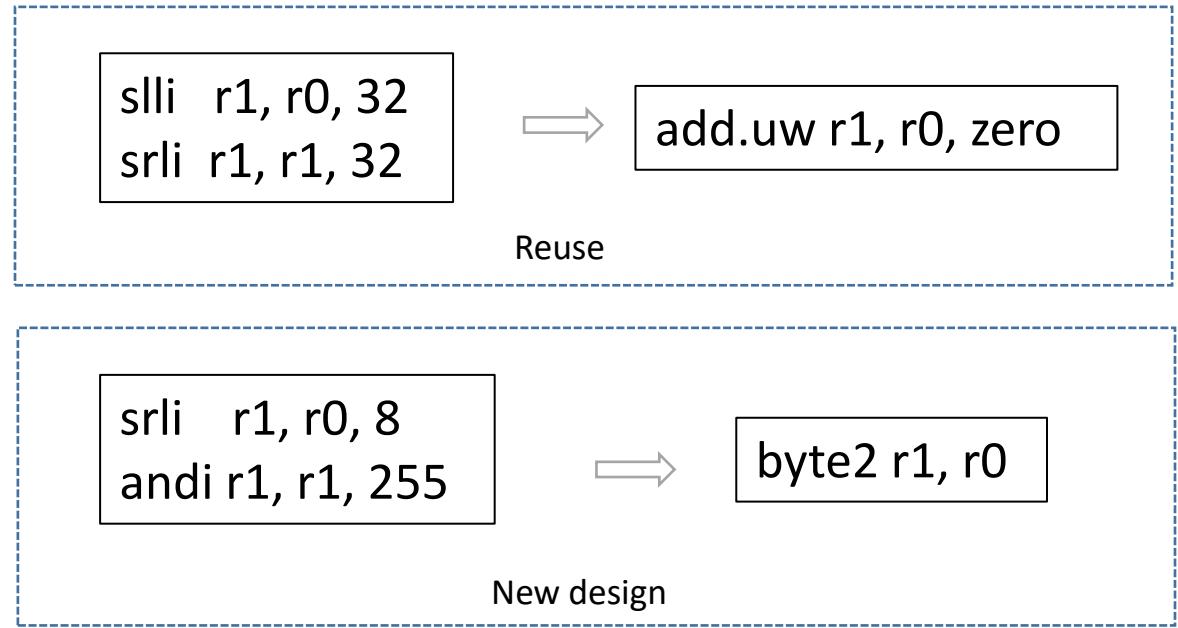
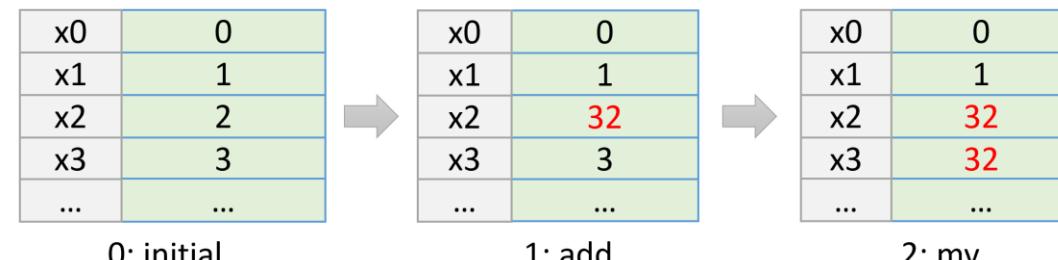
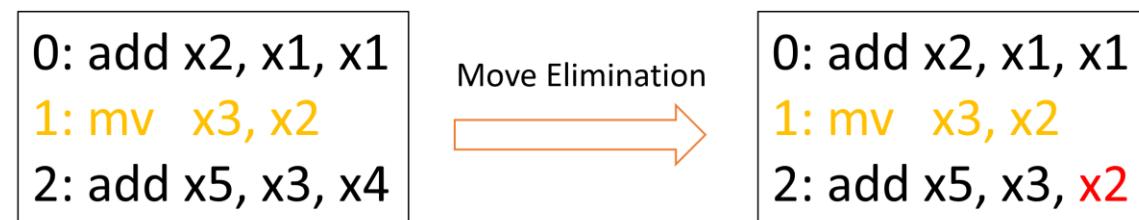


Figure. Normalized Dynamic  
Instruction Count for CoreMark (-O2)



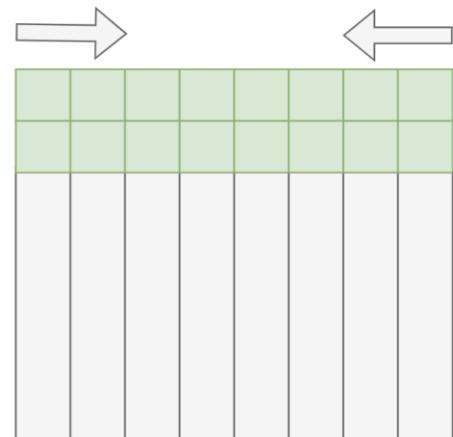
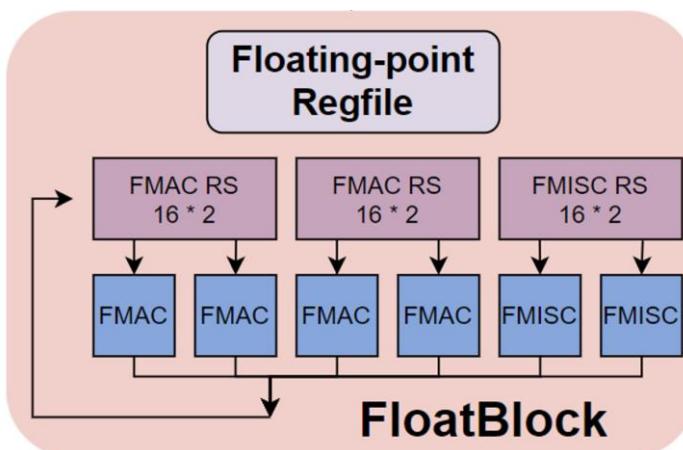
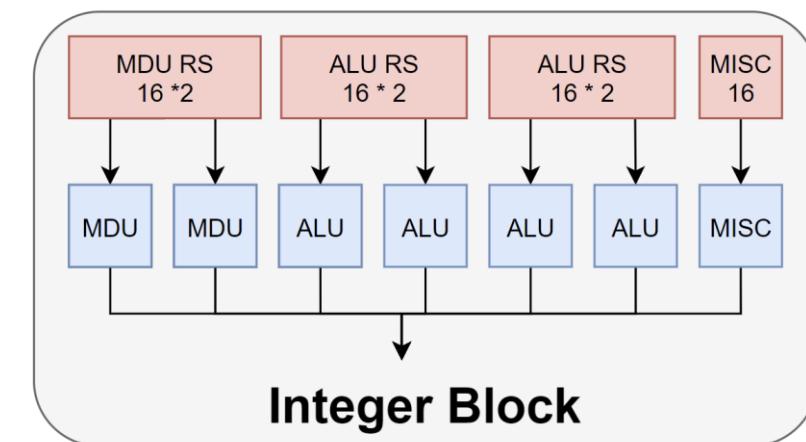
# Rename & Move Elimination

- Maintain the mapping between **arch registers** and **physical registers**
- Move Elimination
  - Mark as **Complete** when Move instruction enters ROB
  - Use **Reference-Counter** to record mapping counts of physical registers



# Reservation Station

- Hybrid reservation station design
  - 2i1o & 1i1o
  - normal selection & age matrix base selection



Normal Select

0	true			
1	true	true		
2	true	false	true	
3	true	false	false	true
	0	1	2	3

Age Matrix



# Floating-Point Units (FPUs)

- Industry competitive Float Point Unit, IEEE 754 compatible

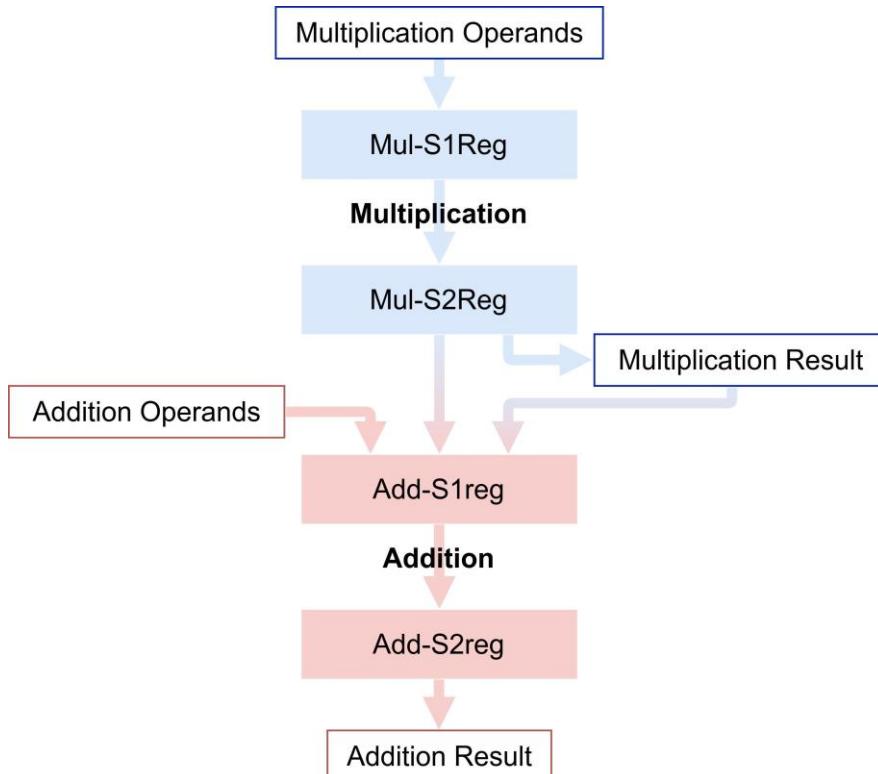
Operations	Latency
<b>FADD</b>	3
<b>FMUL</b>	3
<b>FMA</b>	5
<b>FDIV</b>	$\leq 11$ (float) $\leq 18$ (double)
<b>FSQRT</b>	$\leq 17$ (float) $\leq 31$ (double)
<b>FCVT, FCMP</b>	3

The screenshot shows the GitHub repository page for 'OpenXiangShan/fudian'. The repository is described as an 'Open source high performance IEEE-754 floating unit'. It has 15 branches, 0 tags, and 33 commits. The 'Code' tab is selected. The repository was created on Aug 24, 2022, by 'notlqr'. The 'About' section includes links to 'Readme', '31 stars', '8 watching', and '12 forks'. The 'Releases' section indicates 'No releases published' and 'Create a new release'.

<https://github.com/OpenXiangShan/fudian>

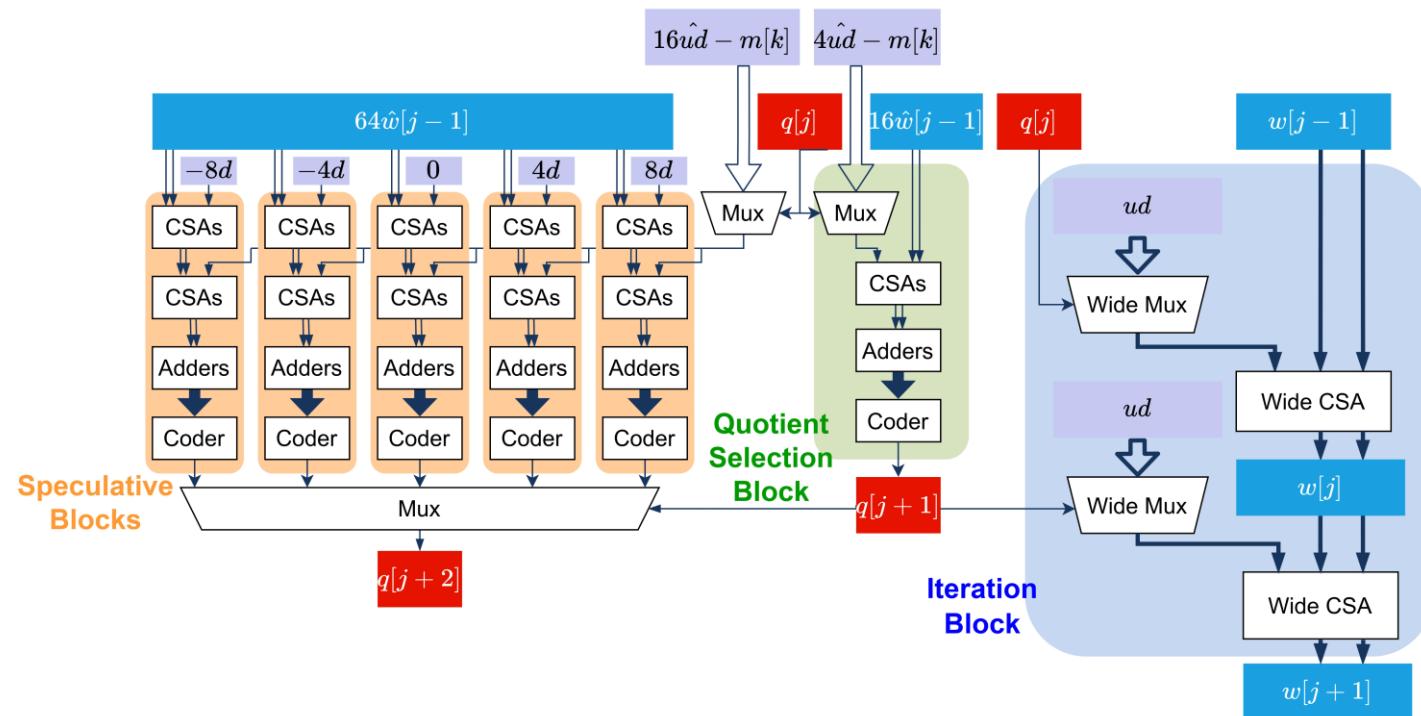
- Cascade FMA

- Reduce FADD delay 5->3



- SRT16 Division units

- $\frac{1}{2}$  delay compared with Yanqihu (SRT4)





# B/K Extension Implemented

- **B:** Bit-Manipulation extension
- **K:** Cryptographic extension
- \* **V(Vector)/H(Hypervisor)** extensions are in development!

1. Extensions . . . . .
1.1. Zba extension . . . . .
1.2. Zbb: Basic bit-manipulation . . . . .
1.2.1. Logical with negate . . . . .
1.2.2. Count leading/trailing zero bits . . . . .
1.2.3. Count population . . . . .
1.2.4. Integer minimum/maximum . . . . .
1.2.5. Sign- and zero-extension . . . . .
1.2.6. Bitwise rotation . . . . .
1.2.7. OR Combine . . . . .
1.2.8. Byte-reverse . . . . .
1.3. Zbc: Carry-less multiplication . . . . .
1.4. Zbs: Single-bit instructions . . . . .

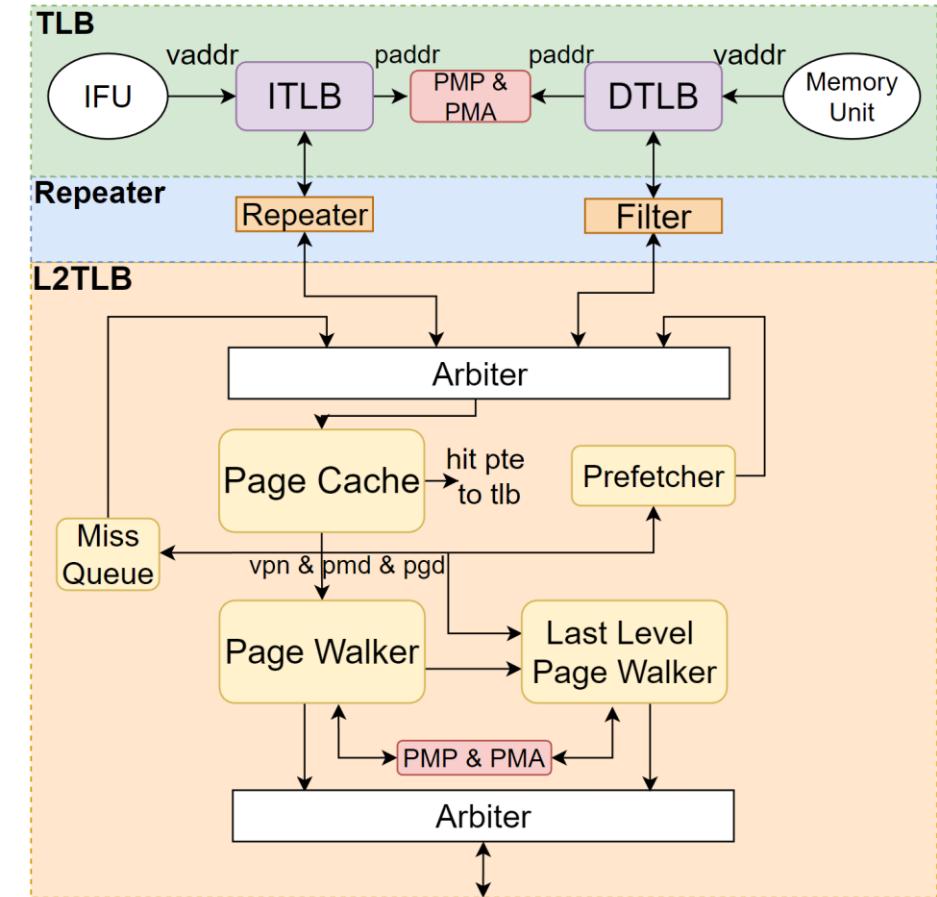
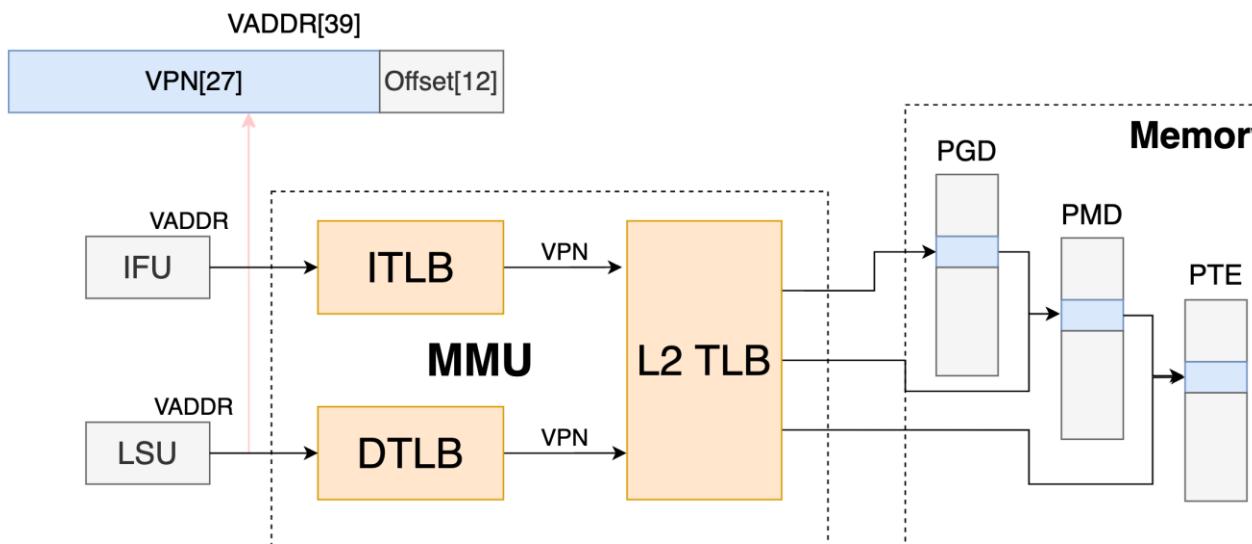
## B Extension

2. Extensions Overview . . . . .
2.1. <a href="#">Zbkb</a> - Bitmanip instructions for Cryptography . . . . .
2.2. <a href="#">Zbkc</a> - Carry-less multiply instructions . . . . .
2.3. <a href="#">Zbkx</a> - Crossbar permutation instructions . . . . .
2.4. <a href="#">Zknd</a> - NIST Suite: AES Decryption . . . . .
2.5. <a href="#">Zkne</a> - NIST Suite: AES Encryption . . . . .
2.6. <a href="#">Zknh</a> - NIST Suite: Hash Function Instructions . . . . .
2.7. <a href="#">Zksed</a> - ShangMi Suite: SM4 Block Cipher Instructions . . . . .
2.8. <a href="#">Zksh</a> - ShangMi Suite: SM3 Hash Function Instructions . . . . .
2.9. <a href="#">Zkr</a> - Entropy Source Extension . . . . .
2.10. <a href="#">Zkn</a> - NIST Algorithm Suite . . . . .
2.11. <a href="#">Zks</a> - ShangMi Algorithm Suite . . . . .
2.12. <a href="#">Zk</a> - Standard scalar cryptography extension . . . . .
2.13. <a href="#">Zkt</a> - Data Independent Execution Latency . . . . .

## K Extension

# Memory Management Unit

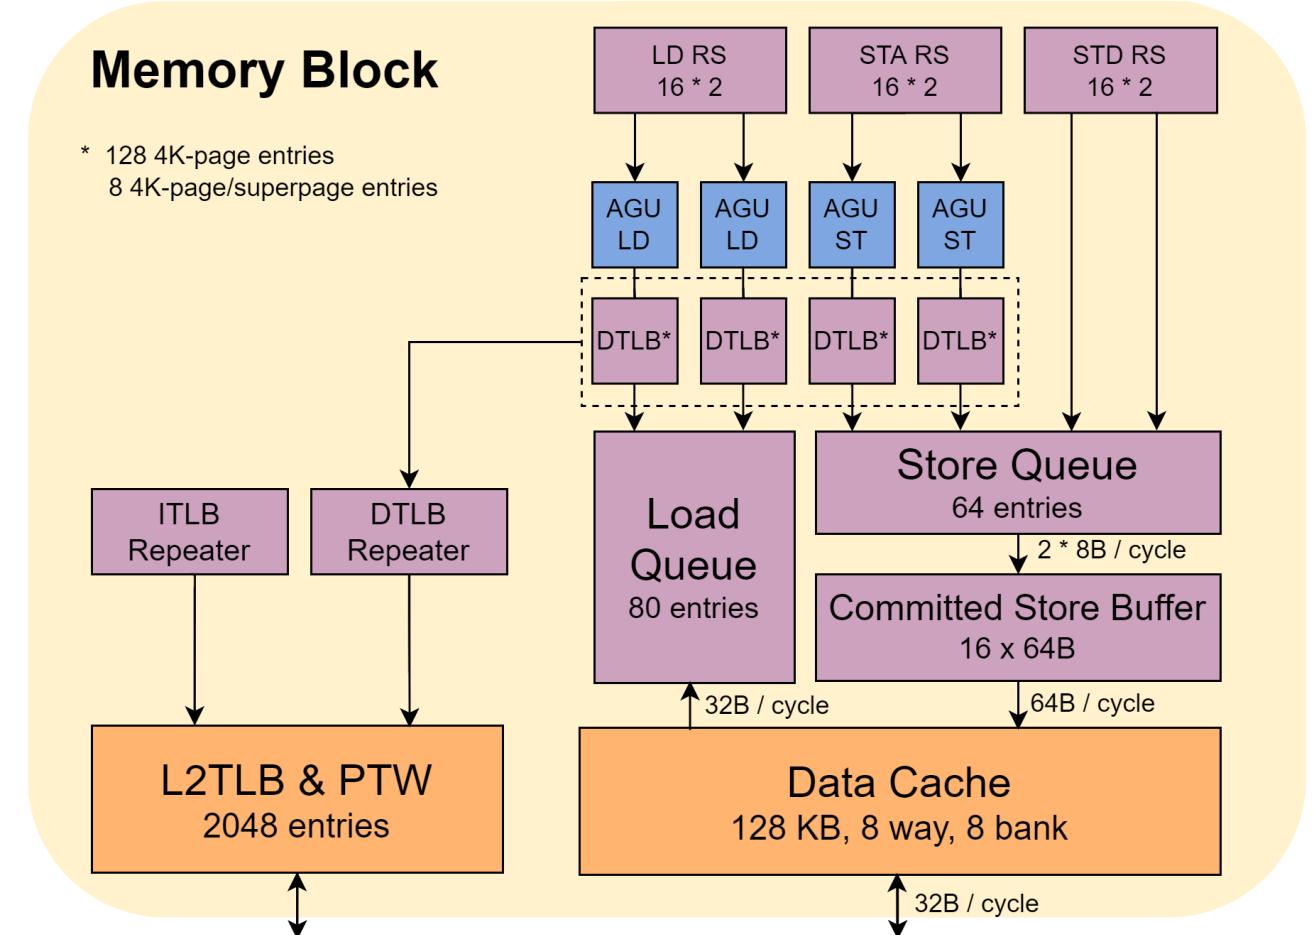
- Translate from virtual address to physical address
- Privilege checking
- Support Sv39



# Memory Block

- Datasheets

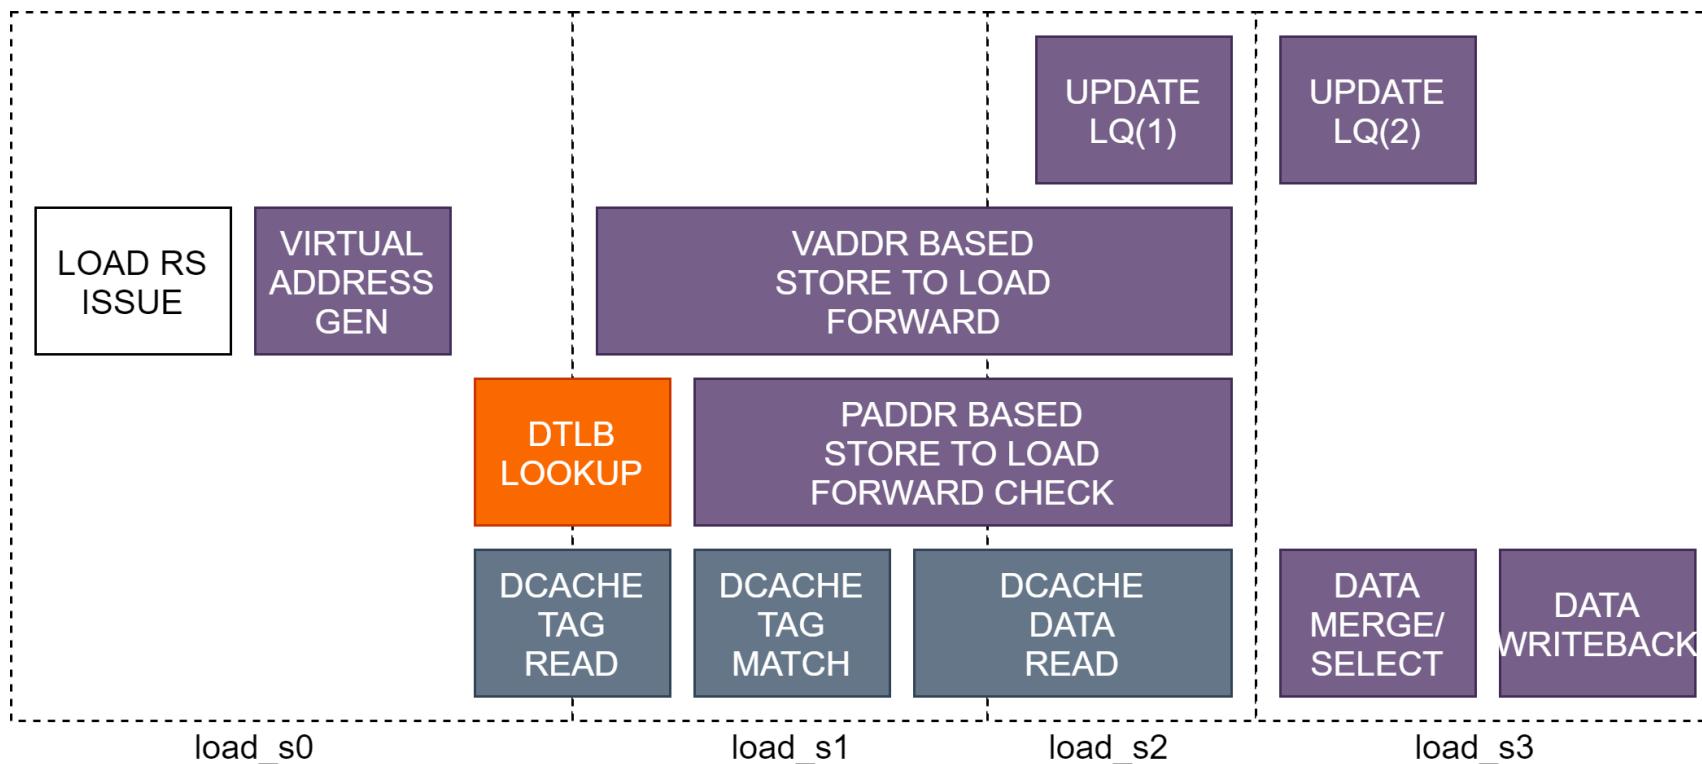
Load pipeline	2
Store address pipeline	2
Store data pipeline	2
L1 LD hit latency	4
L1 LD hit bandwidth	2x8B/cycle
Store data bandwidth	2x8B/cycle
Load Queue Entry	80
Store Queue Entry	64
Merged Store Buffer	16 cachelines





# Load Pipeline

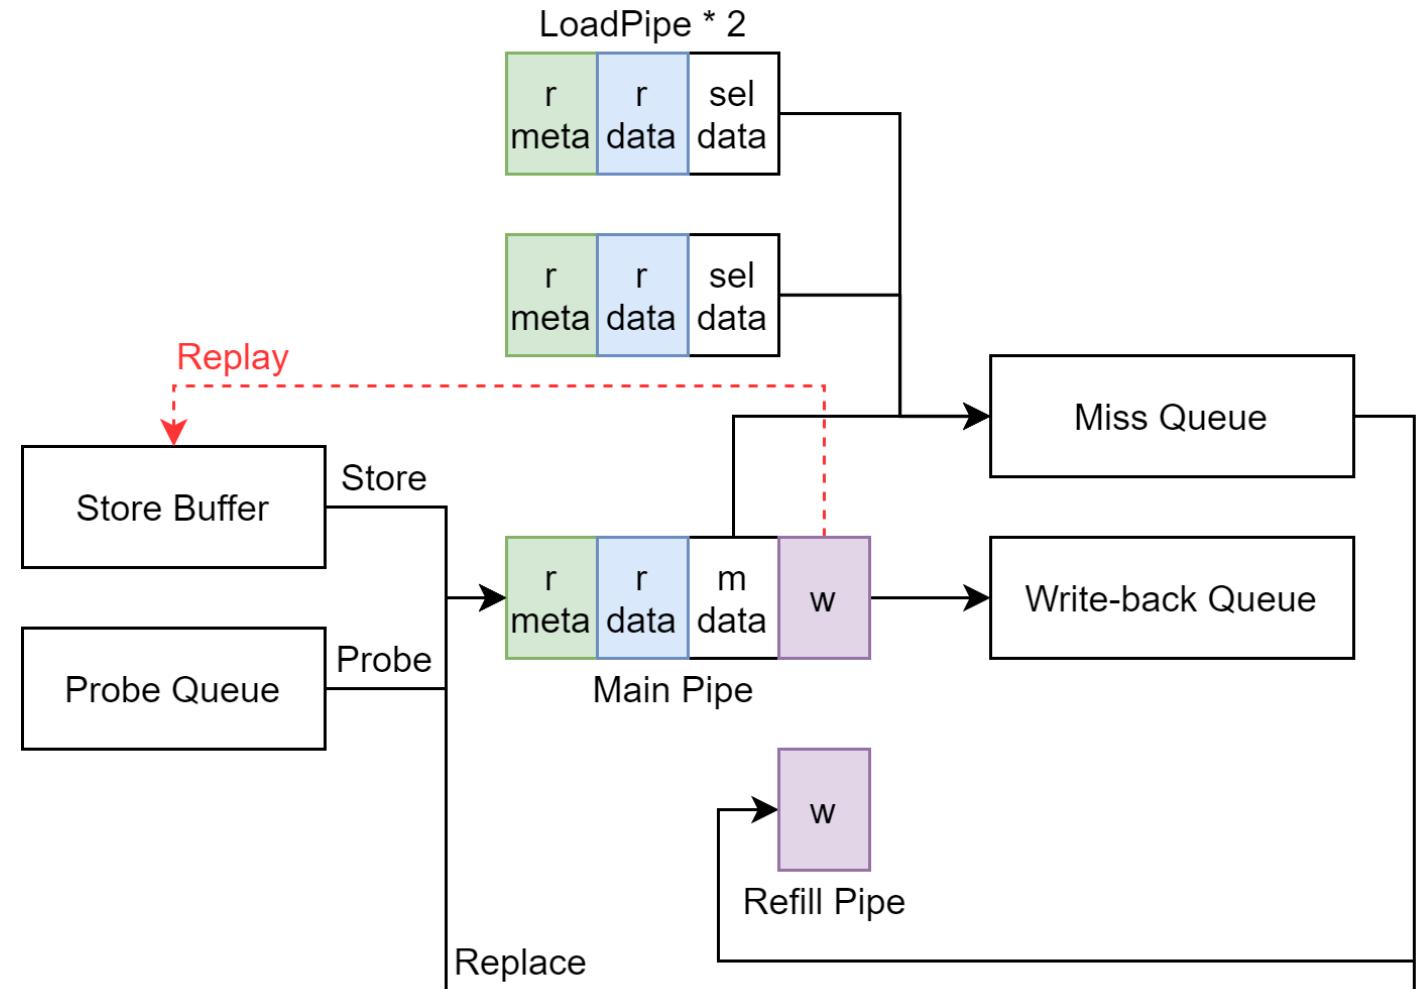
- Four stage
- Optimized for load-to-load & load-store bypass



# L1 DCache

- Coupled with load pipeline
- TileLink protocol

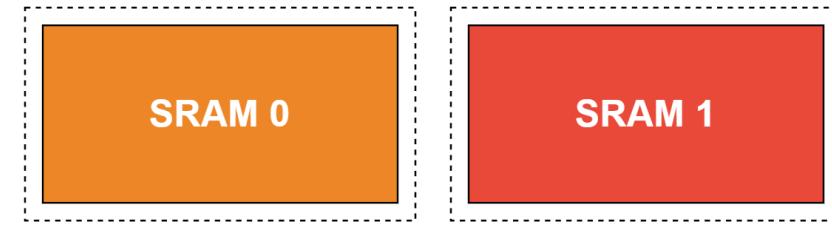
<b>L1 Data cache size</b>	<b>64KB</b>
L1/L2 Bus Width	256bit
Store Buffer	16
Probe Queue	8
Miss Queue	16
Write-back Queue	18
<b>Prefetch Instruction</b>	<b>Support</b>
<b>ECC</b>	<b>Support</b>



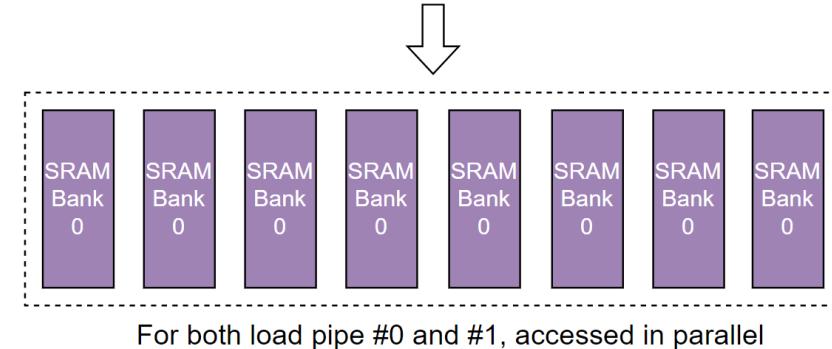
# L1 DCache

- Noteworthy features
  - Multi-bank
  - Anti-alias
  - MSHR request merge
  - Optional L1D prefetcher

Yanqihu



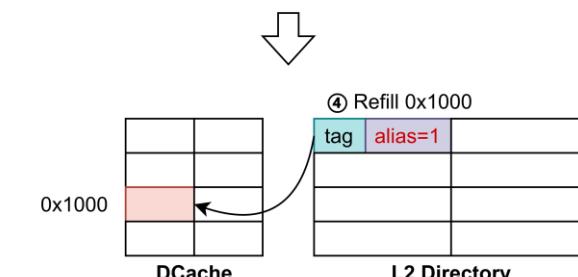
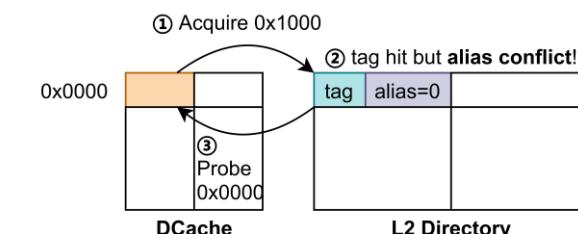
Nanhu



**Alias Bit transferred to L2 cache**

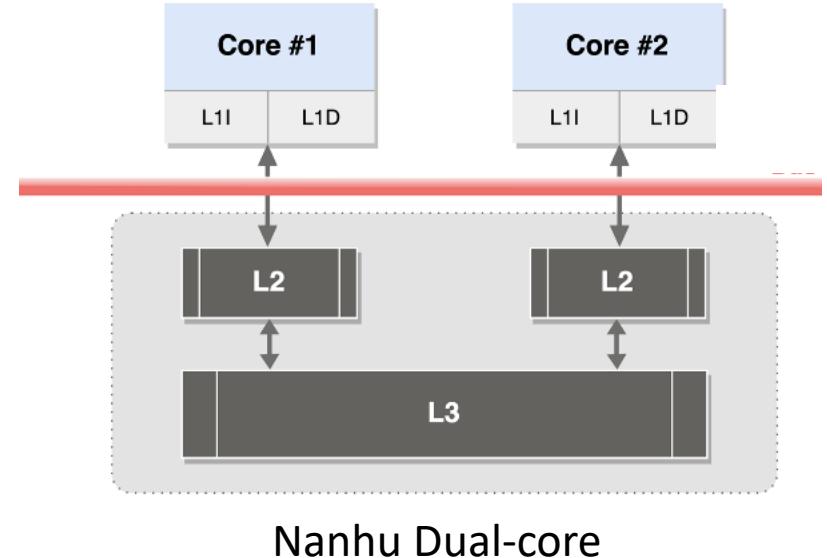


**Use paddr tag for data cache tag comparison**



# L2/L3 Cache

- High performance non-blocking **L2/L3** cache
- Design features
  - **Directory based coherence**
  - **Inclusive/Non-inclusive**
  - **TileLink protocol**
- Highly configurable parameters
  - Slice, Size, #Ways, #MSHRs
  - Replacement: **Random/PLRU/RRIP**
  - Prefetch: **BOP/SMS**
- Maintain coherence with **both** L1I & L1D



Nanhua Dual-core

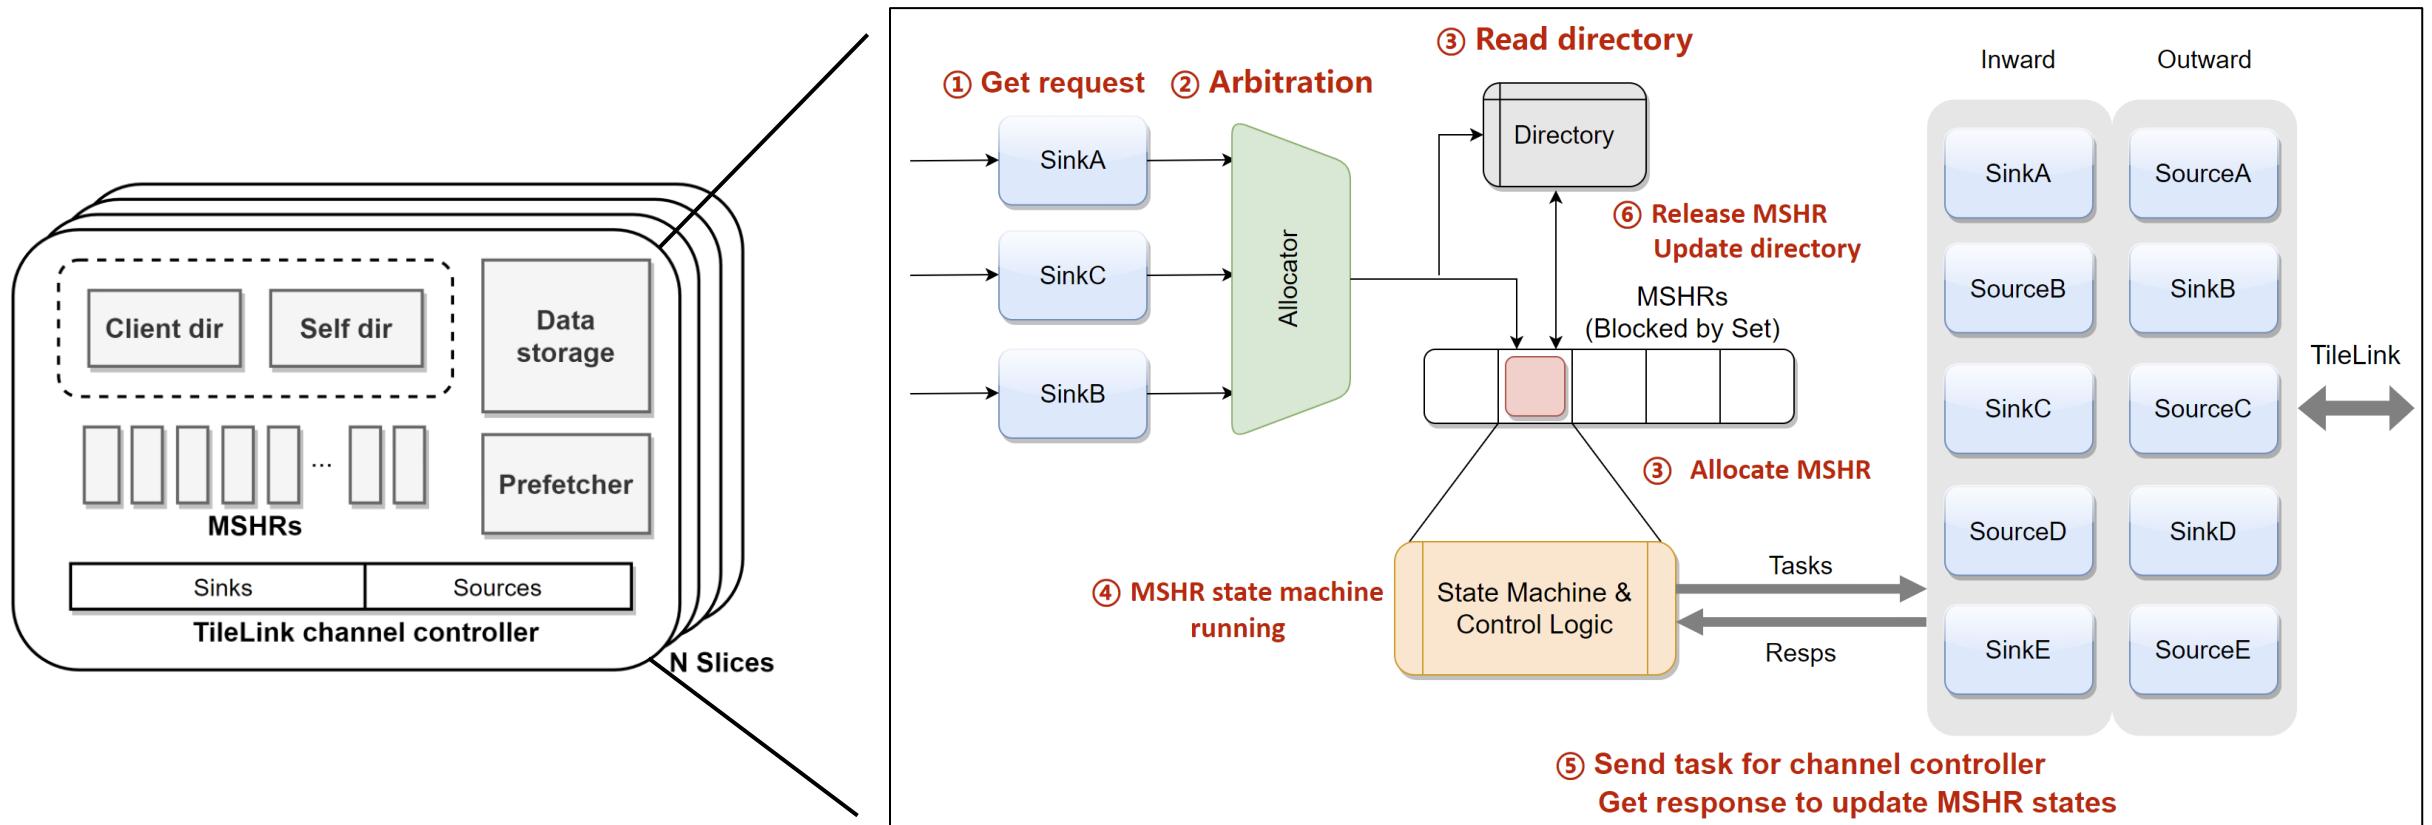
A screenshot of the GitHub repository page for OpenXiangShan/HuanCun. The repository has 36 branches and 613 commits. The commit history shows various contributions from different users, including Lemover, GitHub workflows, rocket-chip, scripts, src, .gitignore, .gitmodules, .mill-version, scalafmt.conf, LICENSE, LICENSE.SiFive, Makefile, README.md, and build.sc. The commits range from 8 days ago to 13 months ago, with many commits from Lemover and GitHub workflows.

Commit	Author	Date	Message
dd9e2e28	dd9e2e28	8 days ago	Merge pull request #91 from OpenXiangShan/sram-old-hazard
Turn on chisel3-plugin	github/workflows	3 months ago	
Bump rocketchip	rocket-chip @ 85f319c	3 months ago	
script: refine tl filter(no tip on tip & must have tip on trunk)	scripts	9 months ago	
util.sram: rm a r/w hazard mux which is not needed	src	10 days ago	
gitignore: add build/	.gitignore	8 months ago	
misc: use https url for git submodule	.gitmodules	6 months ago	
Turn on chisel3-plugin	.mill-version	3 months ago	
Add scalafmt check	scalafmt.conf	13 months ago	
add Mulan PSL2 license (#15)	LICENSE	10 months ago	
add Mulan PSL2 license (#15)	LICENSE.SiFive	10 months ago	
test: enable prefetch during tl-test	Makefile	5 months ago	
Create README.md	README.md	13 months ago	
Turn on chisel3-plugin	build.sc	3 months ago	

<https://github.com/OpenXiangShan/HuanCun>

# L2/L3 Cache

- Concurrent request processing flow





# Microarchitecture Implementation

- Implement by **Chisel** HDL

- High readability
- High modifiability



- Beyond built-ins, we designed some open-source high-performance utilities/components

 BinaryArbiterNode.scala  
 BitUtils.scala  
 ChiselDB.scala  
 CircularQueuePtr.scala  
 Constantin.scala  
 DataModuleTemplate.scala  
 ECC.scala  
 ExcitingUtils.scala

 ExtractVerilogModules.scala  
 FastArbiter.scala  
 FileRegisters.scala  
 GTimer.scala  
 Hold.scala  
 IntBuffer.scala  
 LFSR64.scala  
 LatencyPipe.scala

 LookupTree.scala  
 MIMOQueue.scala  
 Misc.scala  
 ParallelMux.scala  
 Pipeline.scala  
 PipelineConnect.scala  
 PriorityMuxDefault.scala  
 PriorityMuxGen.scala

 RegMap.scala  
 Replacement.scala  
 ResetGen.scala  
 SRAMTemplate.scala  
 StopWatch.scala  
 TLClientsMerger.scala  
 TLEdgeBuffer.scala

**Check it out!**





# Microarchitecture Implementation

- Highly configurable
  - You can setup parameters as you wish

```
IBufSize: Int = 48,  
DecodeWidth: Int = 6,  
RenameWidth: Int = 6,  
CommitWidth: Int = 6,  
FtqSize: Int = 64,  
EnableLoadFastWakeUp: Boolean = true,  
TssQueSize: Int = 16,  
NRPhyRegs: Int = 192,  
LoadQueueSize: Int = 80,  
LoadQueueNWriteBanks: Int = 8,  
StoreQueueSize: Int = 64,  
StoreQueueNWriteBanks: Int = 8,  
RobSize: Int = 256,  
dpParams: DispatchParameters =  
  DispatchParameters(  
    IntDqSize = 16,  
    FpDqSize = 16,  
    LsDqSize = 16,  
    IntDqDeqWidth = 4,  
    FpDqDeqWidth = 4,  
    LsDqDeqWidth = 4  
  ),  
  
  exuParameters: ExuParameters =  
    ExuParameters(  
      JmpCnt = 1,  
      AluCnt = 4,  
      MulCnt = 0,  
      MduCnt = 2,  
      FmacCnt = 4,  
      FmiscCnt = 2,  
      FmiscDivSqrtCnt = 0,  
      LduCnt = 2,  
      StuCnt = 2  
    ),  
  
  prefetcher: Option[PrefetcherParams] =  
    Some(SMSParams()),  
  LoadPipelineWidth: Int = 2,  
  StorePipelineWidth: Int = 2,  
  StoreBufferSize: Int = 16,  
  StoreBufferThreshold: Int = 7,  
  EnableLoadToLoadForward: Boolean = true,  
  EnableFastForward: Boolean = false,  
  EnableLdVioCheckAfterReset: Boolean = true,  
  EnableSoftPrefetchAfterReset: Boolean = true,  
  EnableCacheErrorAfterReset: Boolean = true,  
  EnablePTWPreferCache: Boolean = true,  
  EnableAccurateLoadError: Boolean = true,  
),
```

Some key parameters

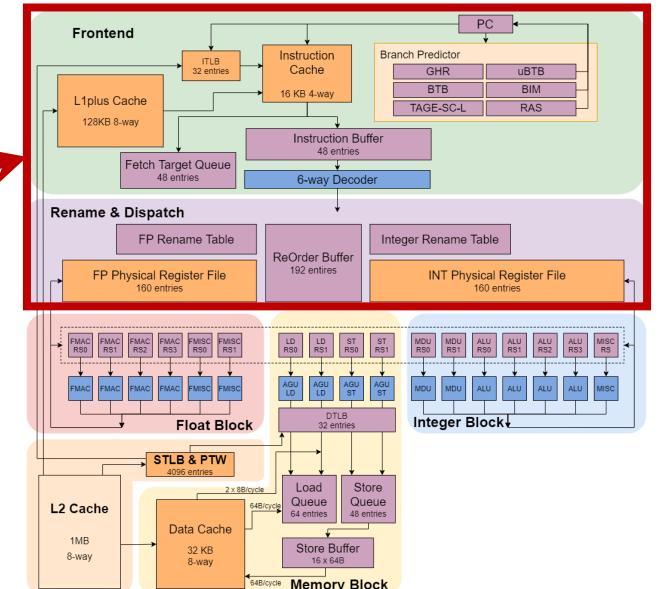
e.g.

How to change rename/decode/issue width?

All you need is to change a number

```
23 // Synthesizable minimal XiangShan  
24 // * It is still an out-of-order, super-scalar arch  
25 // * L1 cache included  
26 // * L2 cache NOT included  
27 // * L3 cache included  
28 class MinimalConfig(n: Int = 1) extends Config(  
29   new DefaultConfig().alter(site, here, up) => {  
30     case SoCParamsKey => up(SoCParamsKey).copy(  
31       cores = up(SoCParamsKey).cores.map(_.copy(  
32         DecodeWidth = 2,  
33         RenameWidth = 2,  
34         FetchWidth = 4,  
35         IssueSize = 8,  
36         NRPhyRegs = 80,  
37         LoadQueueSize = 16,  
38         StoreQueueSize = 16,  
39         RqSize = 32,  
40         BrqSize = 8,  
41         FtqSize = 16,  
42         IBufSize = 16,  
43         StoreBufferSize = 4,  
44         StoreBufferThreshold = 3,  
45         dpParams = DispatchParameters(  
46           IntDqSize = 8,  
47           FpDqSize = 8,  
48           LsDqSize = 8  
        ),  
      ))  
    })  
})
```

All affected  
But it works

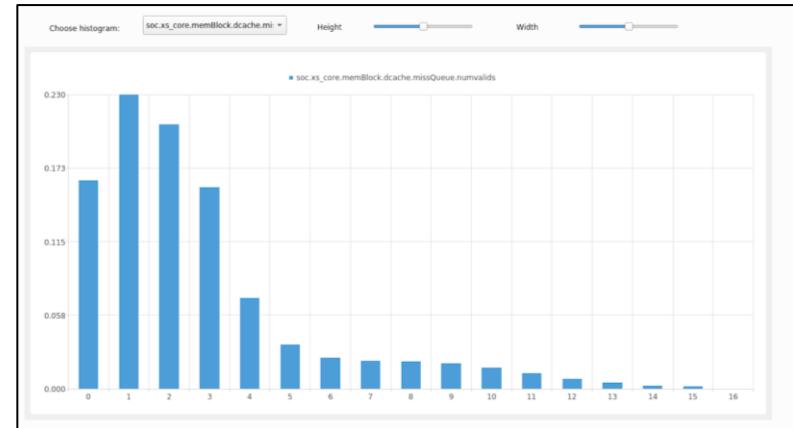




# Microarchitecture Implementation

- Easy-to-use performance counter system
- Multiple collection types
  - Cumulative counter
  - Delay counter
  - Transaction counter
- Multiple analysis style
  - Key-value list
  - Histogram
  - Database
  - Time series visualization

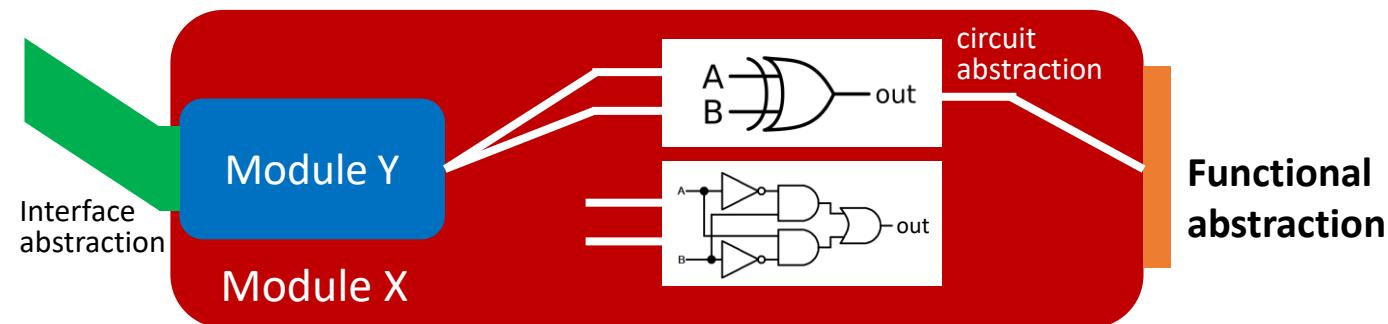
TOP.XSSimSoC.soc_xs_core.ctrlBlock.ftq: BpBInstr,	57899
TOP.XSSimSoC.soc_xs_core.ctrlBlock.ftq: BpBRight,	54695
TOP.XSSimSoC.soc_xs_core.ctrlBlock.ftq: BpBWrong,	3204
TOP.XSSimSoC.soc_xs_core.ctrlBlock.ftq: BpCRight,	3615
TOP.XSSimSoC.soc_xs_core.ctrlBlock.ftq: BpCWrong,	34
TOP.XSSimSoC.soc_xs_core.ctrlBlock.ftq: BpInstr,	78176
TOP.XSSimSoC.soc_xs_core.ctrlBlock.ftq: BpIRight,	5471
TOP.XSSimSoC.soc_xs_core.ctrlBlock.ftq: BpIWrong,	6099
TOP.XSSimSoC.soc_xs_core.ctrlBlock.ftq: BpJRight,	8707
TOP.XSSimSoC.soc_xs_core.ctrlBlock.ftq: BpJWrong,	0
TOP.XSSimSoC.soc_xs_core.ctrlBlock.ftq: BpRight,	68873





# Microarchitecture Implementation

- Modules are designed according to object-oriented programming
- Support multi-level abstraction



- e.g., to implement a new prefetcher
  - We provide a clean set of interfaces that allow the developer to focus on the implementation of the algorithm itself



# KUNMINGHU: 3<sup>rd</sup> generation microarchitecture

- Work with *Beijing Institute of Open Source Chip (BOSC)*
  - Non-profit organization with 18 founding members from the industry
- RISC-V Vector Extension 1.0 Support
- RISC-V Hypervisor Extension 1.0 Support
- Loop predictor, loop buffer
- Refactored issue queues, execution units, load/store units/queues
- L1D prefetcher, cross-level cache optimizations
- AMBA CHI compatible
- Comprehensive verification plan
- Advanced process node and EDA flows
- GEM5 simulator with aligned microarchitecture



The screenshot shows a GitHub repository interface with a sidebar for navigating between branches. The main area displays a timeline of commits grouped by date. Each commit entry includes the author, commit message, date, and a 'Verified' badge. The commits are related to various microarchitectural components like dcache, MMU, and ftq.

- Commits on Mar 19, 2023:
  - Fix reply logic in unified load queue (#1966) - happy-lx committed 2 days ago ✓ Verified
  - util: change ElaborationArtifacts to FileRegisters (#1973) - Mapleza-U committed 2 days ago ✓ Verified
- Commits on Mar 18, 2023:
  - dcache: optimize the ready signal of missqueue (#1965) - happy-lx committed 5 days ago ✓ Verified
  - bump diffest, track master branch (#1967) - Lemover committed 5 days ago ✓ Verified
- Commits on Mar 15, 2023:
  - MMU: Add sector tlb for larger capacity (#1964) - good-circle committed last week ✓ Verified
- Commits on Mar 13, 2023:
  - dcache: fix plru update logic (#1921) - AugustinWillWing committed last week ✓ Verified
- Commits on Feb 27, 2023:
  - c: use checkout@v3 instead of v2 (#1942) - Tang-Haigin committed 3 weeks ago ✓ Verified
- Commits on Feb 22, 2023:
  - ftq: revert #1875, #1920 (#1931) - stencema and lyn committed last month ✓ Verified

Active development on GitHub



# Roadmap of Kunminghu

- **Functional Improvement**
  - Support RISC-V **Vector** extension
  - Support RISC-V **Hypervisor** extension
- **Performance Exploration**
  - Performance upgrade of front-end, back-end, load-store unit and cache
  - Performance model calibrated with RTL
  - Workflow: **DSE on perf Model => Impl. & fine tuning on RTL**
- **Functional Verification**
  - **Hierarchical veri. flow** including ST、IT、UT + FPGA
  - Industrial-grade verification process
- **Physical Design**
  - Experienced physical design team
  - Simultaneous iteration of RTL coding with physical design timing evaluation



# Estimated Performance of Kuminghu

- Estimated **SPECint 2006 40.95, SPECfp 2006 43.96@3GHz**

- Compile with GCC 12.0, -O3, RV64GCB
- SPEC CPU2006 checkpoints selected by Simpoint
- RTL simulation, 1MB L2 + 16MB L3, DDR4-3200MHz under DRAMsim3 (66 cycle access padding)

SPECint 2006	
400.perlbench	37.587
401.bzip2	23.673
403.gcc	46.215
429.mcf	53.832
445.gobmk	29.61
456.hmmmer	40.236
458.sjeng	30.219
462.libquantum	123.249
464.h264ref	57.426
471.omnetpp	32.748
473.astar	28.965
483.xalancbmk	41.511
SPECint2006@3GHz	40.9462

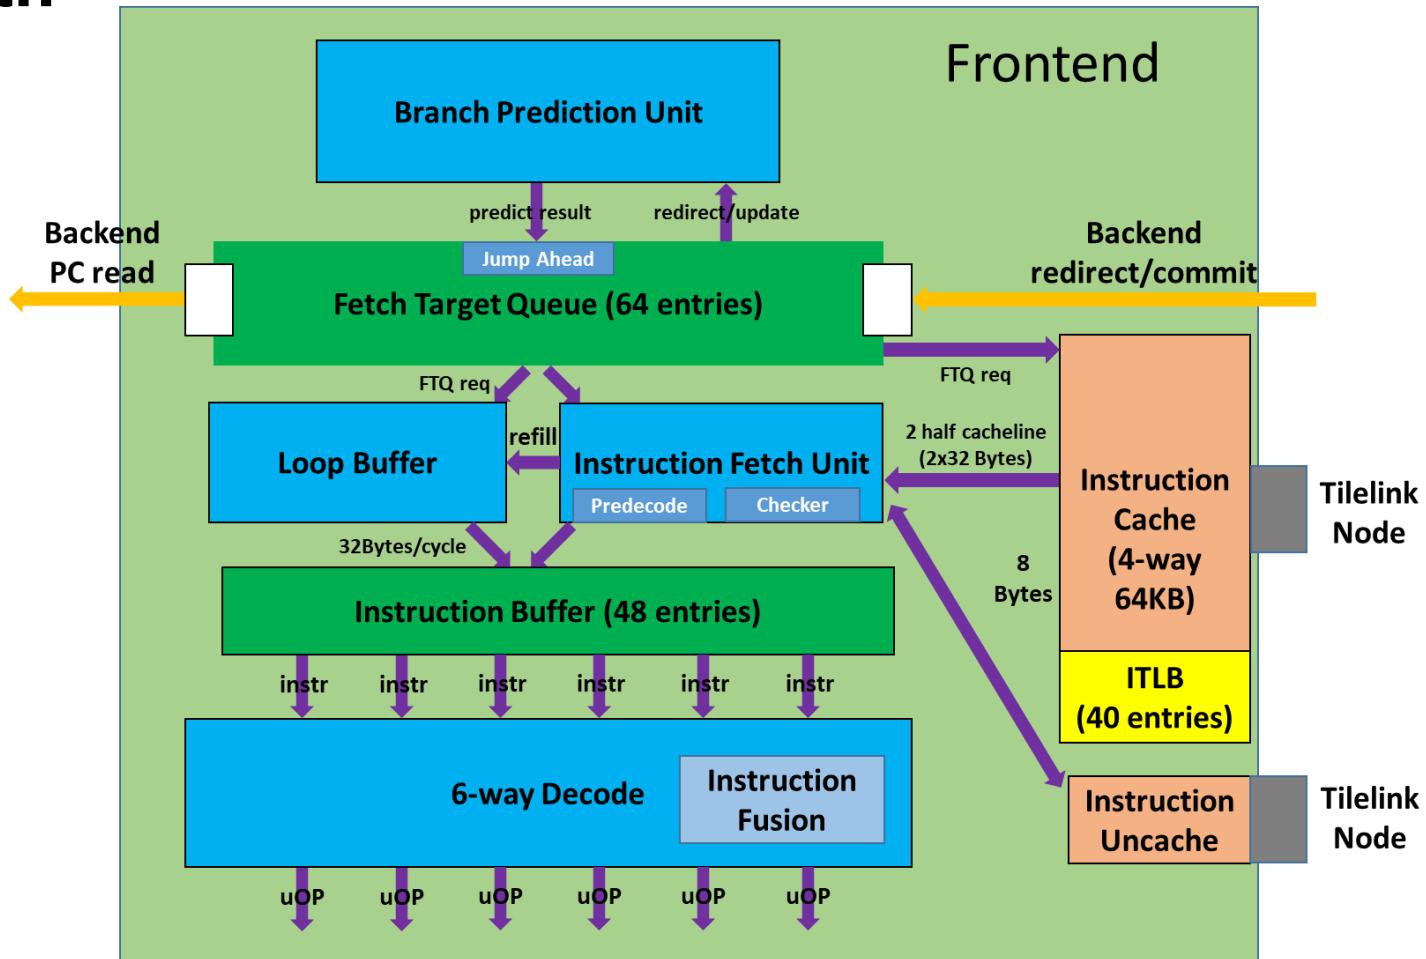
SPECfp 2006	
410.bwaves	76.941
416.gamess	44.178
433.milc	33.402
434.zeusmp	45.684
435.gromacs	31.035
436.cactusADM	47.286
437.leslie3d	42.138
444.namd	38.427
447.dealII	57.177
450.soplex	55.869
453.povray	58.761
454.Calculix	15.135
459.GemsFDTD	37.728
465.tonto	33.228
470.lbm	91.608
481.wrf	37.389
482.sphinx3	56.82
SPECfp2006@3GHz	43.9578

More feature to be merged yet



# ① Frontend Upgrade

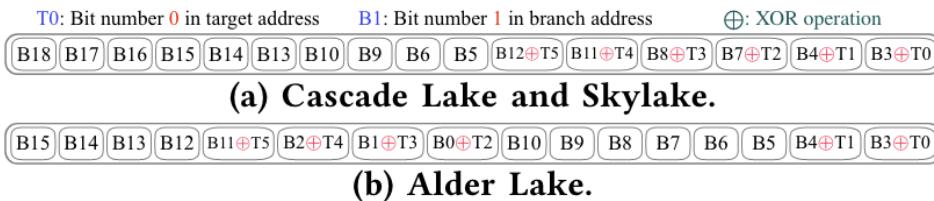
- Larger fetch/prediction width
- Lower prediction miss rate
- Higher frequency





# Feature 1: TAGE Predictor Fine Tuning

- ① Branch History Hashing



Hash algorithm affects performance

Genetic and particle swarm algorithms are employed to adjust parameters automatically.

Param of Nanhu → Param tuned: IPC +1.44%

Bonus!

- ② #Table & History Length

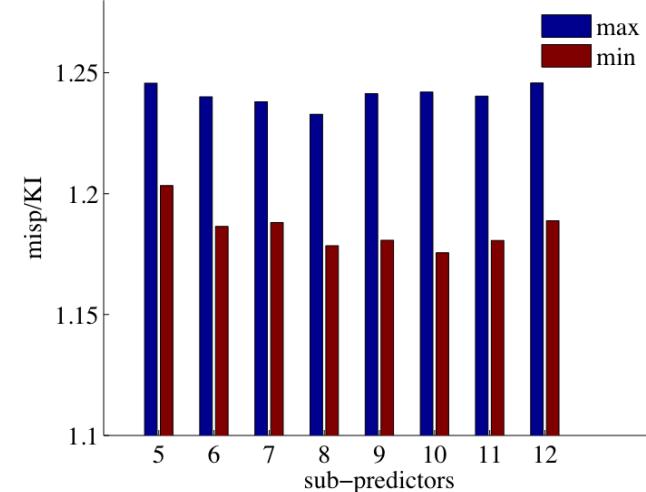


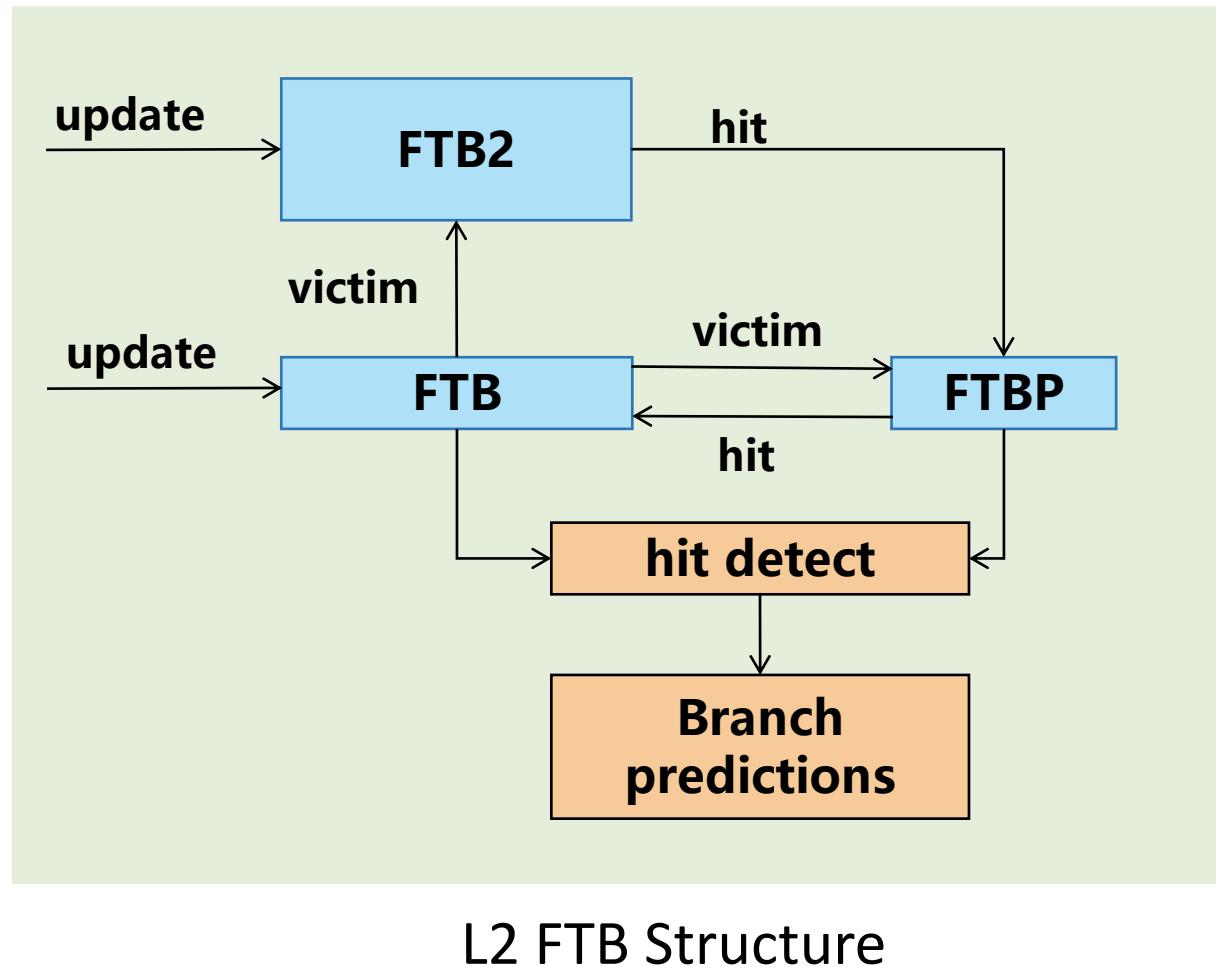
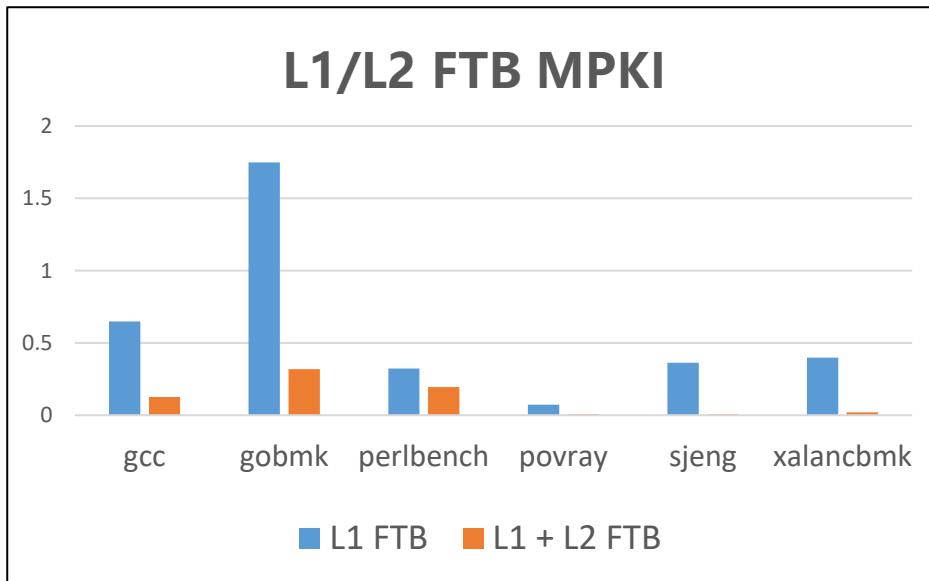
Figure 2: Performance of random parameters

Previous studies have shown that the #tables and history length have a significant effect on TAGE performance, up to 5.64% in MPKI [1]

[1] C. Zhou, L. Huang, Z. Li, T. Zhang, and Q. Dou, “Design Space Exploration of TAGE Branch Predictor with Ultra-Small RAM,” in *Proceedings of the on Great Lakes Symposium on VLSI 2017*, Banff Alberta Canada: ACM, May 2017, pp. 281–286.

# Feature 2: L2 FTB (BTB)

- Hierarchical FTB
  - Support 4K L2 FTB
  - Semi-exclusive
  - FTBP acts as a buffer between L1 & L2 FTB



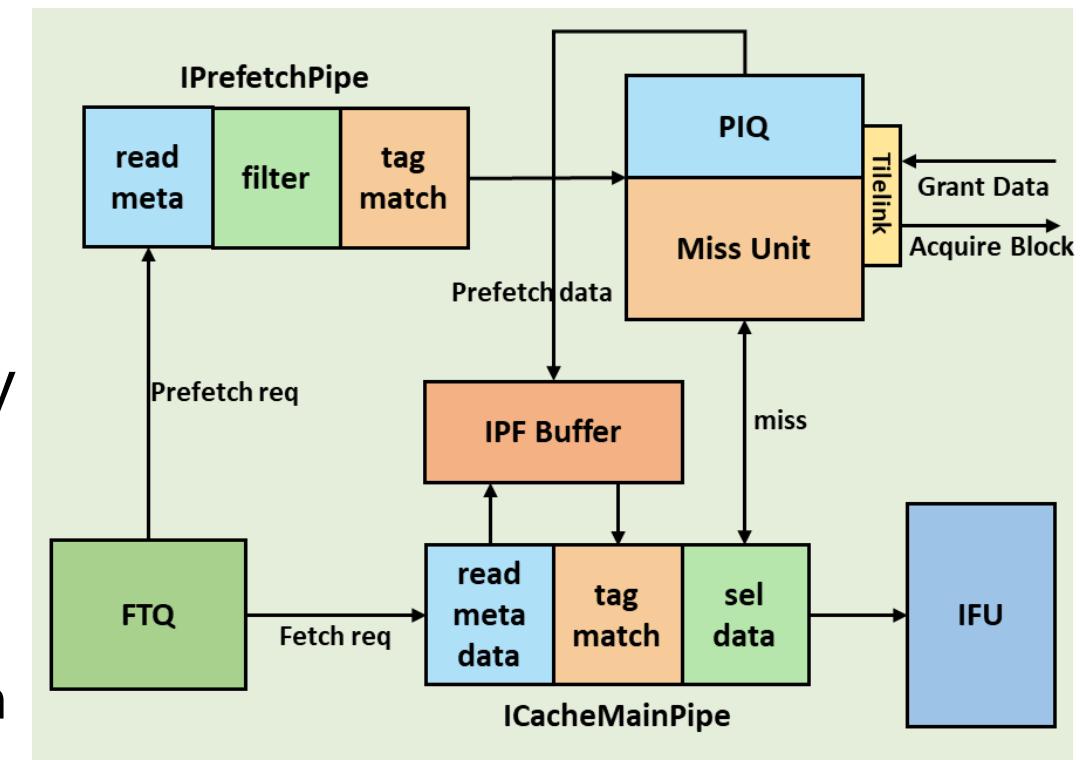
# Feature 3: Smart ICache Prefetcher

- Based decoupled front-end, FDIP prefetching algorithm is implemented

## Prefetcher workflow

1. Select prefetch request from PF-queue
2. Write refilled data into prefetch buffer
3. Query icache & prefetch buffer parallelly

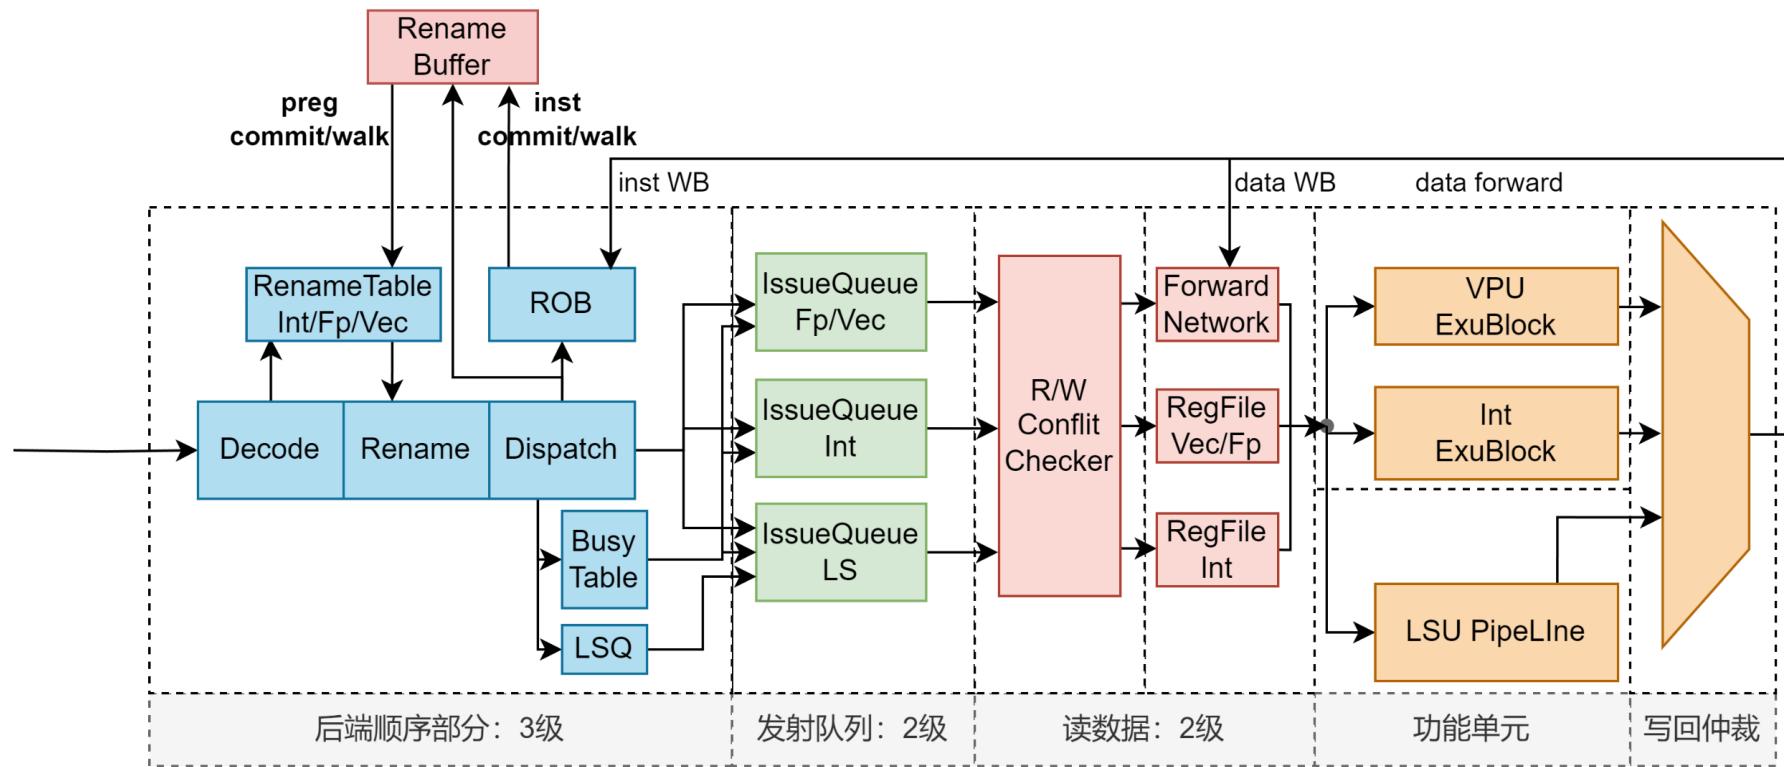
- Prefetch position: L2 Cache → L1 ICache
- Dual prefetching pipelines to enlarge bandwidth





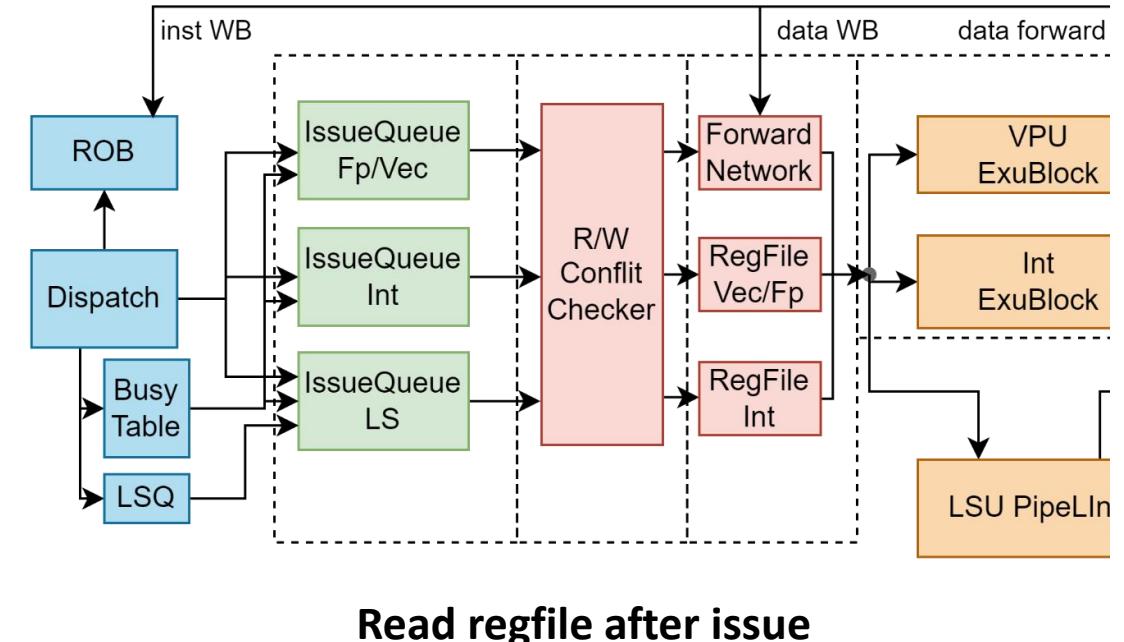
## ② OoO Backend Upgrade

- Read regfile after issue
- Faster μop commit and redirect
- Vector execution unit



# Feature 1: Read regfile after issue

- Advantage
  - Bottleneck is shifted back to  $\mu$ op issue
  - Reduce RS storage to save SRAM
  - Reduce latency and increase RS capacity
- Optimization
  - Efficient arbitration of regfile reading ports
  - Accurate speculative wakeup and cancelling



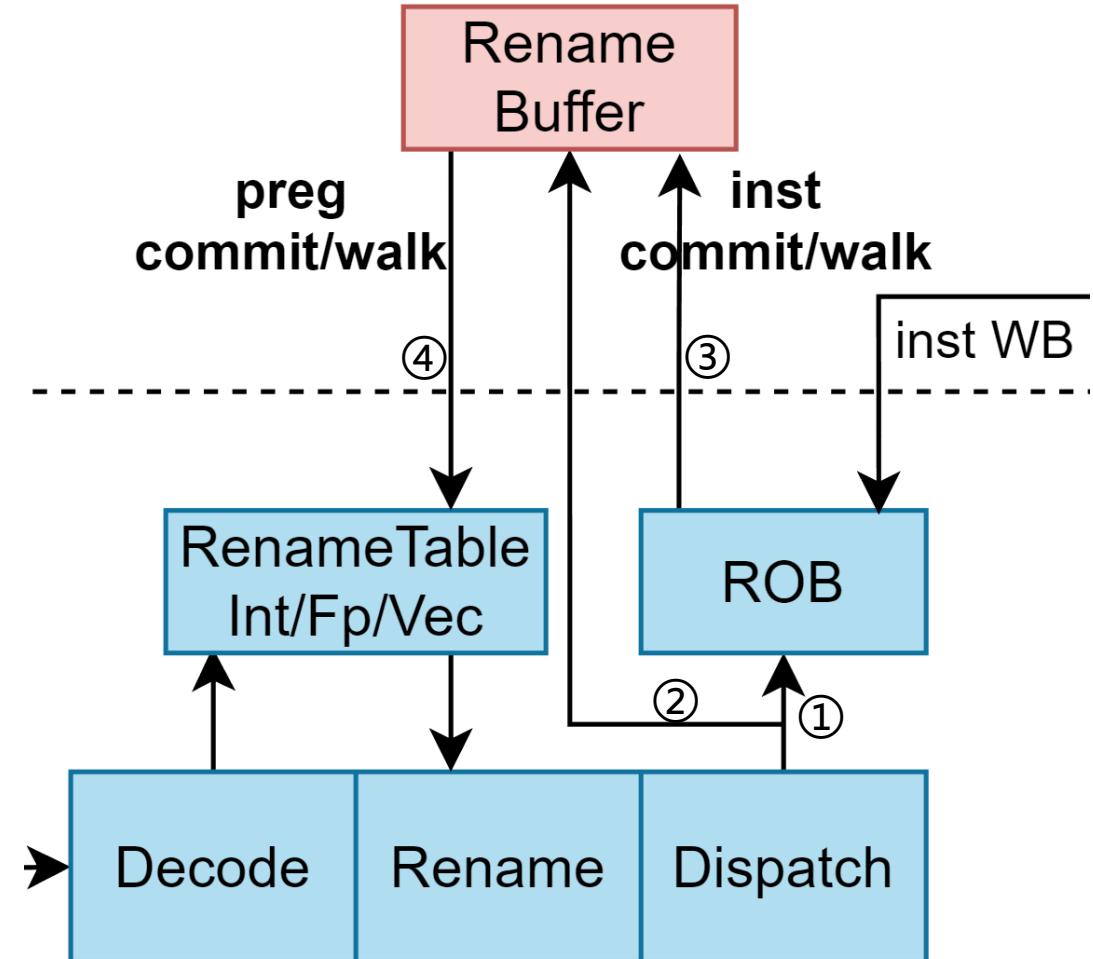


## Feature 2: Decouple μop commit & rename-table update

- New Rename Buffer (**RAB**) design: record μops that have been committed but not updated the RAT

- ① Save regfile mapping info to RenameBuffer (**RAB**)
- ② On instruction commit/redirect, wake RAB up
- ③ RAB is responsible to update RenameTable (**RAT**)

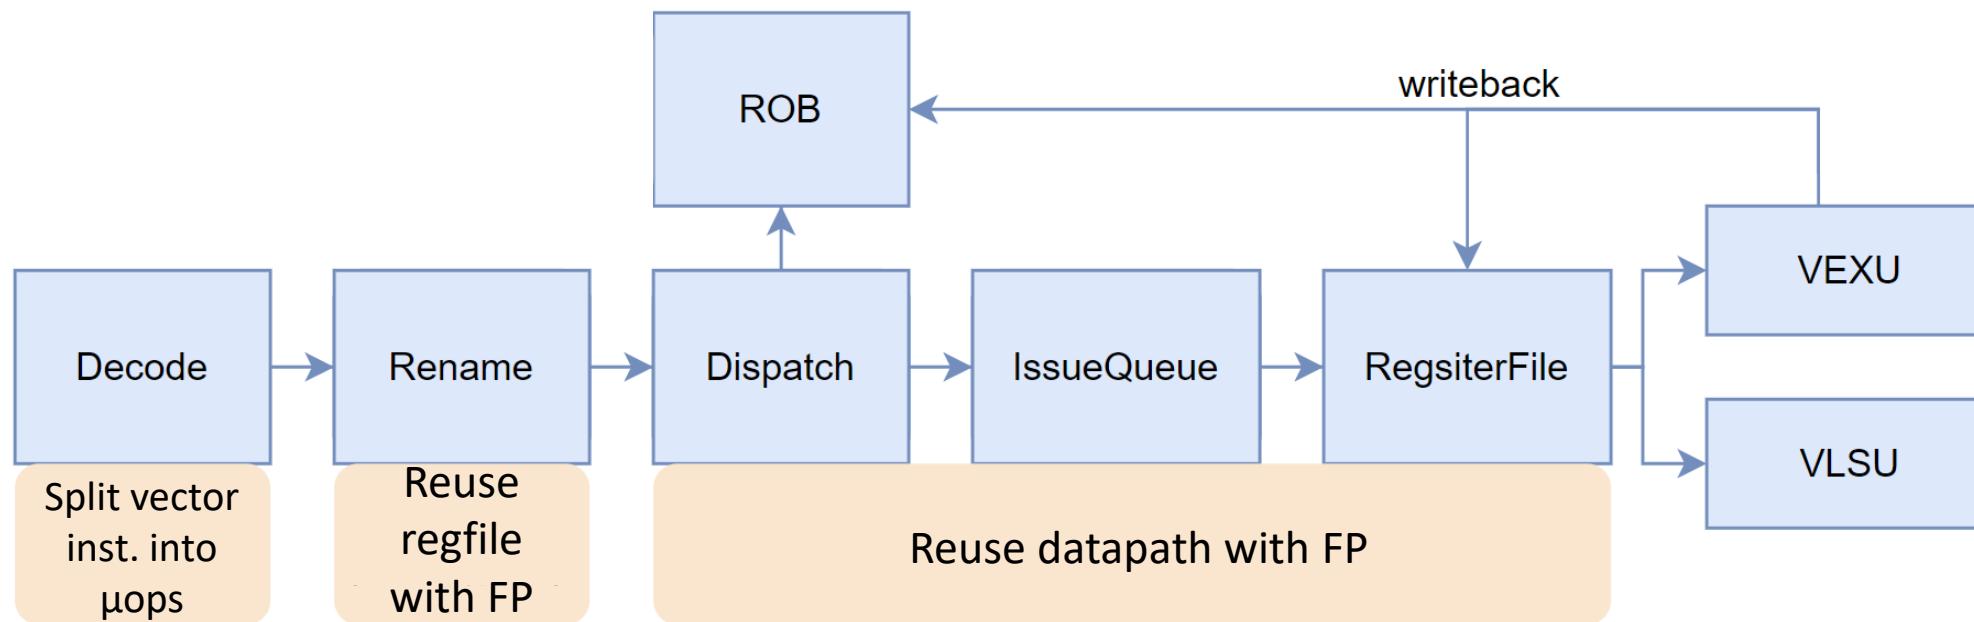
- Benefits
  - Reduce ROB entry and optimize timing
  - Support μop split of vector instructions
  - Support ROB compression





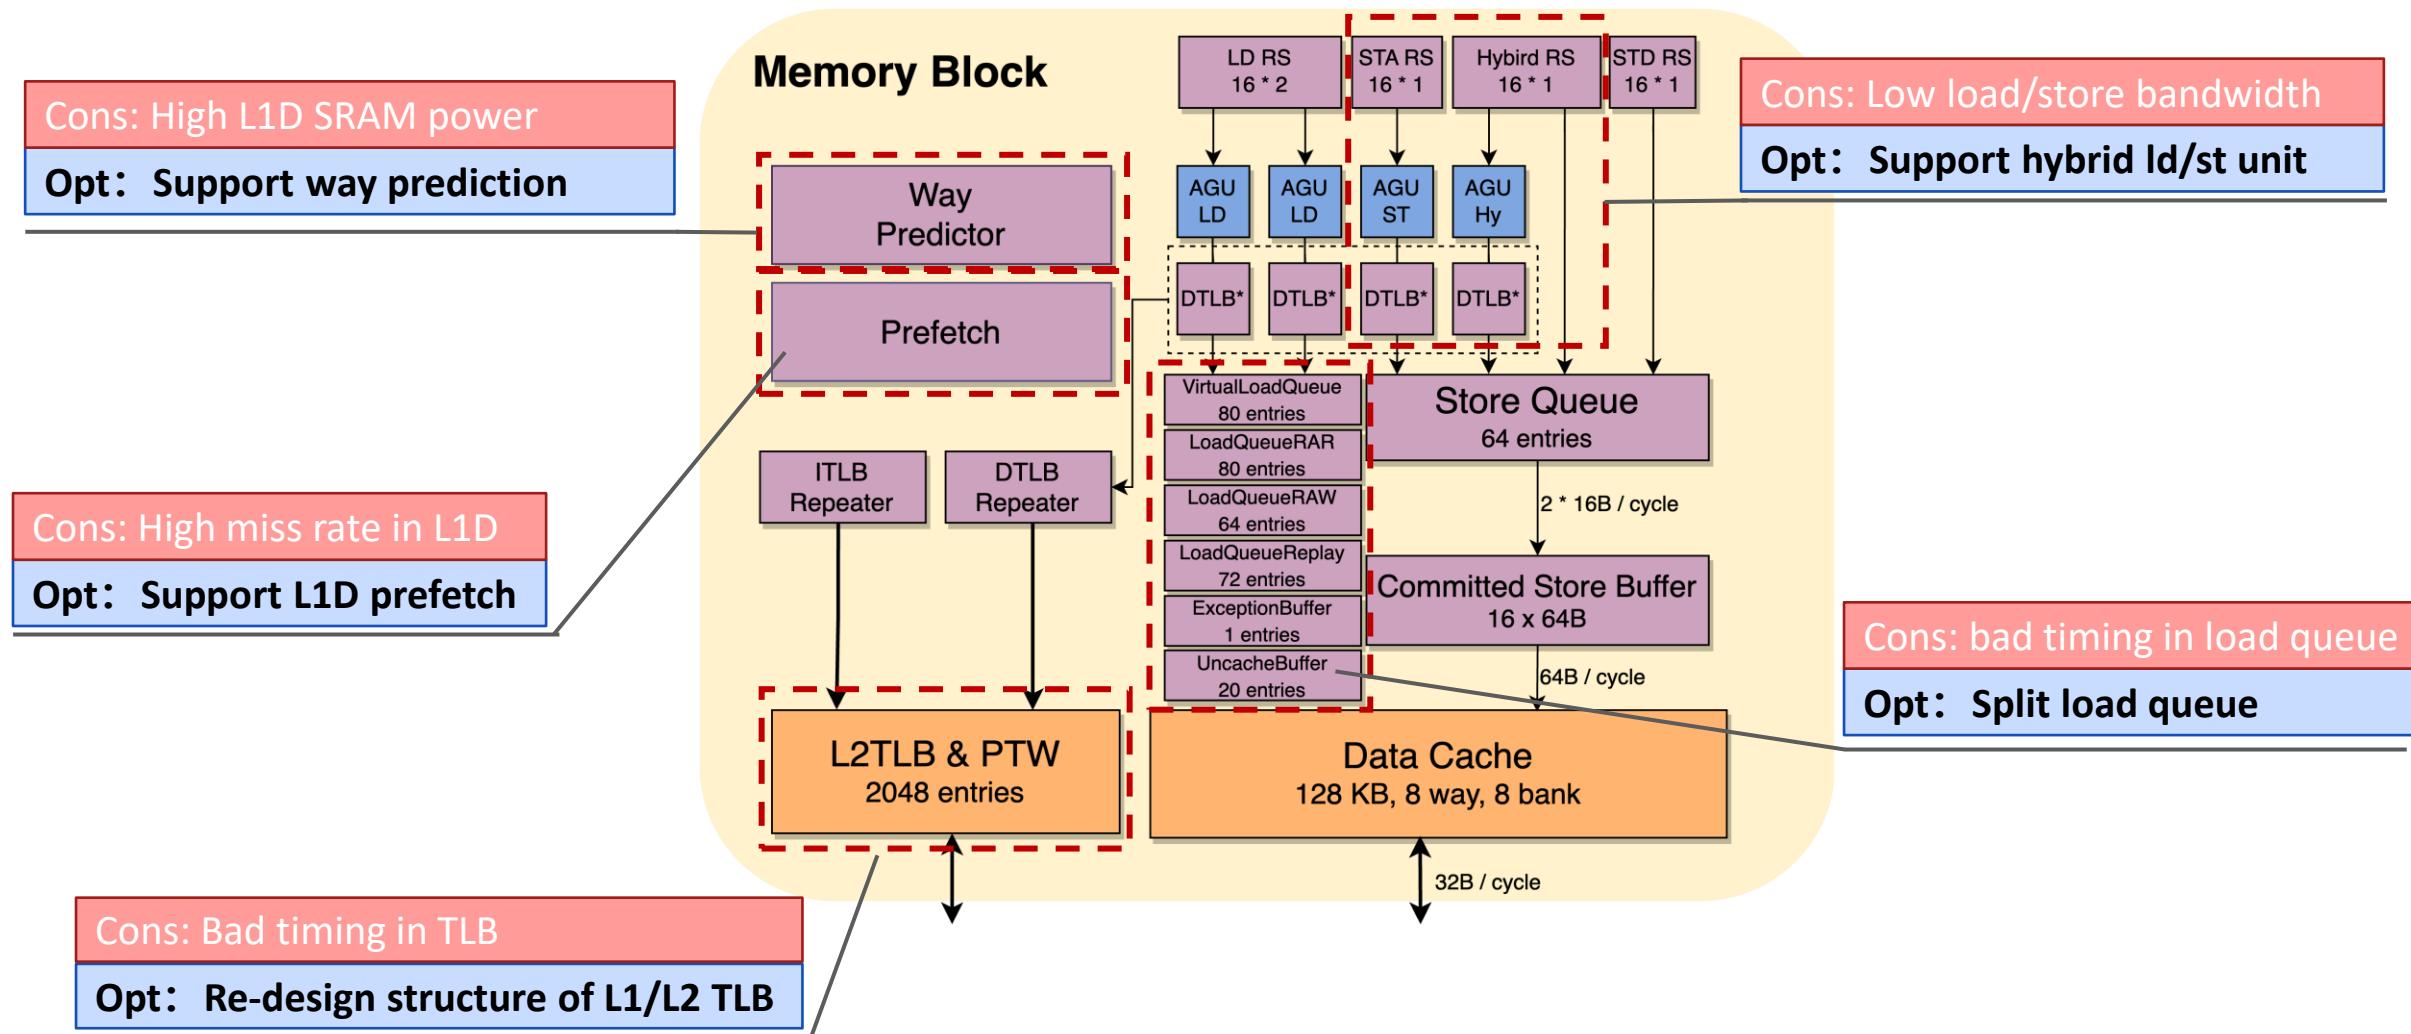
# Feature 3: Support Vector Extension

- Decoupled vector unit design, Compatible with RISC-V Vector 1.0
  - VLEN = 128
- Reuse backend pipeline: Decode/Dispatch/OoO/Execute





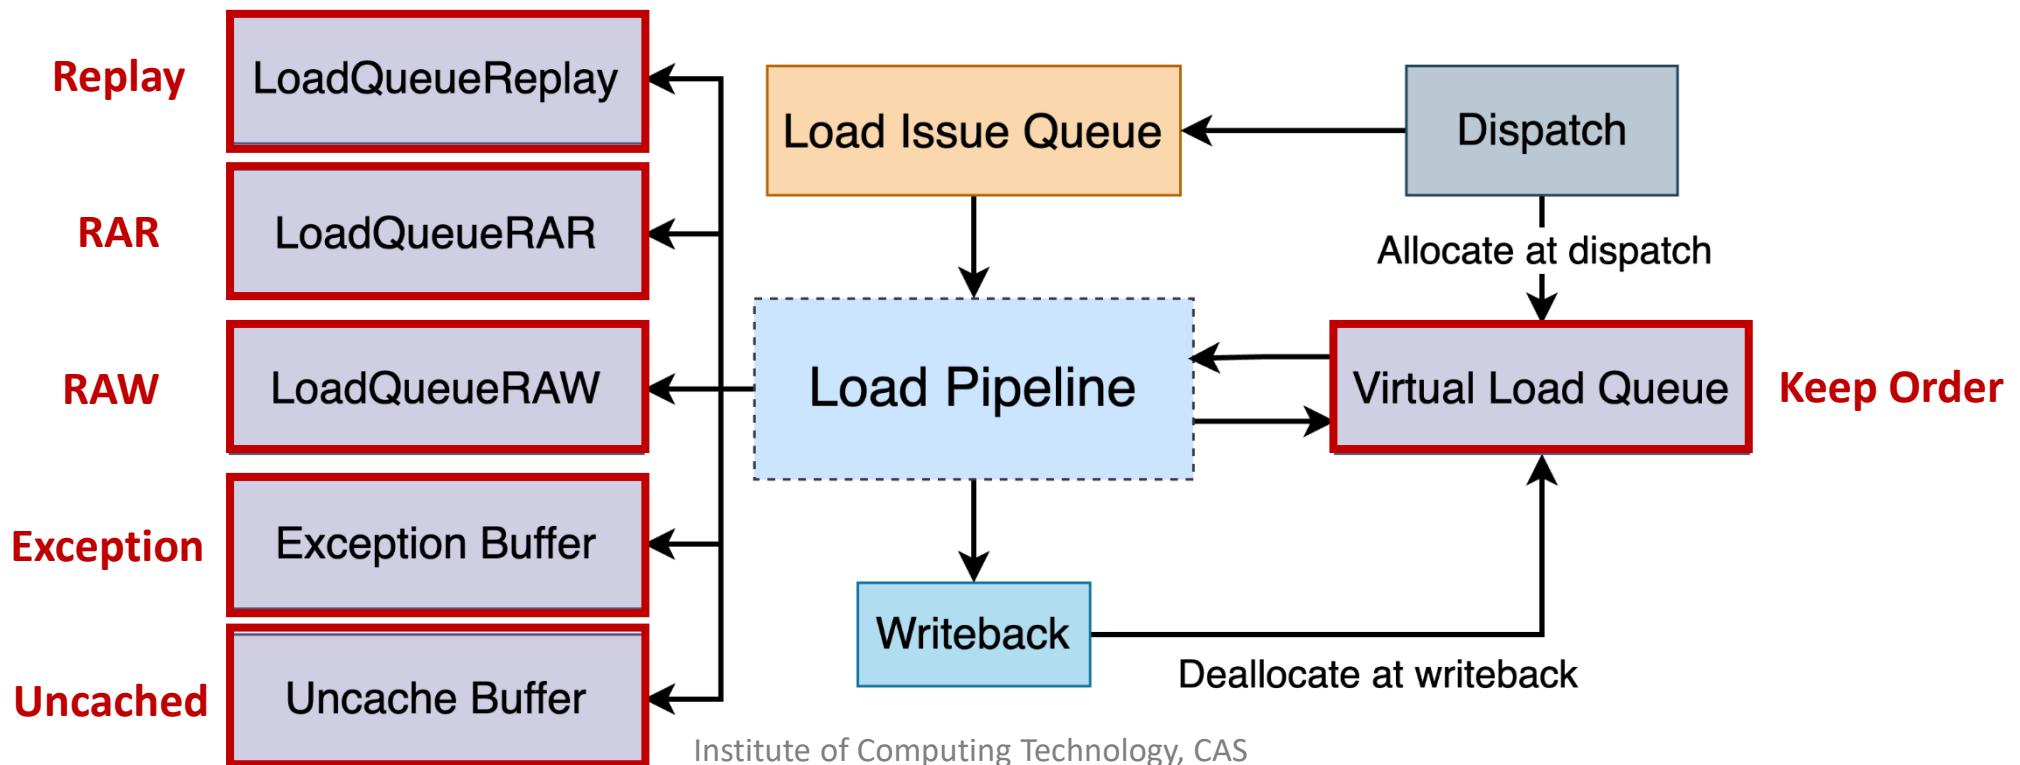
# ③ Load-Store Unit Upgrade





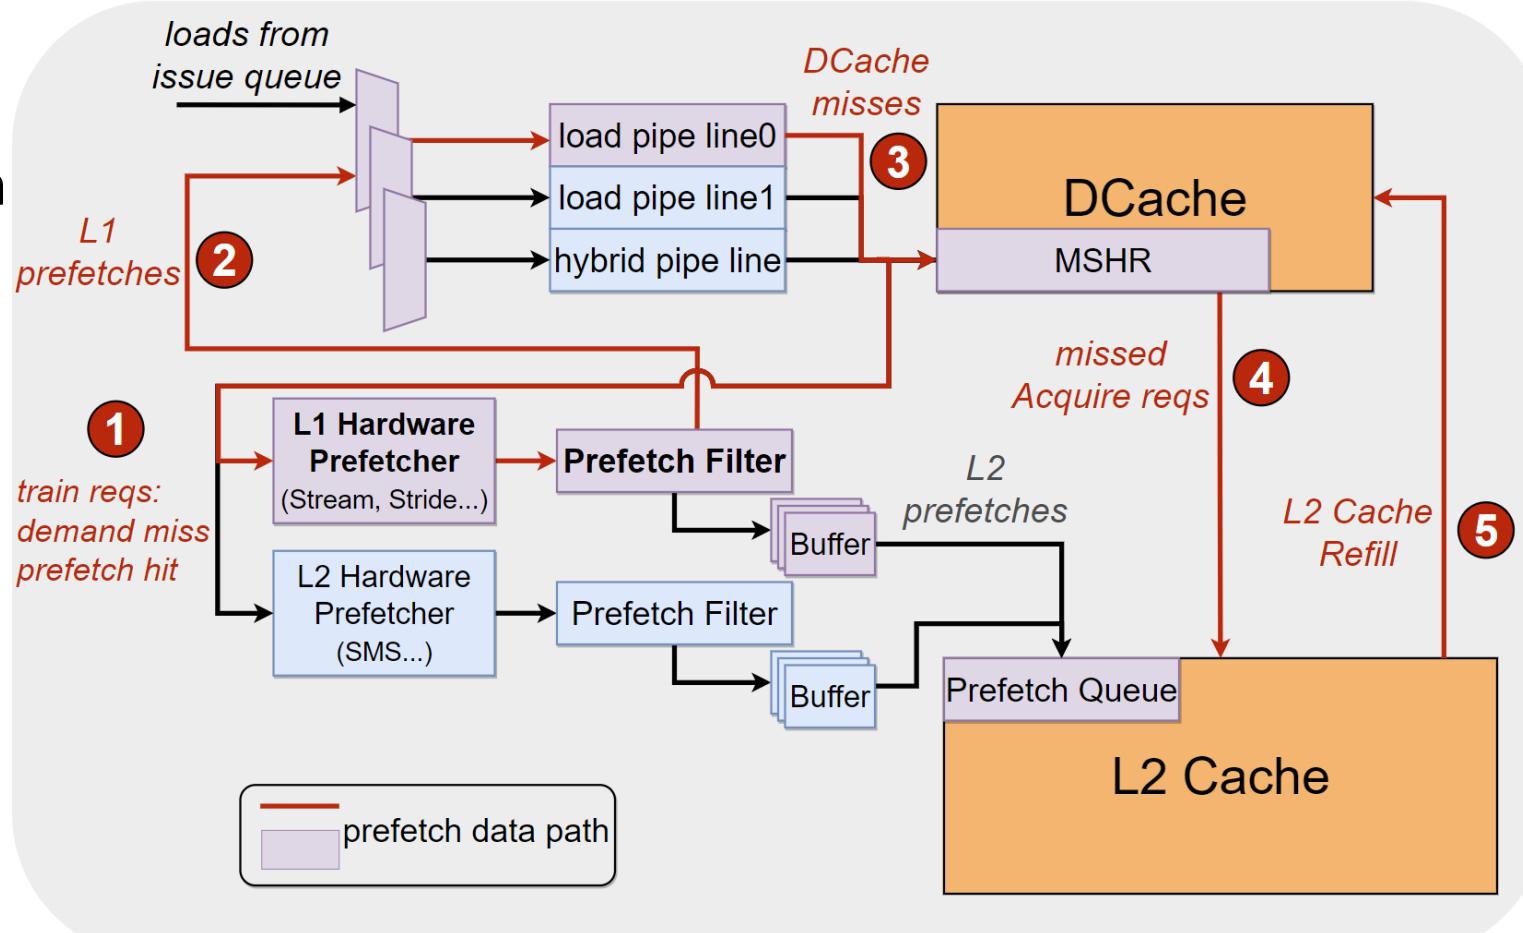
# Feature 1: Split Load Queue

- Split load queue **based on different responsibilities**
  - **Early commit:** load instruction can be retired sequentially from virtual load queue after write-back
  - Logic simplified, timing friendly and area acceptable



# Feature 2: L1D Prefetch

- L1D prefetch workflow
  - training, pattern detection
  - **Reuse load pipelines**
  - Fetch data from D\$ MSHR
- Prefetch algorithm
  - **Stream**<sup>[1]</sup>
    - $(x, x+1, x+2, x+3\dots)$
  - **Stride**<sup>[2]</sup>
    - $(x, x+n, x+2n, x+3n\dots)$

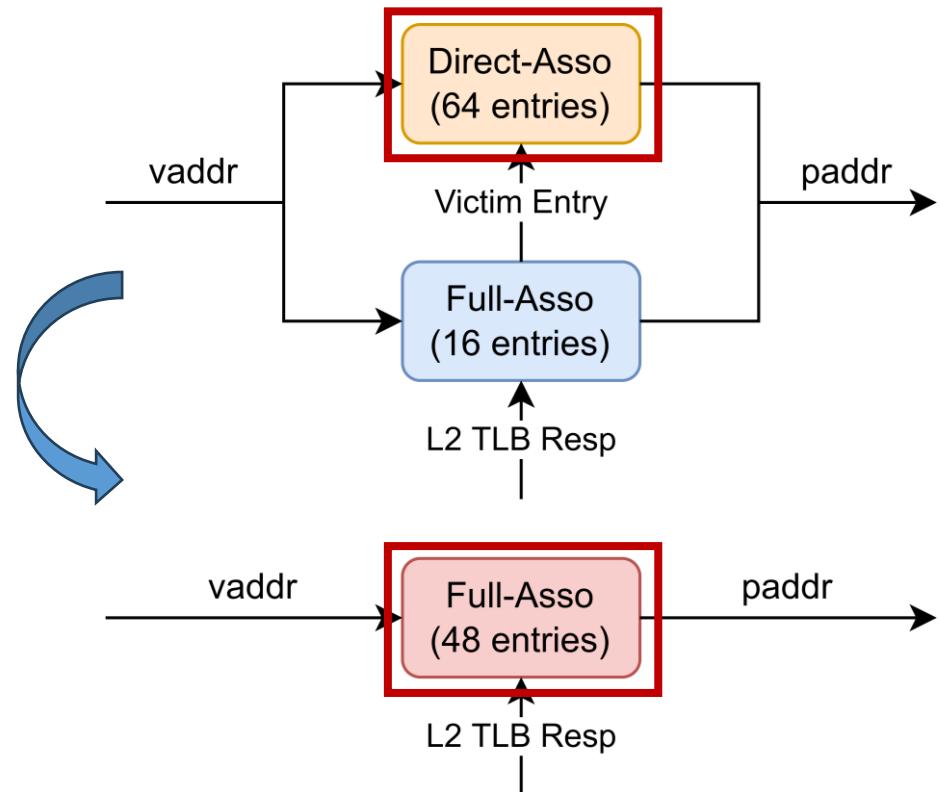


[1] S. Srinath, O. Mutlu, H. Kim and Y. N. Patt, "Feedback Directed Prefetching: Improving the Performance and Bandwidth-Efficiency of Hardware Prefetchers," 2007 IEEE 13th International Symposium on High Performance Computer Architecture, Scottsdale, AZ, USA, 2007

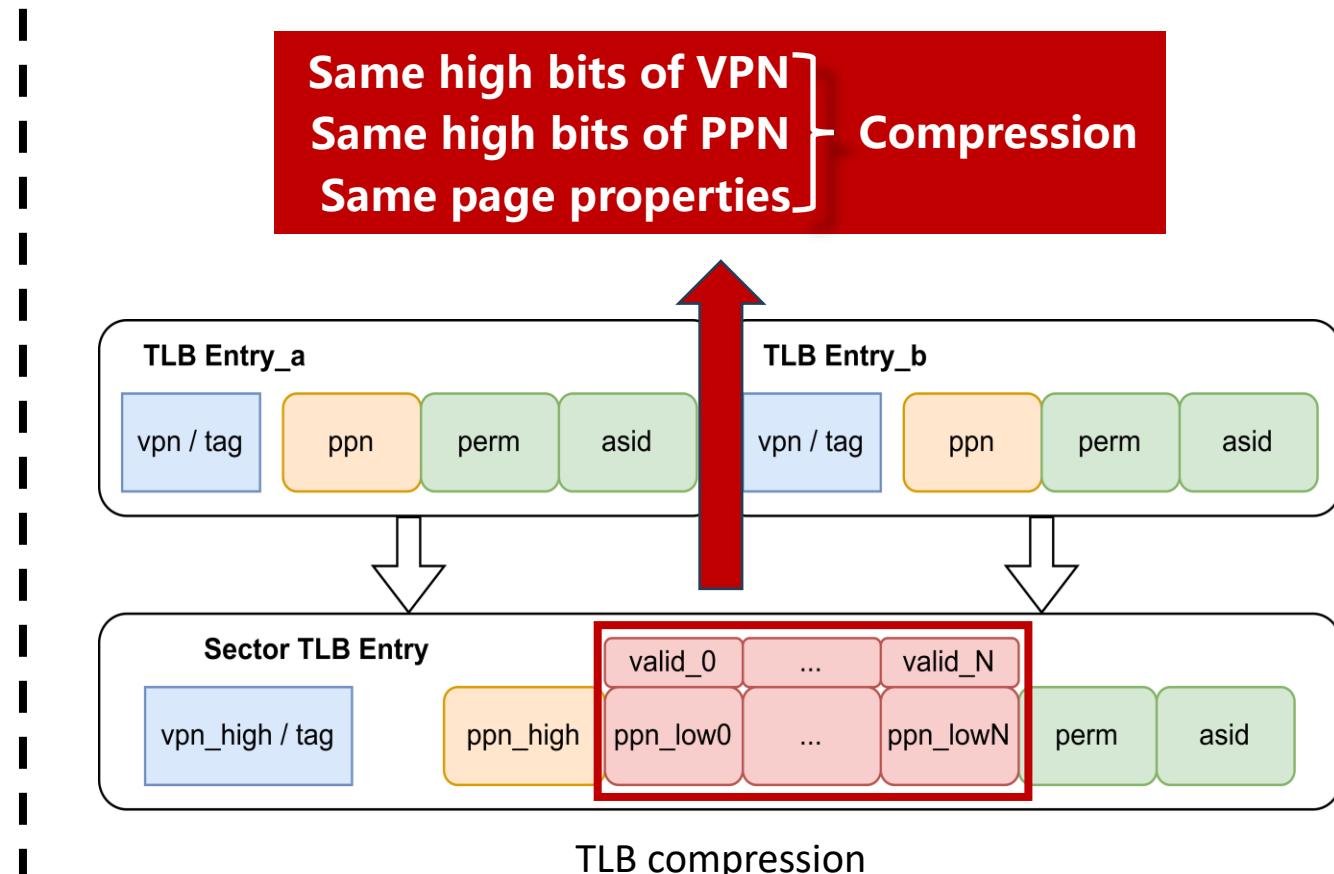
[2] J. -L. Baer and T. -F. Chen, "An effective on-chip preloading scheme to reduce data access penalty," Supercomputing '91:Proceedings of the 1991 ACM/IEEE Conference on Supercomputing, Albuquerque, NM, USA, 1991

# Feature 3: Powerful MMU

- 16 fully associative entry + 64 direct mapping entry → **48 fully associative entry**
- Support TLB compression, Merges contiguous page table entries



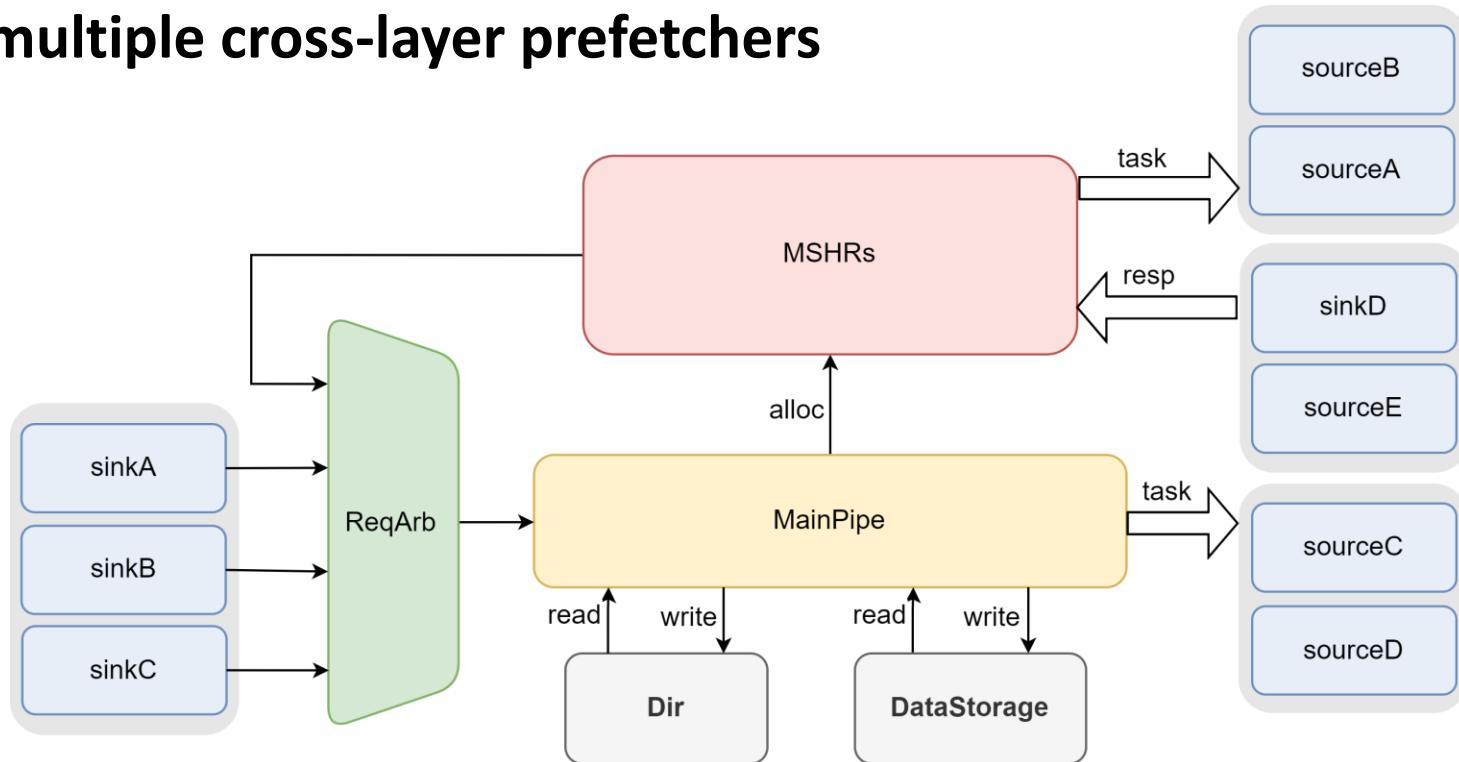
DTLB organization in Nanhu and Kunminghu





## ④ L2 Cache Upgrade

- Multiple concurrent pipelines → Non-blocking main pipeline
- L1-L2 Refill wake-up collaboratively
- Eliminate transactions serialization in the same set & Request merge
- Aggressive multiple cross-layer prefetchers



# Feature 1: L1-L2 Collaborative optimization

- **Observation:**

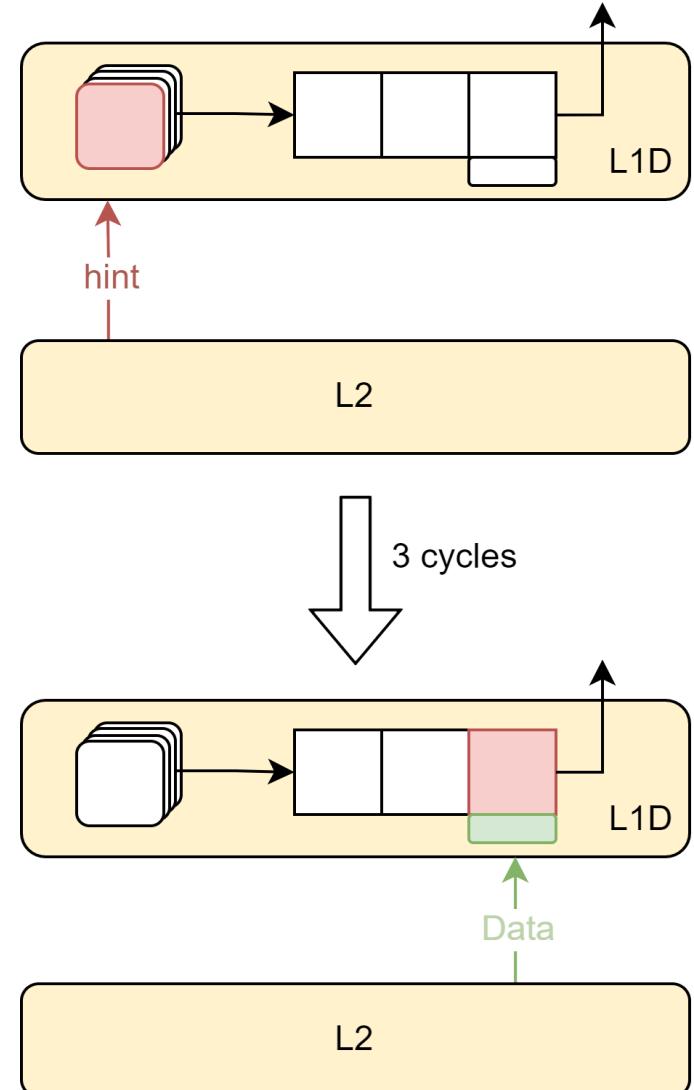
- There's chance to hint LSU to replay before refill
- Interval between hint and refill signal is 3 cycles

- **New Design:**

- Let L2 give the Hint signal in advance
- **Speedup the wake and replay of load replay queue**

- **Implementation:**

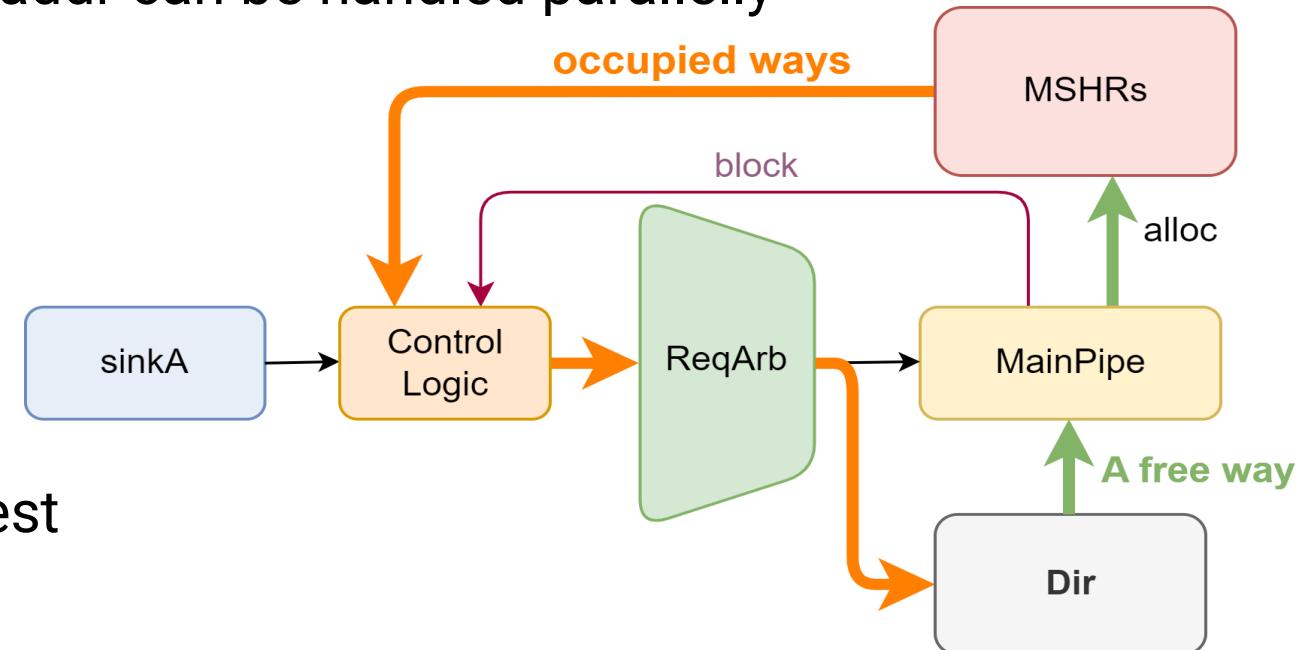
- Set up monitor to track info from main pipeline
- Send hint signal to L1 3 cycles before refill
- Need to calculate the exact Hint timing considering possible pipeline delays





# Feature 2: Eliminate txns serialization in the same set

- **Background:** Txns to the same set were handled serially, which reduces parallelism
- **Difficulty:** Txns to the same set can affect each other due to replacement
- **New Design**
  - Txns to the same set but different addr can be handled parallelly
  - Block txns to same addr only
- **Implementation**
  - Collect the ways in active MSHRs
  - Assign a free way to the new request



# Feature 3: Evict on Refill

- Observation:**

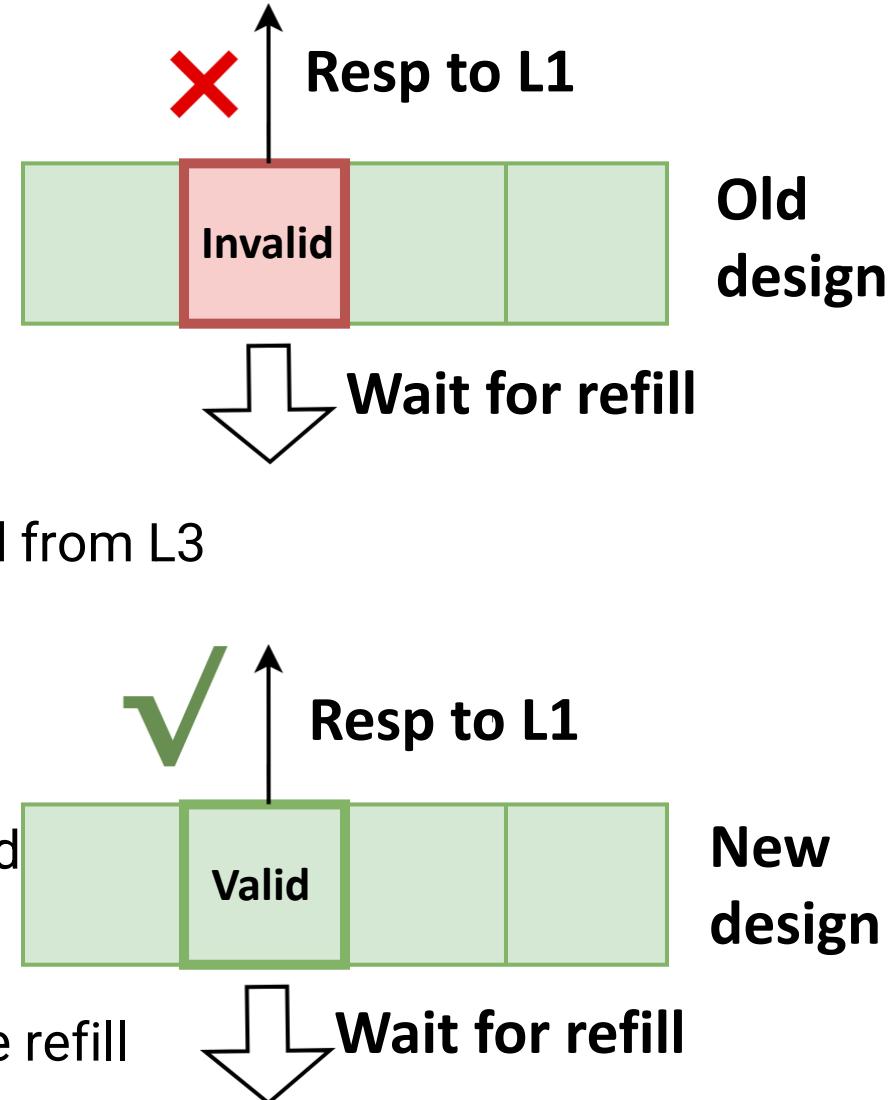
- On L1 acquire, L2 assigns one way by replacement policy
- But this way is occupied until refill, unable to serve L1

- Improvement:**

- New design makes the way be assigned until data is refilled from L3
- So, it can still serve L1 when waiting for refill

- Implementation:**

- Transfer it to L3 on acquire miss, make replacer untouched
- Wait for L3 to refill
- Read directory and assign one way when MSHR handles the refill



# Thanks!