

XiangShan: An Open-Source High-Performance RISC-V Processor and Infrastructure for Architecture Research

The XiangShan Team

Institute of Computing Technology (ICT)
Chinese Academy of Sciences (CAS)

HPCA'24@Edinburgh, Scotland

March 2, 2024



Tutorial@HPCA'24 Schedule

Time (AM)	Topic
8:30 - 9:00	Introduction of the XiangShan Project
9:05 - 10:00	<i>Microarchitecture Design and Implementation</i>
10:05 - 10:50	Hands-on Development
	Coffee Break
11:20 - 12:00	Hands-on Development & Discussions (Cont.)



XiangShan: Open-Source High-Performance Processor

- **1st generation: YQH**
 - 2020/6: first commit of design RTL
 - 2021/7: **28nm tape-out, 1.3GHz**
 - Performance: **SPEC CPU2006 7.01@1GHz, DDR4-1600**
- **2nd generation: NH**
 - 2021/5: starting design exploration and RTL design
 - 2023/4: GDSII delivery
 - **14-nm tape-out, estimated SPEC CPU2006 20@2GHz**
- **3rd generation: KMH**
 - ISA feature: **Vector (V) extension, Hypervisor (H) extension**
 - Comprehensive performance improvement
- Open-sourced at <https://github.com/OpenXiangShan/XiangShan>



Fragrant Hills in Beijing

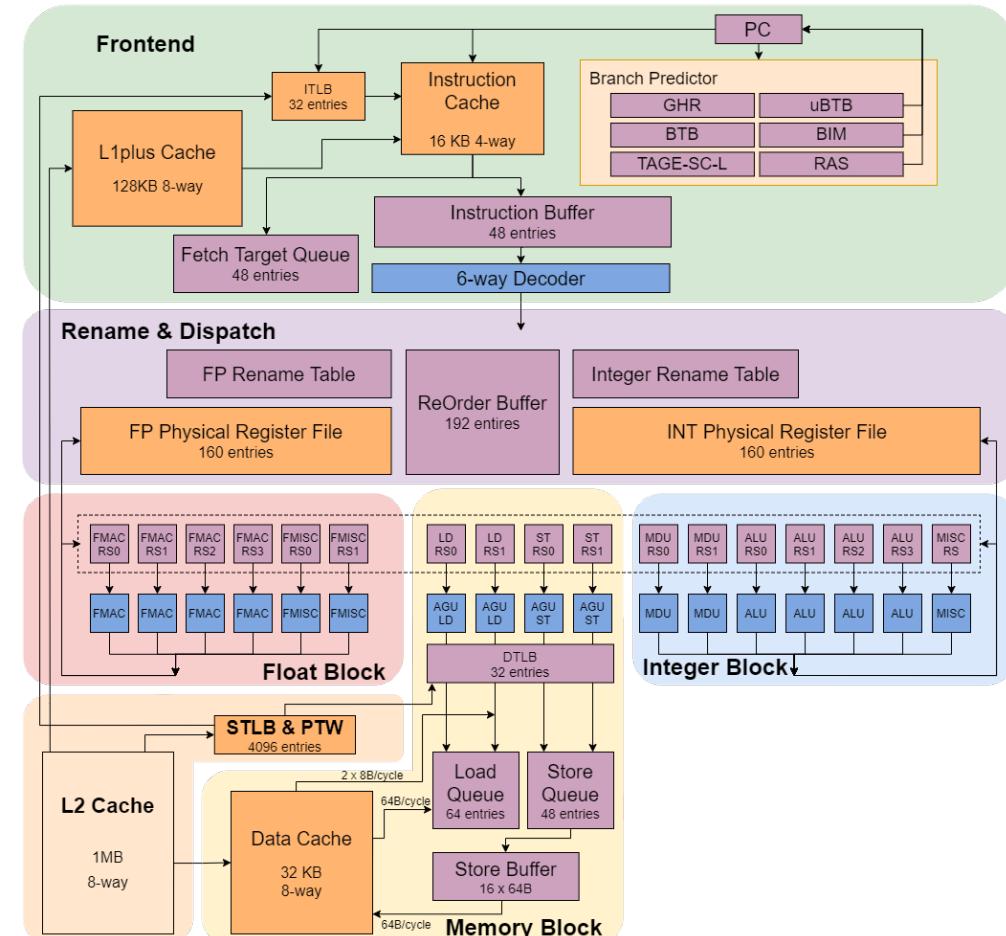
The screenshot shows the GitHub repository page for XiangShan. It includes the repository name, a list of branches (master, 205), tags (4), and recent commits. The commits are listed with their authors, messages, and timestamps. The repository description is "Open-source high-performance RISC-V processor" and it includes links to chisel3, risc-v, and microarchitecture. The sidebar shows activity, custom properties, and statistics like 4.2K stars, 86 watching, and 563 forks.

>4.2K stars, >500 forks on GitHub



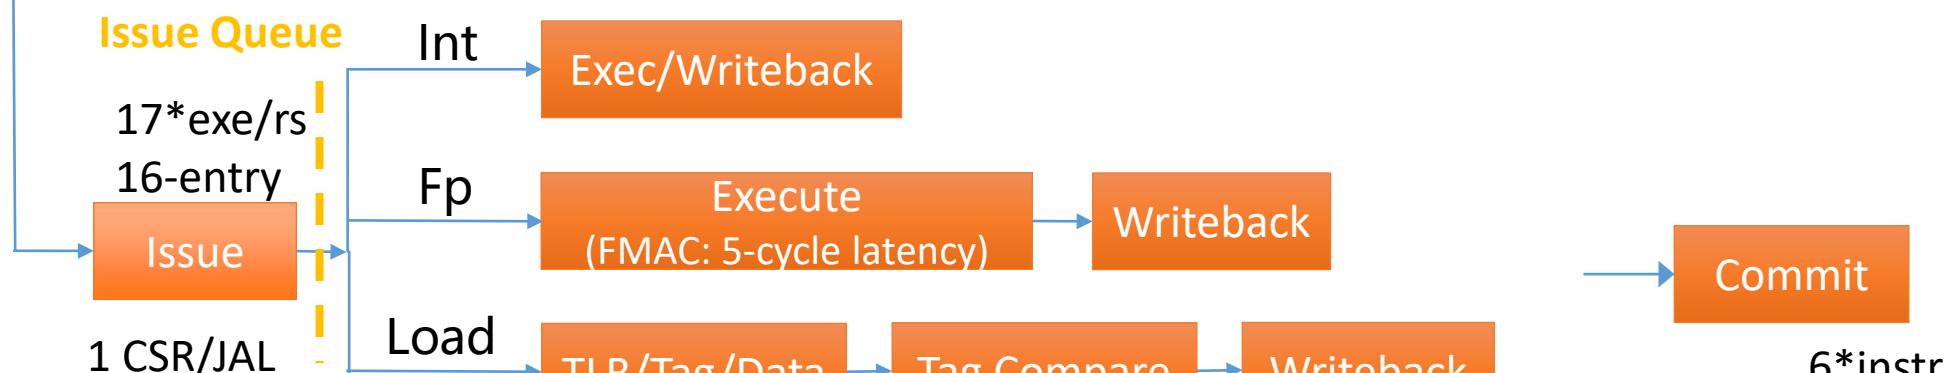
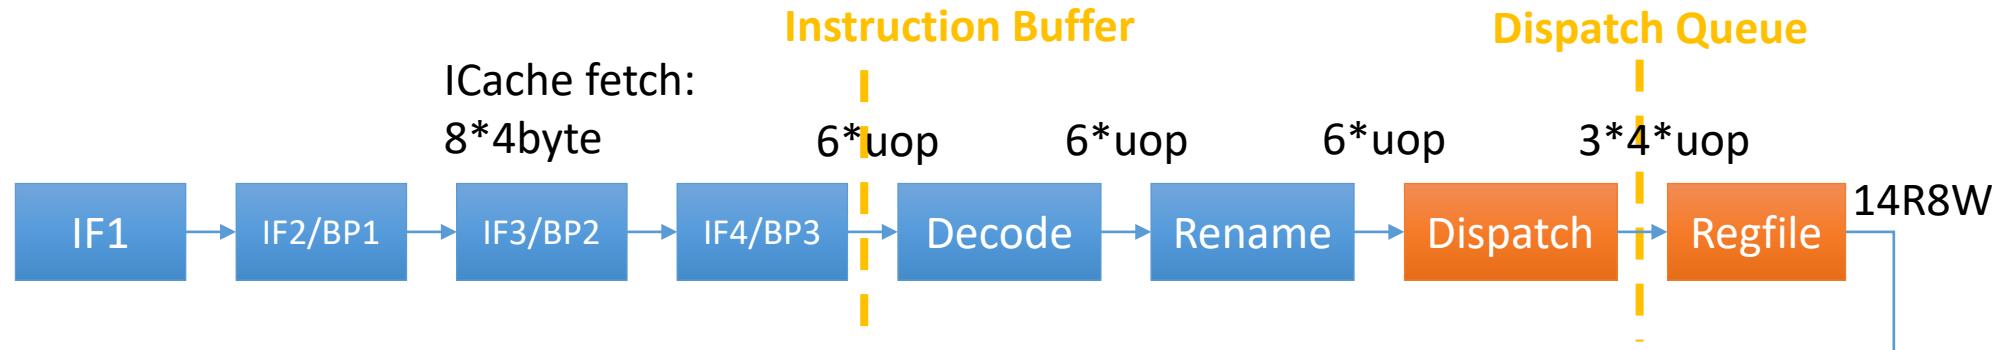
XiangShan Microarchitecture (V1-YQH)

- **11**-stage, **6**-wide decode/ rename
- **TAGE-SC-L** branch prediction
- **160** Int PRF + **160** FP PRF
- **192**-entry ROB, **64**-entry LQ, **48**-entry SQ
- **16**-entry RS for each FU
- **16KB** L1 Cache, **128KB** L1plus Cache for instruction
- **32KB** L1 Data Cache
- **32**-entry ITLB/DTLB, **4K**-entry STLB
- **1MB** inclusive L2 Cache





XiangShan Microarchitecture (V1-YQH)

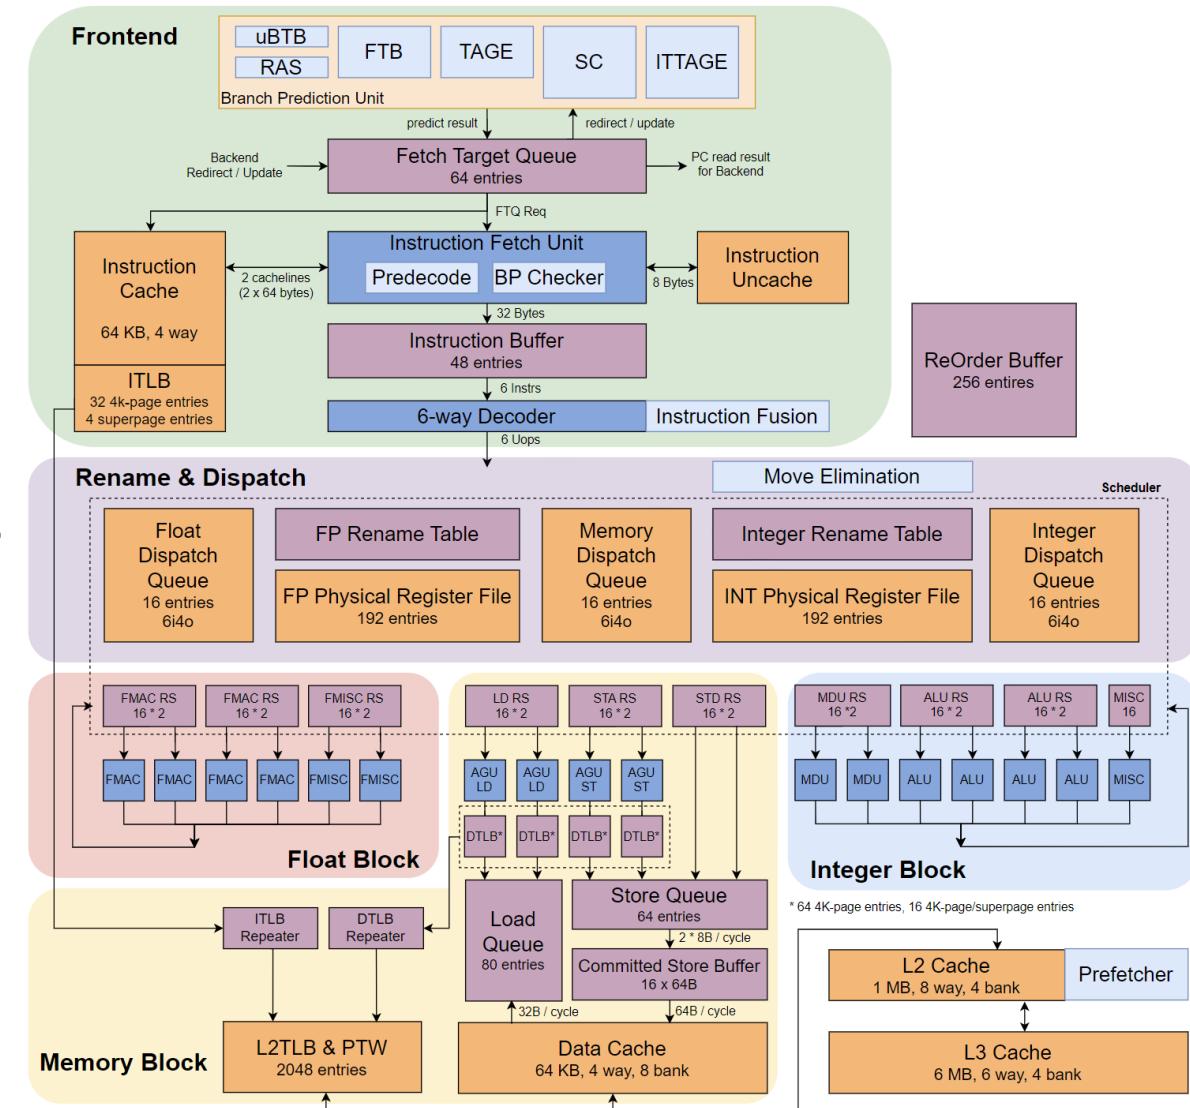


	In-order
	Out-of-order
	Queues



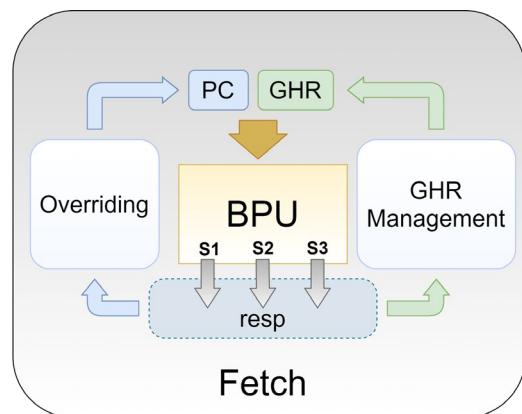
XiangShan Microarchitecture (V2-NH)

- **192** Int PRF + **192** FP PRF
- **256**-entry ROB, **80**-entry LQ, **64**-entry SQ
- **32**-entry RS for two FUs
- **64KB** L1 Instr. Cache, **64KB** L1 Data Cache
- **80**-entry DTLB, **36**-entry ITLB, **2K**-entry STLB
- **1MB** non-inclusive L2 Cache
- **6MB** non-inclusive L3 Cache (LLC)

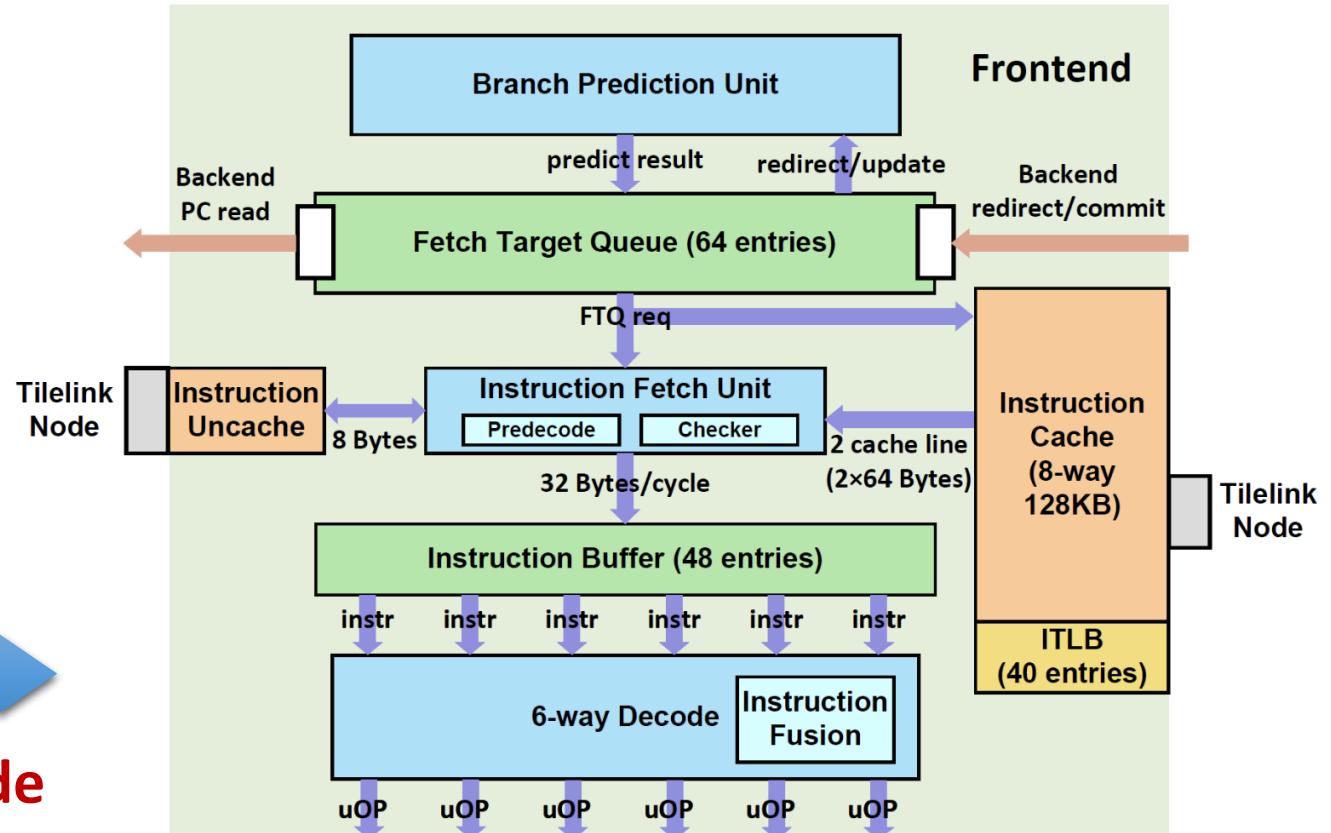


Frontend

- Decoupled frontend
 - **Producer**: branch predictor unit
 - **Consumer**: instruction fetch unit
- Merit
 - Fewer fetch bubbles
 - Fetch directed instr. prefetching



Upgrade



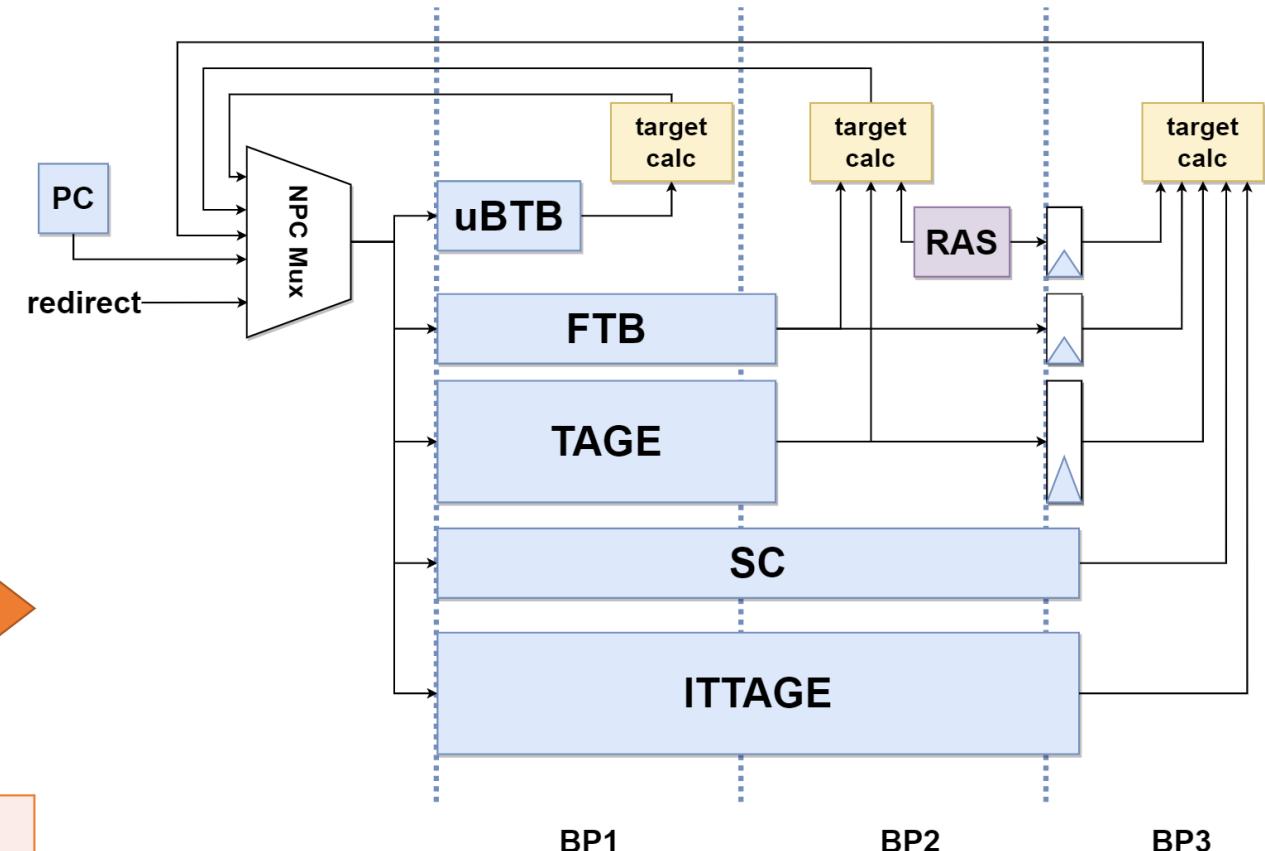
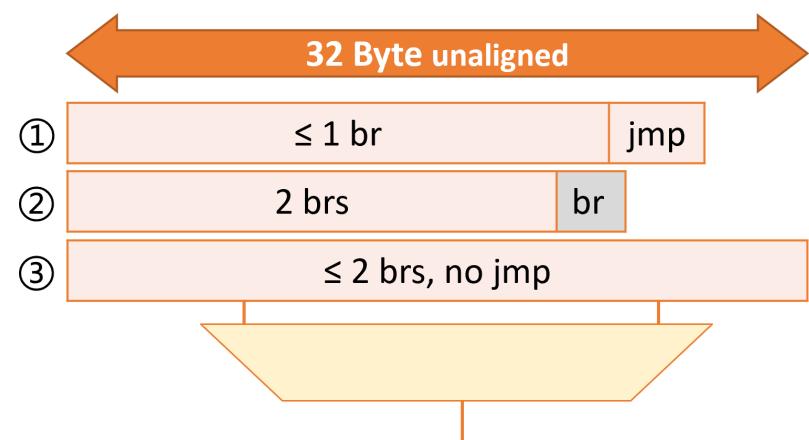
Yanqihu

Nanhu

Branch Predictor

Three-stage overriding prediction

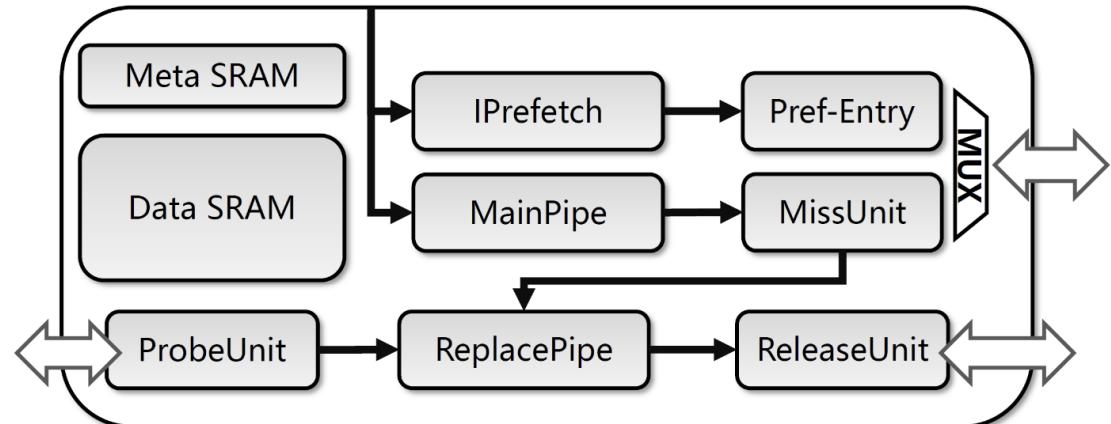
- Stage 1: uBTB
- Stage 2: FTB + TAGE
- Stage 3: SC + ITTAGE + RAS
- FTB: Fetch Target Buffer



Instruction Cache

- **Main feature**

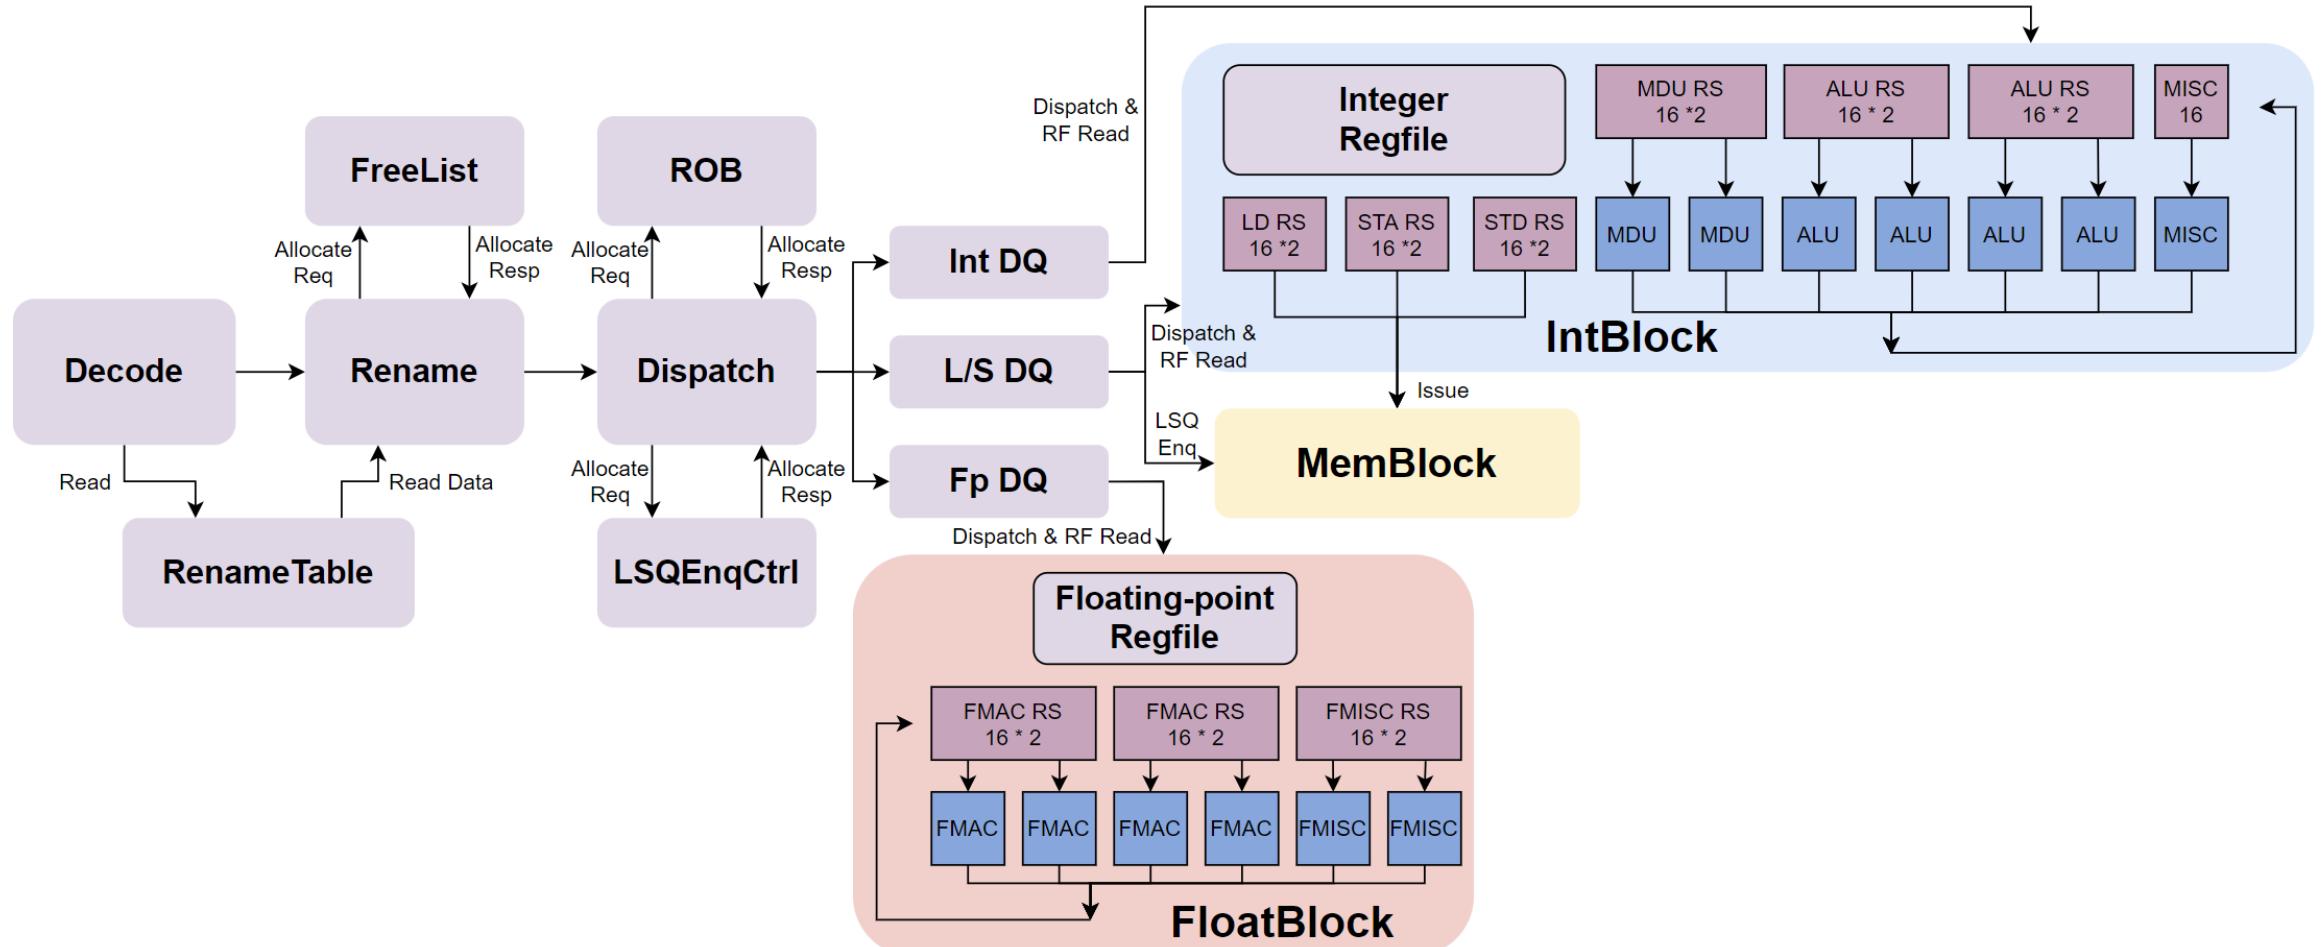
- 64KB 4-w VIPT blocking cache
- Even/odd bank interleaving read
- **Support TileLink protocol**
- PLRU replacement



	Yanqihu	Nanhu	
L1 ICache	16KB, 4-w	64KB, 4-w	4×
L1plus Cache	128KB, 8-w	-	Reduce
Read bandwidth	64 B	128 B	2×
TileLink Support	No	Yes	New
Instruction prefetch	Stream	L2 FDIP	New



Backend



Instruction Fusion

- Fusion of adjacent UOPs
 - Improve backend efficiency
 - Reduce data bypass delay
- Fusion Types
 1. Reuse of the original calculation
 2. Add new type of calculation
- **Reduce dynamic instr. greatly**

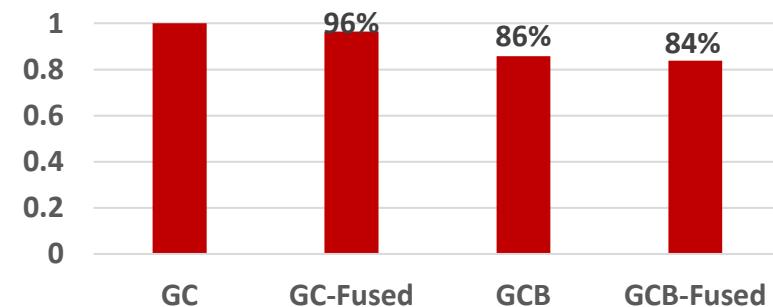
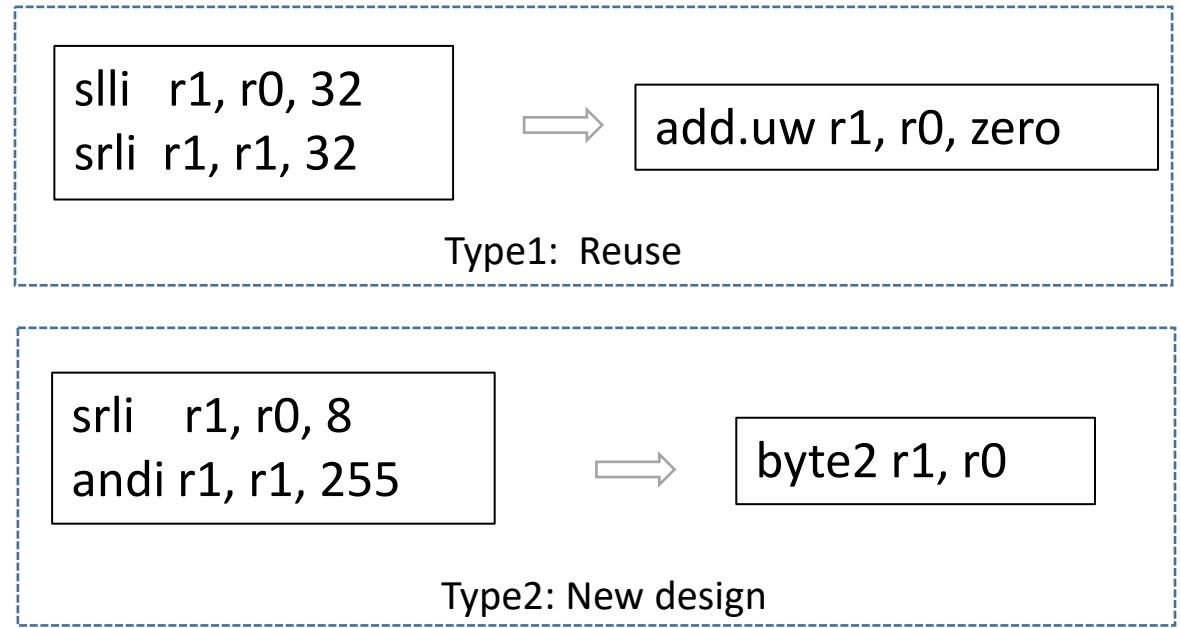
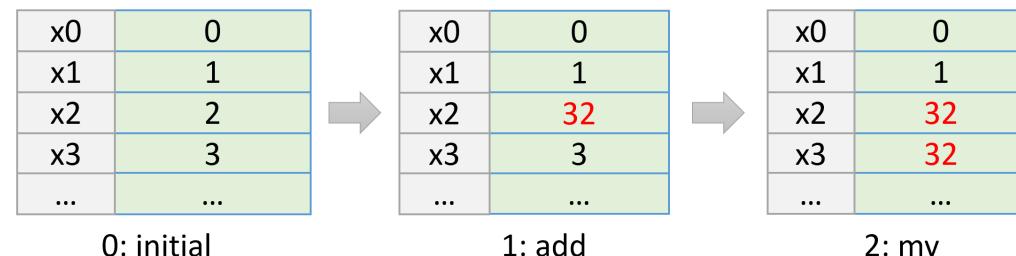
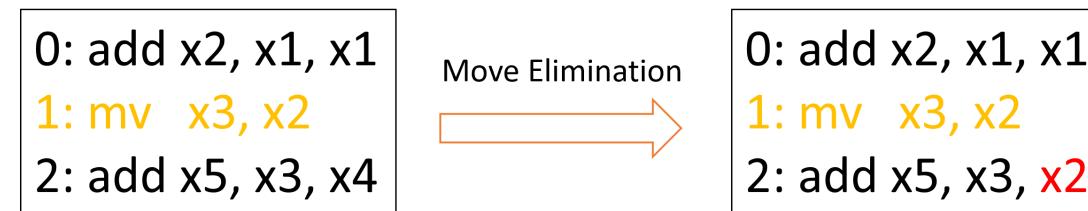


Figure. Normalized Dynamic Instruction Count for CoreMark (-O2)



Rename & Move Elimination

- RenameMap: the mapping between **arch registers** and **physical registers**
- Move Elimination
 - No need to allocate new physical register by just changing the RenameMap
 - Mark as **Complete** when Move instruction enters ROB
 - Use **Reference-Counter** to record mapping counts of physical registers



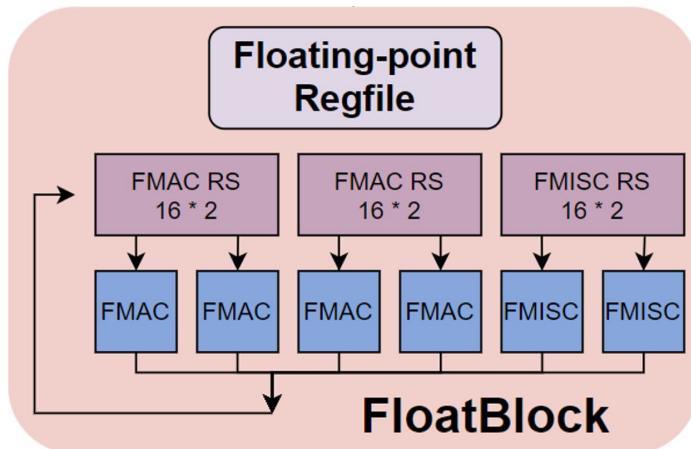
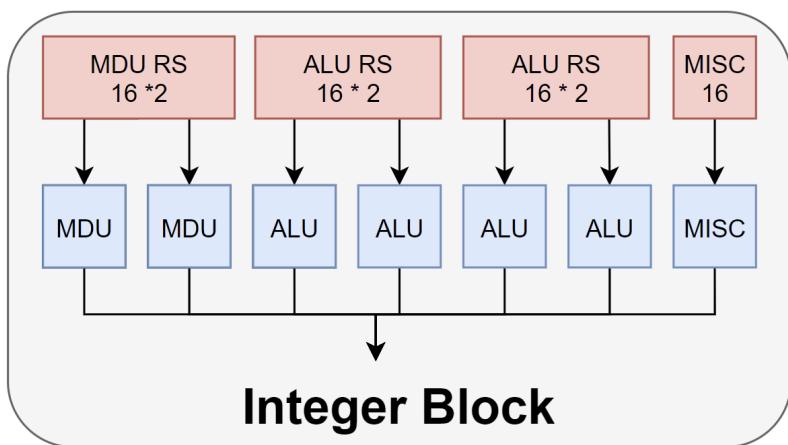


Reservation Station

- Main Changes:
 - 2i2o & 1i1o
 - Age Matrix based selection

0	true			
1	true	true		
2	true	false	true	
3	true	false	false	true
	0	1	2	3

Age Matrix





Floating Point Units (FPUs)

- Industry competitive Floating Point Unit, IEEE 754 compatible

Operations	Latency
FADD	3
FMUL	3
FMA	5
FDIV	≤ 11 (float) ≤ 18 (double)
FSQRT	≤ 17 (float) ≤ 31 (double)
FCVT, FCMP	3

The screenshot shows the GitHub repository page for 'OpenXiangShan/fudian'. The repository is public, has 8 watchers, 12 forks, and 31 stars. It contains 15 branches and 0 tags. The main branch has 33 commits. The commits are listed as follows:

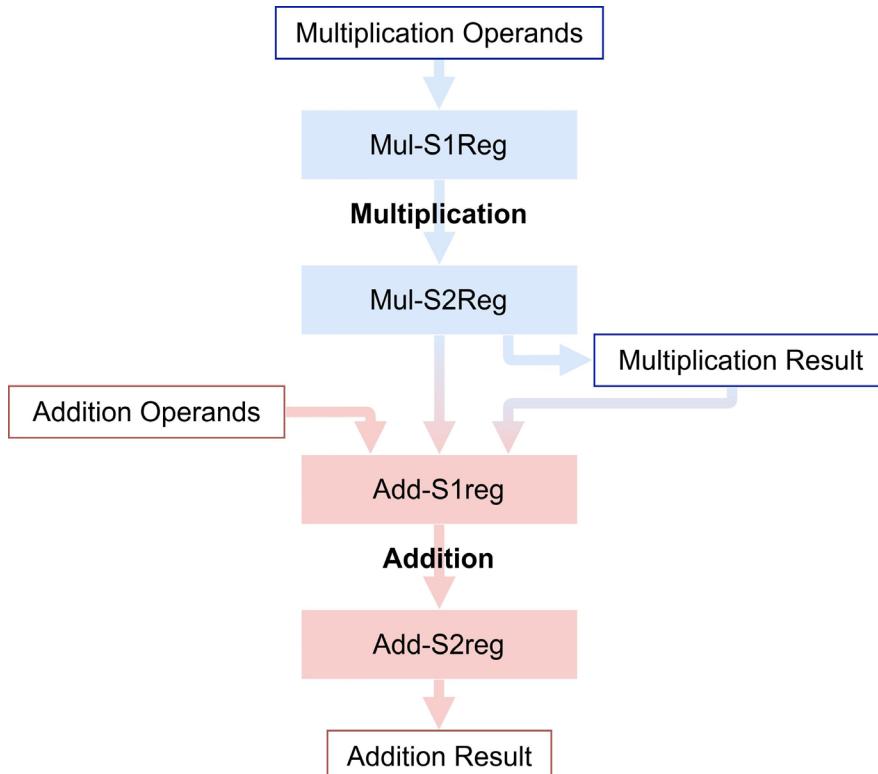
- notlqr Add instructions to run unit tests (1bad625 on Aug 24, 2022)
- berkeley-softfloat-3 ... Initial commit (2 years ago)
- berkeley-testfloat-3 ... Initial commit (2 years ago)
- src FADD: avoid use decoder (2 years ago)
- .gitignore Initial commit (2 years ago)
- .gitmodules Initial commit (2 years ago)
- .mill-version Initial commit (2 years ago)
- .scalafmt.conf Initial commit (2 years ago)

The repository description states: "Open source high performance IEEE-754 floating unit".

<https://github.com/OpenXiangShan/fudian>

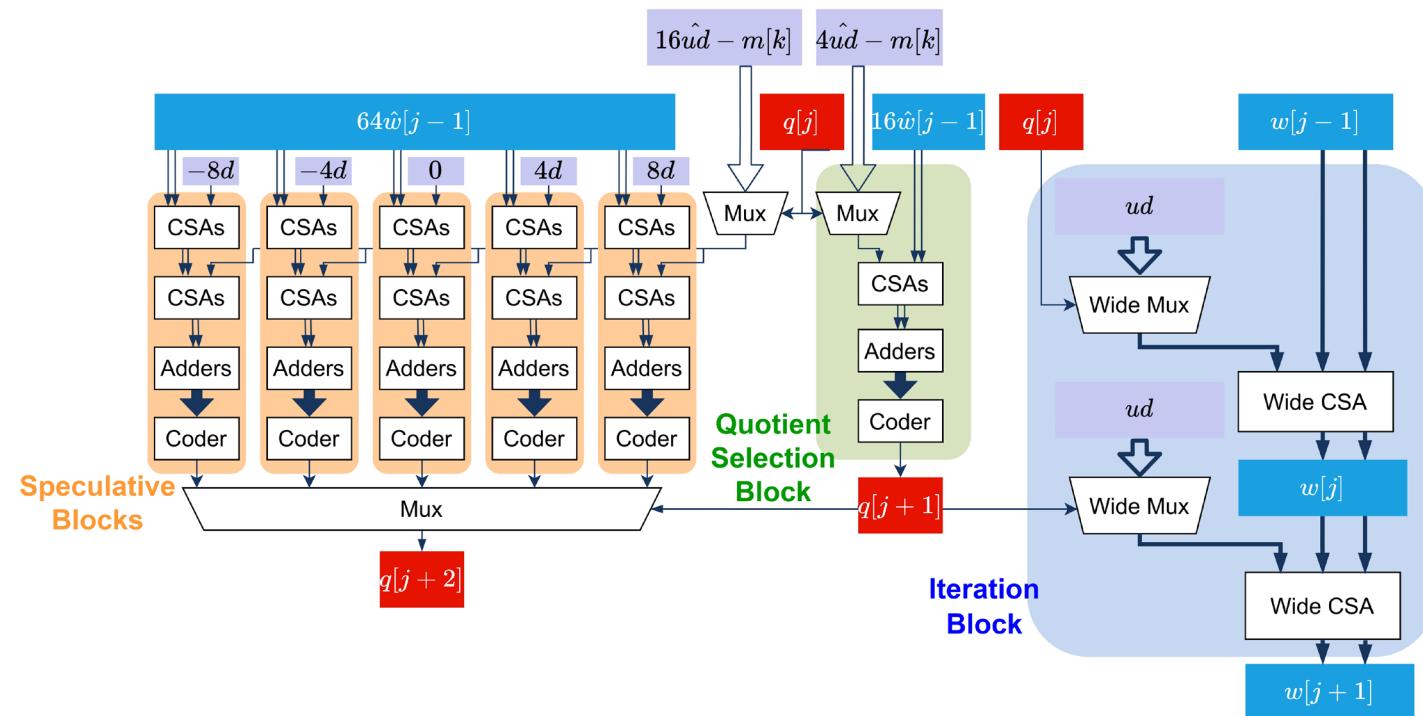
- Cascade FMA

- Reduce FADD delay 5->3



- SRT16 Division units

- $\frac{1}{2}$ delay compared with YQH (SRT4)





B/K Extension Implemented

- B: Bit-Manipulation extension
- K: Cryptographic extension

1. Extensions
1.1. Zba extension
1.2. Zbb: Basic bit-manipulation
1.2.1. Logical with negate
1.2.2. Count leading/trailing zero bits
1.2.3. Count population
1.2.4. Integer minimum/maximum
1.2.5. Sign- and zero-extension
1.2.6. Bitwise rotation
1.2.7. OR Combine
1.2.8. Byte-reverse
1.3. Zbc: Carry-less multiplication
1.4. Zbs: Single-bit instructions

B Extension

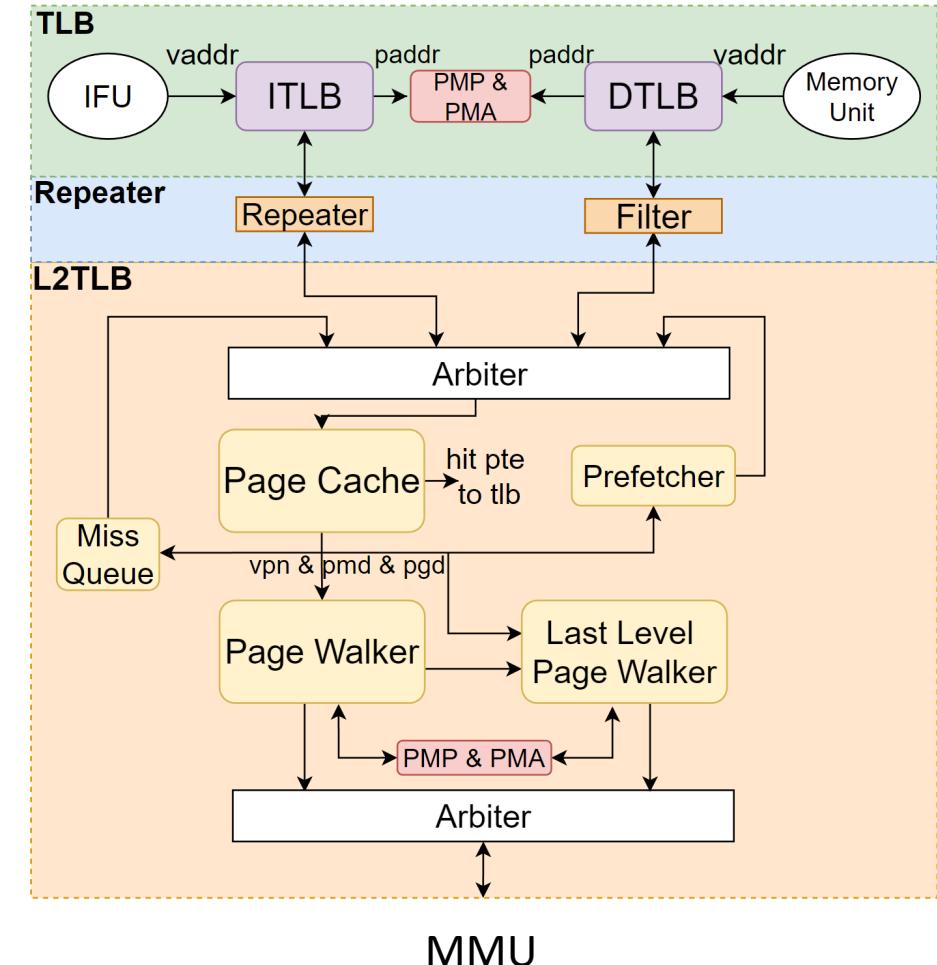
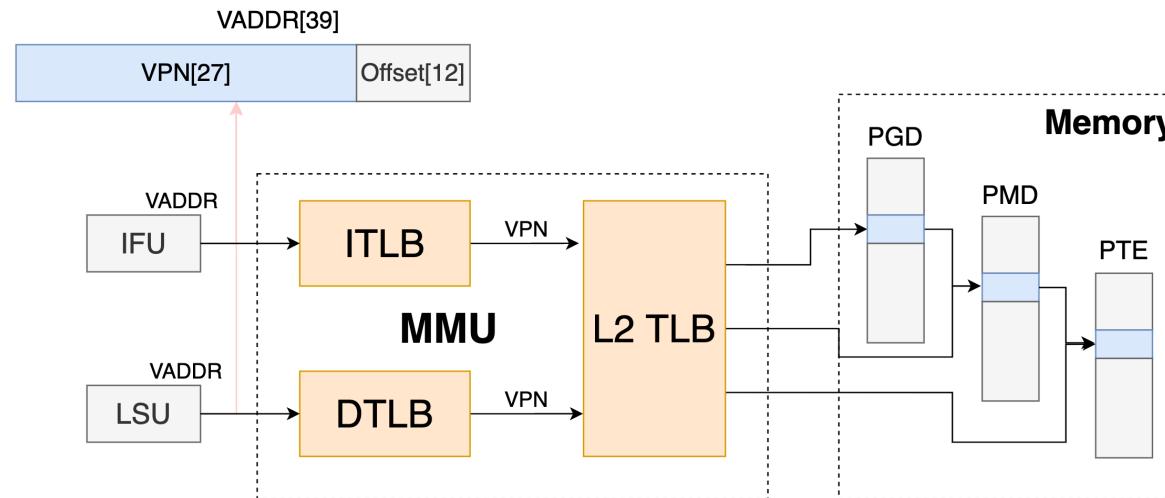
2. Extensions Overview
2.1. Zbkb - Bitmanip instructions for Cryptography
2.2. Zbkc - Carry-less multiply instructions
2.3. Zbkx - Crossbar permutation instructions
2.4. Zknd - NIST Suite: AES Decryption
2.5. Zkne - NIST Suite: AES Encryption
2.6. Zknh - NIST Suite: Hash Function Instructions
2.7. Zksed - ShangMi Suite: SM4 Block Cipher Instructions
2.8. Zksh - ShangMi Suite: SM3 Hash Function Instructions
2.9. Zkr - Entropy Source Extension
2.10. Zkn - NIST Algorithm Suite
2.11. Zks - ShangMi Algorithm Suite
2.12. Zk - Standard scalar cryptography extension
2.13. Zkt - Data Independent Execution Latency

K Extension

Memory Management Unit

- Improvement:

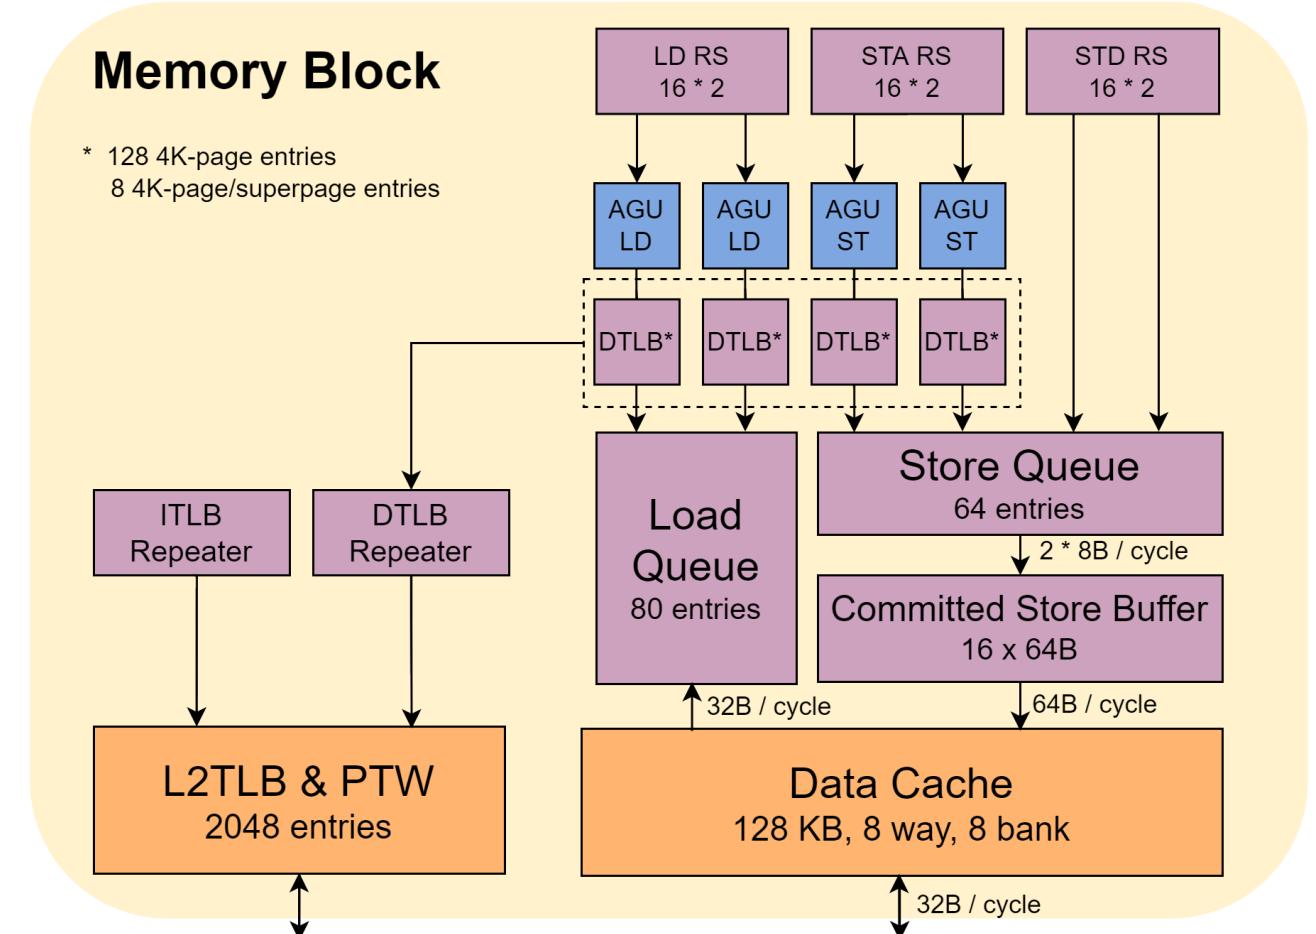
- Larger Data TLB size
- Next-Line Prefetch for STLB
- Parallel Page Table Walks
- PMP/PMA support



Memory Block

- Datasheets

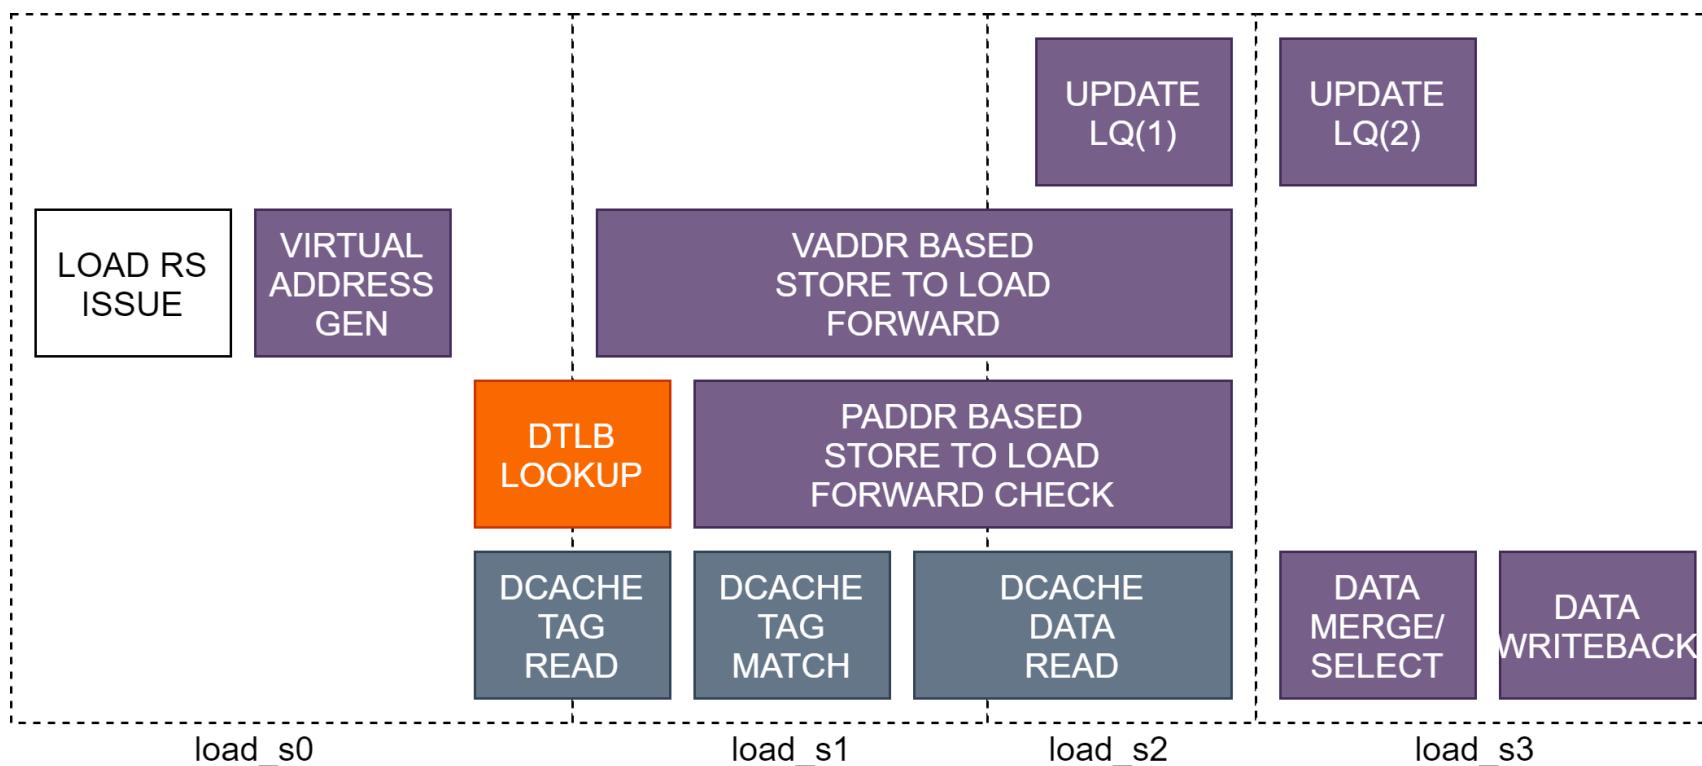
Load pipeline	2
L1 LD hit bandwidth	2x8B/cycle
L1 LD hit latency	4
Store address pipeline	2
Store data pipeline	2
Store data bandwidth	2x8B/cycle
Load Queue Entry	80
Store Queue Entry	64
Merged Store Buffer	16 cachelines





Load Pipeline

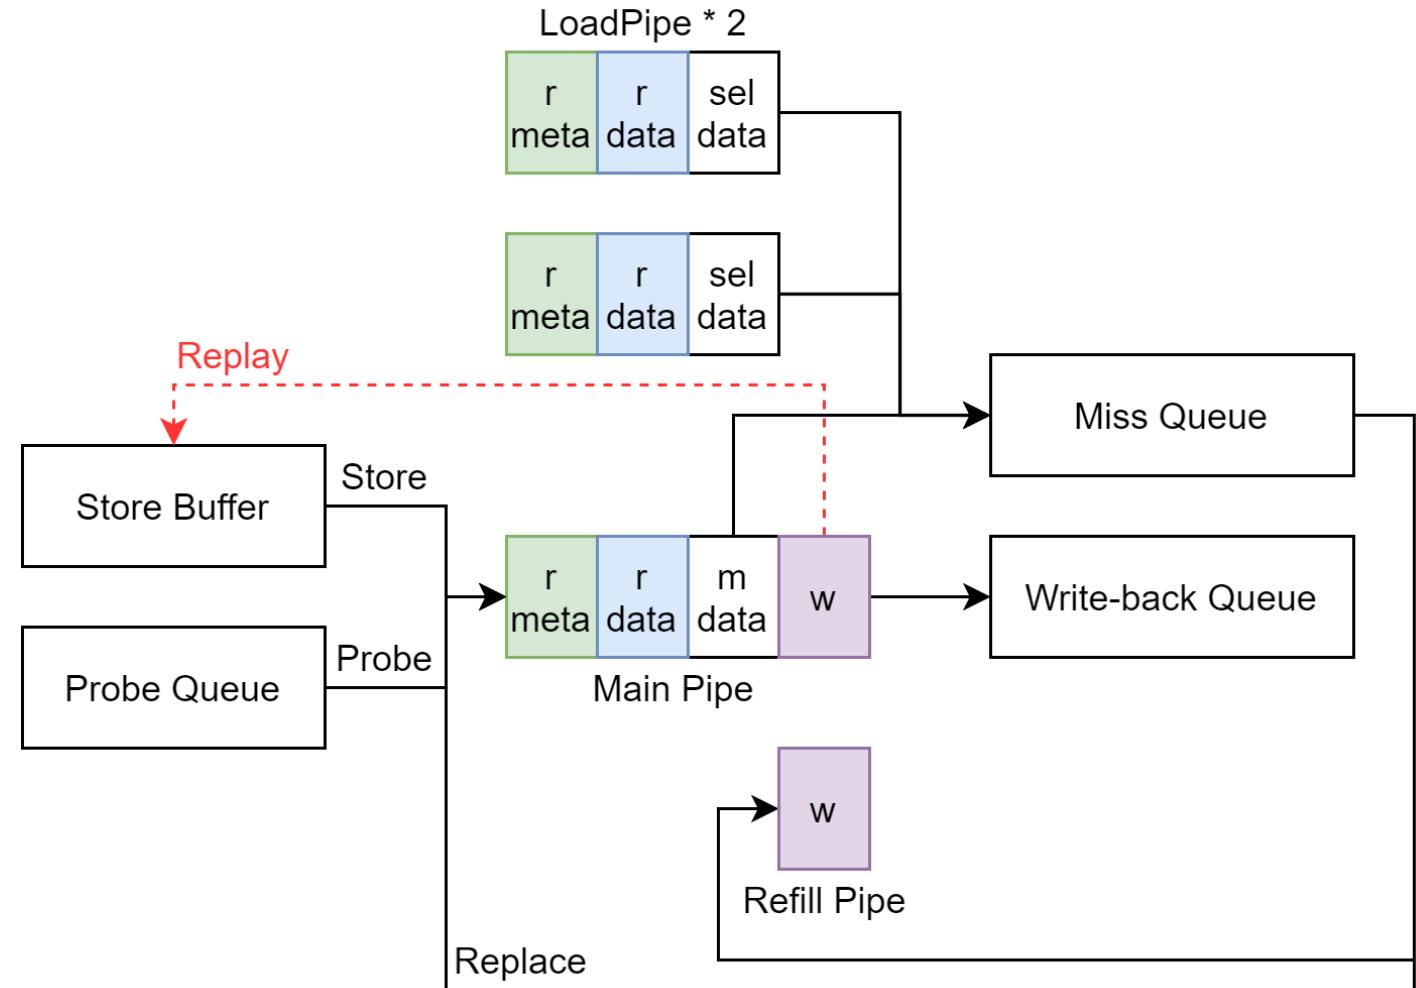
- Four stage
- Optimized for load-to-load & load-store bypass



L1 DCache

- Coupled with load pipeline
- TileLink protocol

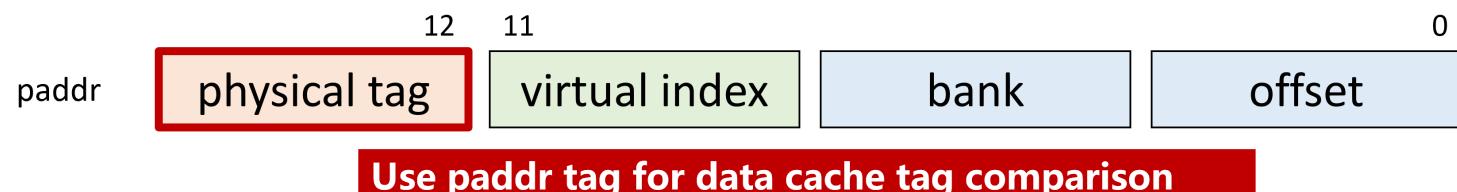
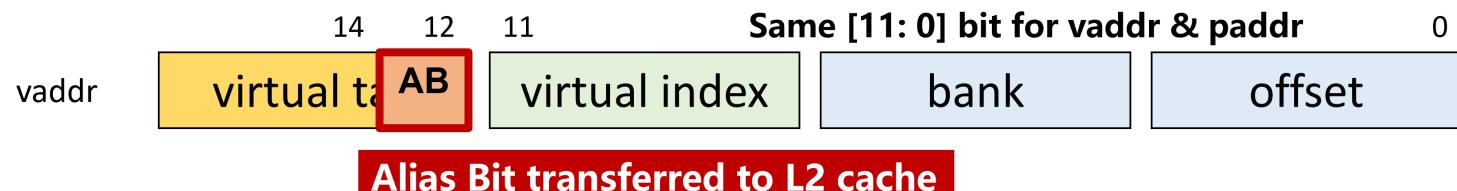
L1 Data cache size	64KB-8w
L1/L2 Bus Width	256bit
Store Buffer	16
Probe Queue	8
Miss Queue	16
Write-back Queue	18
Software Prefetch	Support
ECC	Support



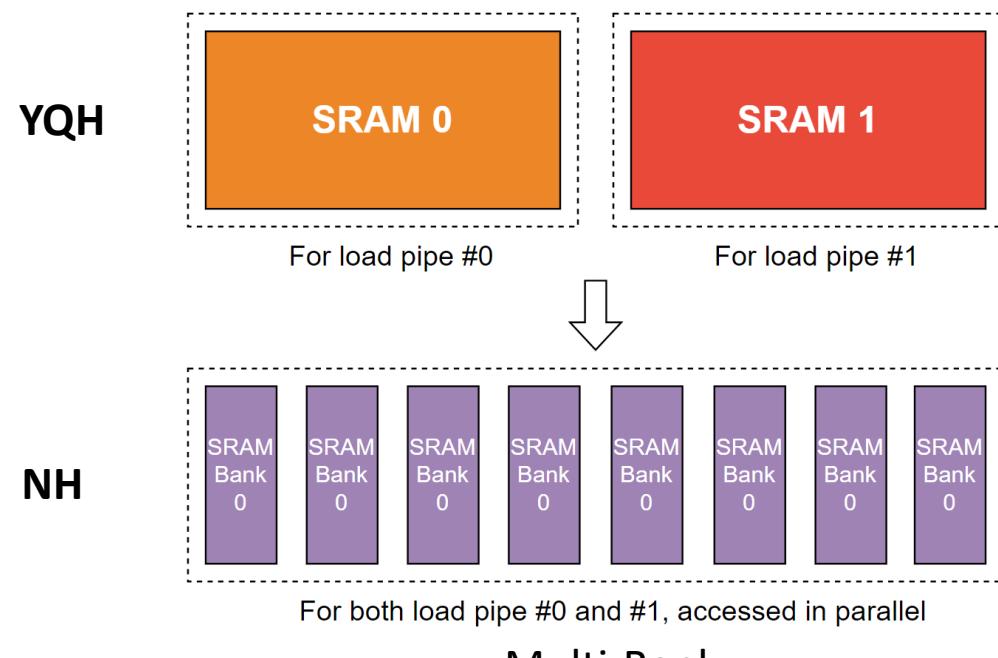
L1 DCache

- Noteworthy features

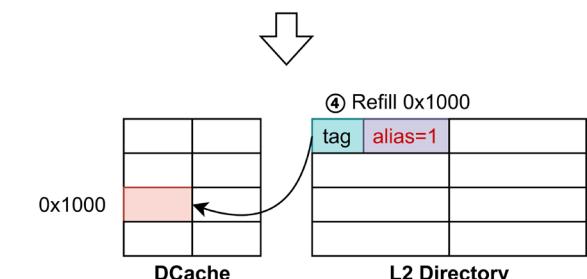
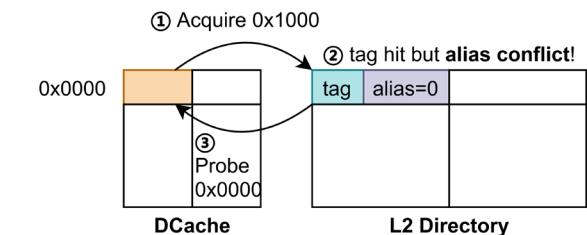
- Multi-bank for multi read ports
- Anti-alias for larger size
- MSHR request merge
- Hardware L1 data prefetcher



Anti-Alias processing flow

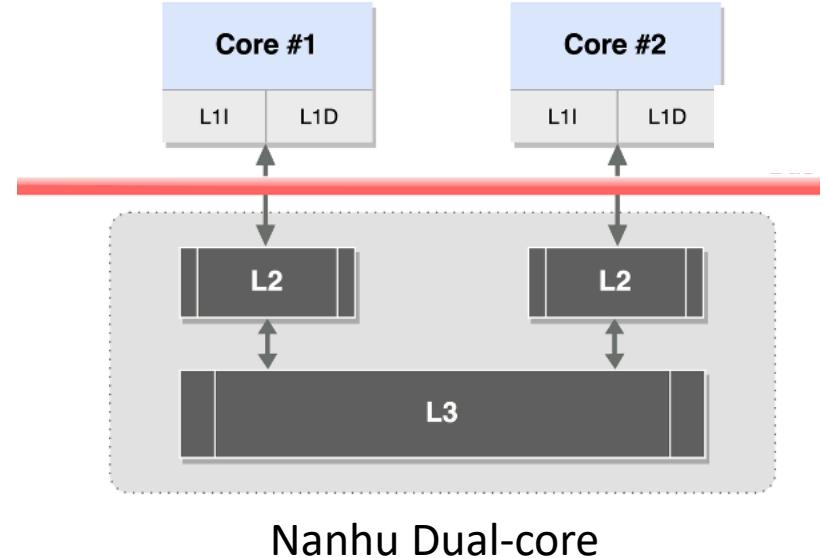


Multi-Bank



L2/L3 Cache

- High performance non-blocking **L2/L3** cache
- Design features
 - **Directory based coherence**
 - **Inclusive/Non-inclusive**
 - **TileLink protocol**
- Highly configurable parameters
 - Slice, Size, #Ways, #MSHRs
 - Replacement: **Random/PLRU/RRIP**
 - Prefetch: **BOP/SMS**



Nanhua Dual-core

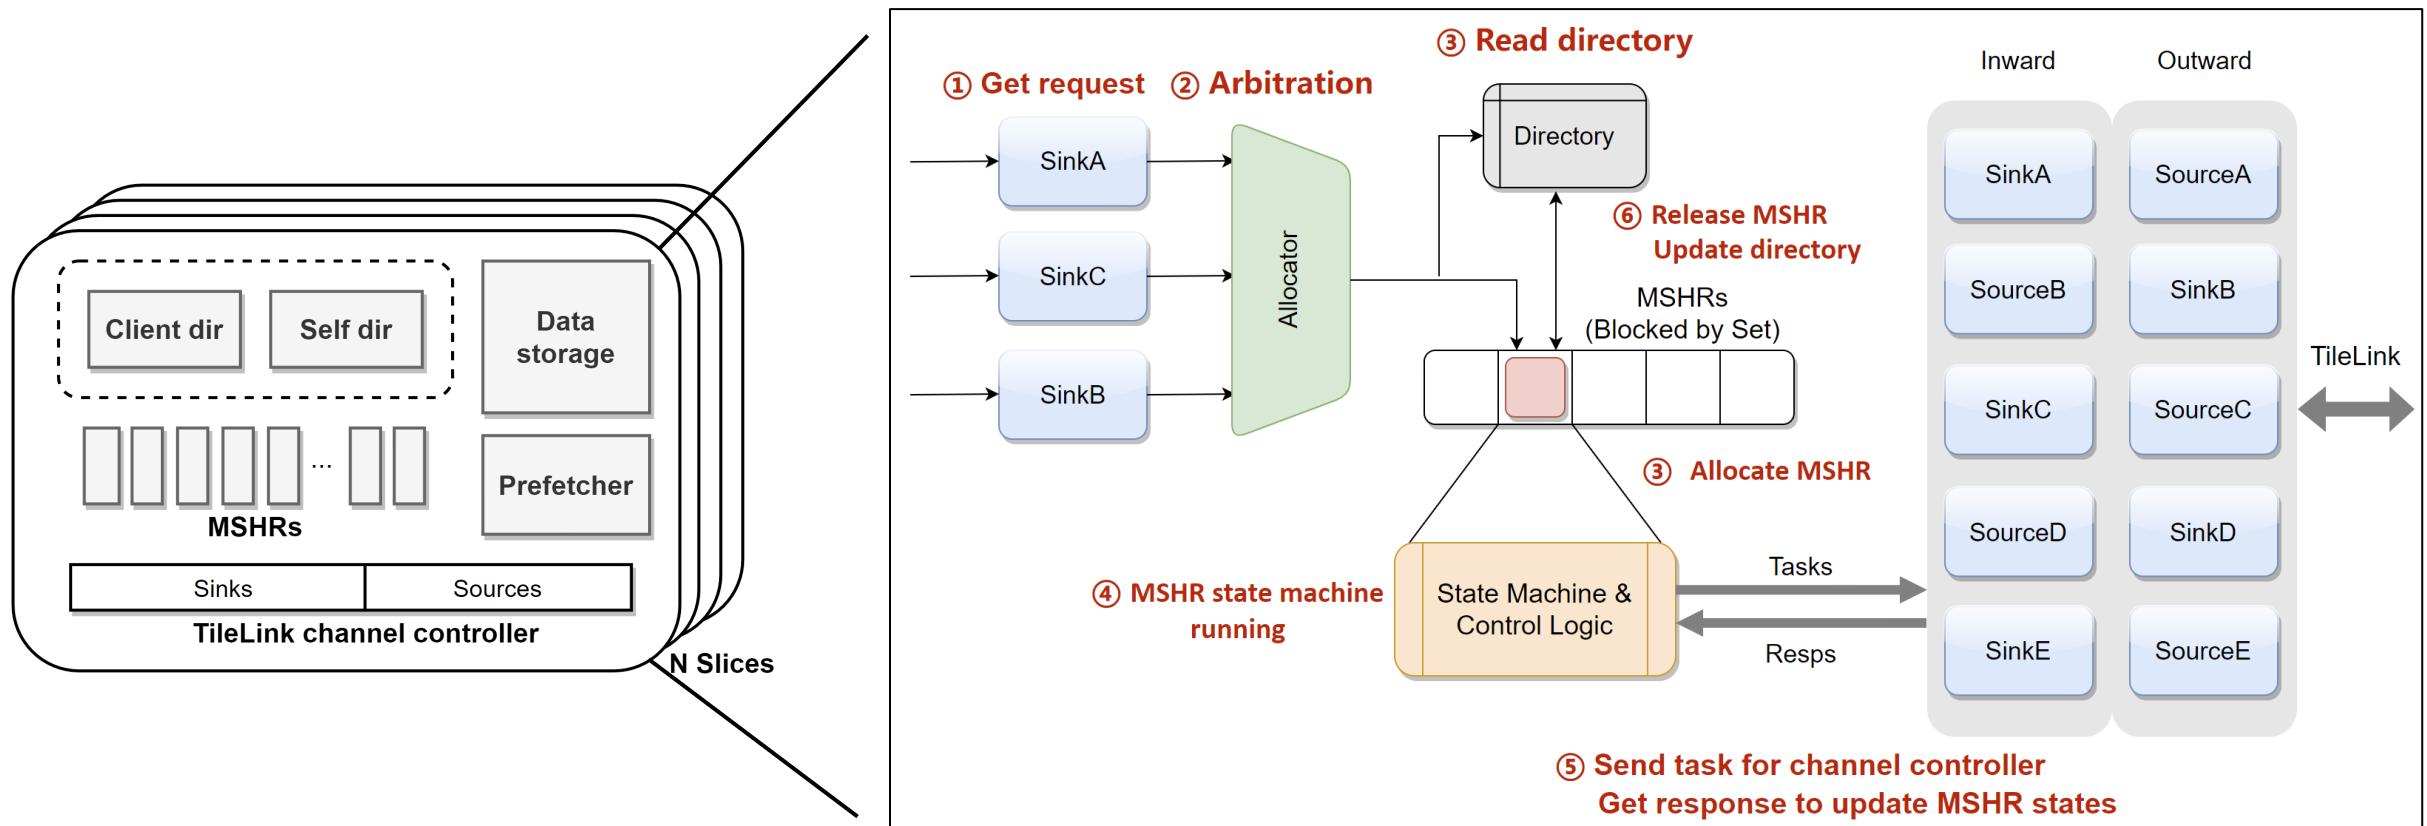
A screenshot of the GitHub repository page for OpenXiangShan/HuanCun. The repository is public and has 36 branches and 0 tags. The codebase is non-inclusive. Recent commits include a merge pull request from Lemover, changes to .github/workflows, rocket-chip, scripts, and src, and updates to .gitignore, .gitmodules, .mill-version, scalafmt.conf, LICENSE, LICENSE.SiFive, Makefile, README.md, and build.sc. The commits are dated from 8 days ago to 3 months ago.

Commit	Date	Message
Merge pull request #91 from OpenXiangShan/sram-old-hazard ...	8 days ago	dd9e202
.github/workflows	3 months ago	Turn on chisel3-plugin
rocket-chip @ 85f319c	3 months ago	Bump rocketchip
scripts	9 months ago	script: refine tl filter(no tip on tip & must have tip on trunk)
src	10 days ago	util.sram: rm a r/w hazard mux which is not needed
.gitignore	8 months ago	gitignore: add build/
.gitmodules	6 months ago	misc: use https url for git submodule
.mill-version	3 months ago	Turn on chisel3-plugin
scalafmt.conf	13 months ago	Add scalafmt check
LICENSE	10 months ago	add Mulan PSL2 license (#15)
LICENSE.SiFive	10 months ago	add Mulan PSL2 license (#15)
Makefile	5 months ago	test: enable prefetch during tl-test
README.md	13 months ago	Create README.md
build.sc	3 months ago	Turn on chisel3-plugin

<https://github.com/OpenXiangShan/HuanCun>

L2/L3 Cache

- Concurrent request processing flow





Microarchitecture Implementation

- Implement by **Chisel** HDL
 - High readability
 - High modifiability
- Beyond built-ins, we designed some open-source high-performance utilities/components



BinaryArbiterNode.scala
BitUtils.scala
ChiselDB.scala
CircularQueuePtr.scala
Constantin.scala
DataModuleTemplate.scala
ECC.scala
ExcitingUtils.scala

ExtractVerilogModules.scala
FastArbiter.scala
FileRegisters.scala
GTimer.scala
Hold.scala
IntBuffer.scala
LFSR64.scala
LatencyPipe.scala

LookupTree.scala
MIMOQueue.scala
Misc.scala
ParallelMux.scala
Pipeline.scala
PipelineConnect.scala
PriorityMuxDefault.scala
PriorityMuxGen.scala

RegMap.scala
Replacement.scala
ResetGen.scala
SRAMTemplate.scala
StopWatch.scala
TLClientsMerger.scala
TLEdgeBuffer.scala

Check it out! <https://github.com/OpenXiangShan/Utility>

Institute of Computing Technology, CAS

Microarchitecture Implementation

- **Highly configurable**
 - You can setup parameters as you wish

```
IBufSize: Int = 48,
DecodeWidth: Int = 6,
RenameWidth: Int = 6,
CommitWidth: Int = 6,
FtqSize: Int = 64,
EnableLoadFastWakeUp: Boolean = true,
IssQueSize: Int = 16,
NRPhyRegs: Int = 192,
LoadQueueSize: Int = 80,
LoadQueueNWriteBanks: Int = 8,
StoreQueueSize: Int = 64,
StoreQueueNWriteBanks: Int = 8,
RobSize: Int = 256,
dpParams: DispatchParameters =
  DispatchParameters(
    IntDqSize = 16,
    FpDqSize = 16,
    LsDqSize = 16,
    IntDqDeqWidth = 4,
    FpDqDeqWidth = 4,
    LsDqDeqWidth = 4
  ),
exuParameters: ExuParameters =
  ExuParameters(
    JmpCnt = 1,
    AluCnt = 4,
    MulCnt = 0,
    MduCnt = 2,
    FmacCnt = 4,
    FmiscCnt = 2,
    FmiscDivSqrtCnt = 0,
    LduCnt = 2,
    StuCnt = 2
  ),
prefetcher: Option[PrefetcherParams] =
  Some(SMSParams()),
LoadPipelineWidth: Int = 2,
StorePipelineWidth: Int = 2,
StoreBufferSize: Int = 16,
StoreBufferThreshold: Int = 7,
EnableLoadToLoadForward: Boolean = true,
EnableFastForward: Boolean = false,
EnableLdVioCheckAfterReset: Boolean = true,
EnableSoftPrefetchAfterReset: Boolean = true,
EnableCacheErrorAfterReset: Boolean = true,
EnablePTWPreferCache: Boolean = true,
EnableAccurateLoadError: Boolean = true,
```

Some key parameters

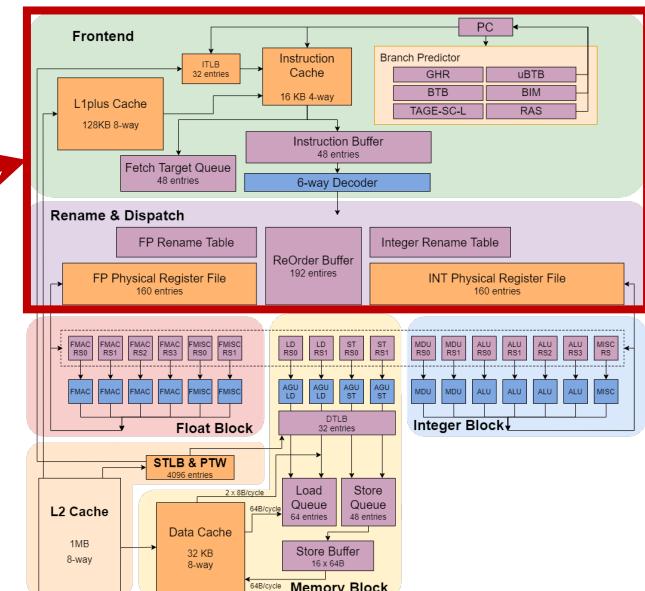
e.g.

How to change rename/decode/issue width?

All you need is to change a number

```
23 // Synthesizable minimal XiangShan
24 // * It is still an out-of-order, super-scalar arch
25 // * L1 cache included
26 // * L2 cache NOT included
27 // * L3 cache included
28 class MinimalConfig(n: Int = 1) extends Config(
29   new DefaultConfig(n).alter((site, here, up) => {
30     case SocParamsKey => up(SocParamsKey).copy(
31       cores = up(SocParamsKey).cores.map(_.copy(
32         DecodeWidth = 2,
33         RenameWidth = 2,
34         FetchWidth = 4,
35         IssQueSize = 8,
36         NRPhysRegs = 80,
37         LoadQueueSize = 16,
38         StoreQueueSize = 16,
39         RoqSize = 32,
40         BrqSize = 8,
41         FtqSize = 16,
42         IBufSize = 16,
43         StoreBufferSize = 4,
44         StoreBufferThreshold = 3,
45         dpParams = DispatchParameters(
46           IntDqSize = 8,
47           FpDqSize = 8,
48           LsDqSize = 8,
```

All affected
But it works

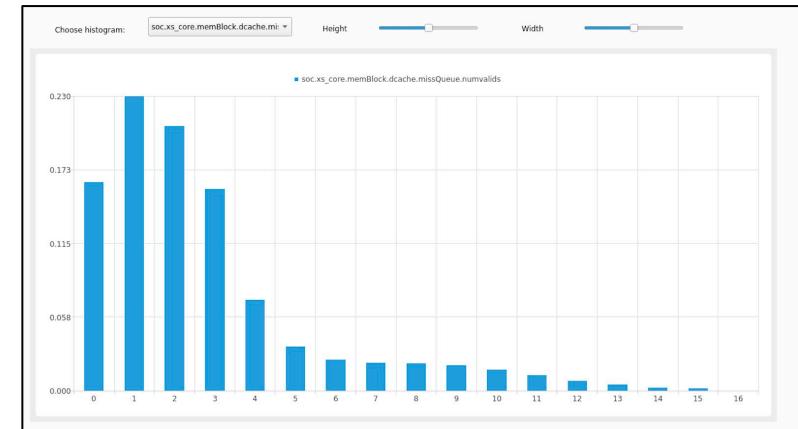




Microarchitecture Implementation

- Easy-to-use performance counter system
- Multiple collection types
 - Cumulative counter
 - Delay counter
 - Transaction counter
 - Queue counter
- Multiple analysis style
 - Key-value list
 - Histogram
 - Database
 - Time series visualization

TOP.XSSimSoC.soc.xs_core.ctrlBlock.ftq: BpBInstr,	57899
TOP.XSSimSoC.soc.xs_core.ctrlBlock.ftq: BpBRight,	54695
TOP.XSSimSoC.soc.xs_core.ctrlBlock.ftq: BpBWrong,	3204
TOP.XSSimSoC.soc.xs_core.ctrlBlock.ftq: BpCRight,	3615
TOP.XSSimSoC.soc.xs_core.ctrlBlock.ftq: BpCWrong,	34
TOP.XSSimSoC.soc.xs_core.ctrlBlock.ftq: BpInstr,	78176
TOP.XSSimSoC.soc.xs_core.ctrlBlock.ftq: BpIRight,	5471
TOP.XSSimSoC.soc.xs_core.ctrlBlock.ftq: BpIWrong,	6099
TOP.XSSimSoC.soc.xs_core.ctrlBlock.ftq: BpJRight,	8707
TOP.XSSimSoC.soc.xs_core.ctrlBlock.ftq: BpJWrong,	0
TOP.XSSimSoC.soc.xs_core.ctrlBlock.ftq: BpRight,	68873





KUNMINGHU: 3rd generation microarchitecture

- Work with *Beijing Institute of Open Source Chip (BOSC)*
 - Non-profit organization with 18 founding members from the industry
- RISC-V Vector Extension 1.0 Support
- RISC-V Hypervisor Extension 1.0 Support
- Loop predictor, loop buffer
- Refactored issue queues, execution units, load/store units/queues
- L1D prefetcher, cross-level cache optimizations
- AMBA CHI compatible
- Comprehensive verification plan
- Advanced process node and EDA flows
- GEM5 simulator with aligned microarchitecture



The screenshot shows a GitHub repository's commit history. It displays multiple groups of commits, each with a date header (e.g., 'Commits on Mar 19, 2023', 'Commits on Mar 18, 2023', etc.). Each group contains several commits with details like author, commit message, and status. The commits are related to various microarchitecture components like dcache, MMU, and ftq.

Date	Commit Details
Commits on Mar 19, 2023	Fix reply logic in unified load queue (#1966) (happy-lx committed 2 days ago), util: change ElaborationArtifacts to FileRegisters (#1973) (Mapleze-U committed 2 days ago)
Commits on Mar 18, 2023	dcache: optimize the ready signal of missqueue (#1965) (happy-lx committed 5 days ago), bump diffest, track master branch (#1967) (Lemover committed 5 days ago)
Commits on Mar 15, 2023	MMU: Add sector tlb for larger capacity (#1964) (good-circle committed last week)
Commits on Mar 13, 2023	dcache: fix plru update logic (#1921) (AugustinWillWing committed last week)
Commits on Feb 27, 2023	c: use checkout@v3 instead of v2 (#1942) (Tang-Haijin committed 3 weeks ago)
Commits on Feb 22, 2023	ftq: revert #1875, #1920 (#1931) (stencema and lyn committed last month)

Active development on GitHub



Roadmap of Kunminghu

- **Functional Improvement**
 - Support RISC-V **Vector** extension
 - Support RISC-V **Hypervisor** extension
- **Performance Exploration**
 - Performance upgrade of front-end, back-end, load-store unit and cache
 - Performance model calibrated with RTL
 - Workflow: **DSE on perf Model => Impl. & fine tuning on RTL**
- **Functional Verification**
 - **Hierarchical veri. flow** including ST、IT、UT + FPGA
 - Industrial-grade verification process
- **Physical Design**
 - Experienced physical design team
 - Simultaneous iteration of RTL coding with physical design timing evaluation



Performance Evaluation of Kuminghu

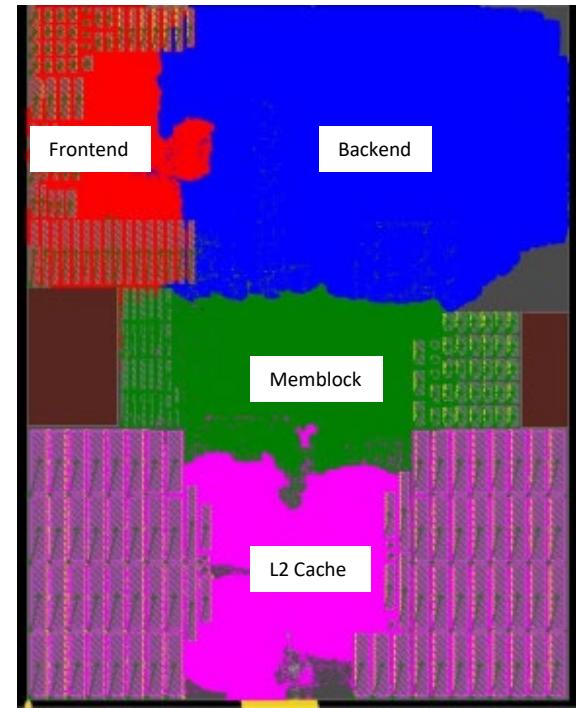
- Method: SPEC CPU checkpoints selected by Simpoint
 - Base: GCC 12 -O3, RV64GCB, jemalloc
 - 1MB L2 and 16MB L3
 - Simulated@3GHz with DRAMsim3 DDR4-3200

SPECint 2006	
400.perlbench	36.648
401.bzip2	24.283
403.gcc	47.692
429.mcf	57.852
445.gobmk	31.711
456.hmmer	39.567
458.sjeng	31.500
462.libquantum	125.487
464.h264ref	57.376
471.omnetpp	42.243
473.astar	30.738
483.xalancbmk	75.535
SPECint2006@3GHz	44.977

Estimated SPECint 2006 Base
44.98@3GHz

SPECint 2006	
400.perlbench	42.143
401.bzip2	24.283
403.gcc	47.692
429.mcf	57.772
445.gobmk	31.711
456.hmmer	60.826
458.sjeng	31.500
462.libquantum	222.412
464.h264ref	60.243
471.omnetpp	42.243
473.astar	30.738
483.xalancbmk	81.213
SPECint2006@3GHz	49.964

Estimated SPECint 2006 Peak
49.96@3GHz

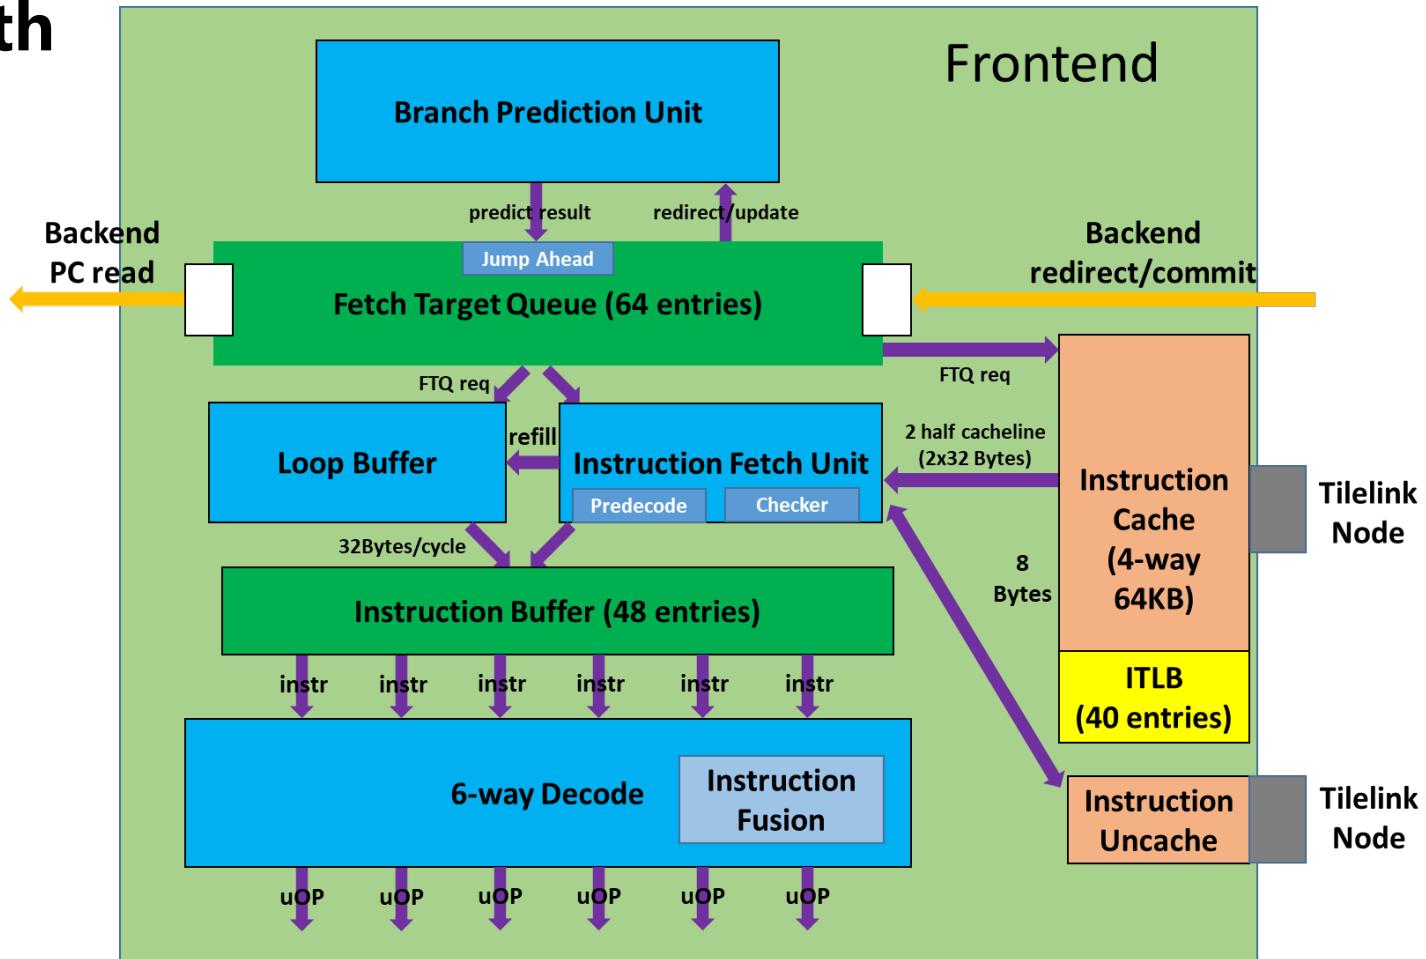


Floorplan of V3 (single core)



① Frontend Upgrade

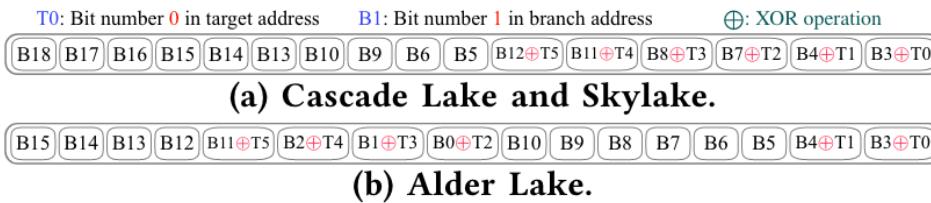
- Larger fetch/prediction width
- Lower prediction miss rate
- Higher frequency





Feature 1: TAGE Predictor Fine Tuning

- ① Branch History Hashing
- ② #Table & History Length



Hash algorithm affects performance

Genetic and particle swarm algorithms are employed to adjust parameters automatically.

Param of Nanhu → Param tuned: IPC +1.44%

Bonus!

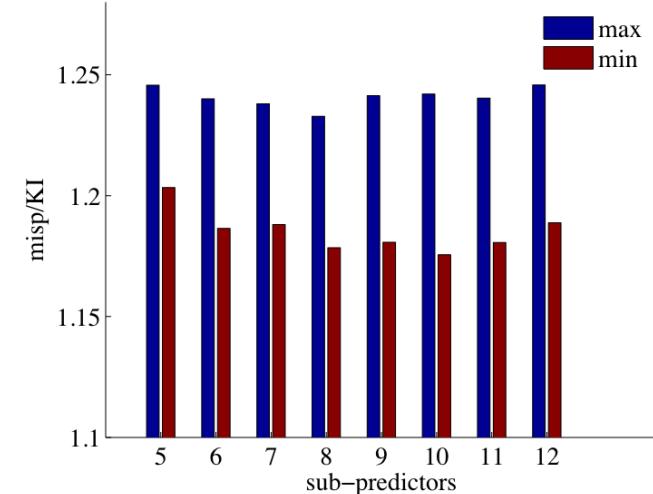


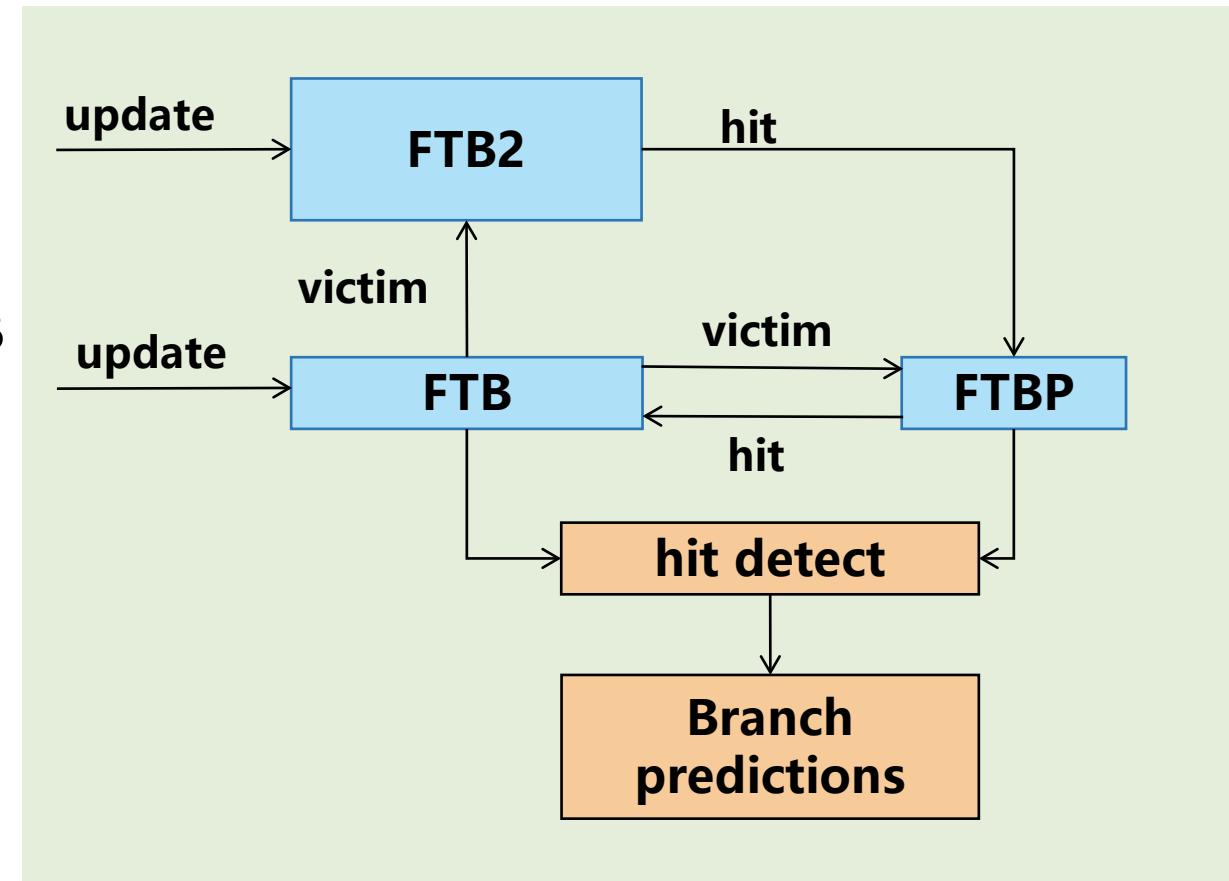
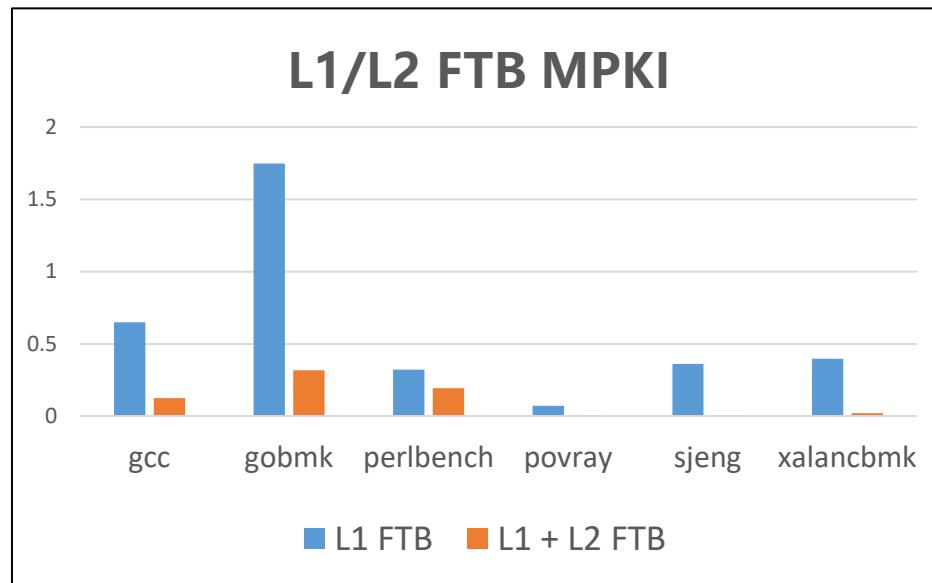
Figure 2: Performance of random parameters

Previous studies have shown that the #tables and history length have a significant effect on TAGE performance, up to 5.64% in MPKI [1]

[1] C. Zhou, L. Huang, Z. Li, T. Zhang, and Q. Dou, “Design Space Exploration of TAGE Branch Predictor with Ultra-Small RAM,” in *Proceedings of the on Great Lakes Symposium on VLSI 2017*, Banff Alberta Canada: ACM, May 2017, pp. 281–286.

Feature 2: L2 FTB (BTB)

- Hierarchical FTB
 - Support 4K L2 FTB
 - Semi-exclusive
 - FTBP acts as a buffer between L1 & L2 FTB



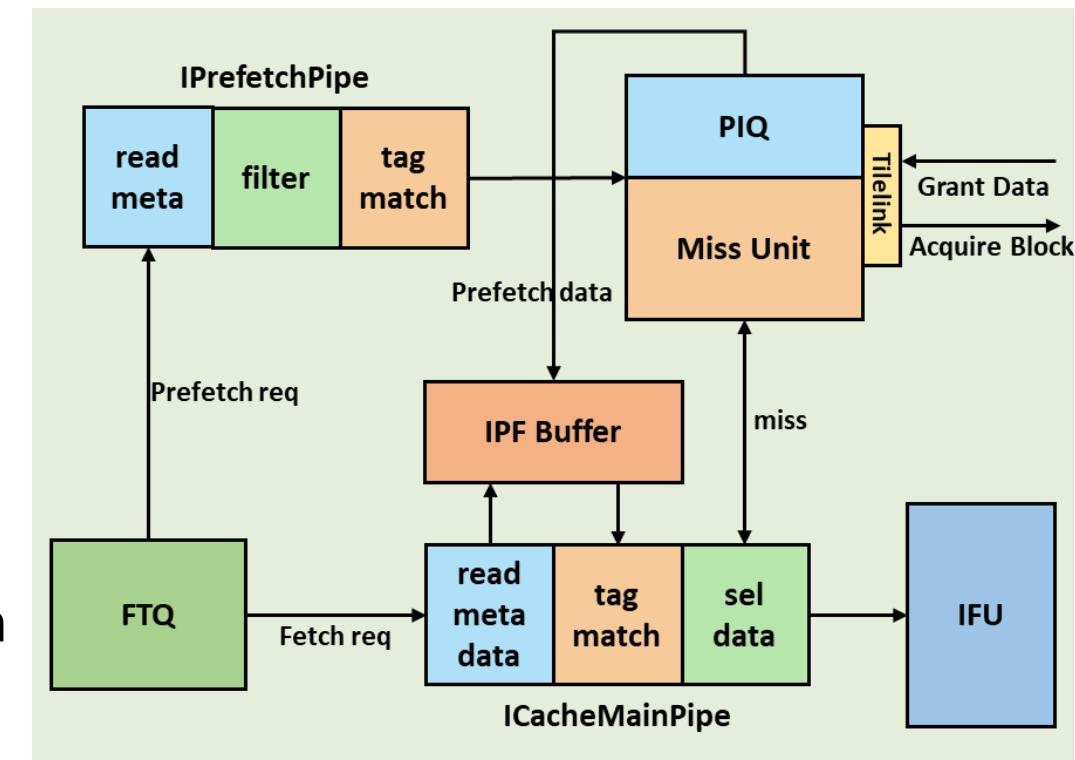
L2 FTB Structure

Feature 3: Smart ICache Prefetcher

- Based on decoupled frontend, FDIP prefetching algorithm is implemented

Prefetcher workflow

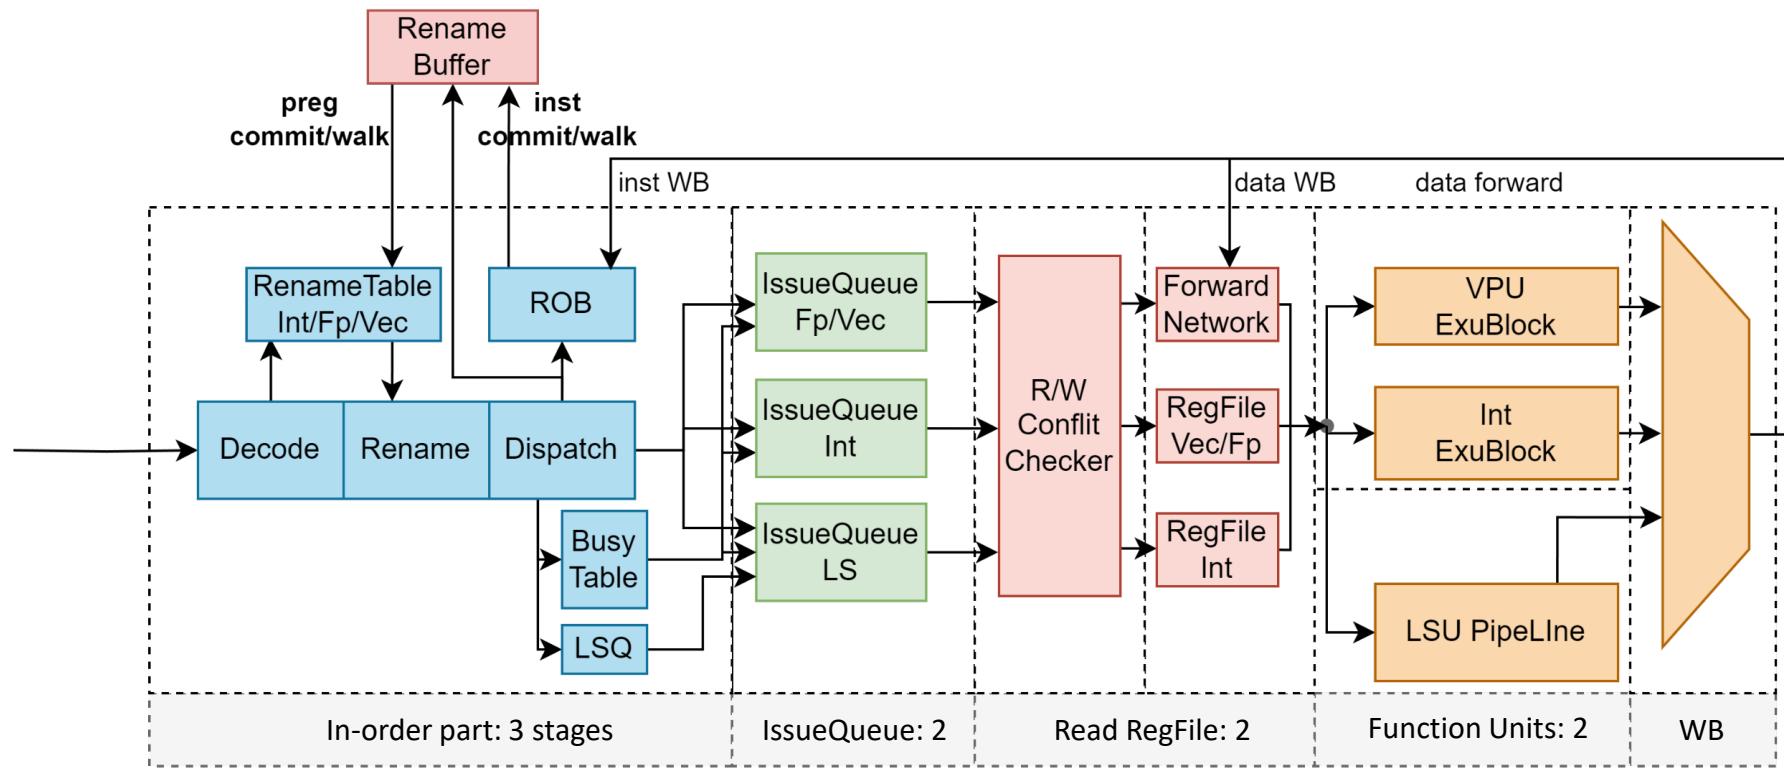
1. Select prefetch request from PF-queue
 2. Write refilled data into prefetch buffer
 3. Fetch query icache & prefetch buffer parallelly
- Prefetch position: L2 Cache or L1 ICache
 - Dual prefetching pipelines to enlarge bandwidth





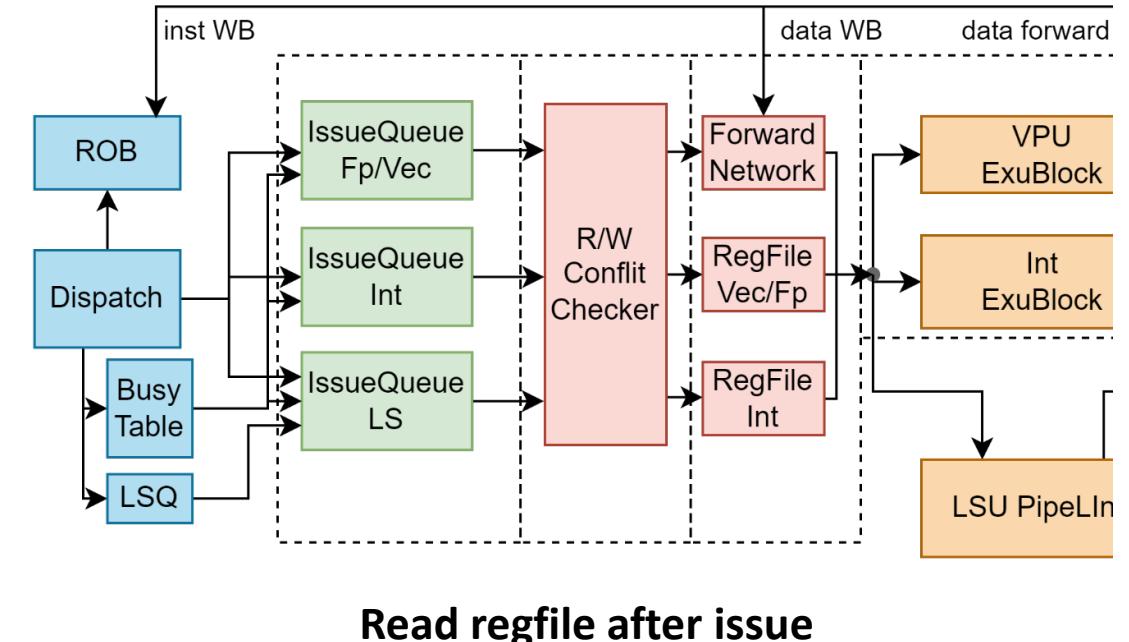
② OoO Backend Upgrade

- Read regfile after issue
- Faster μ op commit and redirect
- Vector execution unit



Feature 1: Read Regfile after issue

- Advantage
 - Bottleneck is shifted to after μ op issue
 - Reduce RS storage to save SRAM
 - Reduce latency and increase RS capacity
- Optimization
 - Efficient arbitration of regfile reading ports
 - Accurate speculative wakeup and cancelling

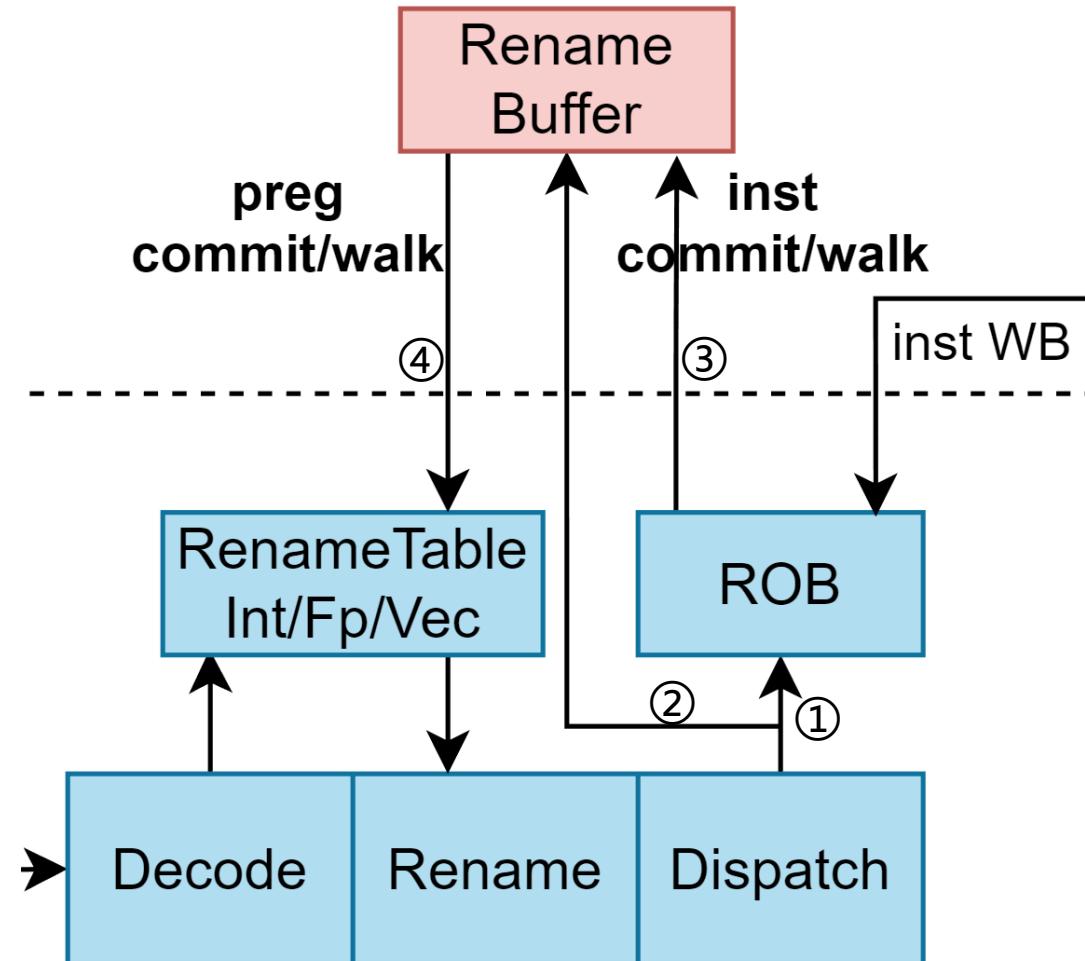


Feature 2: Decouple `pop commit` & `regfile commit`

- New Rename Buffer (**RAB**) design: record regfile changing history

- ① Save regfile mapping info to RenameBuffer (**RAB**)
 - ② On instruction commit/redirect, wake RAB up
 - ③ RAB is responsible to update RenameTable (**RAT**)

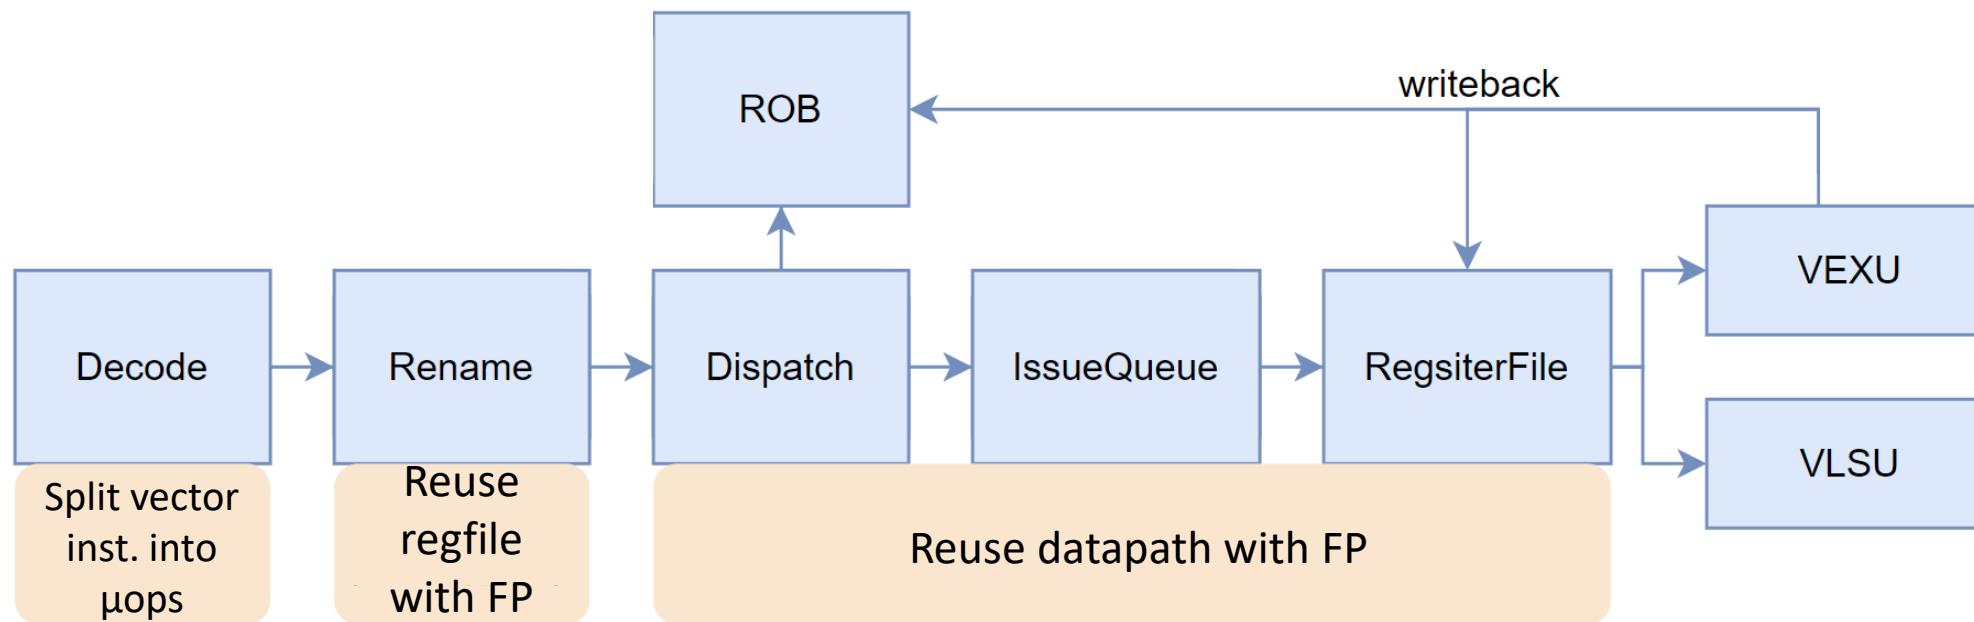
- Benefits
 - Faster ROB entry releasing
 - Optimize timing
 - Support μop split of vector instructions
 - Support ROB compression





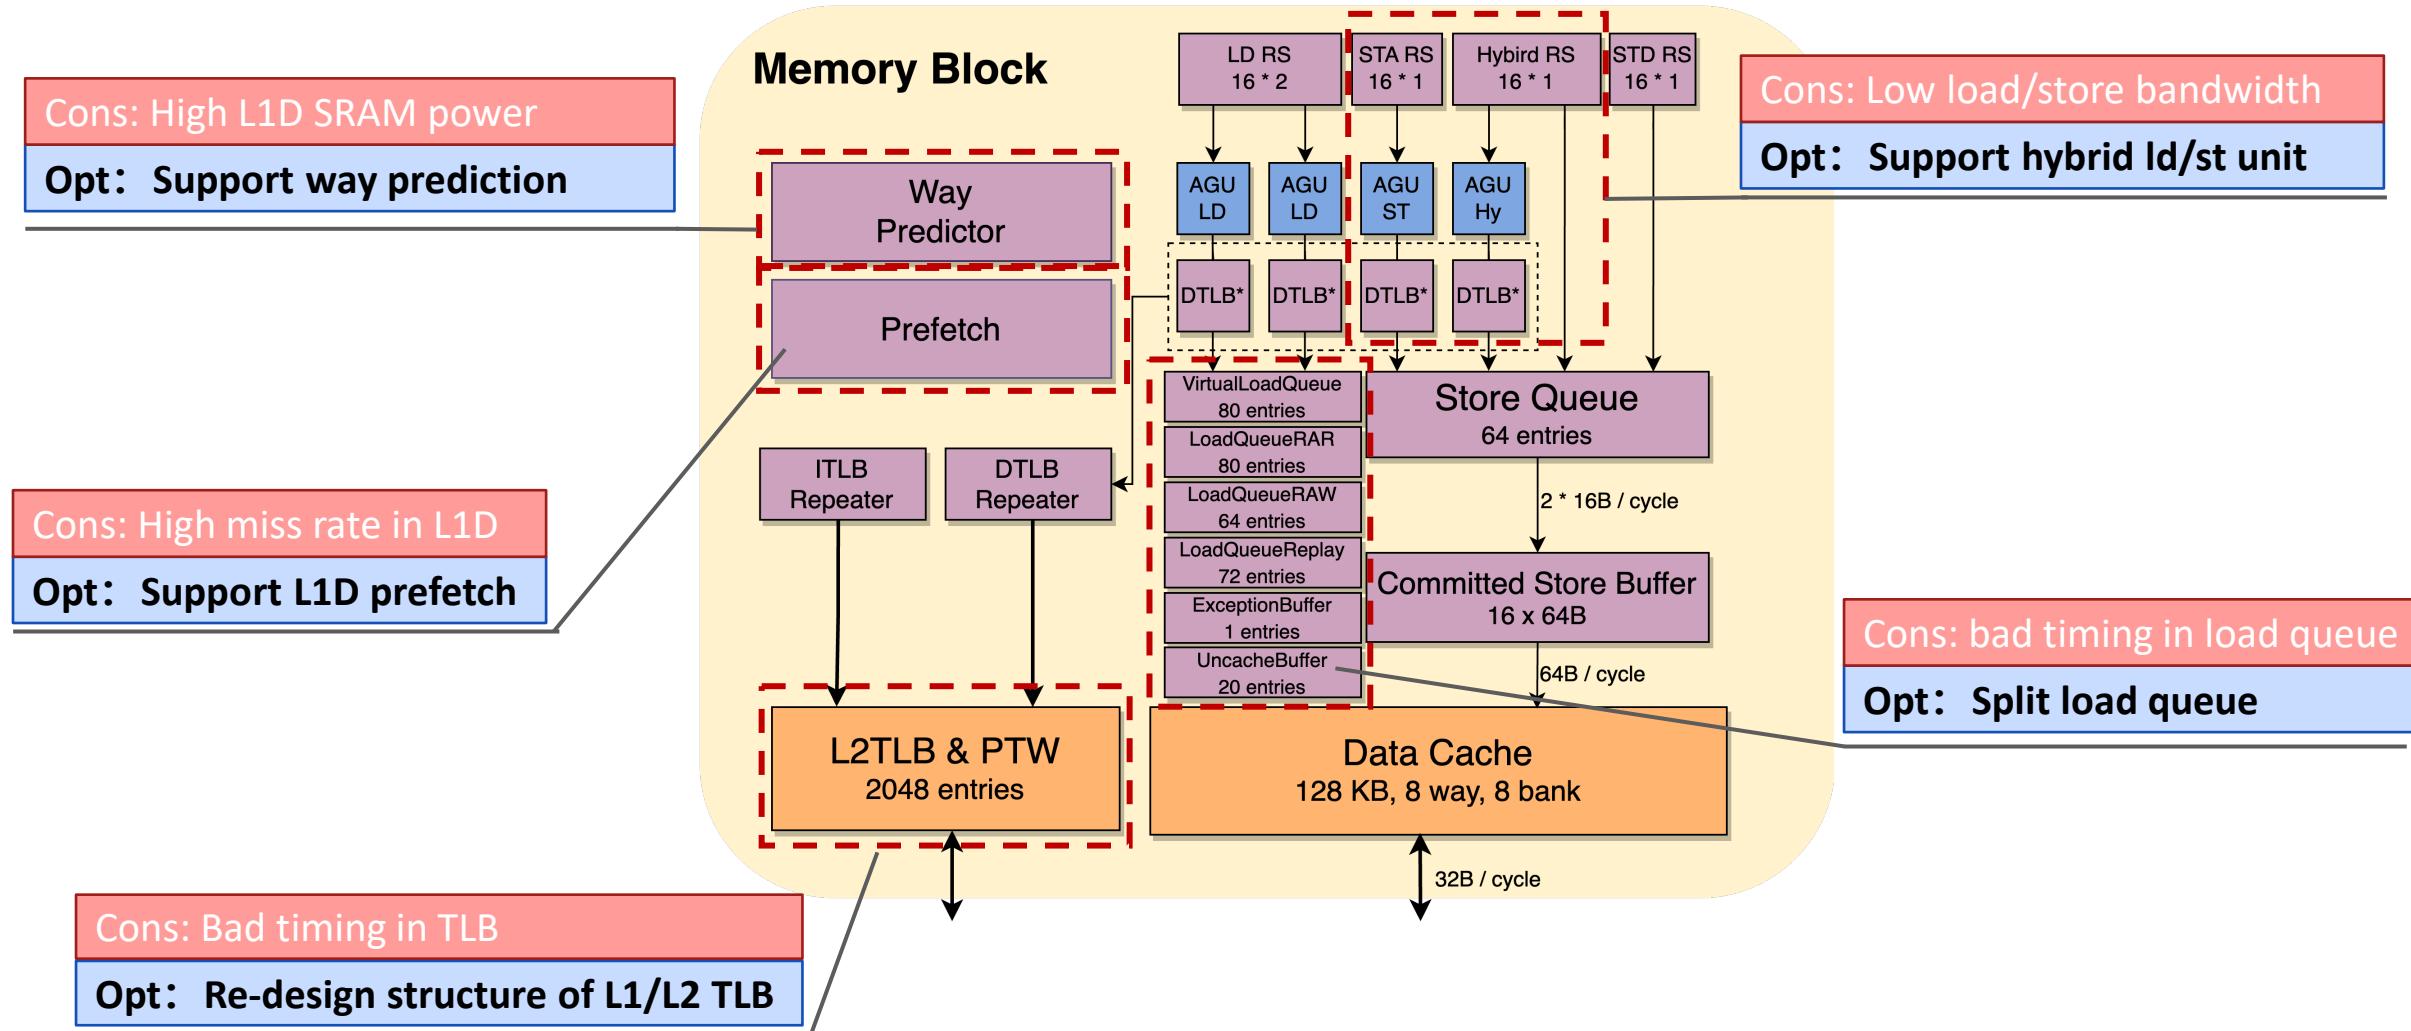
Feature 3: Support Vector Extension

- Coupled vector unit design, compatible with RISC-V Vector 1.0
 - VLEN = 128
- Reuse backend pipeline: Decode/Dispatch/OoO/Execute





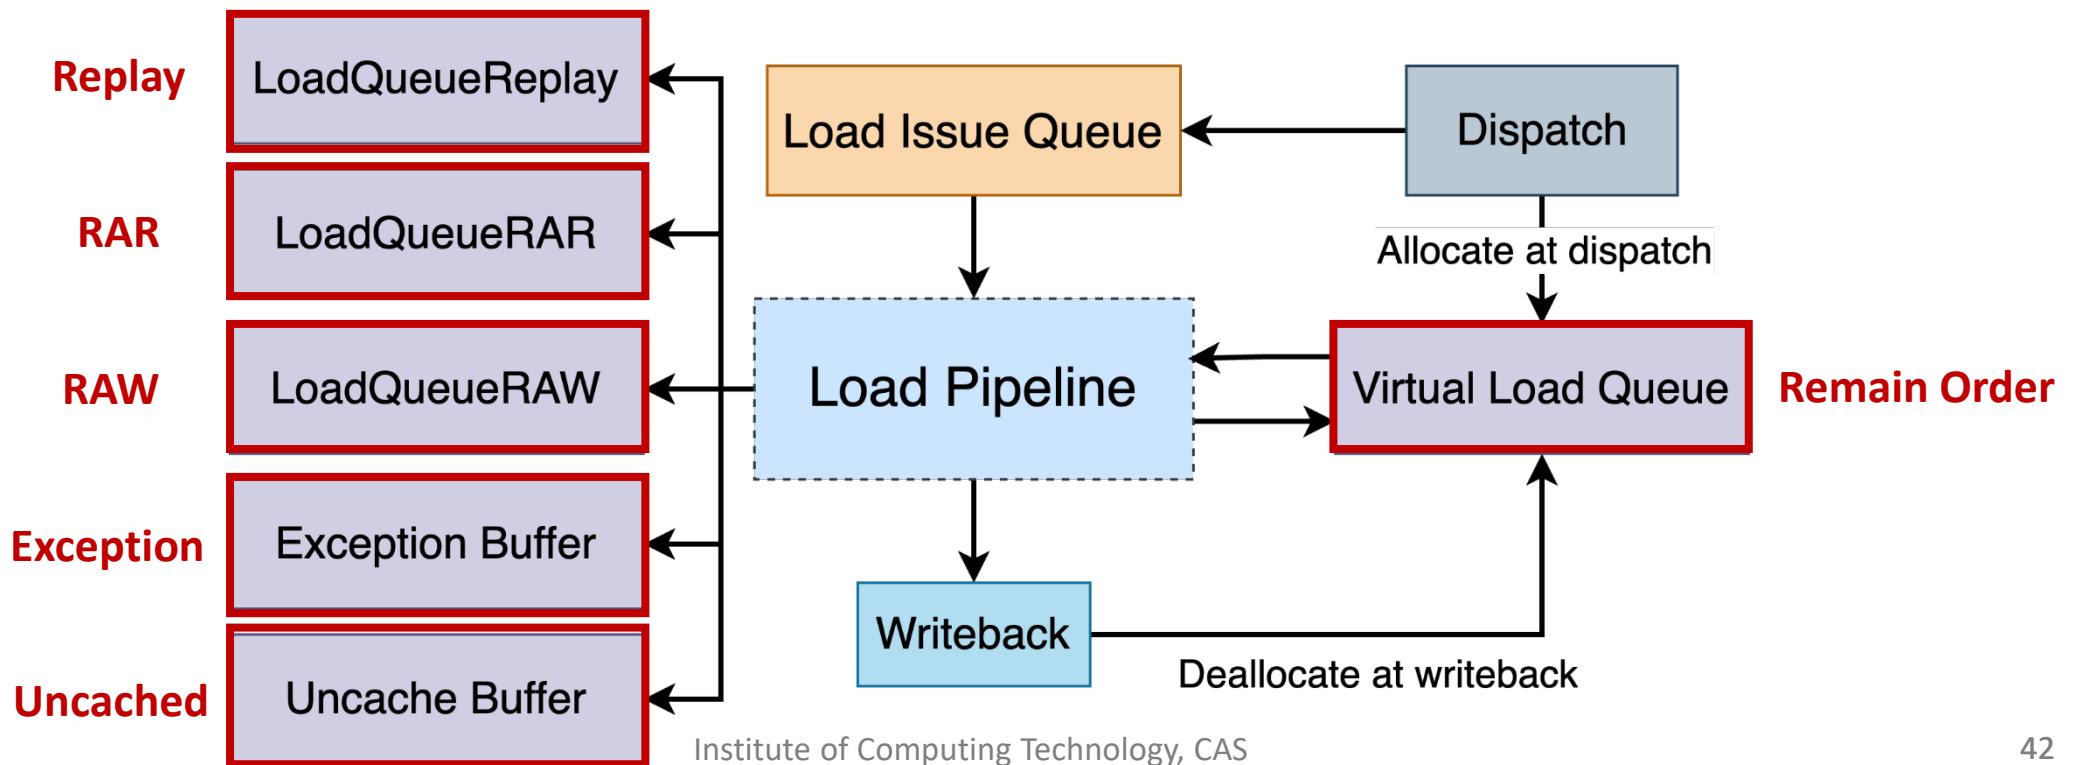
③ Load-Store Unit Upgrade





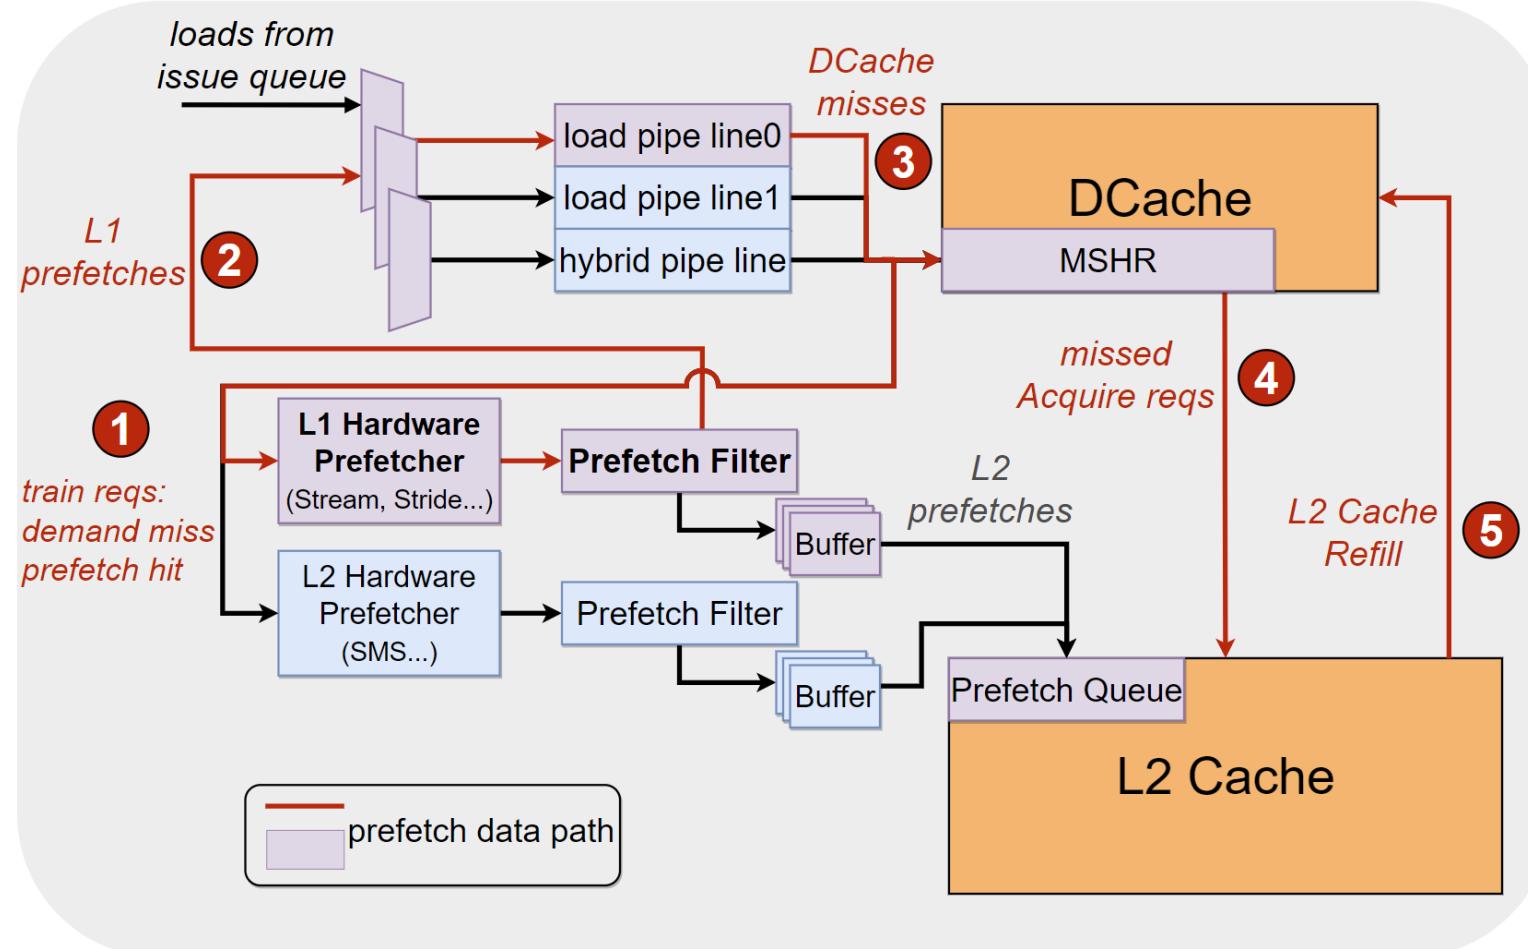
Feature 1: Split Load Queue

- Split load queue **based on different responsibilities**
 - **Early commit:** load instruction can be retired sequentially from virtual load queue after write-back
 - Logic simplified, timing friendly and area acceptable



Feature 2: L1D Prefetch

- L1D prefetch workflow
 - Training pattern detection
 - **Reuse load pipelines**
 - Fetch data from D\$ MSHR
- Prefetch algorithm
 - **Stream**^[1]
 - $(x, x+1, x+2, x+3 \dots)$
 - **Stride**^[2]
 - $(x, x+n, x+2n, x+3n \dots)$

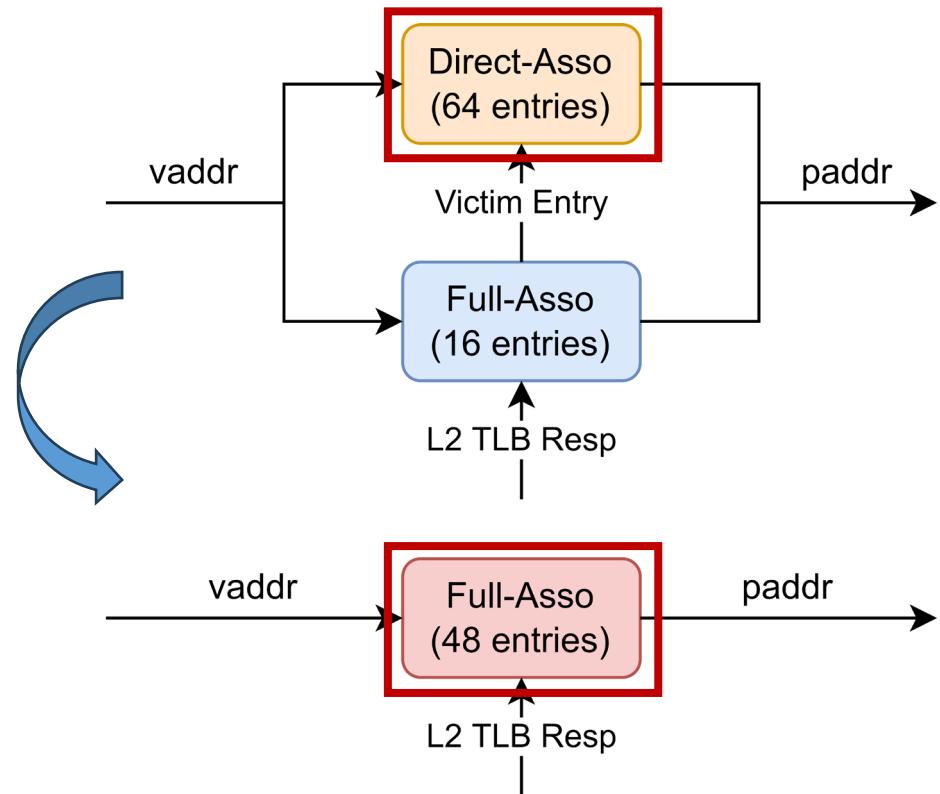


[1] S. Srinath, O. Mutlu, H. Kim and Y. N. Patt, "Feedback Directed Prefetching: Improving the Performance and Bandwidth-Efficiency of Hardware Prefetchers," 2007 IEEE 13th International Symposium on High Performance Computer Architecture, Scottsdale, AZ, USA, 2007

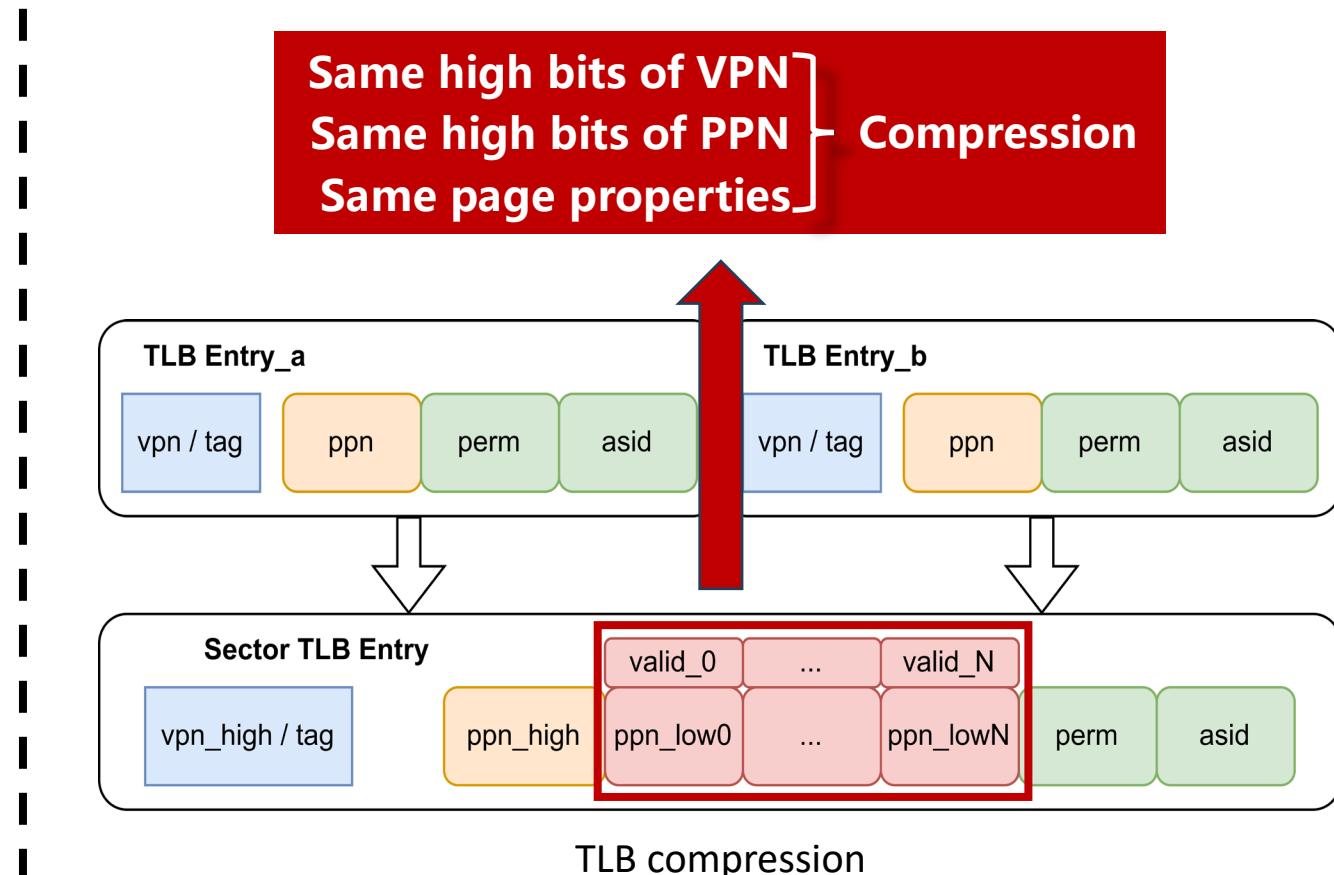
[2] J. -L. Baer and T. -F. Chen, "An effective on-chip preloading scheme to reduce data access penalty," Supercomputing '91:Proceedings of the 1991 ACM/IEEE Conference on Supercomputing, Albuquerque, NM, USA, 1991

Feature 3: Powerful MMU

- 16 fully associative entry + 64 direct mapping entry → **48 fully associative entry**
- Support **TLB compression**, merging contiguous page table entries



DTLB organization in Nanhu and Kunminghu

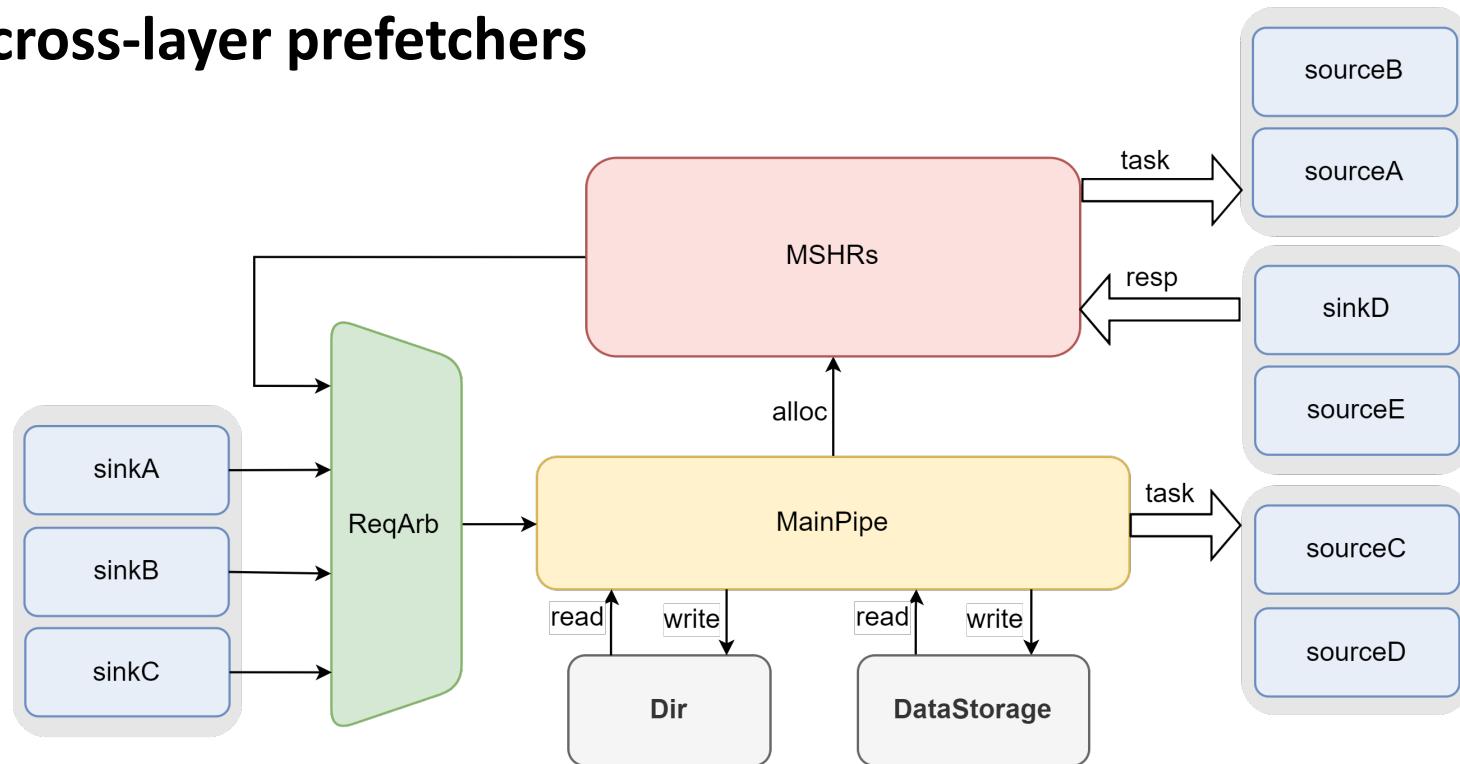


TLB compression



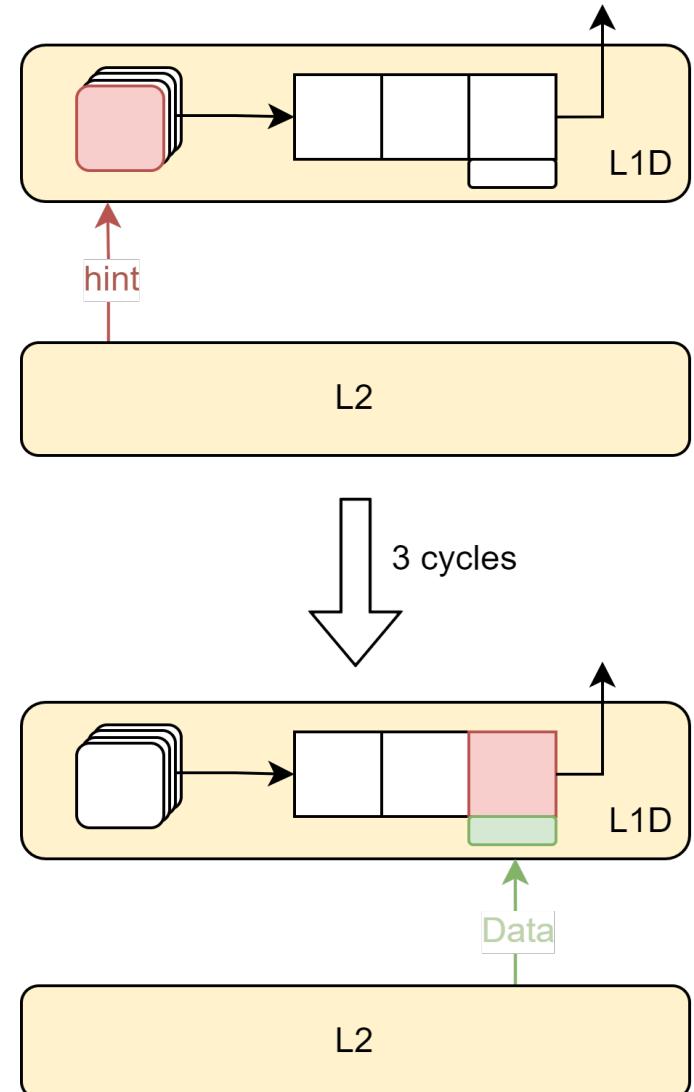
④ L2 Cache Upgrade

- Multiple concurrent pipelines → Non-blocking main pipeline
- L1-L2 Refill wake-up collaboratively
- Eliminate transaction serialization in the same set & Request merge
- Aggressive cross-layer prefetchers



Feature 1: L1-L2 Collaborative Optimization

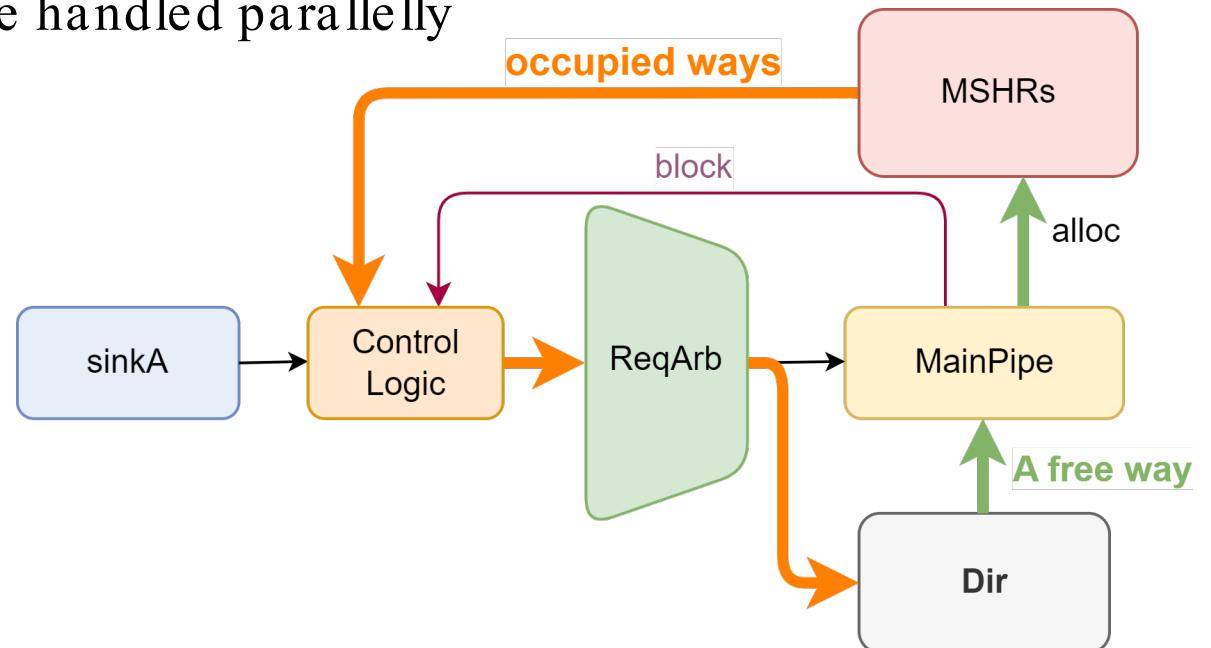
- Observation:
 - Interval between load wake-up and use-refill-data is 3 cycles
 - There's chance to hint LSU to replay before refill
- New Design:
 - Let L2 give Hint signal in advance
 - Speedup the wake and replay of load replay queue
- Implementation:
 - Set up monitor to track info from main pipeline
 - Send hint signal to L1 3 cycles before refill
 - Need to calculate the exact Hint timing considering possible pipeline delays





Feature 2: Eliminate Txns Serialization in the same set

- Background: Txns to the same set were handled serially, which reduces parallelism
- Difficulty: Same-set txns can affect each other, due to replacement
- New Design
 - Txns of same-set but different-addr can be handled parallelly
 - Txns of same-addr are handled serially
- Implementation
 - Record occupied ways in active MSHRs
 - Assign a free way to the new request



Feature 3: Evict on Refill

- Observation:

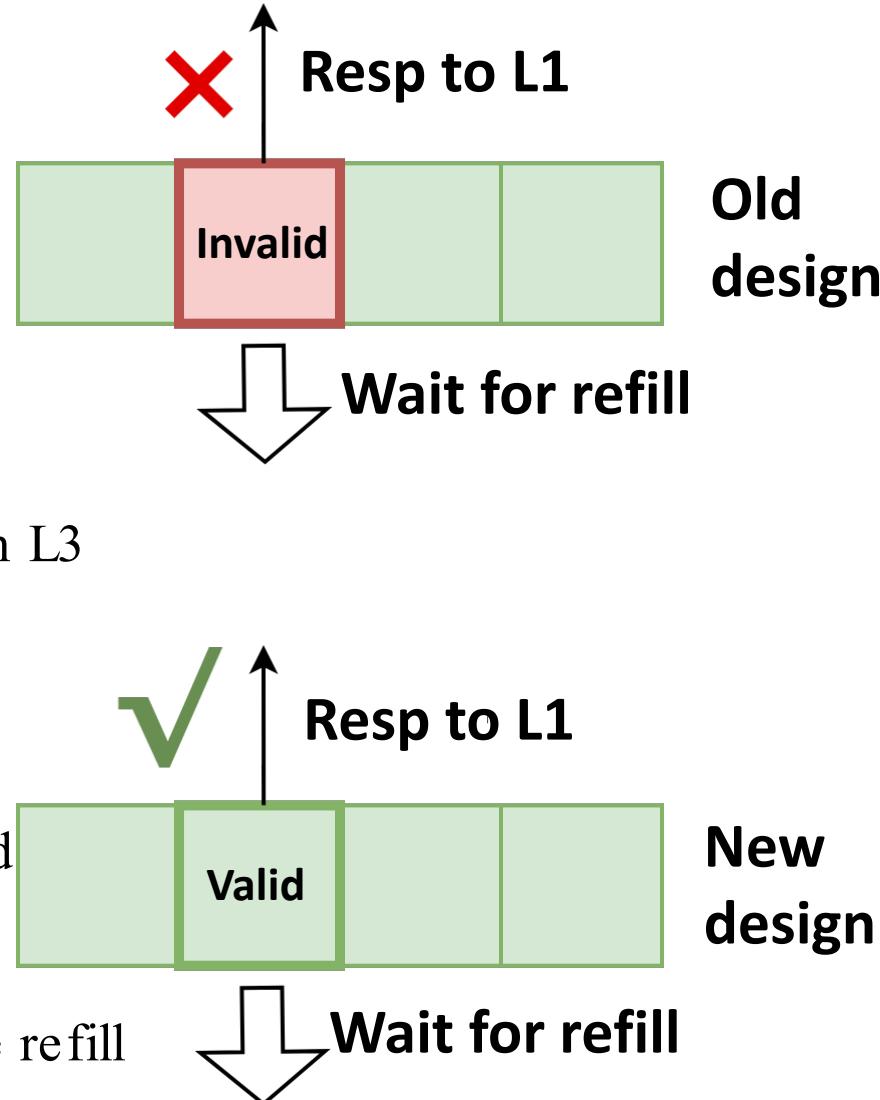
- On L1 acquire, L2 chooses a way for replacement
- But this way is occupied until refill, unable to serve L1

- Improvement:

- New design chooses replaced-way after data is refilled from L3
- So, it can still serve L1 when waiting for refill

- Implementation:

- Transfer it to L3 on acquire miss, make replacer untouched
- Wait for L3 to refill
- Read directory and assign one way when MSHR handles the refill





XiangShan: Open-Source High-Performance Processor

- **1st generation: YANQIHU**
 - 2020/6: first commit of design RTL
 - 2021/7: **28nm tape-out, 1.3GHz**
 - Performance: **SPEC CPU2006 7.01@1GHz, DDR4-1600**
- **2nd generation: NANHU**
 - 2021/5: starting design exploration and RTL design
 - 2023/4: GDSII delivery
 - **14-nm tape-out, estimated SPEC CPU2006 20@2GHz**
- **3rd generation: KUNMINGHU**
 - ISA feature: **Vector (V) extension, Hypervisor (H) extension**
 - Comprehensive performance improvement
- Open-sourced at <https://github.com/OpenXiangShan/XiangShan>



Fragrant Hills in Beijing

The screenshot shows the GitHub repository page for XiangShan. It displays the following information:

- Repository name: XiangShan
- Branch: master
- Commits: 7,609
- Issues: 205
- Tags: 4
- Last commit: b15efc0 · 18 hours ago
- Recent commits (partial list):
 - Tang-Haojin chore: bump to chisel 6.1.0 (#2710)
 - .github MISC: update issue template (#2692)
 - coupledL2 @ 54b7e76 chore: bump chisel 6.0.0 (#2654)
 - debug bump diffest & mkdir for wave/perf for local-ci script's run~...
 - diffest @ f63f678 Makefile,diffest: Support palladium simulation (#2662)
 - fudian @ e1bd469 chore: bump chisel 6.0.0 (#2654)
 - huancong @ e4b0c8a chore: bump chisel 6.0.0 (#2654)
- Activity sidebar:
 - chisel3 risc-v microarchitecture
 - Readme View license
 - Activity Custom properties
 - 4.2k stars 86 watching
 - 86 forks 563 forks

>4.2K stars, >500 forks on GitHub

Thanks!