

XiangShan: An Open-Source High-Performance RISC-V Processor and Infrastructure for Architecture Research

The XiangShan Team

MICRO'24@Austin, USA

November 02, 2024



What we will cover in this tutorial

- Introduce of XiangShan (13:00 – 13:40, 40 minutes)
- CPU Microarchitecture (13:50 – 15:00, 70 minutes)
 - Design and implementation – How to implement novel ideas on XiangShan
 - Frontend: branch prediction and instruction fetch
 - Backend: out-of-order scheduler, execution units
 - Load/Store Unit: LSQ, pipelines, TLBs, data caches
 - L2/L3 caches and prefetchers
- Development workflows (15:30 – 17:00, 90 minutes)
 - Introduction of their usages – How to develop on XiangShan with MinJie
 - Simulation and Debugging
 - Research Demo



In this tutorial, we will

- Provide access to cloud servers
- Prepare the XiangShan development environment for you
- Go through the development workflows including:

Simulation

- Server login
- Environment
- RTL Generation
- RTL Simulation
-

Func. Verification

- Nexus-AM
- NEMU
- Difftest
- TL-Test
-

Perf. Verification

- Checkpoint
- Gem5
- Top-Down
- Constantin
-

- ***Required: machines with network and SSH***



Demo Instructions

- Shell commands are highlighted in **brown box** with prefix \$

```
$ echo "Hello, XiangShan"  
$ echo "Have a nice day"
```

- Description and notes are presented in **grey box** with prefix #

```
# Please prepare a laptop with an SSH client.  
# Next, let's start the demo session !
```

Let's Start!



Prerequisites

- Login to the provided cloud server
 - Windows (Windows Terminal / PowerShell is recommended)

```
PS > ssh guest@tutorial.xiangshan.cc  
# Password: xiangshan-2024
```

- Linux / Mac

```
$ ssh guest@tutorial.xiangshan.cc  
# Password: xiangshan-2024  
$ guest@xiangshan-tutorial:~$ pwd  
/home/guest
```

For offline users, please refer to <https://github.com/OpenXiangShan/xs-env/tree/tutorial-new>



Prerequisites

```
# Copy tutorial environment to your dir based on your name
$ cp -r /opt/xs-env ~/<YOUR_NAME>

# Enter your dir
$ cd ~/<YOUR_NAME>

# Set up environment variables
# DO IT AGAIN when opening a new terminal
$ source env.sh

# SET XS_PROJECT_ROOT:           /home/guest/YOUR_NAME
# SET NOOP_HOME (XiangShan RTL Home): $XS_PROJECT_ROOT/XiangShan
# SET NEMU_HOME:                 $XS_PROJECT_ROOT/NEMU
# SET AM_HOME:                   $XS_PROJECT_ROOT/nexus-am
# SET TLT_HOME:                  $XS_PROJECT_ROOT/tl-test-new
# SET gem5_home:                 $XS_PROJECT_ROOT/gem5
```



Prerequisites

Project Structure

```
$ tree -d -L 1
```

```
.
├── DRAMsim3
├── gem5
├── NEMU
├── nexus-am
├── NutShell
├── tl-test-new
├── tutorial
└── XiangShan
```

Enter XiangShan directory

```
$ cd XiangShan
```



Chisel Compilation

- **Compile RTL and build simulator with Verilator**



```
$ make emu -j4
```

Options:

CONFIG=MinimalConfig

Configuration of XiangShan

// EMU_THREADS=4

Simulation threads

// EMU_TRACE=1

Enable waveform dump

// WITH_DRAMSIM=1

Enable DRAMSim3 for DRAM simulation

// WITH_CHISELDB = 1

Enable ChiselDB feature

// WITH_CONSTANTIN = 1

Enable Constantin feature

- **Compilation might take ~20 mins**



Open Another Terminal

```
# login to the cloud server again  
$ ssh guest@tutorial.xiangshan.cc  
# Password: xiangshan-2024  
  
# Enter your dir and set up  
$ cd ~/<YOUR_NAME> && source env.sh  
  
# Enter tutorial for following operations  
$ cd tutorial
```



Run RTL Simulation with Verilator

- After building, we can run the simulator

```
# run pre-built simulator
$ cd p1-basic-func
$ ./emu -i hello.bin --no-diff 2>hello.err

# Some key options:
-i                                     # Workload to run
-C / -I                                # Max cycles /Max Insts
--diff=PATH / --no-diff      # Path of Reference Model / disable difftest
```

Great!

We have learned the basic simulation process of Xiangshan.

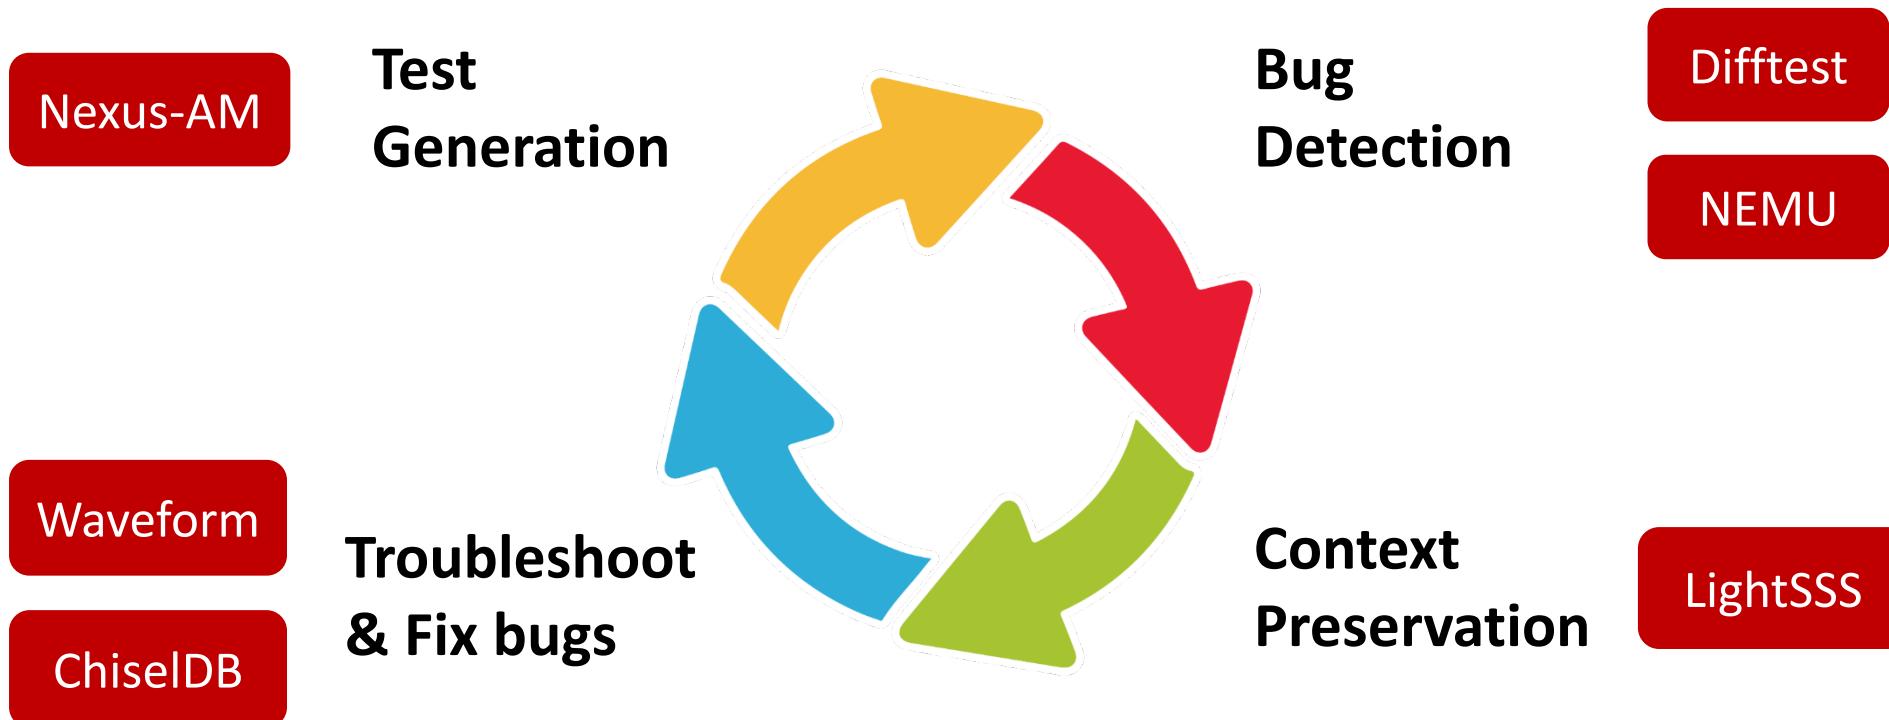
Once we get the RTL simulator ready, what's next?

How can we

- **generate desired workloads**
- **find and fix functional bugs**
- **do performance analysis**
- **do research on micro-architecture**



MinJie Functional Verification





Nexus-AM: Bare Metal Runtime

- Purpose
 - Generate workloads **agilely** without an OS
 - Provide runtime framework for **bare metal machines** like XiangShan
- We present a bare metal runtime called **Nexus-AM**
 - Light-weight, easy to use
 - Implement basic **system call** interfaces and exception handlers
 - Support multiple **ISAs and configurations**

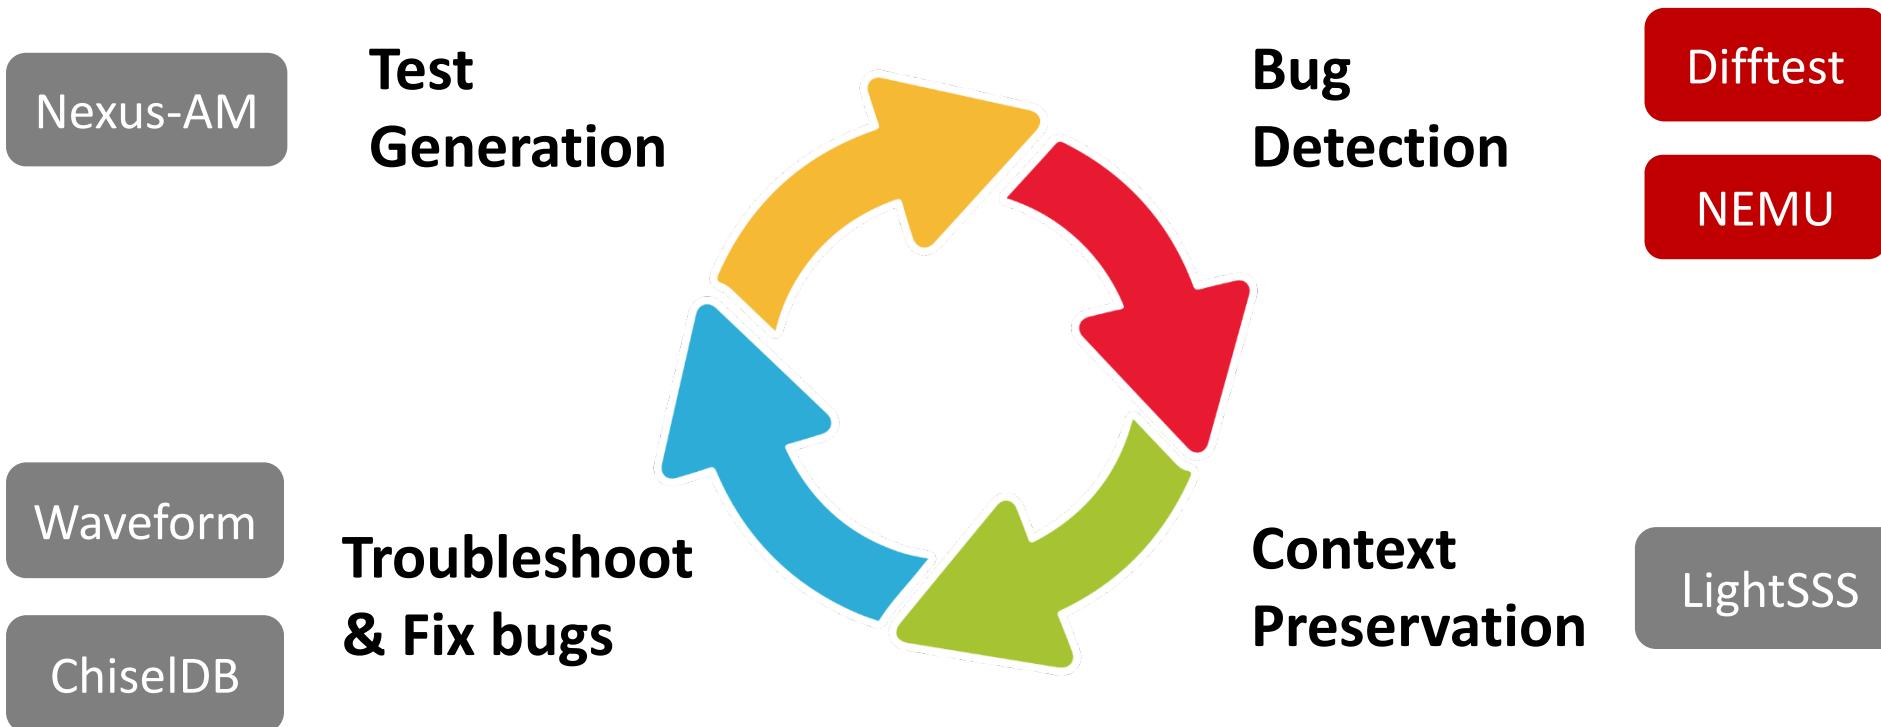
- **Hands-on:** build Coremark workload
- **Showcase**

```
$ bash build-am.sh
```

```
# cd $AM_HOME/apps/coremark
# make ARCH=riscv64-xs \
    ITERATIONS=1 LINUX_GNU_TOOLCHAIN=1 TOTAL_DATA_SIZE=400
# ls -l build
#   coremark-riscv64-xs.bin      Binary image of the program
#   coremark-riscv64-xs.elf      ELF of the program
#   coremark-riscv64-xs.txt     Disassembly of the program
```



MinJie Functional Verification





NEMU: ISA REF

- Purpose
 - Golden model for verification
 - Simple yet high performance
- We present **NEMU**
 - An **instruction set simulator**, similar to Spike
 - Through **optimization techniques**, achieve performance similar to QEMU.
 - A **set of APIs** is provided to assist XiangShan in comparing and verifying the **architectural states**

- Showcase

```
$ bash build-nemu.sh
```

```
# cd $NEMU_HOME  
# make clean  
# make riscv64-xs_defconfig  
# make -j      build NEMU as the bare metal machine, can run the Coremark from the previous step  
# make clean-softfloat  
# make riscv64-xs-ref_defconfig  
# make -j      build NEMU as the reference model for XiangShan
```

- **Hands-on:** run Coremark workload on NEMU
- **Showcase**

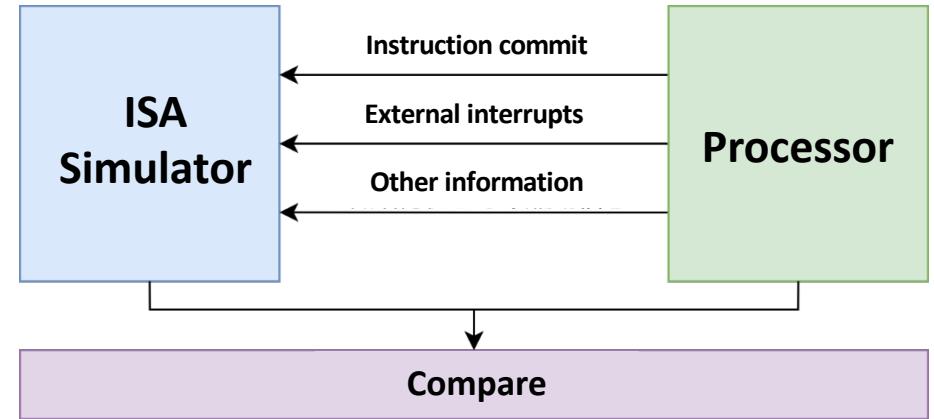
```
$ bash run-nemu.sh
```

```
# cd $NEMU_HOME
# ./build/riscv64-nemu-interpreter -b \run in batch mode, faster
$AM_HOME/apps/coremark/build/coremark-riscv64-xs.bin set workspace to Coremark
```



Difftest: ISA Co-simulation Framework

- **Basic flows**
 - Instructions commit/other states update
 - The simulator executes the same instructions
 - Compare the architectural states between DUT and REF
 - Abort or continue
- **Provided as verification infrastructure for processors**
 - APIs for HDLs such as Chisel/Verilog
 - RTL simulators such as Verilator, VCS
 - RISC-V ISS such as **NEMU**, **Spike**
- **SMP-Difftest: co-simulation on an SMP processor**
 - SMP Linux kernel and multi-threading programs
 - Online checking of cache coherency and memory consistency



Basic architecture of DiffTest

```
while (1) {  
    icnt = cpu_step();  
    nemu_step(icnt);  
    r1s = cpu_getregs();  
    r2s = nemu_getregs();  
    if (r1s != r2s) { abort();  
    }  
}
```

Online checking

- **Hands-on:** run Coremark workload on XiangShan, difftest with NEMU
- **Showcase**

```
# Running Coremark takes about 1 minutes
```

```
$ bash run-emu.sh
```

```
# ./emu \
-i $AM_HOME/apps/coremark/build/coremark-riscv64-xs.bin \ use coremark.bin
--diff $NEMU_HOME/build/riscv64-nemu-interpreter-so \ difftest with NEMU
2>perf.out      redirect stderr to file
```



Difftest

```
./emu -i $AM_HOME/apps/coremark/build/coremark-riscv64-xs.bin --diff $NEMU_HOME/build/riscv64-nemu-interpreter
-so 2>perf.err
emu compiled at Aug 19 2024, 22:33:45
[WARNING] /home/guest/linzhida/XiangShan/build/constantin.txt does not exist, so all constants default to initialized values.
Using simulated 32768B flash
Using simulated 8192MB RAM
The image is /home/guest/linzhida/nexus-am/apps/coremark/build/coremark-riscv64-xs.bin
The reference model is /home/guest/linzhida/NEMU/build/riscv64-nemu-interpreter-so
The first instruction of core 0 has commited. Difftest enabled.
Running CoreMark for 1 iterations
400 performance run parameters for coremark.
CoreMark Size      : 133
Total time (ms)   : 200
Iterations        : 1
Compiler version  : GCC11.4.0
seedcrc           : 0xe97b
[0]crclist        : 0x31be
[0]crcmatrix      : 0x0000
[0]crcstate       : 0x7bd3
[0]crcfinal       : 0x31be
Finished in 200 ms.
=====
CoreMark Iterations/Sec 5000
Core 0: HIT GOOD TRAP at pc = 0x80001eb4
Core-0 instrCnt = 39606, cycleCnt = 68161, IPC = 0.581065
Seed=0 Guest cycle spent: 68165 (this will be different from cycleCnt if emu loads a snapshot)
Host time spent: 35280ms
```

- **Hands-on:**

- We injected a bug in a pre-built XiangShan processor
- Now, let's trigger it by Difftest

- **Showcase**

```
$ bash run-emu-diff.sh
```

```
# ./emu-bug \
-i $AM_HOME/apps/coremark/build/coremark-riscv64-xs.bin \
--diff $NEMU_HOME/build/riscv64-nemu-interpreter-so    # difftest with NEMU
```

- Difftest will report the error once failed
 - Dump info like PC, GPRs, etc.

```
===== Commit Group Trace (Core 0) =====
commit group [00]: pc 0080000872 cmtcnt 9
commit group [01]: pc 0080000838 cmtcnt 9
commit group [02]: pc 00800015de cmtcnt 10
commit group [03]: pc 0080001614 cmtcnt 13
commit group [04]: pc 008000163e cmtcnt 19
commit group [05]: pc 0080001aaa cmtcnt 11 <-
commit group [06]: pc 00800028d8 cmtcnt 8
commit group [07]: pc 0080002a96 cmtcnt 8
commit group [08]: pc 0080002aa6 cmtcnt 10
commit group [09]: pc 0080000162 cmtcnt 8
commit group [10]: pc 0080000838 cmtcnt 10
commit group [11]: pc 00800015b0 cmtcnt 9
commit group [12]: pc 008000083e cmtcnt 9
commit group [13]: pc 008000082a cmtcnt 9
commit group [14]: pc 0080000860 cmtcnt 9
commit group [15]: pc 0080000830 cmtcnt 9
```

```
===== REF Regs =====
$0: 0x0000000000000000      ra: 0x00000000800015da      sp: 0x000000008000bf00      gp: 0x0000000000000000
tp: 0x0000000000000000      t0: 0x0000000000000000      t1: 0x0000000000000001      t2: 0x0000000000000009
s0: 0x0000000000000000      s1: 0x0000000000000000      a0: 0x0000000000000085      a1: 0x000000008000bf18
a2: 0x0000000000000002      a3: 0x0000000000000085      a4: 0x0000000000000002      a5: 0x0000000000000007
a6: 0x0000000000000001      a7: 0x0000000000000003      s2: 0x0000000000000000      s3: 0x0000000000000000
s4: 0x0000000000000000      s5: 0x0000000000000000      s6: 0x0000000000000000      s7: 0x0000000000000000
s8: 0x0000000000000000      s9: 0x0000000000000000      s10: 0x0000000000000000      s11: 0x0000000000000000
t3: 0x0000000080003b10     t4: 0x0000000000000001      t5: 0x000000008000bda1      t6: 0x0000000000000031
ft0: 0xfffffffff00000000      ft1: 0xfffffffff00000000      ft2: 0xfffffffff00000000      ft3: 0xfffffffff00000000
ft4: 0xfffffffff00000000      ft5: 0xfffffffff00000000      ft6: 0xfffffffff00000000      ft7: 0xfffffffff00000000
fs0: 0xfffffffff00000000      fs1: 0xfffffffff00000000      fa0: 0xfffffffff00000000      fa1: 0xfffffffff00000000
fa2: 0xfffffffff00000000      fa3: 0xfffffffff00000000      fa4: 0xfffffffff00000000      fa5: 0xfffffffff00000000
fa6: 0xfffffffff00000000      fa7: 0xfffffffff00000000      fs2: 0xfffffffff00000000      fs3: 0xfffffffff00000000
fs4: 0xfffffffff00000000      fs5: 0xfffffffff00000000      fs6: 0xfffffffff00000000      fs7: 0xfffffffff00000000
fs8: 0xfffffffff00000000      fs9: 0xfffffffff00000000      fs10: 0xfffffffff00000000      fs11: 0xfffffffff00000000
ft8: 0xfffffffff00000000      ft9: 0xfffffffff00000000      ft10: 0xfffffffff00000000      ft11: 0xfffffffff00000000
pc: 0x0000000080001aa2      mstatus: 0x8000000a00006000      mcause: 0x0000000000000000      mepc: 0x0000000000000000
                                         sstatus: 0x8000000200006000      scause: 0x0000000000000000      sepc: 0x0000000000000000
```

```
privilegeMode: 3
    a1 different at pc = 0x0080001aaa, right= 0x000000008000bf18, wrong = 0x000000008000bf20
    a3 different at pc = 0x0080001aaa, right= 0x0000000000000085, wrong = 0x0000000000000000
Core 0: ABORT at pc = 0x80000f24
Core-0 instrCnt = 1,296, cycleCnt = 11,810, IPC = 0.109738
Seed=0 Guest cycle spent: 11,813 (this will be different from cycleCnt if emu loads a snapshot)
Host time spent: 5,376ms
```



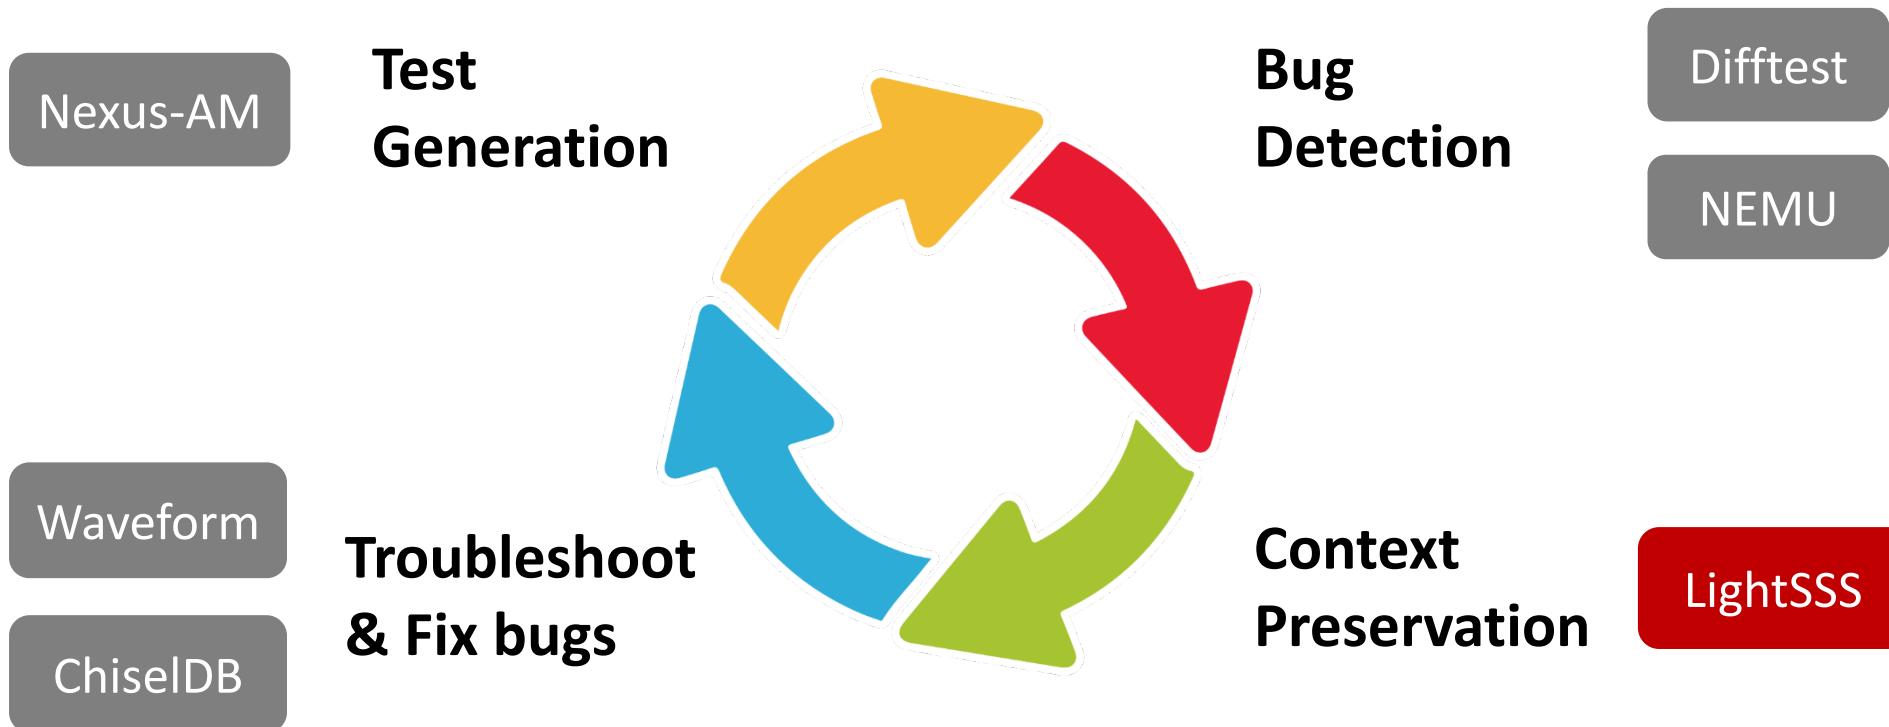
- **Hands-on:** Dump waveform after Difftest detects the bug
- **Showcase**

```
$ bash run-emu-dumpwave.sh  
$ ls $NOOP_HOME/build/*.vcd -lh # There should be a .vcd waveform
```

```
# ./emu-bug \  
-i $AM_HOME/apps/coremark/build/coremark-riscv64-xs.bin \  
--diff $NEMU_HOME/build/riscv64-nemu-interpreter-so \  
-b 11000 \      the begin cycle of the wave, you may find XiangShan trap at 12000 cycles in previous step  
-e 13000 \      the end cycle of the wave  
--dump-wave \   set this option to dump wave, you will find *.vcd on $NOOP_HOME/build
```



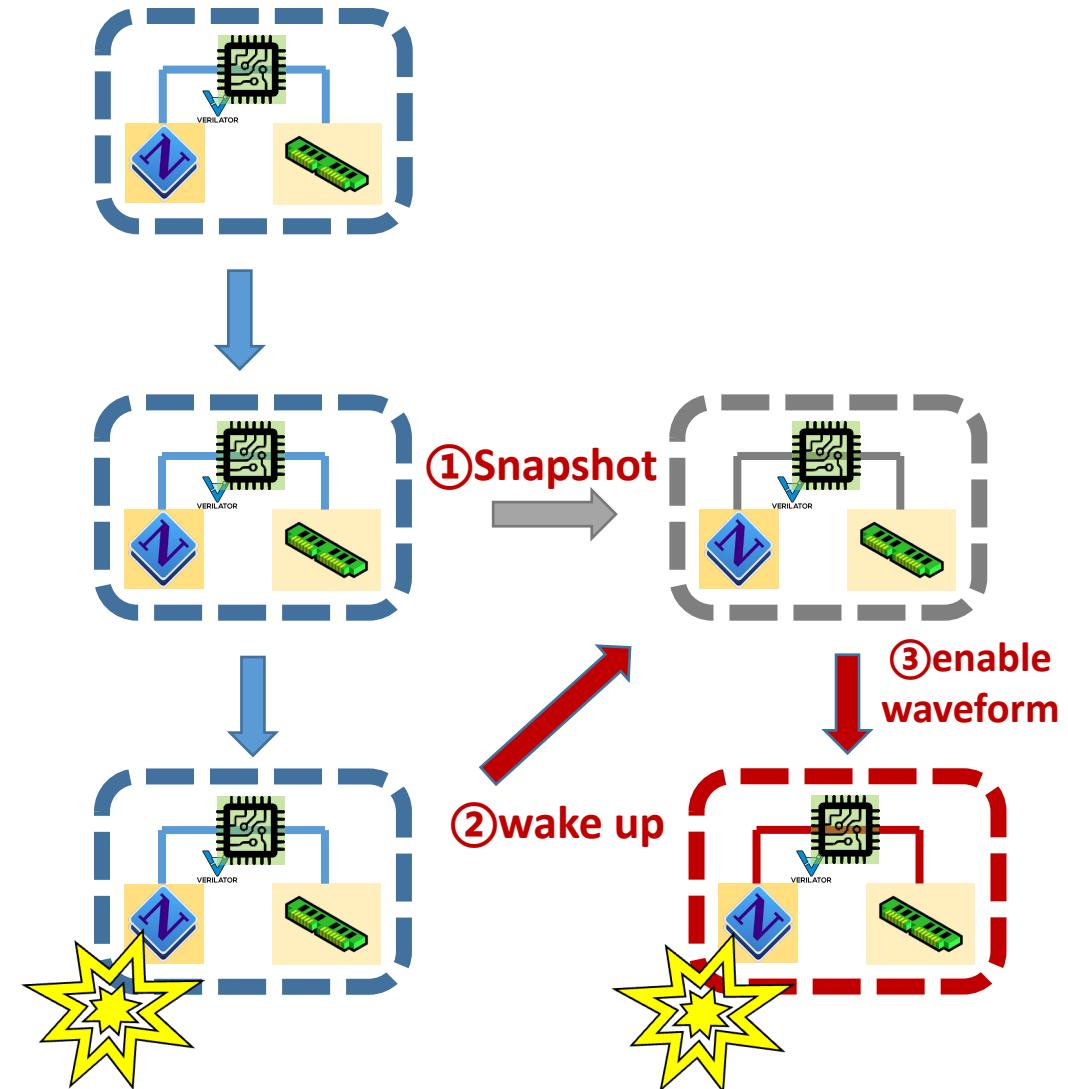
MinJie Functional Verification





LightSSS: Light-weight Simulation SnapShot

- Re-run simulation is time-consuming!
 - Snapshot is the way out
- LightSSS: snapshots of the process with fork()
 - Copy-on-write on Linux Kernel
- Advantage 1: Good portability and scalability
 - Snapshots for any external models (such as C++)
 - No need to understand details of external models
- Advantage 2: Low overhead of snapshots
 - ~500us for taking a snapshot
 - Far less overhead than RTL snapshots by Verilator





- **Hands-on:** use lightSSS to dump waveform
- **Showcase**

```
$ bash run-emu-lightsss.sh
```

```
# ./emu-bug \
-i $AM_HOME/apps/coremark/build/coremark-riscv64-xs.bin \
--diff $NEMU_HOME/build/riscv64-nemu-interpreter-so \
--enable-fork \      No longer need to use complex parameters
2 > lightsss.err
```



- LightSSS is working if you see:

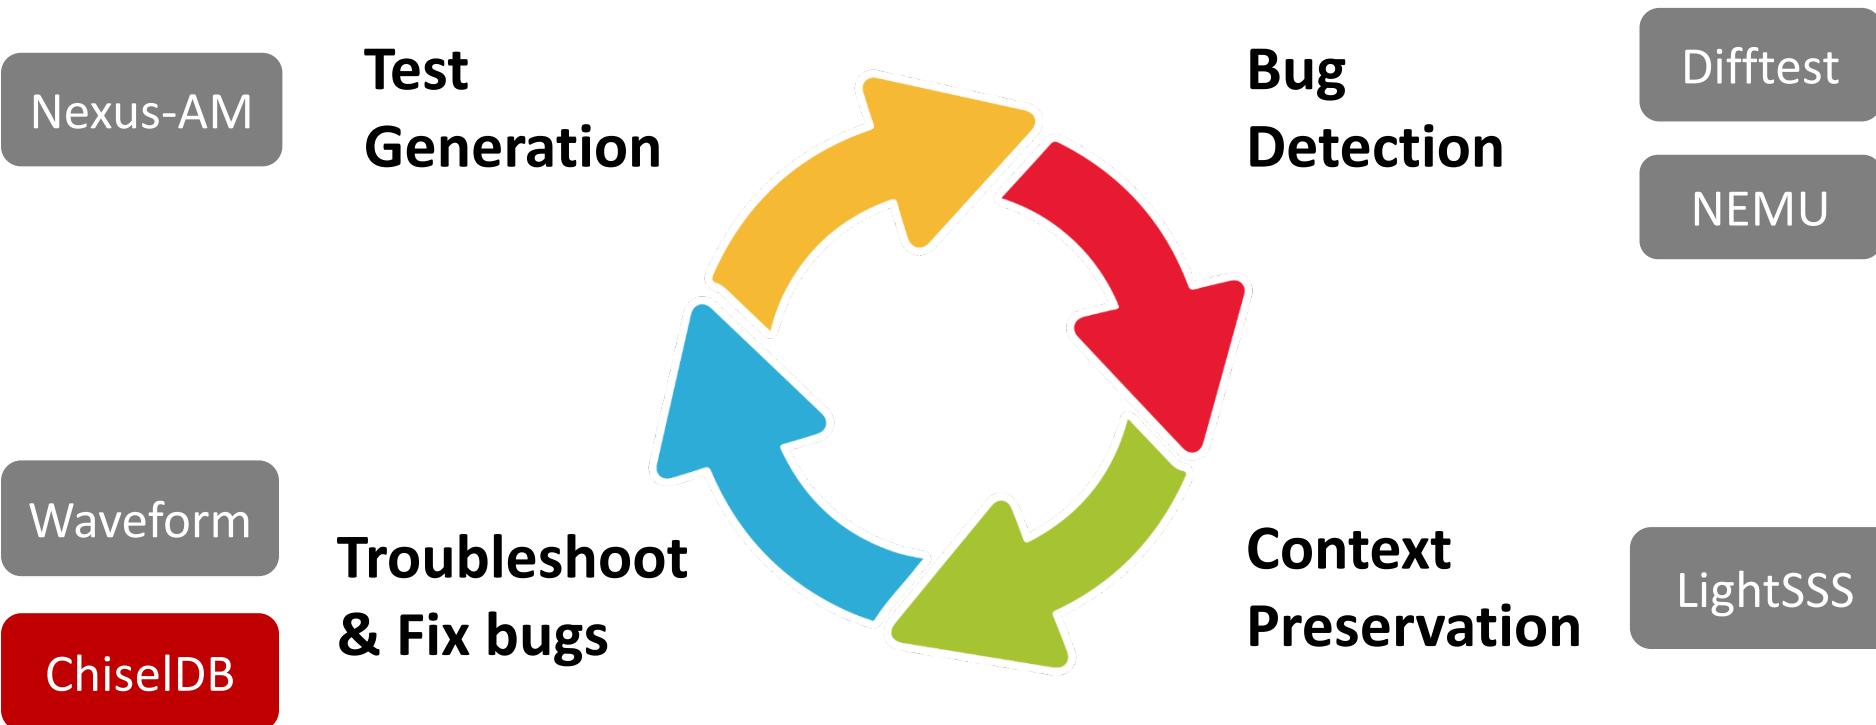
```
privilegeMode: 3
    a1 different at pc = 0x0080001aaa, right= 0x000000008000bf18, wrong = 0x000000008000bf20
    a3 different at pc = 0x0080001aaa, right= 0x00000000000000085, wrong = 0x000000000000000000000000
[FORK_INFO pid(29770)] the oldest checkpoint start to dump wave and dump nemu log...
[FORK_INFO pid(29770)] dump wave to /home/guest/liyanqin-tmp/XiangShan/build/2024-11-02@23:40:18_9208.vcd...
ing CoreMark for i iterations

===== Commit Group Trace (Core 0) =====
commit group [00]: pc 0080000872 cmtcnt 9
commit group [01]: pc 0080000838 cmtcnt 9
commit group [02]: pc 00800015de cmtcnt 10
commit group [03]: pc 0080001614 cmtcnt 13
commit group [04]: pc 008000163e cmtcnt 19
commit group [05]: pc 0080001aaa cmtcnt 11 <--
```

- After that, simulation will start again from the existing latest snapshot



MinJie Functional Verification





ChiselDB: Debug-friendly Structured Database

- **Motivation:**
 - Waveforms are **large** in size and **hard to apply** further analysis
 - Need to analyze structured data like memory transaction trace
- **We present ChiselDB**
- **Highlights:**
 - Inserting probes between module interfaces in hardware
 - DPI-C: Using C++ function in Chisel code to transfer data
 - Persist in database, SQL queries supported

- **Design source code:** *XiangShan/utility/src/main/scala/utility/ChiselDB.scala*
- **Usage:** Create table

```
// API: def createTable[T <: Record](tableName: String, hw: T): Table[T]
import huancun.utils.ChiselDB

class MyBundle extends Bundle {
    val fieldA = UInt(10.W)
    val fieldB = UInt(20.W)
}
val table = ChiselDB.createTable("MyTableName", new MyBundle)
```

- **Usage:** Register probes

```
/* APIs
def log(data: T, en: Bool, site: String = "", clock: Clock, reset: Reset)
def log(data: Valid[T], site: String, clock: Clock, reset: Reset): Unit
def log(data: DecoupledIO[T], site: String, clock: Clock, reset: Reset): Unit
*/
table.log(
    data = my_data /* hardware of type T */,
    en = my_cond,    site = "MyCallSite",
    clock = clock,   reset = reset
)
```

- Example: *XiangShan/src/main/scala/xiangshan/cache/mmu/L2TLB.scala*

```
class L2TlbMissQueueDB(implicit p: Parameters) extends TlbBundle {  
    val vpn = UInt(vpnLen.W)  
}  
  
val L2TlbMissQueueInDB, L2TlbMissQueueOutDB = Wire(new L2TlbMissQueueDB)  
L2TlbMissQueueInDB.vpn := missQueue.io.in.bits.vpn  
L2TlbMissQueueOutDB.vpn := missQueue.io.out.bits.vpn  
  
val L2TlbMissQueueTable = ChiselDB.createTable(  
    "L2TlbMissQueue_hart" + p(XSCoreParamsKey).HartId.toString, new L2TlbMissQueueDB)  
  
L2TlbMissQueueTable.log(L2TlbMissQueueInDB, missQueue.io.in.fire, "L2TlbMissQueueIn", clock, reset)  
L2TlbMissQueueTable.log(L2TlbMissQueueOutDB, missQueue.io.out.fire, "L2TlbMissQueueOut", clock, reset)
```

- **TL-Log : Persistence Transactions/Database**

- Record bus transactions to SQLite3 database using ChiselDB
- Support offline bug detect and performance analysis

Time	Name	Channel	Opcode	Permission	State	SourceID	SinkID	Address	Data			
116854134	L2_L1I_0	A	AcquireBlock	Grow	NtoB	1	0	803d3680	0000000000000000	0000000000000000	0000000000000000	0000000000000000
116854146	L2_L1I_0	D	GrantData	Cap	toT	1	16	803d3680	fa8437839043903	06090063843c23	84382300093783	378300093023c3bd
116854147	L2_L1I_0	D	GrantData	Cap	toT	1	16	803d3680	03e32e07bb030089	f0ef8526c881f49b	b60300893783bbdf	855a010925832e07
123191840	L2_L1I_0	C	ReleaseData	Shrink	TtON	0	0	803d3680	fa8437839043903	06090063843c23	84382300093783	378300093023c3bd
123191841	L2_L1I_0	B	Probe	Cap	toN	0	0	803d3680	0000000000000000	0000000000000000	0000000000000000	0000000000000000
123191841	L2_L1I_0	C	ReleaseData	Shrink	TtoN	0	0	803d3680	03e32e07bb030089	f0ef8526c881f49b	b60300893783bbdf	855a010925832e07
123191848	L2_L1I_0	D	ReleaseAck	Cap	toT	0	31	803d3680	0000000020ab05b	0000000020fab45b	0000000020fab85b	0000000020fabc5b
123191851	L3_L2_0	C	ReleaseData	Shrink	TtoN	31	0	803d3680	fa8437839043903	06090063843c23	84382300093783	378300093023c3bd
123191852	L3_L2_0	C	ReleaseData	Shrink	TtoN	31	0	803d3680	03e32e07bb030089	f0ef8526c881f49b	b60300893783bbdf	855a010925832e07
123191862	L3_L2_0	D	ReleaseAck	Cap	toT	31	16	803d3680	86a6f7043583dc5	f0ef856a865a8762	3583b7654481e17	865a86a68762f704
123191863	L2_L1I_0	C	ProbeAck	Shrink	TtoN	0	0	803d3680	0000000000000000	0000000000000000	0000000000000000	0000000000000000
123191878	L3_L2_0	C	Release	Report	NtoN	30	0	803d3680	5597bf494981aa09	8522cc2585930000	f84ef15dda0f00ef	001947036775e84a

Example: requests sequence in database that violate cache coherence

- Hands-on: Analysis of Cache Coherence Violation

```
# Inject bug – We set all Release Data (L2->L3) as a constant value
```

```
$ cd ../../p2-chiseldb && cat cdb_err.patch
```

```
--- a/src/main/scala/huancun/noninclusive/SinkC.scala
+++ b/src/main/scala/huancun/noninclusive/SinkC.scala

@@ -99,7 +99,7 @@ class SinkC(implicit p: Parameters) extends
BaseSinkC {
-    buffer(insertIdx)(count) := c.bits.data
+    buffer(insertIdx)(count) := 0xABCD.EF.U
```

- Hands-on: Analysis of Cache Coherence Violation

```
# simulate using pre-built emu (ChiselDB enabled)
$ bash chiseldb_step1-run.sh
```

```
# ./emu-cdb-err -i $NOOP_HOME/ready-to-run/Linux.bin \
--diff $NOOP_HOME/ready-to-run/riscv64-nemu-interpreter-so \
--dump-db    # set this argument to dump data base
2>linux.err
```



- Hands-on: Analysis of Cache Coherence Violation
 - Difftest Info
 - Error Address : 0x800419c0

- **Hands-on: Analysis of Cache Coherence Violation**

```
# Analyze  
$ bash chiseldb_step2-analyze.sh
```

```
# 1. sqlite query for all transactions on Eaddr  
# 2. use script to parse TLLog  
# sqlite3 $NOOP_HOME/build/*.db \  
"select * from TLLOG where ADDRESS=0x800419c0" | \  
sh $NOOP_HOME/scripts/utils/parseTLLog.sh
```



ChiselDB:

Result: [Time | To_From | Channel | Opcode | Permission | Address | Data]

Data successfully transferred from L1D to L2

28189 L2_L1D_0 C ReleaseData Shrink TtoN 800419c0

0000000000000000 0000000000000000 000000080041a70 000000080016198

Data successfully transferred from L2 to L3

39637 L3_L2_0 C ProbeAckData Shrink TtoN 800419c0

0000000000000000 0000000000000000 000000080041a70 000000080016198

But when L1D acquires Eaddr again, data loaded from L3 is wrong

47371 L2_L1D_0 A AcquireBlock Grow NtoB 800419c0

47377 L3_L2_0 A AcquireBlock Grow NtoB 0 1 800419c0

47413 L3_L2_0 D GrantData Cap toT 800419c0

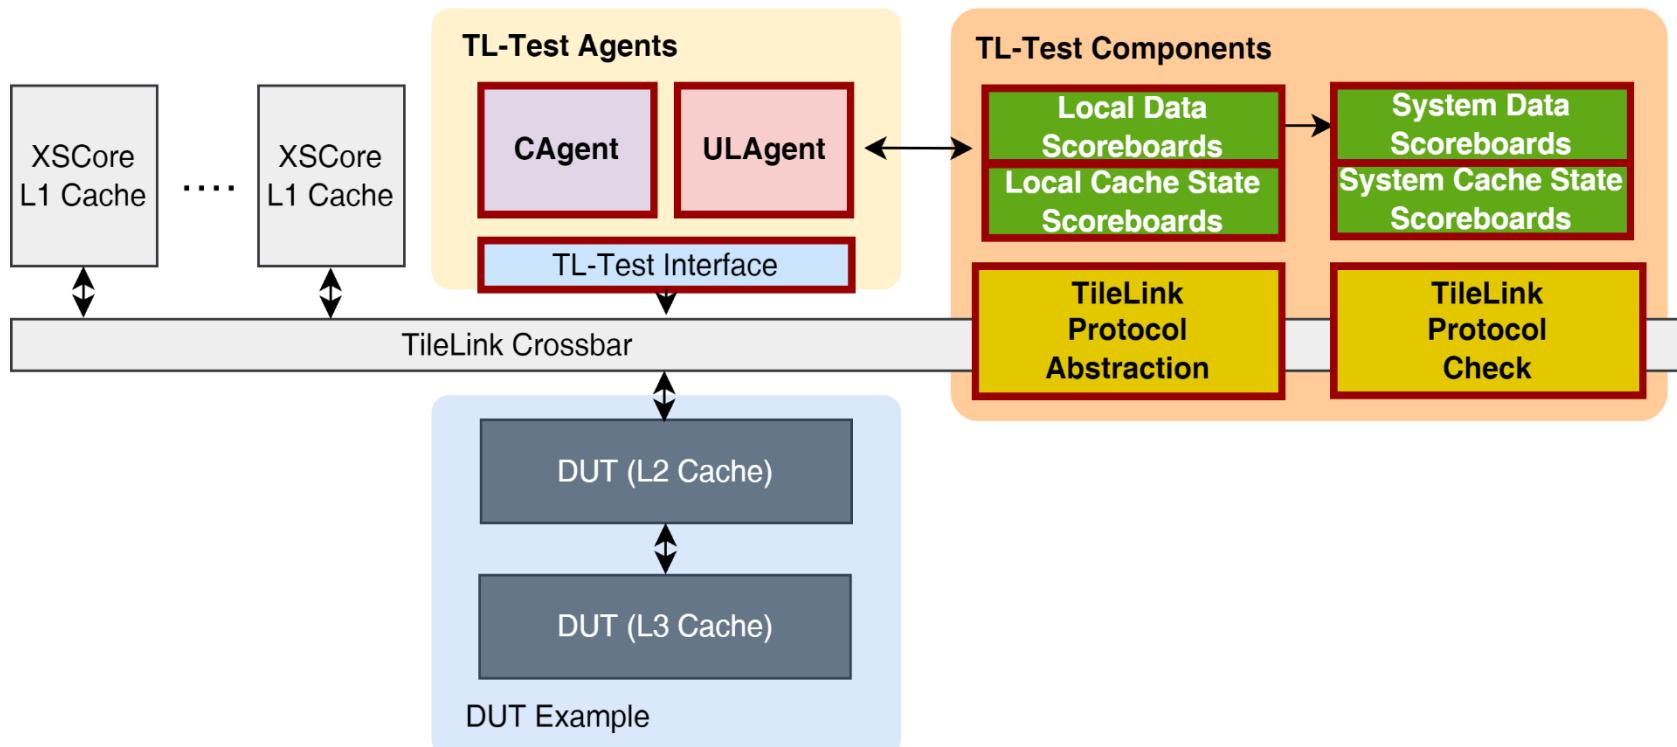
0000000000abcdef 0000000000000000 0000000000000000 0000000000000000

So there must be something wrong when L3 records Release Data

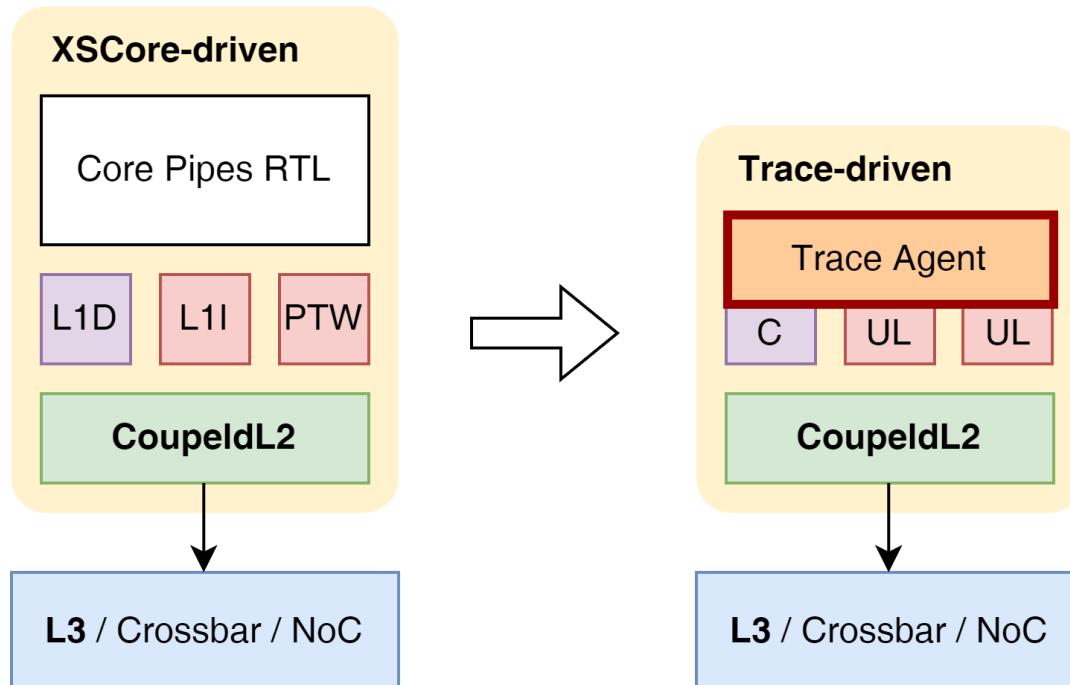


TL-Test: Cache Verification Framework

- As an infrastructure for cache verification
 - Support **Tilelink protocol** and **cache coherence check**
 - **Fast test cases generation** and **constrained random test**



- **TL-Trace: Trace-based testbench generation**
 - Convert cache access trace into testbench
 - Fast reproduction of functional bugs and analysis of performance problems



Project Structure

```
$ cd $TLT_HOME && tree -d -L 1
```

```
.
├── assets
├── configs
├── dut
├── main
├── run
├── scripts
└── tutorial
```

Enter TL-Test tutorial

```
$ cd ../../tutorial/p3-tl-test
```

- Hands-on: Analysis of Cache Coherence Violation

```
# Inject bug – We wrongly shift the Release Data (L1->L2)
```

```
$ cat tlt_err.patch
```

```
--- a/tl-test-new/dut/CoupledL2/src/main/scala/coupledL2/GrantBuffer.scala
+++ b/tl-test-new/dut/CoupledL2/src/main/scala/coupledL2/GrantBuffer.scala
@@ -87,7 +87,7 @@ class GrantBuffer(implicit p: Parameters) extends
LazyModule {
-
d.data := data
+
d.data := data << 8
```



- **Hands-on: Analysis of Cache Coherence Violation**

```
# We provide a pre-compiled TL-Test  
  
# Run the TL-Test  
  
$ bash tltest_step1_run.sh
```

```
# cd $TLT_HOME/run && ./tltest_v3lt 2>&1 | tee tltest_v3lt.log
```



- **Hands-on: Analysis of Cache Coherence Violation**
- **TL-Test Info**
 - Error Address : 0x16000

```
[1740] [tl-test-new-INFO] #0 L2[0].C[0] [data complete D] [GrantData toT] source: 0x1, addr: 0x16000, alias: 0, data: [ 00 d6 ... 00 75 ... ]  
dut: [ 00 d6 ... 00 75 ... ]  
ref: [ d6 26 ... 75 0d ... ]  
[1740] [tl-test-new-ERROR] [tlc_assert failure at int GlobalBoard<unsigned long>::data_check(TLLocalContext *, const uint8_t *, const uint8_t *, std::string) [T = unsigned long]:263]  
[1740] [tl-test-new-ERROR] [tlc_assert failure from system #0]  
[1740] [tl-test-new-ERROR] info: Data mismatch from status SB_VALID!
```



- **Hands-on: Analysis of Cache Coherence Violation**

```
# Analyze  
$ bash tltest_step2_analyze.sh
```

```
# 1. grep all transactions on Eaddr(0x16000)  
# 2. use TL-Log to help debugging  
  
# grep “addr: 0x16000” $TLT_HOME/run/tltest_v3lt.log
```

```
# Result: [Time | Core | Channel | Opcode | Source | Address | Data]

# Data successfully transferred from L1D to L2
[800] L2[0].C[0] [fire C] [ReleaseData TtoN] source: 0xa, addr: 0x16000, data: [ d6 26 ... ]
[802] L2[0].C[0] [fire C] [ReleaseData TtoN] source: 0xa, addr: 0x16000, data: [ 75 0d ... ]

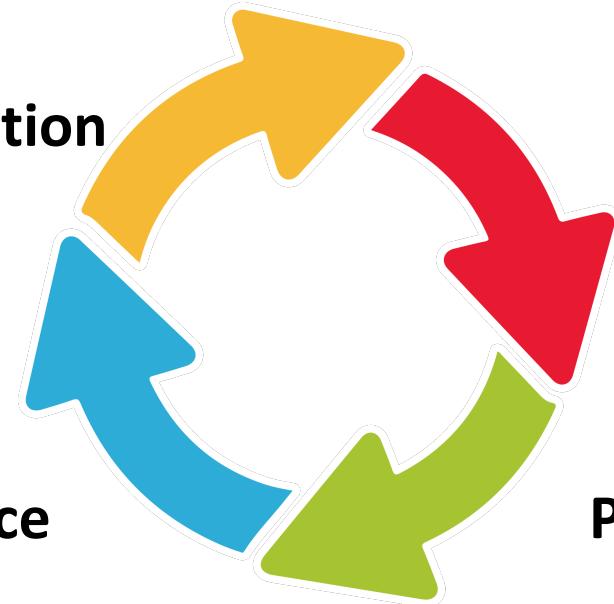
# But when L1D acquires Eaddr again, data loaded from L2 is wrong
[844]  L2[0].C[0] [fire A] [AcquireBlock NtoT] source: 0x1, addr: 0x16000
[1738] L2[0].C[0] [fire D] [GrantData toT] source: 0x1, addr: 0x16000, data: [ 00 d6 ... ]
[1740] L2[0].C[0] [fire D] [GrantData toT] source: 0x1, addr: 0x16000, data: [ 00 75 ... ]

# So there must be something wrong when L2 records Release Data
```

MinJie Performance Verification

Chisel-based
prototypes

Implementation



XSPerf

Constantin

Top-Down

**Performance
Analysis**

Run Tests

**Performance
Evaluation**



umd-memsys/
DRAMsim3

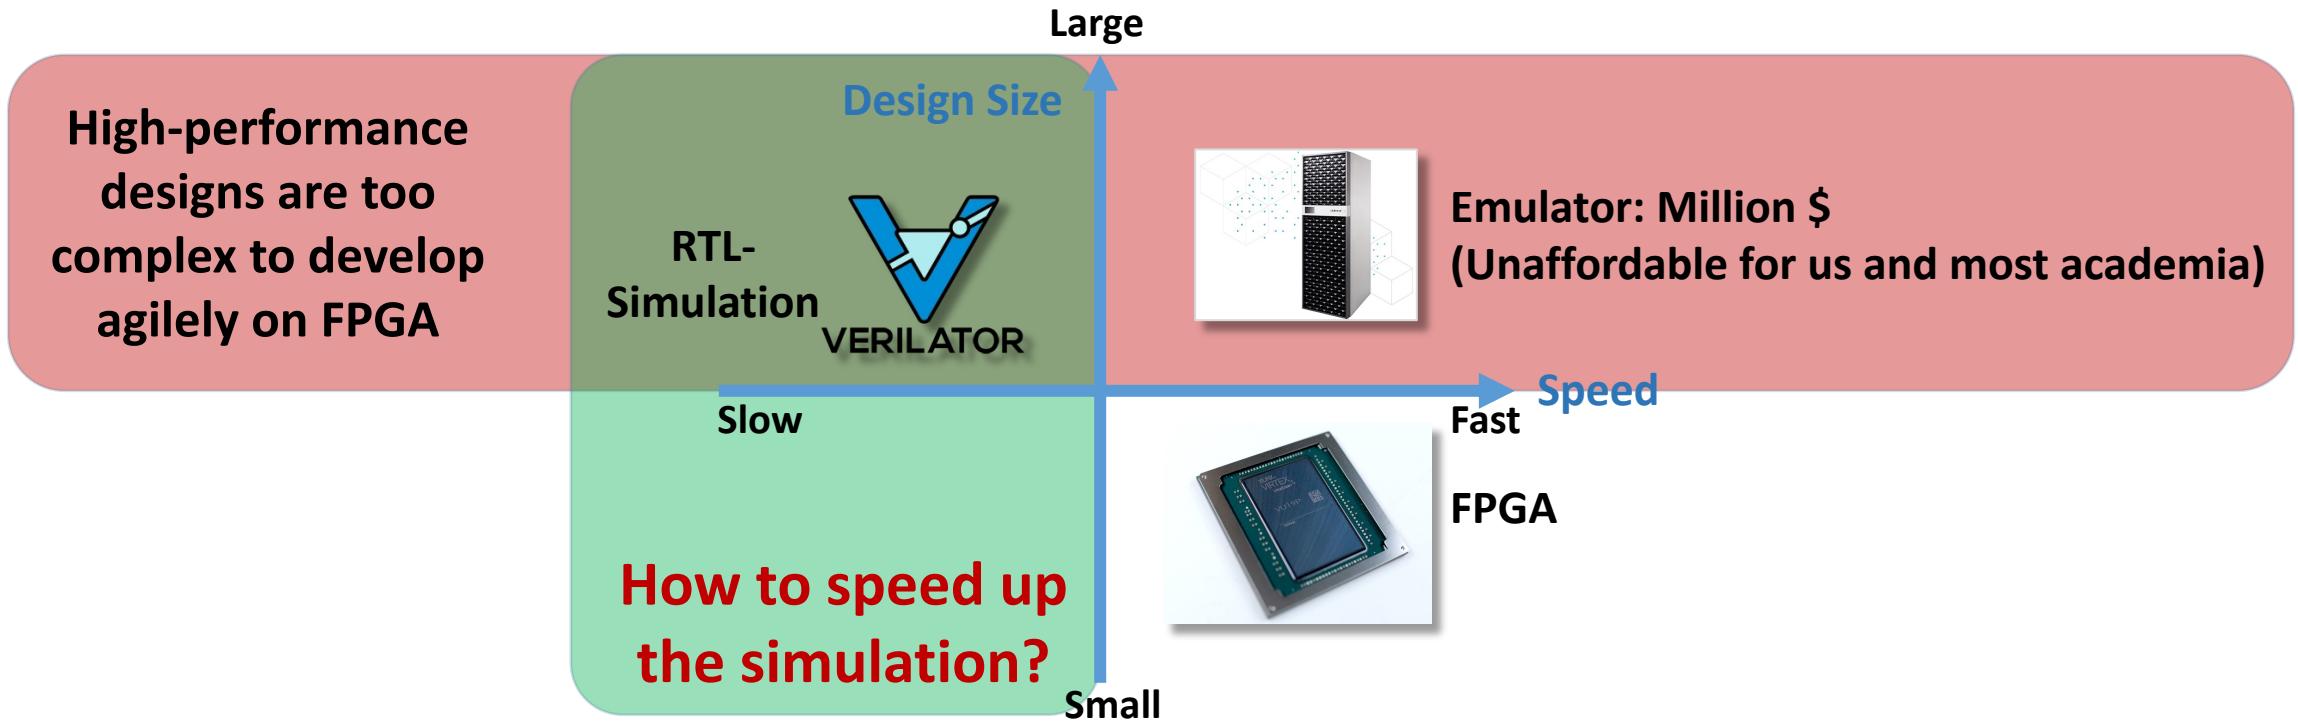
DRAMsim3: a Cycle-accurate, Thermal-Capable
DRAM Simulator

RISC-V
Checkpoint





Checkpoint: Agile Performance Evaluation



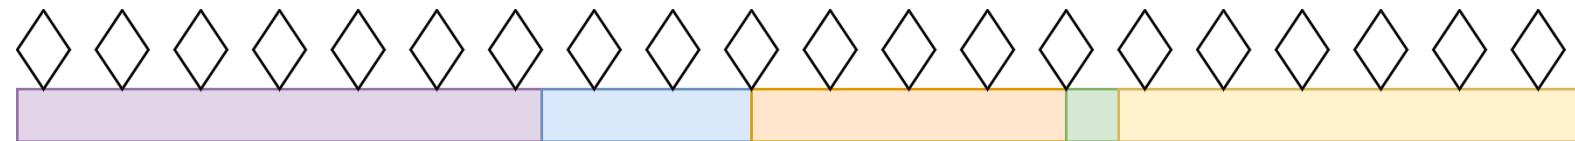
Time for performance modelling of single-core XiangShan on SPEC CPU2006

	RTL-simulation	FPGA	RTL-Simulation w/ Checkpoint
Compile	20 minutes	5 hours	20 minutes
Simulation	958 years@2KHz (2K CPS)	7 days@100MHz	5.5 hours with enough x86 servers Our Approach



Checkpoint: Agile Performance Evaluation

- Uniform Checkpoint
 - Divide a program to equal segments for parallel simulation with uniform sampling



- SimPoint Checkpoint
 - Select representative segments with cluster analysis and weighted sampling





Checkpoint: SimPoint Checkpoint

- Step 0: Prepare NEMU environment for SimPoint checkpoint

```
$ cd ../../p4-checkpoint && bash simpoint_step0-prepare.sh
```

```
# simpoint_step0_prepare.sh (~23s)

# cd $NEMU_HOME
# git submodule update --init
# cd $NEMU_HOME/resource/simpoint/simpoint_repo
# make clean && make           # generate simpoint generator binary

# cd $NEMU_HOME
# make clean
# make riscv64-xs-cpt_defconfig && make -j 8      # compile NEMU

# cd $NEMU_HOME/resource/gcpt_restore
# rm -rf $XS_PROJECT_ROOT/tutorial/part5-checkpoint/gcpt
# make -C $NEMU_HOME/resource/gcpt_restore/ \      # generate gcpt restorer binary
#       O=$XS_PROJECT_ROOT/tutorial/part5-checkpoint/gcpt \ # directory of results
#       GCPT_PAYLOAD_PATH=$XS_PROJECT_ROOT/tutorial/part5-checkpoint
#       /bin/stream_100000.bin
```



Checkpoint: SimPoint Checkpoint

- Step 1: Execute the workload and collect program behavior

```
$ bash simpoint_step1-profiling.sh
```

```
# simpoint_step1_profiling.sh (~18s)
# source simpoint_env.sh                                # configure environment variables

# rm -rf $RESULT

# $NEMU ${BBL_PATH}/${workLoad}.bin \
#       -b                                         \      # specify workload
#       -D $RESULT                                \      # run with batch mode
#       -C $profiling_result_name                 \      # directory where the checkpoint is generated
#       -w stream                                 \      # name of this task
#       --simpoint-profile                         \      # name of workload
#       --cpt-interval ${interval}                \      # is simpoint profiling
#       > >(tee $log/stream-out.txt) 2> >(tee ${Log}/stream-err.txt) # simpoint interval instructions: 50000000
#       >> (tee $log/stream-err.txt) # redirect stdout and stderr
```



Checkpoint: SimPoint Checkpoint

Your clock granularity/precision appears to be 26 microseconds.
Each test below will take on the order of 1023153 microseconds.
(= 39352 clock ticks)

Increase the size of the arrays if this shows that
you are not getting at least 20 clock ticks per test.

WARNING -- The above is only a rough guideline.

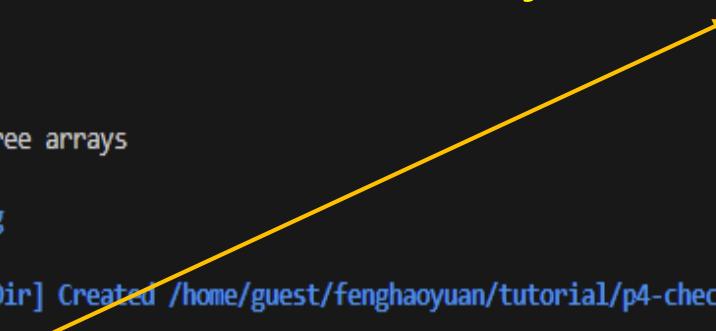
For best results, please be sure you know the
precision of your system timer.

Function	Best Rate MB/s	Avg time	Min time	Max time
Copy:	941.2	0.085167	0.084997	0.085384
Scale:	215.4	0.371746	0.371488	0.371877
Add:	275.9	0.435193	0.434995	0.435358
Triad:	256.4	0.468193	0.467980	0.468407

Solution Validates: avg error less than 1.000000e-06 on all three arrays

```
[src/monitor/monitor.c:195,parse_args] Doing Simpoint Profiling
[src/checkpoint/path_manager.cpp:54,init] Cpt id: -1
[src/checkpoint/path_manager.cpp:72,setSimpointProfilingOutputDir] Created /home/guest/fenghaoyuan/tutorial/p4-checkpoint/simpoint_result/simpoint-profiling/stream/
[src/checkpoint/simpoint.cpp:83,init] Doing simpoint profiling with interval 50000000
[src/checkpoint/serializer.cpp:378,next_index] set next index 0
[src/checkpoint/path_manager.cpp:91,setCheckpointingOutputDir] do not set checkpoint path without Checkpoint mode
[src/isa/riscv64/init.c:191,init_isa] NEMU will start from pc 0x80000000
[src/monitor/image_loader.c:203,load_img] Loading image (checkpoint/bare metal app/bbl) from cmdline: /home/guest/fenghaoyuan/tutorial/p4-checkpoint/gcpt/build/gcpt.l
```

*NEMU will do profiling while
executing workload STREAM at
intervals of 50,000,000 instructions.*





Checkpoint: SimPoint Checkpoint

- Step 2: Cluster and obtain multiple representative slices

```
$ bash simpoint_step2_cluster.sh
```

```
# simpoint_step2_cluster.sh

# export CLUSTER=$RESULT/cluster/${workLoad} && mkdir -p $CLUSTER
# mkdir -p $LOG_PATH/cluster_Logs/cluster

# random1=`head -20 /dev/urandom | cksum | cut -c 1-6`
# random2=`head -20 /dev/urandom | cksum | cut -c 1-6`

# $NEMU_HOME/resource/simpoint/simpoint_repo/bin/simpoint \
#   -LoadFVFile $PROFILING_RES/${workload}/simpoint_bbv.gz \
#   -saveSimpoints $CLUSTER/simpoints0 \
#   -saveSimpointWeights $CLUSTER/weights0 \
#   -inputVectorsGzipped \
#   -maxK 3 \
#   -numInitSeeds 2 \
#   -iters 1000 \
#   -seedkm ${random1} \
#   -seedproj ${random2} \
#   >>($tee $Log/${workLoad}-out.txt) 2>>($tee $Log/${workLoad}-err.txt)
# redirect stdout and stderr
```



Checkpoint: SimPoint Checkpoint

```
Run number 2 of at most 4, k = 3
-----
Initialization seed trial #1 of 2; initialization seed = 634853
-----
Initialized k-means centers using random sampling: 3 centers
Number of k-means iterations performed: 1
BIC score: -27.9287
Distortion: 7.02268
Distortions/cluster: 0.542115 3.14007 3.34049
Variance: 0.702268
Variances/cluster: 0.271058 1.04669 0.668099
```

Compute the information of K-means clustering

```
Initialization seed trial #2 of 2; initialization seed = 634854
-----
Initialized k-means centers using random sampling: 3 centers
Number of k-means iterations performed: 4
BIC score: -10.3726
Distortion: 6.04602
Distortions/cluster: 5.29257 0.753145 0.000309052
Variance: 0.604602
Variances/cluster: 0.661571 0.753145 0.000309052
The best initialization seed trial was #2
```

Obtain multiple program representative checkpoints

```
Post-processing runs
```

```
For the BIC threshold, the best clustering was run 2 (k = 3)
```



Checkpoint: SimPoint Checkpoint

- Step 3: Generate corresponding Checkpoints based on clustering results

```
$ bash simpoint_step3_genspt.sh # It may take about 2 minutes
```

```
# simpoint_step3_genspt.sh

# export CLUSTER=$RESULT/cluster
# mkdir -p $LOG_PATH/checkpoint_Logs

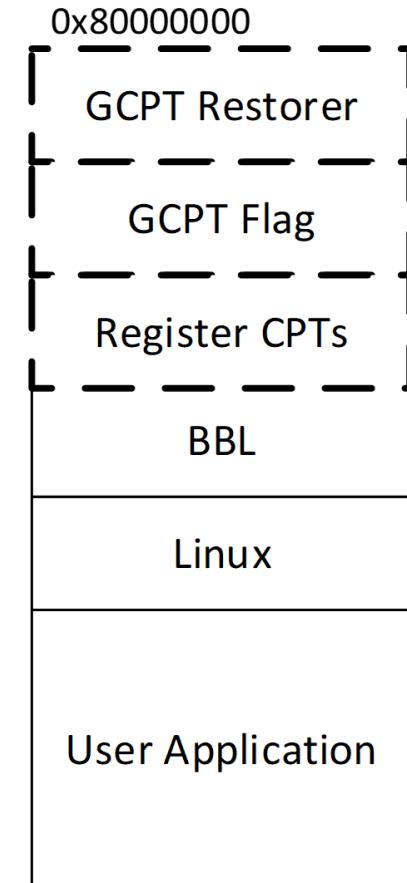
# $NEMU ${BBL_PATH}/${workLoad}.bin \
#   -b                                     \ run with batch mode
#   -D $RESULT                                \ directory where the checkpoint is generated
#   -C checkpoint                            \ name of this task
#   -w stream                                 \ name of workload
#   -S $CLUSTER                               \ simpoints results director
#   --cpt-interval $interval                 \ simpoint interval instructions: 50,000,000
#   >>(tee $Log/stream-out.txt) 2>>(tee $Log/stream-err.txt) redirect stdout and stderr
```



Checkpoint: SimPoint Checkpoint

```
[src/isa/riscv64/exec/special.c,38,exec_nemu_trap] Start profiling, resetting inst count → Start making checkpoint
[src/checkpoint/serializer.cpp,100,serializeRegs] Writing int registers to checkpoint memory @[0x80001000, 0x80001100) [0x1000, 0x1100)
[src/checkpoint/serializer.cpp,110,serializeRegs] Writing float registers to checkpoint memory @[0x80001100, 0x80001200) [0x1100, 0x1200)
[src/checkpoint/serializer.cpp,118,serializeRegs] Writing PC: 0xffffffffe000720f20 at addr 0x80001200
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x105: 0xffffffffe000697cc8
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x106: 0xfffffffffffffff
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x141: 0x200013417e
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x142: 0x8
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x180: 0x80000000000f805d
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x300: 0xa00000022
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x301: 0x800000000014112d
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x302: 0xb109
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x303: 0x222
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x304: 0xzzz → Save int & float registers
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x305: 0x800a0004
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x306: 0xfffffffffffffff
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x340: 0x800abdc0
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x341: 0xffffffffe0008d6eb2
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x342: 0x2
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x3a0: 0x1f0800
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x3b0: 0x20028000
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x3b1: 0x2002d000
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x3b2: 0xfffffffffffffff → Save CSR registers
[src/checkpoint/serializer.cpp,144,serializeRegs] Writing CSR to checkpoint memory @[0x80001300, 0x80009300) [0x1300, 0x9300)
[src/checkpoint/serializer.cpp,152,serializeRegs] Touching Flag: 0xbeef at addr 0x80000f00
[src/checkpoint/serializer.cpp,156,serializeRegs] Record mode flag: 0x1 at addr 0x80000f08
[src/checkpoint/serializer.cpp,160,serializeRegs] Record time: 0x1 at addr 0x80000f10
[src/checkpoint/serializer.cpp,164,serializeRegs] Record time: 0x1 at addr 0x80000f18
[src/checkpoint/serializer.cpp,52,serializePMem] Created tutorial_simpoint/take_cpt/stream_0_0.083333/0/ → Save other information

Opening tutorial_simpoint/take_cpt/stream_0_0.083333/0/_0_.gz as checkpoint output file
[src/checkpoint/serializer.cpp,83,serializePMem] Written 0xffffffff bytes → write checkpoints into .gz
[src/checkpoint/serializer.cpp,83,serializePMem] Written 0xffffffff bytes
[src/checkpoint/serializer.cpp,83,serializePMem] Written 0xffffffff bytes
[src/checkpoint/serializer.cpp,83,serializePMem] Written 0xffffffff bytes
[src/checkpoint/serializer.cpp,83,serializePMem] Written 0x4 bytes
[src/checkpoint/serializer.cpp,89,serializePMem] Checkpoint done!
```



Basic format of
RISC-V Checkpoint



Checkpoint: SimPoint Checkpoint

- Step 4: Run the SimPoint checkpoint on NEMU or XiangShan
- Here we take XiangShan as an example

```
$ bash simpoint_step4_run_xs.sh
```

```
# simpoint_step4_run_xs.sh (~ 1 min)

# ./emu \
#   -i `find $RESULT/checkpoint/stream -type f -name "*_.gz" | tail -1` \
#         get the path of workload
#   --diff $NOOP_HOME/ready-to-run/riscv64-nemu-interpreter-so \
#         Enable the reference standard design path for difftest
#   --max-cycles=50000 \
#         Maximum execution instruction count
#   2>simpoint.err
```



Checkpoint: SimPoint Checkpoint

- Step 5: Provide configuration files for batch running on Gem5/XiangShan

```
$ python3 simpoint_step5_dumpresult.py  
$ ls -l simpoint_result/checkpoint
```

- Collect profiling/checkpoint information to generate a list file for Gem5 execution

```
#checkpoint.lst  
stream_5 stream/5 0 0 20 20  
stream_11 stream/11 0 0 20 20  
stream_1 stream/1 0 0 20 20  
Workload Name | Checkpoint Path | Instructions to Skip (default 0) | Functional Warmup Instructions (default 0) | Detailed Warmup Instructions (default 20M) | Sample Instructions (default 20M)
```

- Collect cluster information to generate a JSON file for XiangShan execution

```
# cluster.json  
{"stream": {"insts": "654228721",  
"points": {"7": "0.615385", "0": "0.153846", "11": "0.230769"}}}  
Workload Name | Instruction Count | Checkpoint Index | Checkpoint Weight
```



Checkpoints: Uniform Checkpoint

- **Step 0: Prepare NEMU environment for checkpoints**

```
$ cd $XS_PROJECT_ROOT/tutorial/p4-checkpoint  
$ bash simpoint_step0_prepare.sh
```

- **Step 1: Generate uniform checkpoints using NEMU**

```
$ bash uniform_cpt.sh
```

- **Step 2: Run uniform checkpoint on NEMU or XiangShan**

- Taking NEMU as an example:

```
$ bash uniform_run_nemu.sh
```



XSim: Simulation Performance Collection

- Purpose
 - Collect simulation performance data under different requirements
- XSim
 - Multiple types for various fine-grained performance analysis
 - **Accumulation**: basic accumulator counter (log with stderr)
 - **Histogram**: count the distribution (log with stderr)
 - **Rolling**: rolling curve collection and visualization (ChiselDB)



XPerf: Accumulate & Histogram

- Example

Accumulation:

```
[PERF][time=10000] ctrlBlock.rob: commitInstr,      1500  
[PERF][time=10000] ctrlBlock.rob: waitLoadCycle,     800
```

Histogram:

```
[PERF][time=10000] l2cache: acquire_period_10_20,    15  
[PERF][time=10000] l2cache: acquire_period_20_30,    200  
[PERF][time=10000] l2cache: acquire_period_30_40,    60
```



XSPerf: Accumulate & Histogram

- **Usage**

- Add `XSPerfAccumulate('name', signal)` or `XSPerfHistogram('name', signal)`
- **By default**, it is printed after simulation finishes
- Set `DebugOptions(EnablePerfDebug = false)` to turn off

- **Showcase**

```
$ cd ../../p5-xs-perf && bash xs-perf-log.sh | head -n 20
```

```
# print the log result of p1-basic-func
# cat ${XS_PROJECT_ROOT}/tutorial/p1-basic-func/perf.err
```



XPerf: Accumulate & Histogram

```
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.ctrlBlock.decode: utilization,           1183
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.bpu.predictors.components_3.tables_0.wrbypass: wrbypass_hit,          0
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.icache.missUnit.entries_1: entryPenalty1,             714
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.ctrlBlock.decode: waitInstr,                414
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.bpu.predictors.components_3.tables_0.wrbypass: wrbypass_miss,          0
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.ctrlBlock.decode: stall_cycle,               393
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.icache.missUnit.entries_1: entryReq1,                 14
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.bpu.predictors.components_3.tables_1.wrbypass: wrbypass_hit,          0
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend: FrontendBubble,                  4085
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.icache.missUnit.entries_0: entryPenalty0,             248
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.bpu.predictors.components_3.tables_1.wrbypass: wrbypass_miss,          2
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: utilization,                1826
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.icache.missUnit.entries_0: entryReq0,                 6
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: util_0_1,                  2029
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: util_1_2,                  541
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.misc.busPMU: dcache_bank_0_A_channel_fire,                   8
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: util_2_3,                  45
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.misc.busPMU: dcache_bank_0_A_channel_stall,                 0
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: util_3_4,                  107
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.exuBlocks.scheduler.rs_4.staRS_0.oldestSelection: oldest_override_0,          0
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.misc.busPMU: dcache_bank_0_A_channel_PutFullData_fire,                 0
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: util_4_5,                  44
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.exuBlocks.scheduler.rs_4.staRS_0.oldestSelection: oldest_same_as_selected_0,      0
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.misc.busPMU: dcache_bank_0_A_channel_PutFullData_stall,                 0
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: util_5_6,                  85
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.exuBlocks.scheduler.rs_4.staRS_0.oldestSelection: oldest_override_1,          0
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.misc.busPMU: dcache_bank_0_A_channel_PutPartialData_fire,                 0
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: util_6_7,                  13
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.exuBlocks.scheduler.rs_4.staRS_0.oldestSelection: oldest_same_as_selected_1,      0
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.misc.busPMU: dcache_bank_0_A_channel_PutPartialData_stall,                 0
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: util_7_8,                  27
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.memBlock.dtlb_st_tlb_st.entries: port0_np_sp_multi_hit,                 0
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.misc.busPMU: dcache_bank_0_A_channel_ArithmeticData_fire,                 0
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: full,                     862
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.memBlock.dtlb_st_tlb_st.entries: port1_np_sp_multi_hit,                 0
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.misc.busPMU: dcache_bank_0_A_channel_ArithmeticData_stall,                 0
[PERF ][time=          2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: exHalf,                    125
```

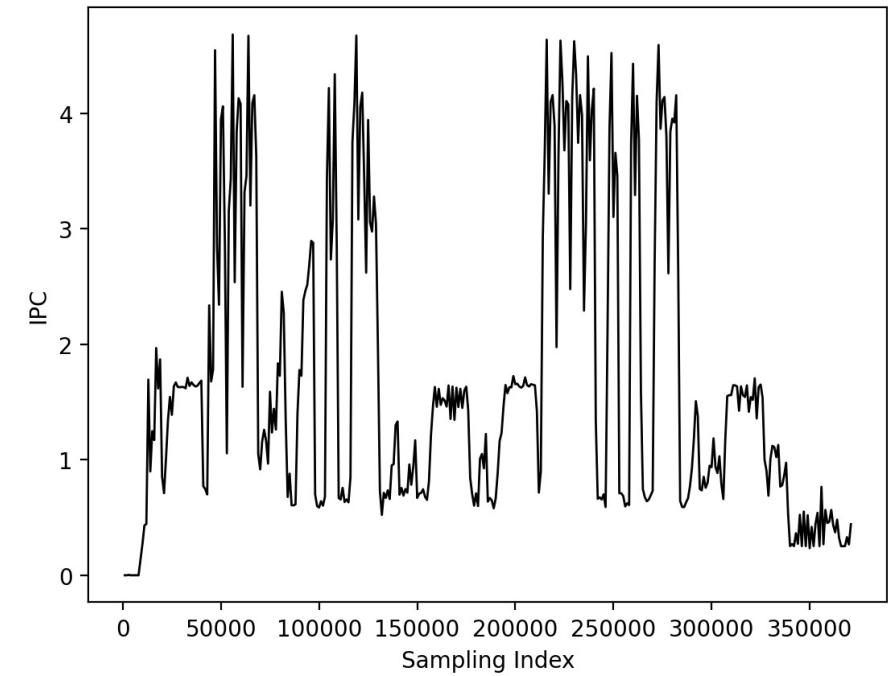
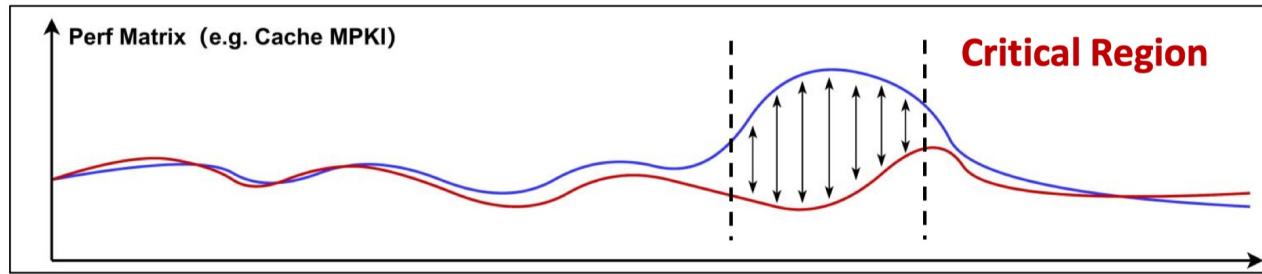
Find more parsing scripts at <https://github.com/OpenXiangShan/env-scripts/blob/main/perf/perf.py>



XSPerf: Rolling

- **Advantages**

- Visualization of performance analysis
- Detailed performance differences among segments
- Help identify parameter-sensitive critical regions





XSPerf: Rolling

- Usage

- Add `XSPERFROLLING('name', perfCnt, granularity, clock, reset)` wherever you want
- Compile with `WITH_CHISELDB=1 WITH_ROLLINGDB=1` , run with `--dump-db`

- Showcase

```
# Build emu with rolling (time consuming, use pre-built emu instead)
# bash xs-perf-prepare.sh
```

```
# cd ${NOOP_HOME}
# make clean
# make emu EMU_THREADS=4 WITH_CHISELDB=1 WITH_ROLLINGDB=1 -j8 \
#     PGO_WORKLOAD=${NOOP_HOME}/ready-to-run/coremark-2-iteration.bin \
#     PGO_MAX_CYCLE=10000 PGO_EMU_ARGS=--no-diff LLVM_PROFDATA=llvm-profdata
# ./build/emu -i ./ready-to-run/coremark-2-iteration.bin \
#     --diff ./ready-to-run/riscv64-nemu-interpreter-so --dump-db
# cp `find ${NOOP_HOME}/build/ -type f -name "*.db" | tail -1` \
#     ${XS_PROJECT_ROOT}/tutorial/p5-xs-perf/xs-perf-rolling.db
```

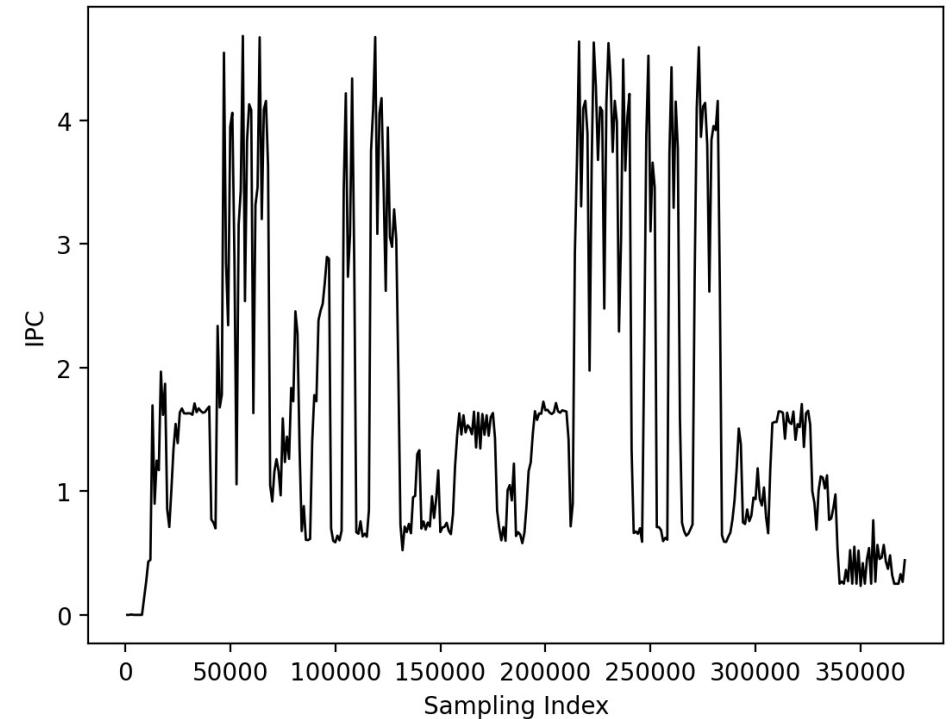


XSPerf: Rolling

- **Hands-on:** check the results and plot the curve of performance segments

```
# tutorial/p5-xs-perf/  
$ bash xs-perf-rolling.sh
```

```
# cd ${NOOP_HOME}/scripts/rolling  
# python3 rollingplot.py  
${XS_PROJECT_ROOT}/tutorial/p6-xs-  
perf/XsPerfRolling-test.db ipc  
# ls ${NOOP_HOME}/scripts/rolling/results  
  
perf.png
```



rolling curve example



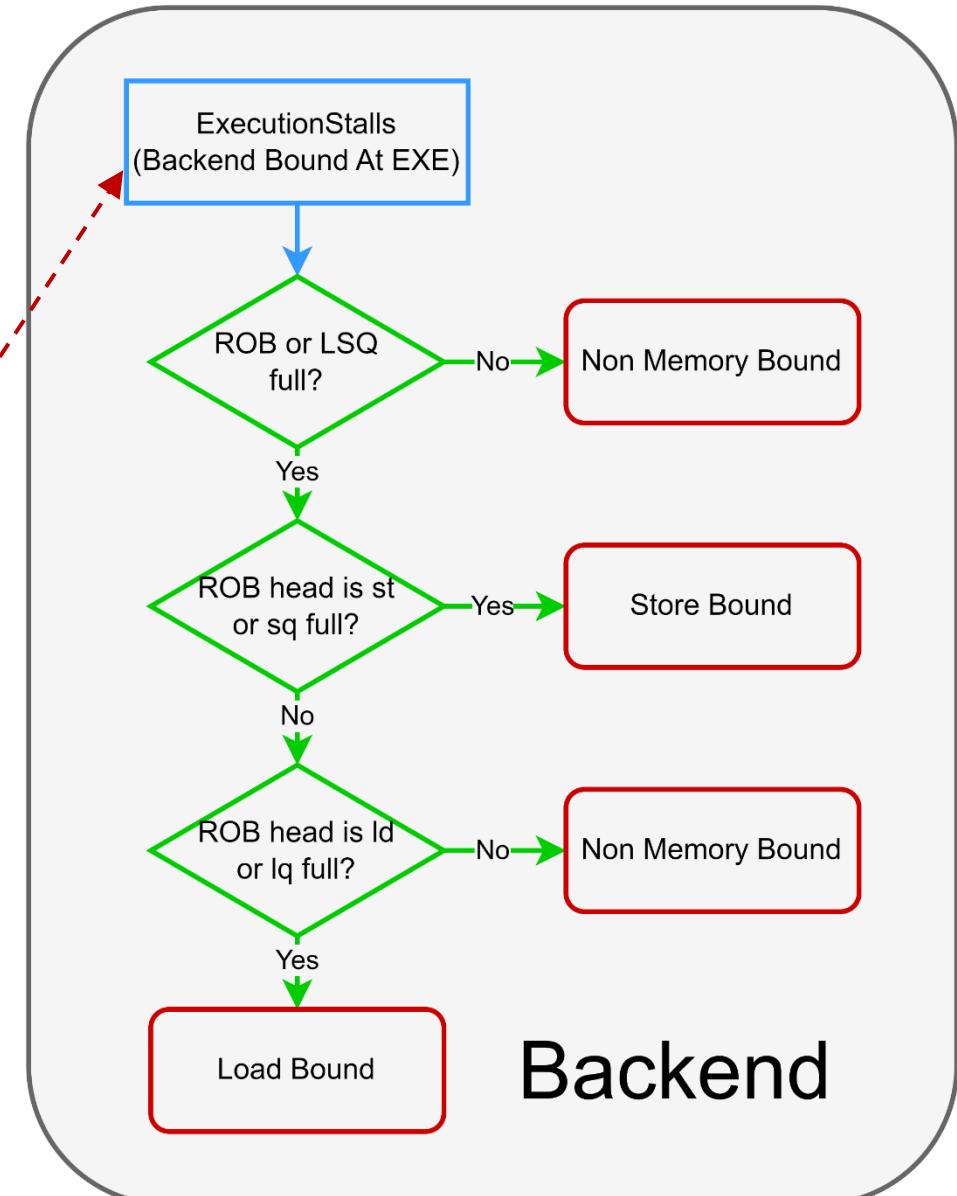
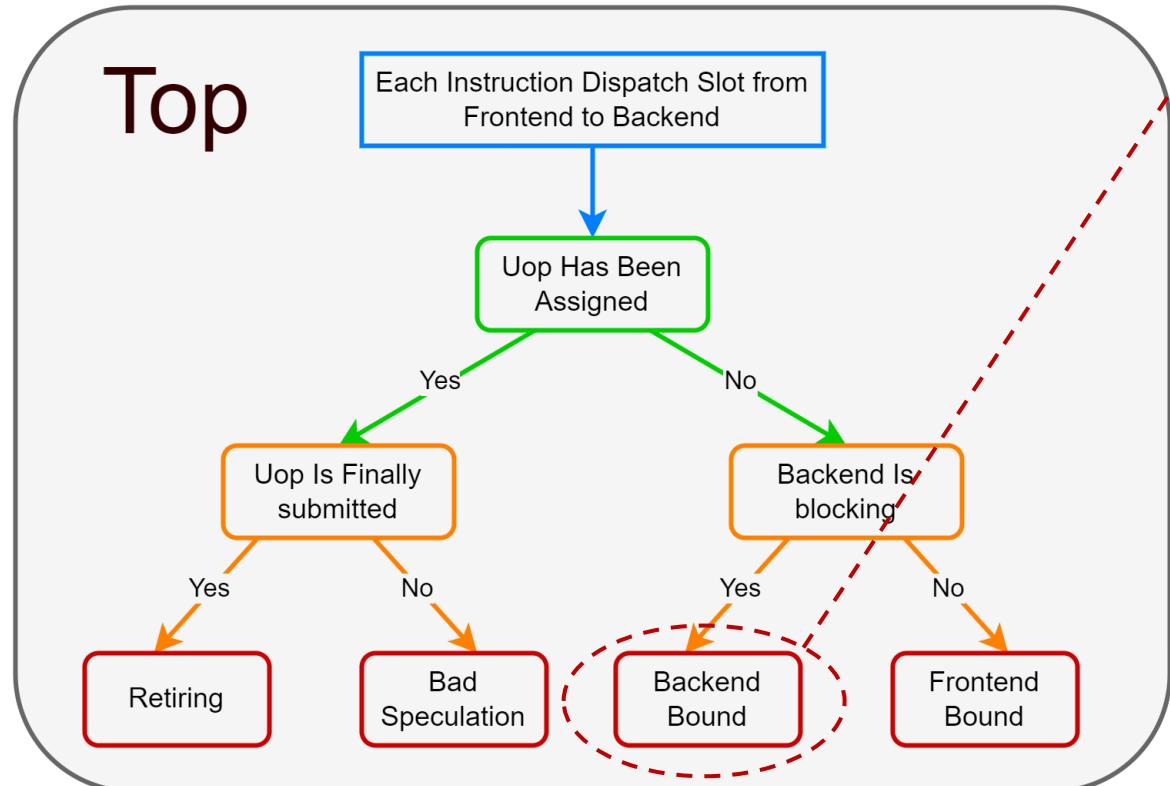
Top-Down: Method for Performance Analysis

- Purpose
 - Organize scattered performance events in a **hierarchical manner**
 - **Accurately** calculate one event's impact on processor performance
- We apply the **Top-Down** approach to XS-Gem5 and XiangShan
 - Complete **targeted optimization adaptation** for RISC-V instruction set
 - Optimize **performance counters** according to XiangShan microarchitecture
 - Further refine the **hierarchical design** of Top-Down model
 - Without missing or duplicating any events
 - Without assuming the blocking cycles of performance events in advance



Top-Down

- Use Top and Backend as examples





Top-Down in RTL

- Setting Performance Counters in RTL Code

```
XiangShan/src/main/scala/xiangshan/backend/Dispatch.scala

val stallReason = Wire(chiselTypeOf(io.stallReason.reason))
// ...
TopDownCounters.values.foreach(ctr =>
    XSPerfAccumulate(ctr.toString(), PopCount(stallReason.map(_ === ctr.id.U)))
)
```

- Do simulation and collect performance counter data



Top-Down in RTL

- Analyze the performance counter data through Top-Down method

```
XiangShan/scripts/top-down/configs.py
```

```
xs_coarse_rename_map = {
    'OverrideBubble': 'MergeFrontend',
    'FtqFullStall': 'MergeFrontend',
    'IntDqStall': 'MergeCoreDQStall', # attribute perf-counters to Top-Down hierarchy
    'FpDqStall': 'MergeCoreDQStall'
}
```

- Obtain analysis results and make targeted optimizations



Top-Down in RTL

- **Hands-on:** analyze the congestion causes in RTL using the Top-Down tool

```
#run the Top-Down analysis scripts (it takes ~2min)
```

```
$ bash rtl-top-down.sh
```

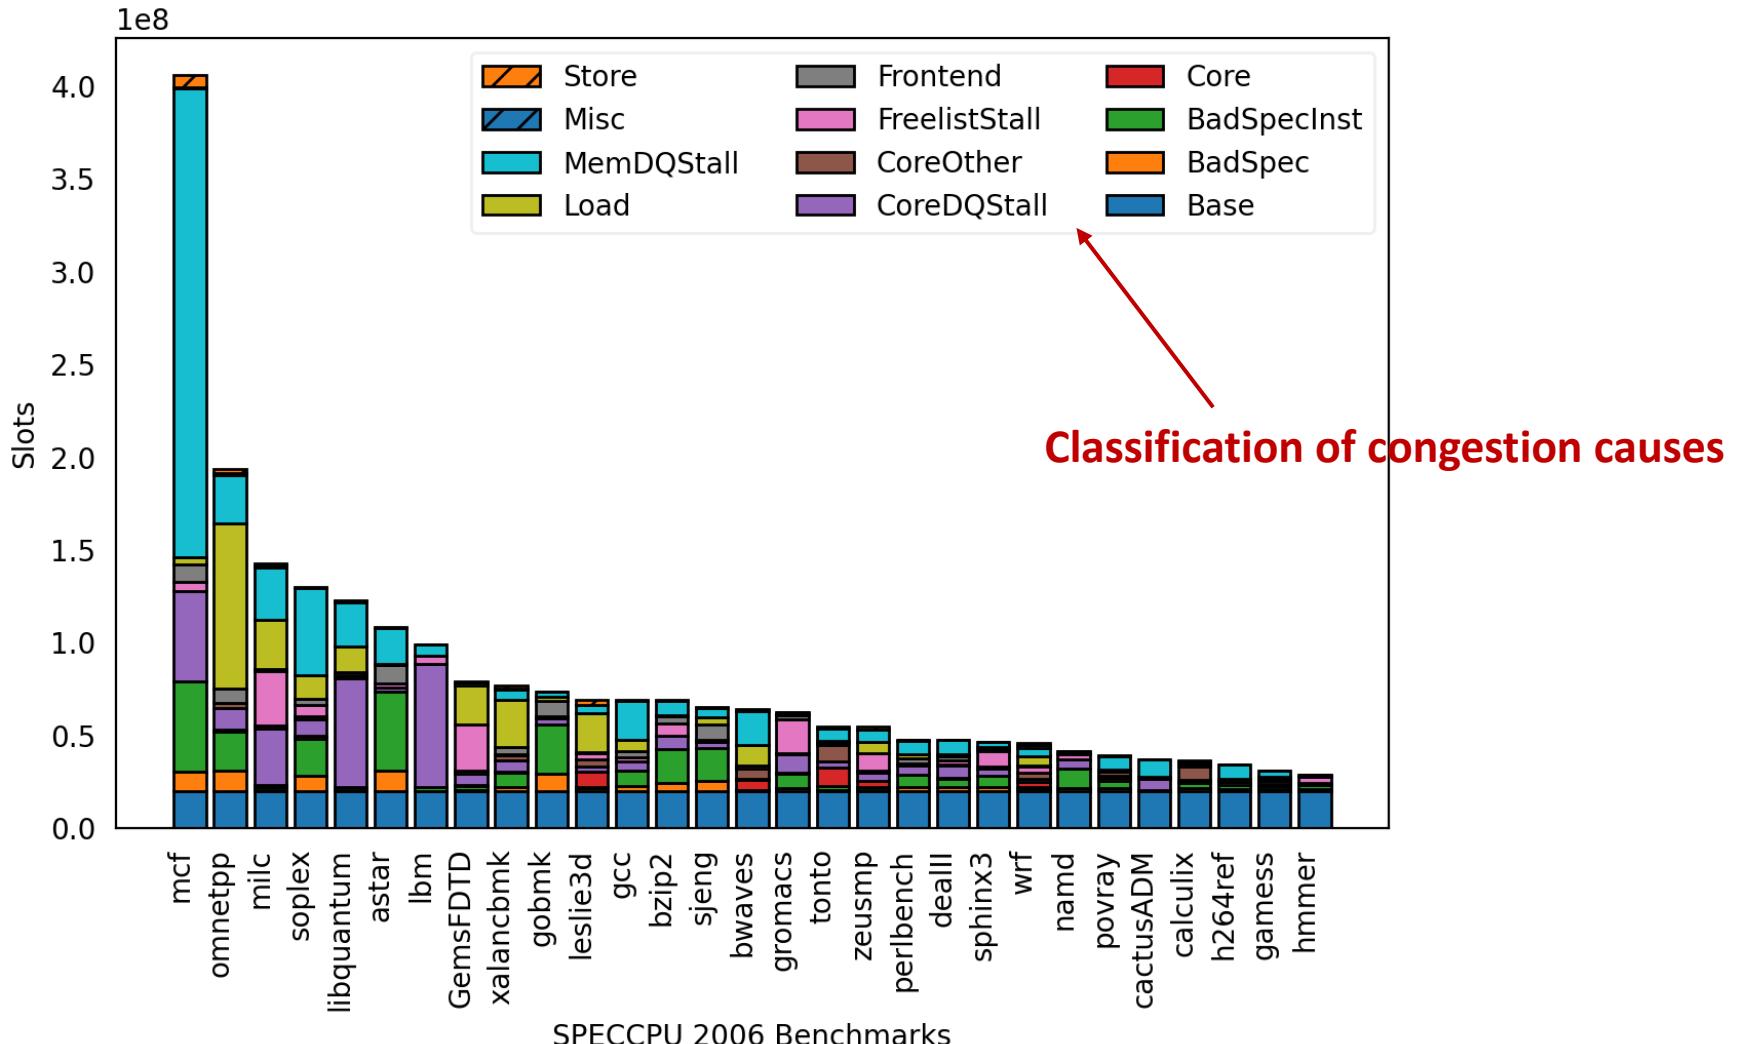
```
# cd ${NOOP_HOME}/scripts/top-down && \
python3 top_down.py \
-s /opt/SPEC06_EmuTasks_topdown \      #path of performance counter results
-j /opt/SPEC06_EmuTasks_topdown.json #json file of checkpoints configuration
```

```
# ls ${NOOP_HOME}/scripts/top-down/results
result.png  results.csv  results-weighted.csv #output of visual analysis results
```



Top-Down in RTL

- Result Example: Top-Down visual analysis results



- **Motivation**

- Want to test the performance under different parameters
- But re-compilation is time consuming
- Can we change parameters (constants) at runtime?

- **ConstantIn**

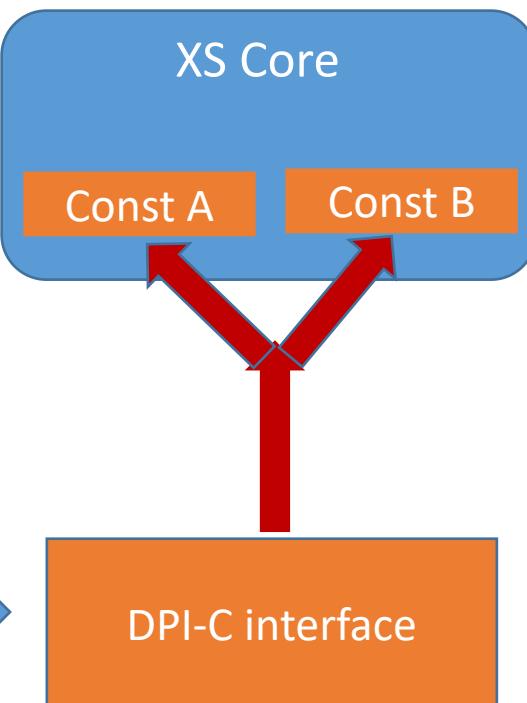
- DPI-C: Using C++ function in Chisel code (through Verilog using BlackBox)
- Signal values configured at runtime rather than compile time



ConstantIn

- Constants can be changed when initializing
- And remain unchanged during runtime

```
block_cycles_cache_0 = 11
block_cycles_cache_1 = 12
block_cycles_cache_2 = 13
block_cycles_cache_3 = 15
```



```
+ XiangShan_another.git:(constantinople) x cat build/constantin.txt
lowerbound 10
predSel 1
+ XiangShan_another.git:(constantinople) x ./build/emu -I ~/oracle/XiangShan_oracle/ready-to-run/microbench.bin -I 10000 -e 0 2> /dev/null
Emu compiled at Feb 6 2023, 11:24:27
Using simulated 32768B flash
No valid flash bin path, use preset flash instead
The image is /nfs/home/chenguokai/oracle/XiangShan_oracle/ready-to-run/microbench.bin
Using simulated 8192MB RAM
--diff is given, try to use $NEMU_HOME/build/riscv64-nemu-interpreter-so by default
NemuProxy using /nfs/home/chenguokai/NEMU_backup/build/riscv64-nemu-interpreter-so
The first instruction of core 0 has compiled. Diffest enabled.
[NEMU] Use built-in image
[src/device/io/mio.c:38,
mmapDeviceMemory, Microbench
{sort} Quick sort Core
total guest instructions
instrCnt = 10,001, cycleCnt
Seed#0 Guest cycle spent: 5,938ms
Host time spent: 5,938ms
Host time spent: 5,938ms
predSel 0
+ XiangShan_another.git:(constantinople) x ./build/emu -I 10000 -e 0 2> /dev/null
Emu compiled at Feb 6 2023, 11:24:27
Using simulated 32768B flash
No valid flash bin path, use preset flash instead
The image is /nfs/home/chenguokai/oracle/XiangShan_oracle/ready-to-run/microbench.bin
Using simulated 8192MB RAM
--diff is not given, try to use $NEMU_HOME/build/riscv64-nemu-interpreter-so by default
NemuProxy using /nfs/home/chenguokai/NEMU_backup/build/riscv64-nemu-interpreter-so
Assertion failed at line 722819.
The simulation stopped. There might be some assertion failed.
Core 0 Guest cycle spent: 6,019ms
total guest instructions = 0
instrCnt = 0, cycleCnt = 3, IPC = 0.000000
Seed#0 Guest cycle spent: 6 (this will be different from cycleCnt if emu loads a snapshot)
Host time spent: 4ms
```

```
+ XiangShan.git:(master) x export NEMU_HOME=/nfs/home/chenguokai/NEMU_backup
+ XiangShan.git:(master) x export NODP_HOME=/nfs/home/chenguokai/tutorial_demo/ConstantIn/XiangShan
+ XiangShan.git:(master) x python3 scripts/constantHelper.py constant.json
iteration 0 begins
[[{"block_cycles_cache_0": 51, "block_cycles_cache_1": 42}, {"block_cycles_cache_2": 10}, {"block_cycles_cache_3": 5}], [{"block_cycles.cache_0": 10}, {"block_cycles.cache_1": 43}, {"block_cycles.cache_2": 49}, {"block_cycles.cache_3": 21}], [{"block_cycles.cache_0": 15}, {"block_cycles.cache_1": 49}, {"block_cycles.cache_2": 43}, {"block_cycles.cache_3": 21}], [{"block_cycles.cache_0": 21}, {"block_cycles.cache_1": 49}, {"block_cycles.cache_2": 43}, {"block_cycles.cache_3": 21}], [{"block_cycles.cache_0": 21}, {"block_cycles.cache_1": 49}, {"block_cycles.cache_2": 41}, {"block_cycles.cache_3": 65}], [{"block_cycles.cache_0": 2}, {"block_cycles.cache_1": 1}, {"block_cycles.cache_2": 1}, {"block_cycles.cache_3": 1}], [{"block_cycles.cache_0": 3}, {"block_cycles.cache_1": 1}, {"block_cycles.cache_2": 1}, {"block_cycles.cache_3": 1}], [{"block_cycles.cache_0": 3}, {"block_cycles.cache_1": 1}, {"block_cycles.cache_2": 1}, {"block_cycles.cache_3": 1}], [{"block_cycles.cache_0": 17}], [{"block_cycles.cache_0": 1}, {"block_cycles.cache_1": 1}, {"block_cycles.cache_2": 1}, {"block_cycles.cache_3": 1}], [{"block_cycles.cache_0": 21}, {"block_cycles.cache_1": 17}, {"block_cycles.cache_2": 45}, {"block_cycles.cache_3": 65}], [{"block_cycles.cache_0": 21}, {"block_cycles.cache_1": 21}, {"block_cycles.cache_2": 19}, {"block_cycles.cache_3": 81}], [{"block_cycles.cache_0": 9}, {"block_cycles.cache_1": 69}, {"block_cycles.cache_2": 22}, {"block_cycles.cache_3": 111}], [{"block_cycles.cache_0": 14}, {"block_cycles.cache_1": 50}, {"block_cycles.cache_2": 29}, {"block_cycles.cache_3": 41}], [{"block_cycles.cache_0": 7}, {"block_cycles.cache_1": 39}, {"block_cycles.cache_2": 19}, {"block_cycles.cache_3": 111}], [{"block_cycles.cache_0": 13}, {"block_cycles.cache_1": 28}, {"block_cycles.cache_2": 6}, {"block_cycles.cache_3": 43}]

opt constant in this round
fitness is -41075

iteration 1 begins
[[{"block_cycles.cache_0": 1}, {"block_cycles.cache_1": 1}, {"block_cycles.cache_2": 1}, {"block_cycles.cache_3": 1}], [{"block_cycles.cache_0": 21}, {"block_cycles.cache_1": 17}, {"block_cycles.cache_2": 45}, {"block_cycles.cache_3": 65}], [{"block_cycles.cache_0": 21}, {"block_cycles.cache_1": 21}, {"block_cycles.cache_2": 19}, {"block_cycles.cache_3": 81}], [{"block_cycles.cache_0": 9}, {"block_cycles.cache_1": 69}, {"block_cycles.cache_2": 22}, {"block_cycles.cache_3": 111}], [{"block_cycles.cache_0": 14}, {"block_cycles.cache_1": 50}, {"block_cycles.cache_2": 29}, {"block_cycles.cache_3": 41}], [{"block_cycles.cache_0": 7}, {"block_cycles.cache_1": 39}, {"block_cycles.cache_2": 19}, {"block_cycles.cache_3": 111}], [{"block_cycles.cache_0": 13}, {"block_cycles.cache_1": 28}, {"block_cycles.cache_2": 6}, {"block_cycles.cache_3": 43}]

opt constant in this round
fitness is -41075

opt constant for gene algorithm is [[{"block_cycles.cache_0": 11}, {"block_cycles.cache_1": 18}, {"block_cycles.cache_2": 127}, {"block_cycles.cache_3": 17}], [{"block_cycles.cache_0": 11}, {"block_cycles.cache_1": 58}, {"block_cycles.cache_2": 127}, {"block_cycles.cache_3": 17}], [{"block_cycles.cache_0": 11}, {"block_cycles.cache_1": 28}, {"block_cycles.cache_2": 6}, {"block_cycles.cache_3": 43}]

+ XiangShan.git:(master) x
```

- Usage: Create signals

```
# XiangShan/coupledL2/src/main/scala/coupledL2/prefetch/TemporalPrefetch.scala
```

```
139  require(cacheParams.hartIds.size == 1)
140  val hartid = cacheParams.hartIds.head
141  // 0 / 1: whether to enable temporal prefetcher
142  private val enableTP = WireInit(Constantin.createRecord("enableTP"+hartid.toString, initialValue = 1.U))
143  // 0 ~ N: throttle cycles for each prefetch request
144  private val tpThrottleCycles = WireInit(Constantin.createRecord("tp_throttleCycles"+hartid.toString, initialValue = 4.U(3.W)))
145  // 0 / 1: whether request to set as trigger on meta hit
146  private val hitAsTrigger = WireInit(Constantin.createRecord("tp_hitAsTrigger"+hartid.toString, initialValue = 1.U))
147  // 1 ~ triggerQueueDepth: enqueue threshold for triggerQueue
148  private val triggerThres = WireInit(Constantin.createRecord("tp_triggerThres"+hartid.toString, initialValue = 1.U(3.W)))
149  // 1 ~ tpEntryMaxLen: record threshold for recorder and sender (storage size will not be affected)
150  private val recordThres = WireInit(Constantin.createRecord("tp_recordThres"+hartid.toString, initialValue = tpEntryMaxLen.U))
151  // 0 / 1: whether to train on vaddr
152  private val trainOnVaddr = WireInit(Constantin.createRecord("tp_trainOnVaddr"+hartid.toString, initialValue = 0.U))
153  // 0 / 1: whether to eliminate L1 prefetch request training
154  private val trainOnL1PF = WireInit(Constantin.createRecord("tp_trainOnL1PF"+hartid.toString, initialValue = 0.U))
```



ConstantIn

- **Hands-on:** Pass constant via standard input stream
 - **By default**, set constant via `${NOOP_HOME}/build/constantin.txt`
 - **For better demonstration**, it is changed to stdin (corresponding modifications can refer to `p6-constantin/utility.patch`)

```
$ cd ../p6-constantin
```

```
# bash step0-build.sh      # Build emu with constantin (time consuming, use pre-built emu instead)
```

```
$ bash step1-basic.sh      # Run emu and pass constant
please input total constant number
2
please input each constant ([constant name] [value])
DelayQueueLatencyvbop 175
DelayQueueLatencycypbop 175
```



ConstantIn

- **Feature:** *AutoSolving* for automatically finding better constant values
- **Process**
 - Enable *AutoSolving* & prepare signals you need
 - Compile the emu with `WITH_CONSTANTIN=1`
 - Run basic demonstration
 - Set parameters for *AutoSolving* in a configuration file
 - Automatically find

- Feature: *AutoSolving* for automatically finding better constant values

```
# Automatical parameter solver configuration file
$ cat my_constantin.json
```

- Parameters defined

- properties of signals
- target of *AutoSolving*
- parameters of genetic algorithm
- parameters of XS simulator

```
{
  "constants": [
    {
      "name": "DelayQueueLatencyvbop",
      "width": 8,
      "guide": 256,
      "init": 177
    },
    {
      "name": "DelayQueueLatencypbop",
      "width": 8,
      "guide": 256,
      "init": 242
    }
  ],
  "opt_target": [
    {"l2cache.topDown: L2MissMatch,": {"policy": "min", "baseline": 409}},
    {"backend.ctrlBlock.rob: clock_cycle,": {"policy": "min", "baseline": 217759}}
  ],
  "population_num": 4,
  "iteration_num": 2,
  "crossover_rate": 50,
  "mutation_rate": 50,
  "emu_threads": 8,
  "concurrent_emu": 4,
  "max_instr": 100000,
  "seed": 3888,
  "work_load": "${XS_PROJECT_ROOT}/tutorial/p7-constantin/maprobe-riscv64-xs.bin"
}
```



ConstantIn

- **Hands-on:** *AutoSolving* example

```
# Run auto solver, output will be the optimal constant currently found  
$ bash step2-solve.sh
```

```
# ~2min  
# The solver generates optimal parameters  
opt constant in this round is [['DelayQueueLatencyvbop', 8], ['DelayQueueLatencypbop', 72]] fitness is 54674  
opt constant for gene algrithom is [['DelayQueueLatencyvbop', 9], ['DelayQueueLatencypbop', 73]] fitness 54674
```



XS-Gem5: Architecture Performance Evaluation Tool

- Purpose
 - Efficiently iterate on microarchitecture features
 - Conduct fast end-to-end performance evaluations
- XS-Gem5
 - Align with XiangShan as much as possible based on Gem5 O3CPU
 - Support Difftest and Checkpoint
 - Add scripts for XiangShan performance analysis

- **Hands-on:** compile and build XS-Gem5

```
$ cd ../../p7-xs-gem5 # Enter XS-Gem5 tutorial directory  
$ bash 0-gem5_prepare.sh # ~8min, time consuming, use pre-built emu instead  
$ cd /home/guest/liyanqin/tutorial/p7-xs-gem5 && export gem5_home=`pwd`
```

```
prepare_gem5() {  
    pushd $gem5_home && \  
    cd ext/dramsim3 && \ # build DRAMSim3  
    git clone git@github.com:umd-memsys/DRAMSim3.git DRAMsim3 && \  
    cd DRAMsim3 && mkdir -p build && cd build && cmake .. && make -j `nproc` && \  
    popd  
}  
build_gem5() {  
    pushd $gem5_home && \  
    scons build/RISCV/gem5.opt --gold-linker -j `nproc` && \  
    popd  
}
```

- **Hands-on:** run bare-metal workload on XS-Gem5

```
$ bash 1-gem5_run_coremark.sh
```

```
#~9s
pushd $gem5_home && \
export GCBV_REF_SO=$NEMU_HOME/build/riscv64-nemu-gem5-ref-so && \
mkdir -p util/xs_scripts/coremark && \ # Setup CoreMark working directory
cd util/xs_scripts/coremark && \
$gem5_home/build/RISCV/gem5.opt $gem5_home/configs/example/xiangshan.py \
--raw-cpt \
--generic-rv-cpt=$NOOP_HOME/ready-to-run/coremark-2-iteration.bin && \
popd
```

```
ecs-user@xiangshan-tutorial:~/chenyangyu/xs-env/tutorial/p5-xs-gem5 ⌂⌘1

coremark-2-iteration.bin to pmem 0x7f72654ee000
build/RISCV/mem/physical.cc:608: info: First 4 bytes are 0x13 0x4 0x0 0x0
build/RISCV/sim/system.cc:561: info: Restored from Xiangshan RISC-V Checkpoint
**** REAL SIMULATION ****
build/RISCV/sim/simulate.cc:194: info: Entering event queue @ 0. Starting simulation...
build/RISCV/cpu/base.cc:1266: warn: Start memcpy to NEMU from 0x7f72654ee000, size=8589934592
build/RISCV/cpu/base.cc:1269: warn: Start regcpy to NEMU
build/RISCV/dev/serial/uartlite.cc:35: warn: Write to other uartlite addr 12 is not implemented
Running CoreMark for 2 iterations
2K performance run parameters for coremark.
CoreMark Size      : 666
Total time (ms)   : 110
Iterations        : 2
Compiler version  : GCC10.2.0
seedcrc           : 0xe9f5
[0]crclist        : 0xe714
[0]crcmatrix      : 0x1fd7
[0]crcstate        : 0x8e3a
[0]crcfinal        : 0x72be
Finised in 110 ms.
=====
CoreMark Iterations/Sec 18181
Exiting @ tick 313394625 because m5_exit instruction encountered when simulating XS
~/chenyangyu/xs-env/tutorial/p5-xs-gem5
→ p5-xs-gem5 git:(tutorial-2024) x █
```

- **Hands-on:** analyze the performance counters

```
$ bash 2-gem5_counter.sh | head -n 20
```

```
cat $gem5_home/util/xs_scripts/coremark/m5out/stats.txt
```



```
ecs-user@xiangshan-tutorial:~/chenyangyu/xs-env/tutorial/p5-xs-gem5 ➜ p5-xs-gem5 git:(tutorial-2024) ✘ ./2-gem5_counter.sh | shuf -n 20 | cut -f1 -d'#'  
system.l3.prefetcher.pfHitInWB          0  
system.cpu.dcache.prefetcher.berti.pfUnused_srcs::11      0  
system.l3.prefetcher.pfHitInCache_srcs::10      0  
system.cpu.dcache.ReadReq.misses::cpu.data      161  
system.cpu.iew.renameStallReason::DTlbStall      10  
system.cpu.mmudtb.l2sptlbUnusedRemove      0  
system.l2_caches.InvalidateReq.missRate::cpu.data      1  
system.cpu.mmudtb.l2_shared.writeL2l3TlbMisses      0  
system.cpu.dcache.prefetcher.spp.pfUseful_srcs::9      0  
system.cpu.branchPred.commitControlSquashLatencyDist::2      0      0.00%      0.00%  
system.l2_caches.ReadExReq.mshrMisses::cpu.dcache.prefetcher      8  
system.cpu.branchPred.tage.bank_0.updateTableMispreds::1      211  
system.cpu.dcache.prefetcher.bop_large.pfUseful_srcs::9      0  
system.cpu.dcache.prefetcher.bop_learned.pfUsefulSpanPage      0  
system.cpu.dcache.prefetcher.pfUseful      0  
system.cpu.dcache.prefetcher.berti.pfHitInWB_srcs::3      0  
system.cpu.dcache.prefetcher.cmc.pfHitInCache_srcs::1      0  
system.mem_ctrls.bwInstRead::cpu.inst      39821998  
system.l2_caches.demandMisses::cpu.data      49  
system.cpu.icache.overallAvgMissLatency::cpu.inst 52512.614493  
➜ p5-xs-gem5 git:(tutorial-2024) ✘
```



XS-Gem5: Performance Analysis Examples

- **Hands-on:** run the Cache MPKI analysis script

```
$ bash 3-gem5_cache.sh
```

```
pushd gem5_data_proc && \
mkdir -p results && \
export PYTHONPATH=`pwd` && \
python3 batch.py \
-s /home/share/xs-model-l1bank \
-o gem5-cache-example.csv \
--cache && \
python3 simpoint_cpt/compute_weighted.py \
-r gem5-cache-example.csv \
-j /home/share/xs-model-l1bank/cluster-0-0.json \
-o weighted.csv
```

Note: here we take the results from SPEC CPU 2006 checkpoints as an example for better showcase.

 Cache MPKI

```
ecs-user@xiangshan-tutorial:~/chenyangyu/xs-env/tutorial/p5-xs-gem5 ⌂⌘1
['GemsFDTD', 'astar', 'bwaves', 'bzip2', 'cactusADM', 'calculix', 'dealII', 'gamess', 'gcc', 'gobmk', 'gromacs', 'h264ref', 'hmmer', 'lmb', 'leslie3d', 'libquantum', 'mcf', 'milc', 'namd', 'omnetpp', 'perlbench', 'povray', 'sjeng', 'soplex', 'sphinx3', 'tonto', 'wrf', 'xalancbmk', 'zeusmp']
   Cycles    Insts   L1D.MPKI   L2.MPKI ...   l3_acc   l3_miss   cpi   coverage
mcf      3.027e+07  2.000e+07   270.673  100.326 ...  2.294e+06  698680.709  1.513   1.0
omnetpp  1.774e+07  2.000e+07    57.713  47.026 ...  1.026e+06  492397.438  0.887   1.0
astar    1.618e+07  2.000e+07   27.252  15.890 ...  3.543e+05  48744.997  0.809   1.0
gobmk    1.155e+07  2.000e+07    5.485   1.525 ...  3.856e+04  14055.195  0.577   1.0
soplex   1.119e+07  2.000e+07   33.391  25.666 ...  9.042e+05  175472.264  0.559   1.0
milc     1.071e+07  2.000e+07   68.772  27.739 ...  9.773e+05  702851.431  0.535   1.0
bzip2    9.521e+06  2.000e+07   12.069  4.226 ...  9.716e+04  4578.537  0.476   1.0
gromacs  9.386e+06  2.000e+07   32.104  0.658 ...  1.976e+04  7220.444  0.469   1.0
gcc      9.261e+06  2.000e+07   13.926  15.001 ...  4.285e+05  114394.979  0.463   1.0
sjeng    9.021e+06  2.000e+07   1.439   1.211 ...  2.422e+04  20958.788  0.451   1.0
bwaves   8.071e+06  2.000e+07   6.480   6.880 ...  5.694e+05  89426.320  0.404   1.0
perlbench 7.967e+06  2.000e+07   4.548   1.431 ...  3.134e+04  19389.243  0.398   1.0
lmb      7.920e+06  2.000e+07   8.029   19.248 ...  6.969e+05  122392.254  0.396   1.0
zeusmp   7.450e+06  2.000e+07   12.428  5.687 ...  1.894e+05  85596.788  0.373   1.0
tonto    6.759e+06  2.000e+07   5.334   2.342 ...  8.166e+04  5793.963  0.338   1.0
calculix 6.720e+06  2.000e+07   0.932   0.340 ...  8.947e+03  3756.312  0.336   1.0
GemsFDTD 6.559e+06  2.000e+07   12.818  14.021 ...  5.250e+05  374770.295  0.328   1.0
leslie3d  6.552e+06  2.000e+07   21.274  10.611 ...  3.842e+05  144669.463  0.328   1.0
xalancbmk 6.515e+06  2.000e+07   19.989  11.652 ...  3.246e+05  38865.836  0.326   1.0
namd     6.494e+06  2.000e+07   14.142  0.212 ...  5.496e+03  4312.841  0.325   1.0
povray   6.162e+06  2.000e+07   17.487  0.195 ...  3.944e+03  3190.261  0.308   1.0
dealII   5.944e+06  2.000e+07    8.144  1.365 ...  9.003e+04  10105.618  0.297   1.0
sphinx3  5.736e+06  2.000e+07   14.614  11.872 ...  3.751e+05  20540.266  0.287   1.0
h264ref  4.941e+06  2.000e+07   3.710   1.137 ...  3.053e+04  10893.001  0.247   1.0
wrf      4.908e+06  2.000e+07   2.839   1.865 ...  1.151e+05  12708.207  0.245   1.0
hmmer   4.489e+06  2.000e+07   1.042   0.818 ...  2.653e+04  494.690  0.224   1.0
cactusADM 4.474e+06  2.000e+07   2.154   2.749 ...  9.884e+04  34987.191  0.224   1.0
libquantum 4.390e+06  2.000e+07   1.185   1.214 ...  5.120e+05  69486.561  0.220   1.0
gamess   4.261e+06  2.000e+07   2.287   0.169 ...  3.496e+03  2324.679  0.213   1.0

[29 rows x 17 columns]
~/chenyangyu/xs-env/tutorial/p5-xs-gem5
→ p5-xs-gem5 git:(tutorial-2024-gem5-reorder) x
```



XS-Gem5: Performance Analysis Example

- **Hands-on:** run the Top-Down analysis script

```
$ bash 4-gem_topdown.sh
```

```
pushd gem5_data_proc && \
python3 batch.py \
-s $gem5_home/util/xs_scripts/coremark \
-t --topdown-raw && \
popd
```



Top-Down

```
ecs-user@xiangshan-tutorial:~/chenyangyu/xs-env/tutorial/p5-xs-gem5 ⌂⌘1
system.l2_caches.demandMisses::cpu.data          49
system.cpu.icache.overallAvgMissLatency::cpu.inst 52512.614493
→ p5-xs-gem5 git:(tutorial-2024) ✘ ./3-gem_topdown.sh
~/chenyangyu/xs-env/tutorial/p5-xs-gem5/gem5_data_proc ~/chenyangyu/xs-env/tutorial/p5-xs-gem5
workload: m5out, point: 0, segments: 1
[('m5out_0', '/home/ecs-user/chenyangyu/xs-env/gem5/util/xs_scripts/coremark/m5out/stats.txt')]
Process finished job: m5out_0
/home/ecs-user/chenyangyu/xs-env/gem5/util/xs_scripts/coremark/m5out/stats.txt
{'m5out_0': {'OtherFetchStall': 0, 'OtherStall': 0, 'CommitSquash': 328812, 'ResumeUnblock': 0, 'OtherMemStall': 0, 'Atomic': 0, 'MemCommitRateLimit': 2058, 'MemNotReady': 1576428, 'MemSquashed': 0, 'StoreMemBound': 0, 'StoreL3Bound': 0, 'StoreL2Bound': 0, 'StoreL1Bound': 930, 'LoadMemBound': 4746, 'LoadL3Bound': 0, 'LoadL2Bound': 0, 'LoadL1Bound': 756, 'InstNotReady': 12, 'SerializeStall': 63390, 'InstSquashed': 130470, 'InstMisPred': 0, 'FetchBufferInvalid': 0, 'SquashStall': 33504, 'FragStall': 1721534, 'TrapStall': 0, 'IntStall': 0, 'BpStall': 186960, 'DTlbStall': 6, 'ITlbStall': 0, 'IcacheStall': 282756, 'NoStall': 951427, 'ipc': 0.66608, 'Insts': 626869, 'Cycles': 941131, 'point': '0', 'workload': 'm5out', 'bmk': 'm5out'}}}
      Atomic  BpStall  CommitSquash  Cycles  ...   bmk    ipc  point  workload
m5out_0        0     186960       328812  941131  ...  m5out  0.666      0     m5out

[1 rows x 37 columns]
~/chenyangyu/xs-env/tutorial/p5-xs-gem5
→ p5-xs-gem5 git:(tutorial-2024) ✘
```



XS-Gem5: Performance Analysis Example

- **Hands-on:** run the SPEC CPU score calculation script

```
$ bash 5-gem5_spec06_score.sh
```

```
pushd gem5_data_proc && \
mkdir -p results && \
export PYTHONPATH=`pwd` && \
python3 batch.py -s ../data/xs-model-l1bank \
-o gem5-score-example.csv && \
python3 simpoint_cpt/compute_weighted.py \
-r gem5-score-example.csv \
-j ../data/xs-model-l1bank/cluster-0-0.json \
--score score.csv
```

Note: here we take the results from SPEC CPU 2006 checkpoints as an example for better showcase.



SPEC CPU score

tmux							
wrf	250.853	11170.0	14.843	[25/2118]	h264ref	428.160	
leslie3d	214.233	9400.0	14.626	1.0	omnetpp	139.743	
gamess	451.597	19580.0	14.452	1.0	astar	229.743	
namd	210.250	8020.0	12.715	1.0	xalancbmk	85.528	
tonto	260.012	9840.0	12.615	1.0	Estimated Int score per GHz: 15.118648674723131		
hmmer	260.274	9330.0	11.949	1.0	Estimated Int score @ 3.0GHz: 45.35594602416939		
gromacs	201.133	7140.0	11.833	1.0	===== FP =====		
perlbench	289.646	9770.0	11.244	1.0	bwaves	175.907	
sjeng	366.934	12100.0	10.992	1.0	gamess	451.597	
gobmk	320.031	10490.0	10.926	1.0	milc	142.358	
astar	229.743	7020.0	10.185	1.0	zeusmp	196.183	
bzip2	410.138	9650.0	7.843	1.0	gromacs	201.133	
calculix	490.428	8250.0	5.607	1.0	cactusADM	252.748	
score.csv					leslie3d	214.233	
===== SPEC06 =====							
===== Int =====							
		time	ref_time	score	coverage		
perlbench	289.646	9770.0	11.244	1.0	dealII	144.375	
bzip2	410.138	9650.0	7.843	1.0	soplex	136.694	
gcc	167.812	8050.0	15.990	1.0	povray	93.311	
mcf	141.107	9120.0	21.544	1.0	calculix	490.428	
gobmk	320.031	10490.0	10.926	1.0	GemsFDTD	176.662	
hmmer	260.274	9330.0	11.949	1.0	tonto	260.012	
sjeng	366.934	12100.0	10.992	1.0	lbm	149.280	
libquantum	148.529	20720.0	46.500	1.0	wrf	250.853	
h264ref	428.160	22130.0	17.229	1.0	sphinx3	350.359	
omnetpp	139.743	6250.0	14.908	1.0	Estimated FP score per GHz: 16.548846674376477		
astar	229.743	7020.0	10.185	1.0	Estimated FP score @ 3.0GHz: 49.64654002312943		
xalancbmk	85.528	6900.0	26.892	1.0	----- Overall -----		
Estimated Int score per GHz: 15.118648674723131							
Estimated Int score @ 3.0GHz: 45.35594602416939							
Estimated overall score per GHz: 15.941323843383495							
Estimated overall score @ 3.0GHz: 47.82397153015049							
~/chenhyangyu/xs-env/tutor-tac/p5-xs-gems							
→ p5-xs-gem5 git:(tutorial-2024-gem5-reorder) ✘							
"xiangshan-tutorial" 21:44 20-Aug-24							

[0] 0:[tmux]*

"xiangshan-tutorial" 21:44 20-Aug-24



Callback: Check your own emu

- Return to the first terminal where you compiled the emu
- Run simulation with your own emulator

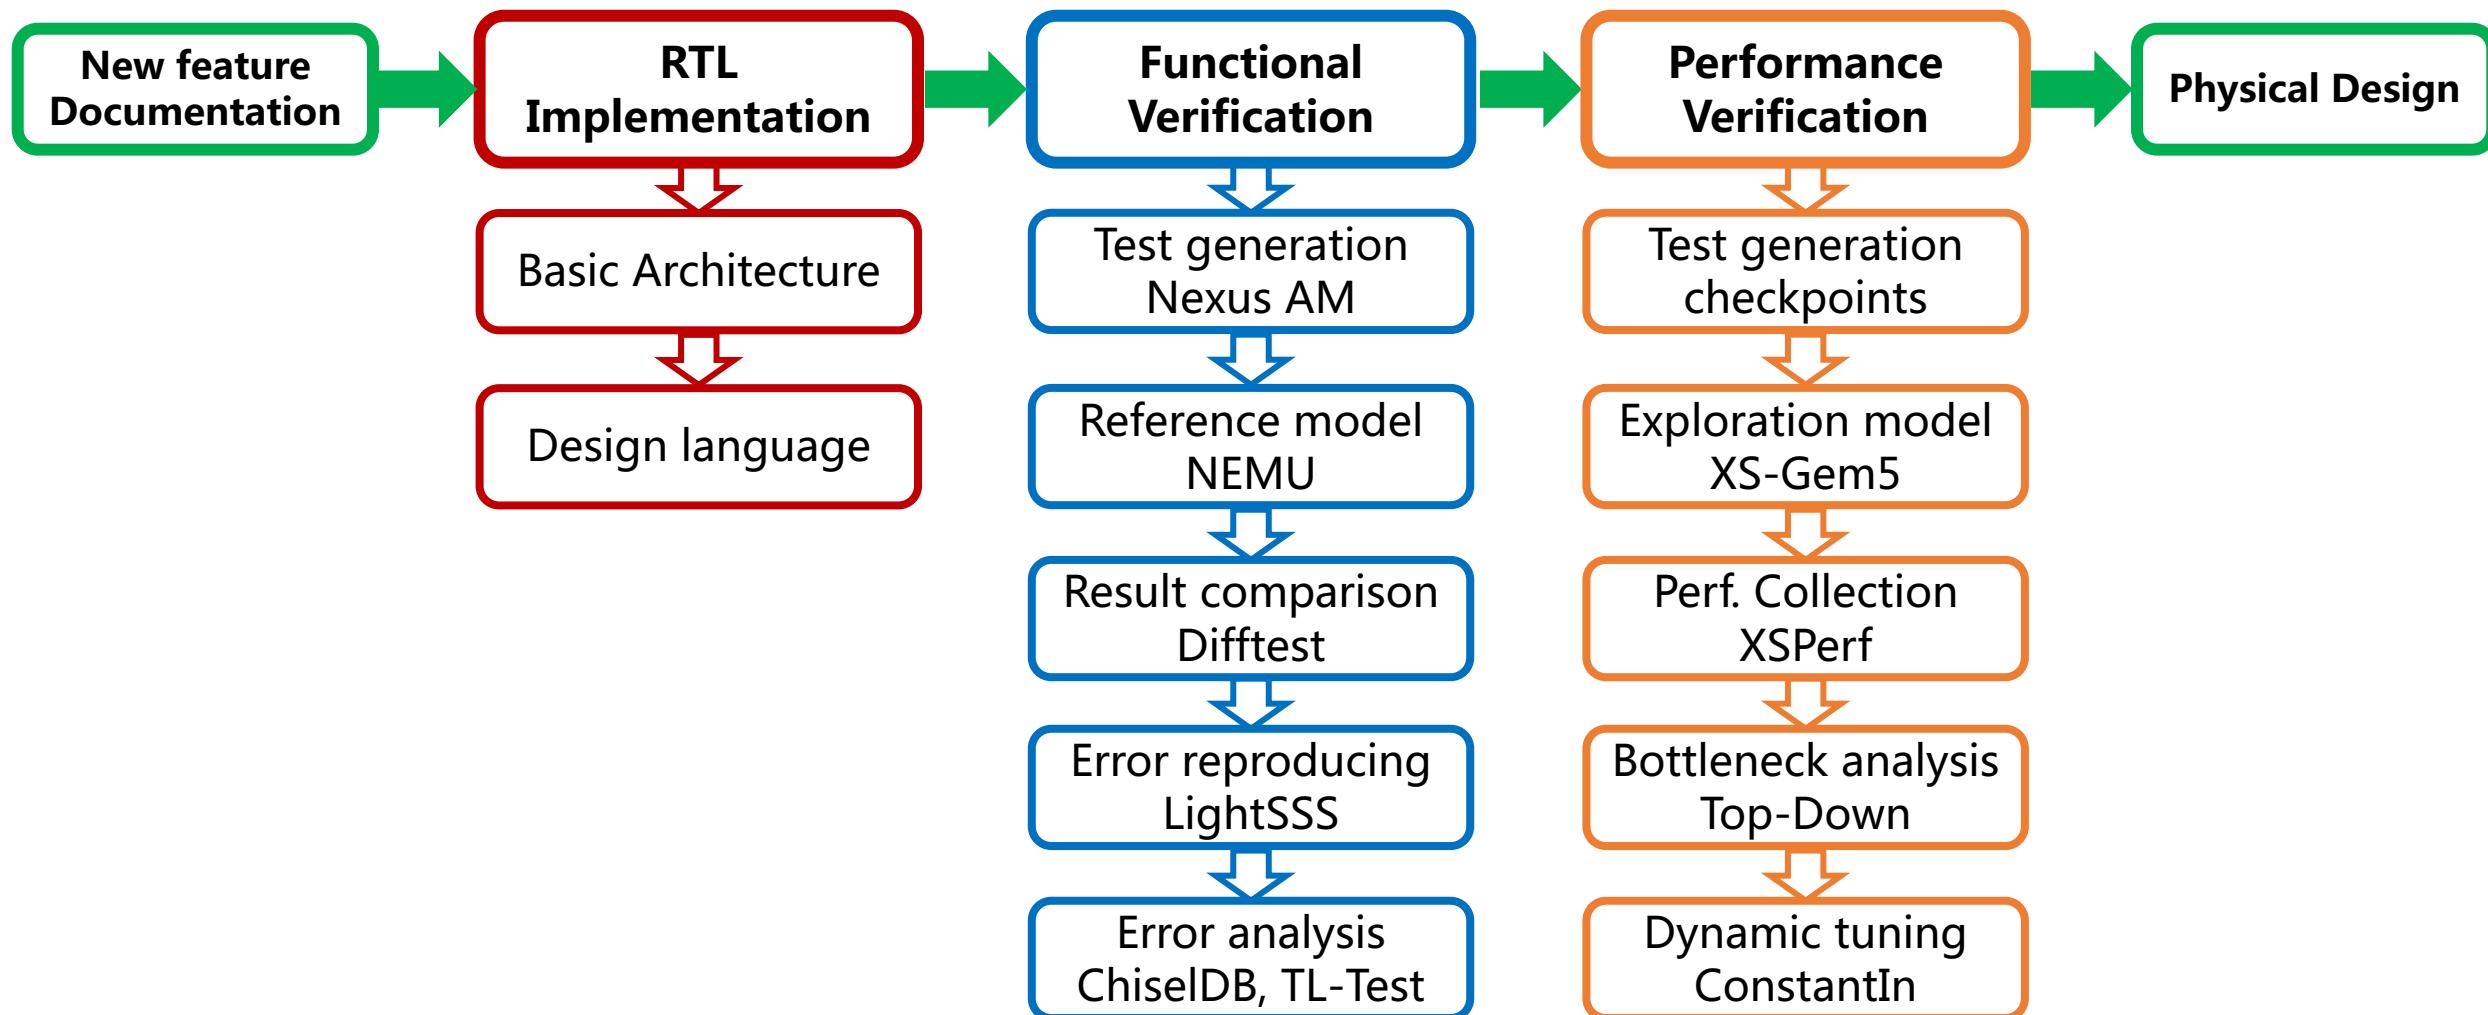
```
$ cd $NOOP_HOME  
$ ./build/emu --help          (automatically runs with EMU_THREADS threads)  
# Some Options:  
  -i                         Workload to run  
  -C / -I / -W                Max cycles / Max Insts / Warmup Insts  
  --diff=PATH / --no-diff     Compare with NEMU for difftest / disable difftest
```

Example: Use the tab key for command completion

```
$ ./build/emu -i ../tutorial/p1-basic-func/hello.bin --diff ./ready-to-  
run/riscv64-nemu-interpreter-so 2> perf.err
```



Summary: MinJie Development Flows and Tools



Thanks!